



NumPyWren

Storage-enabled Scaling of
Serverless Supercomputing



Vaishaal Shankar



Karl Krauth



Qifan Pu



Shivaram
Venkataraman



Ion
Stoica



Ben
Recht

Jonathan
Ragan-Kelly



Berkeley **C**enter for
Computational **I**maging

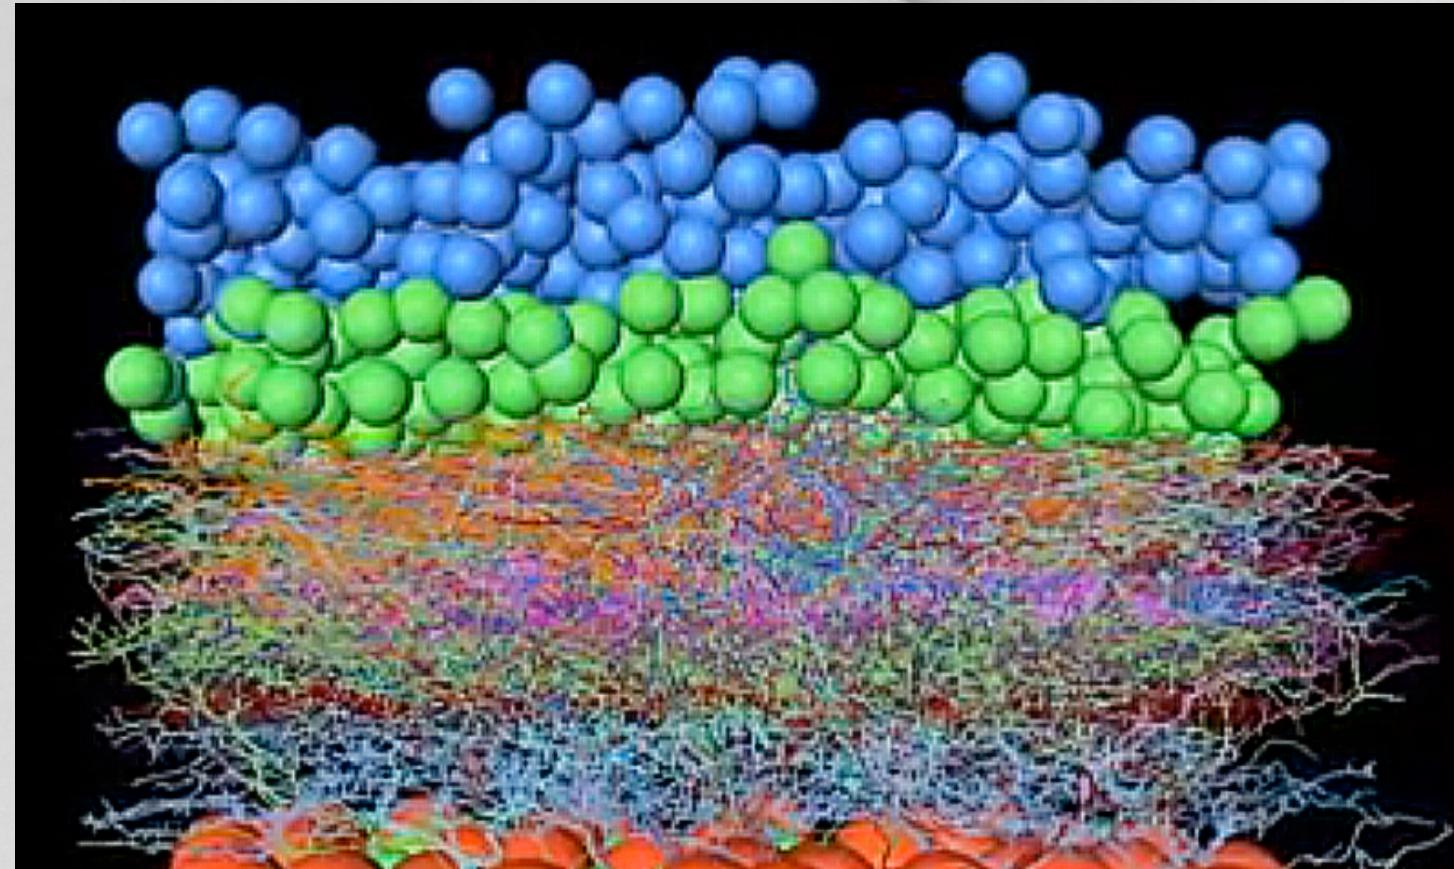
Eric Jonas
Postdoctoral Researcher
jonas@eecs.berkeley.edu
[@stochastician](https://twitter.com/stochastician)

PyWren: Scale For Everyone

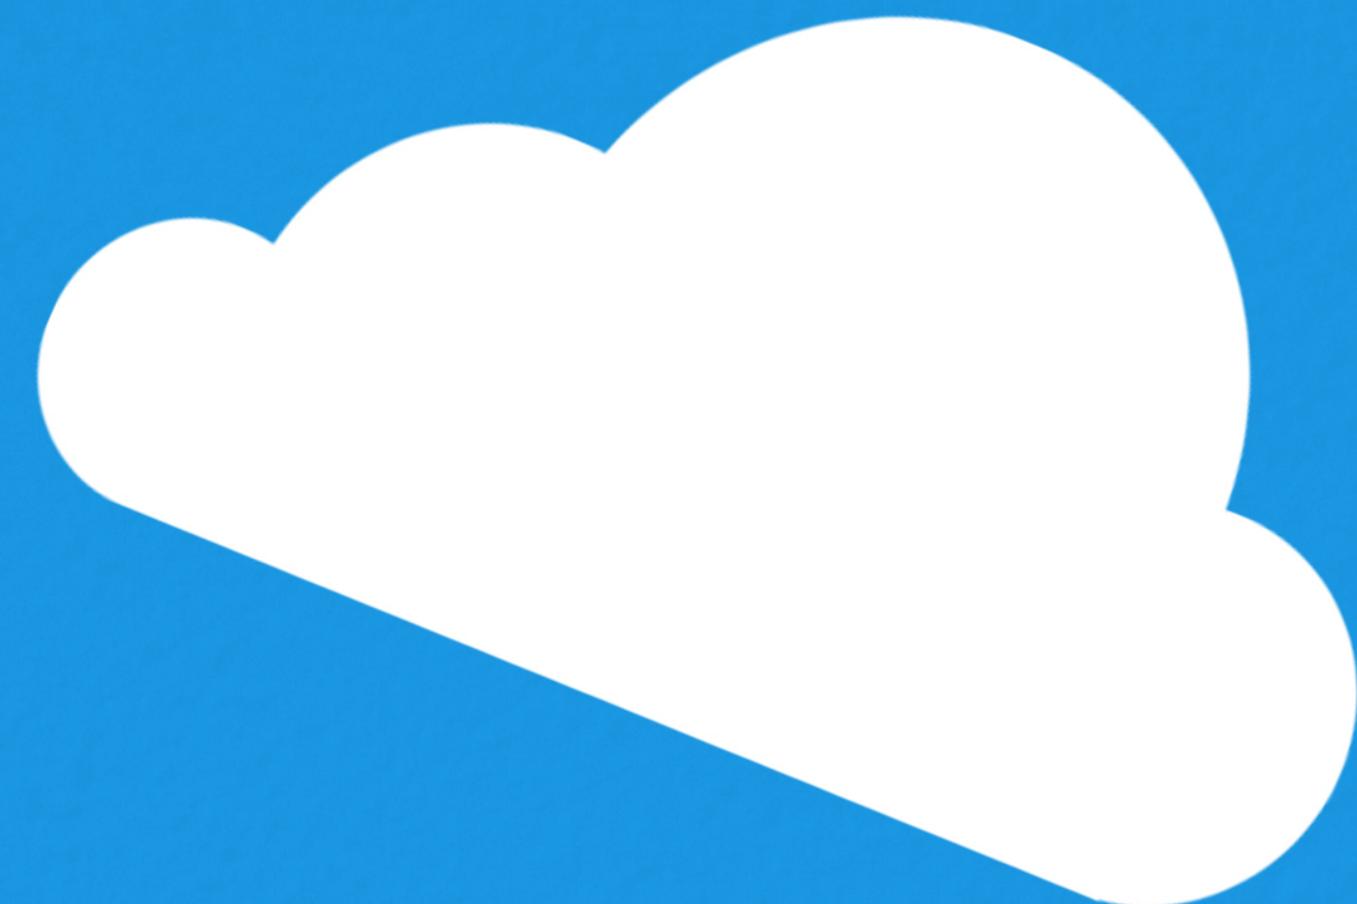
Not just computer scientists

PyWren: Scale For Everyone

Not just computer scientists

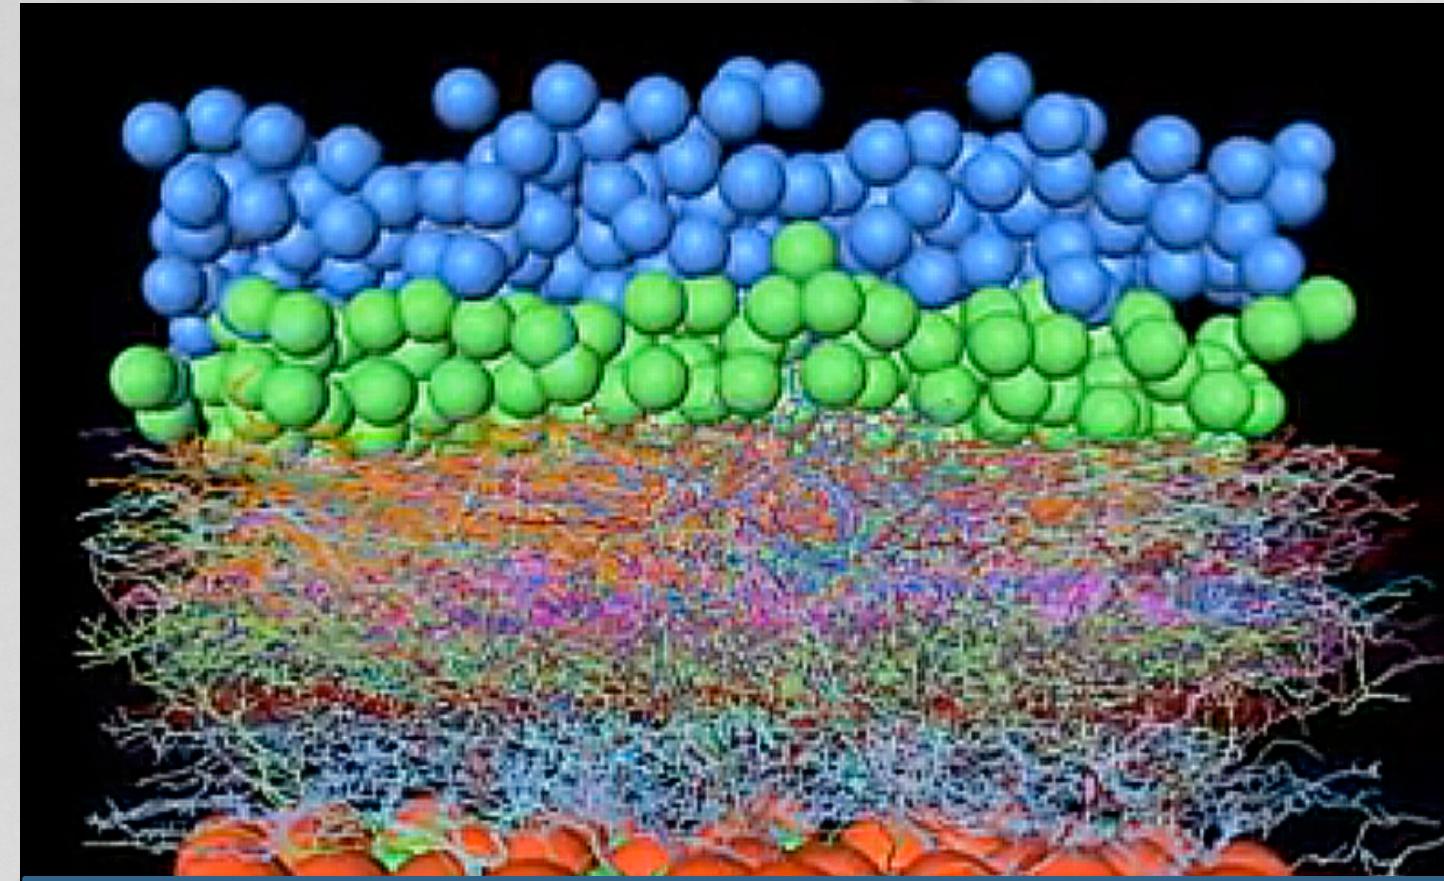


Neuroscientists

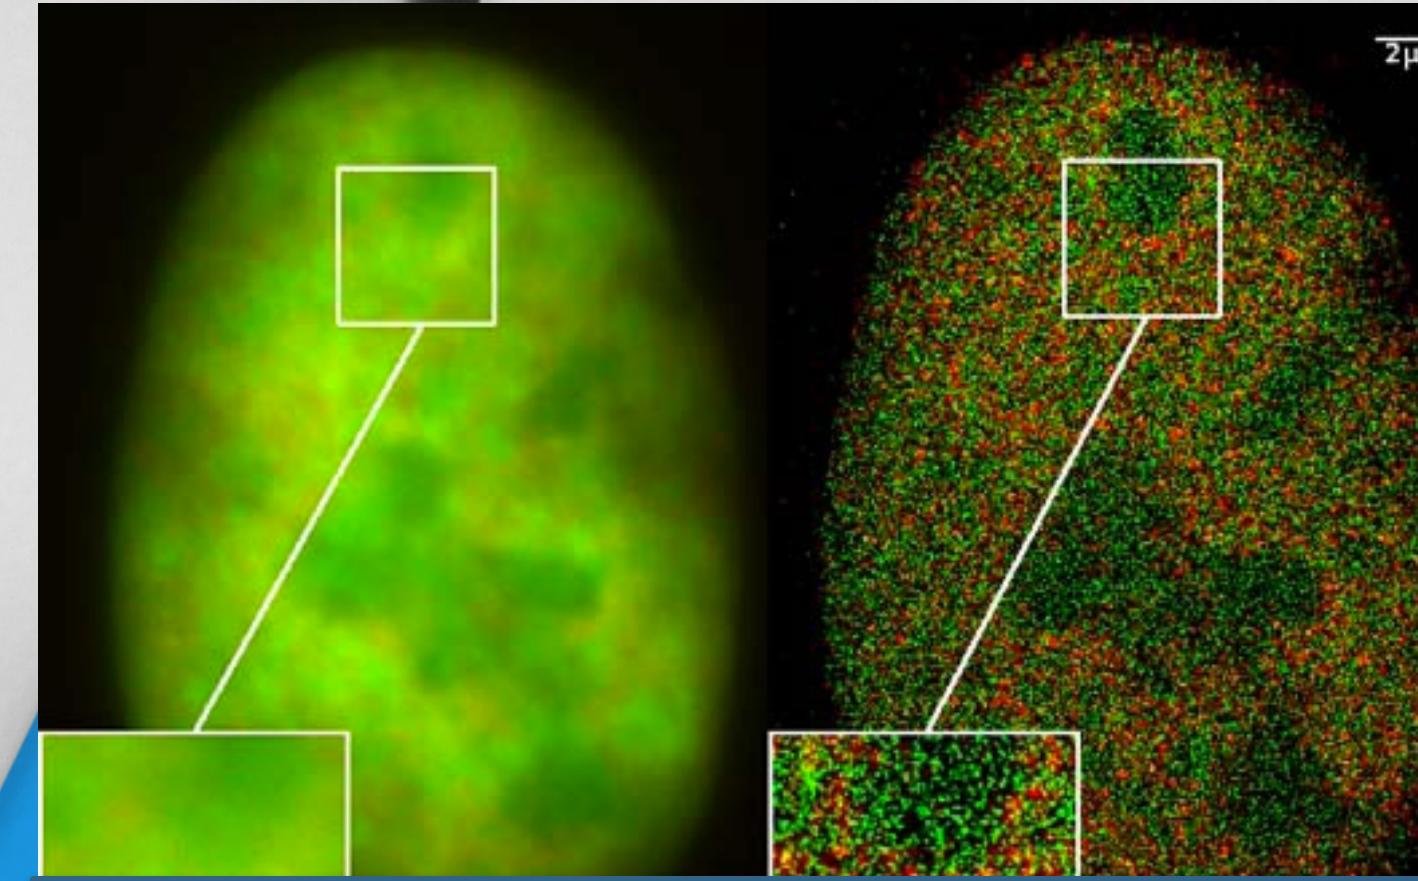


PyWren: Scale For Everyone

Not just computer scientists



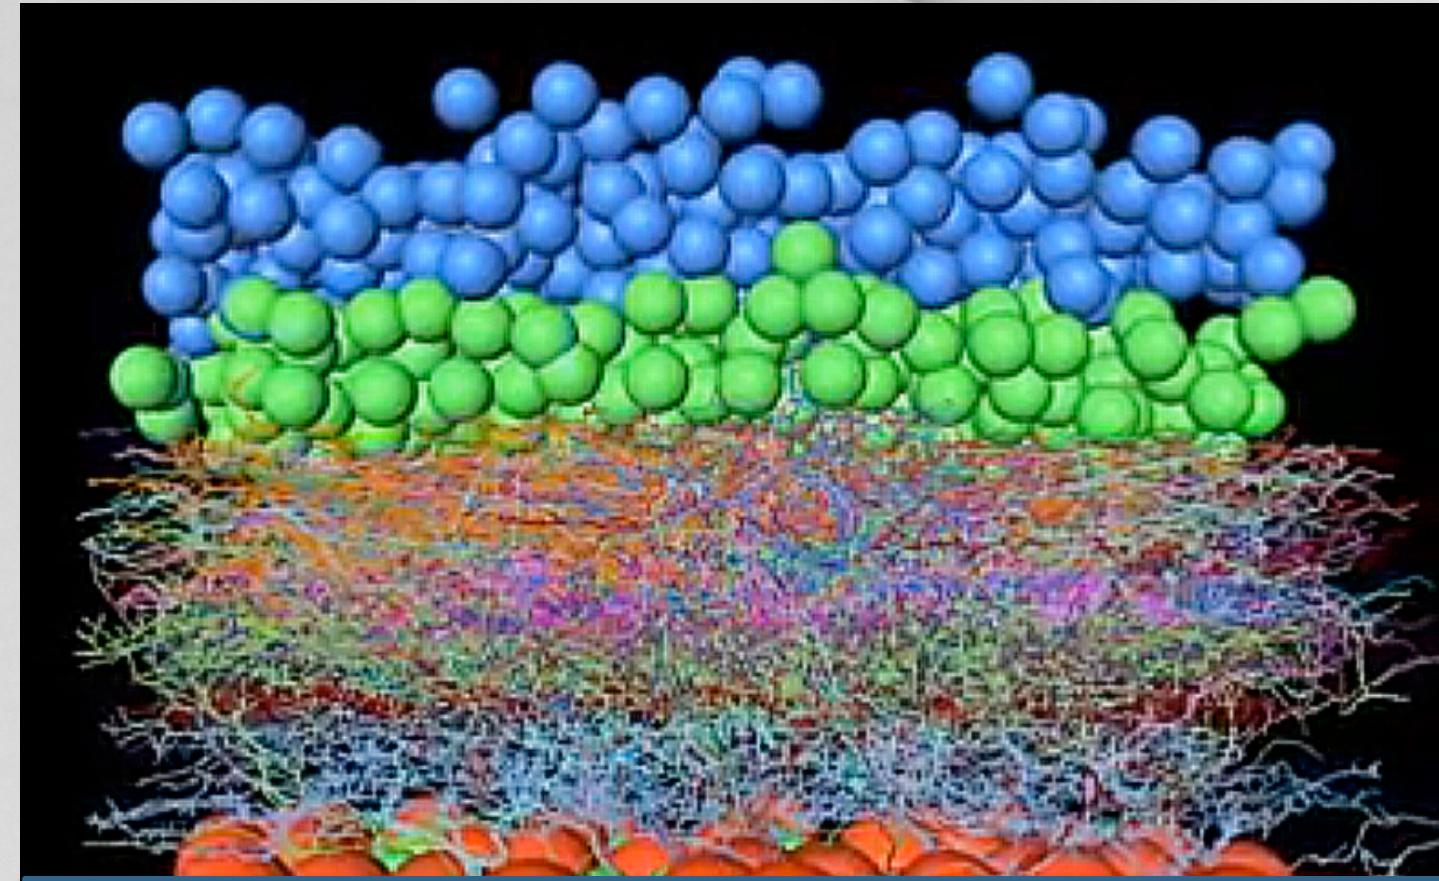
Neuroscientists



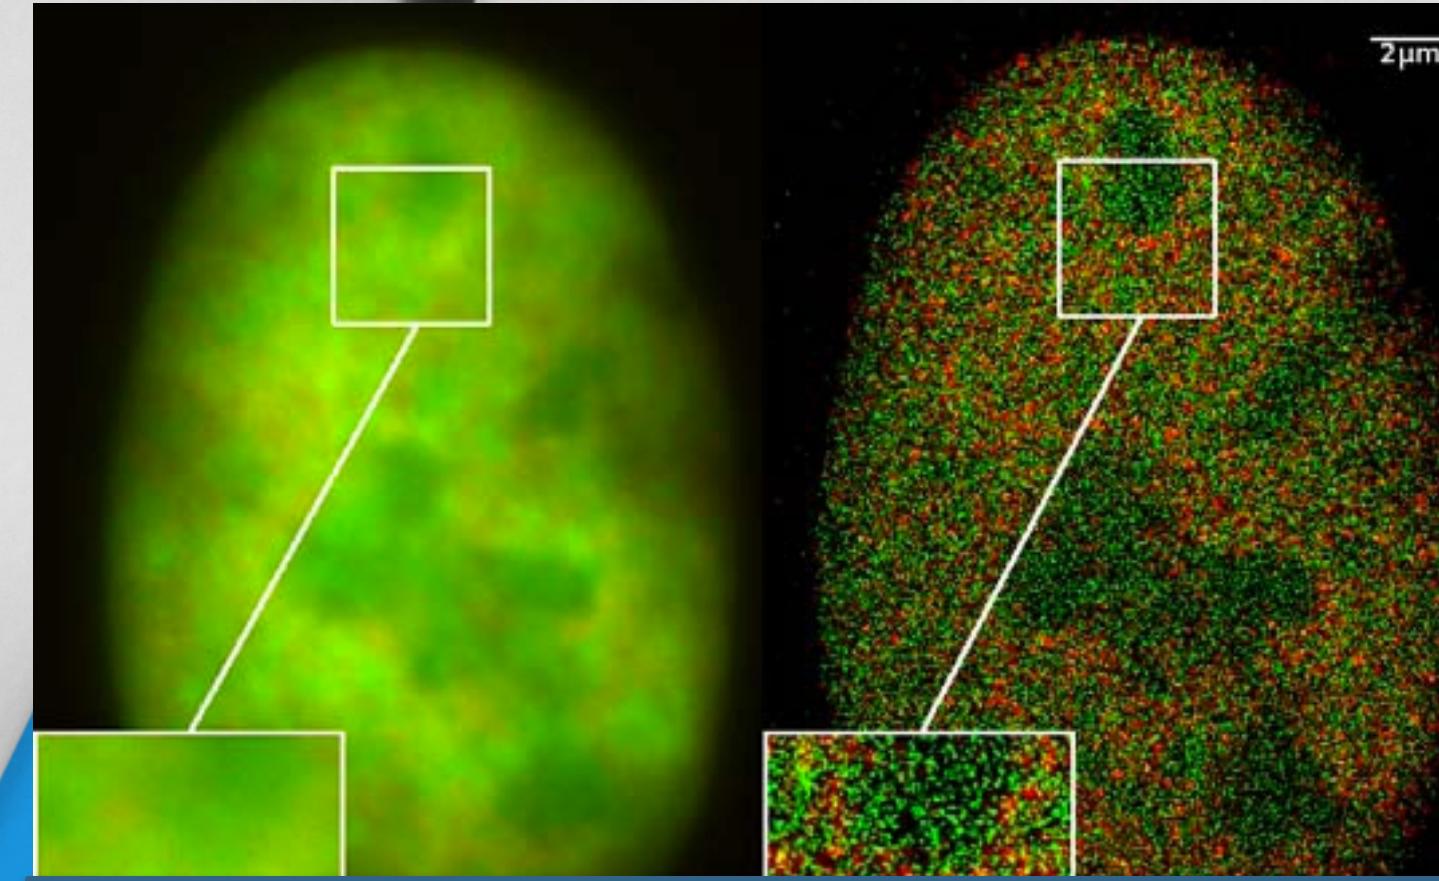
Microscopy and Optics

PyWren: Scale For Everyone

Not just computer scientists



Neuroscientists

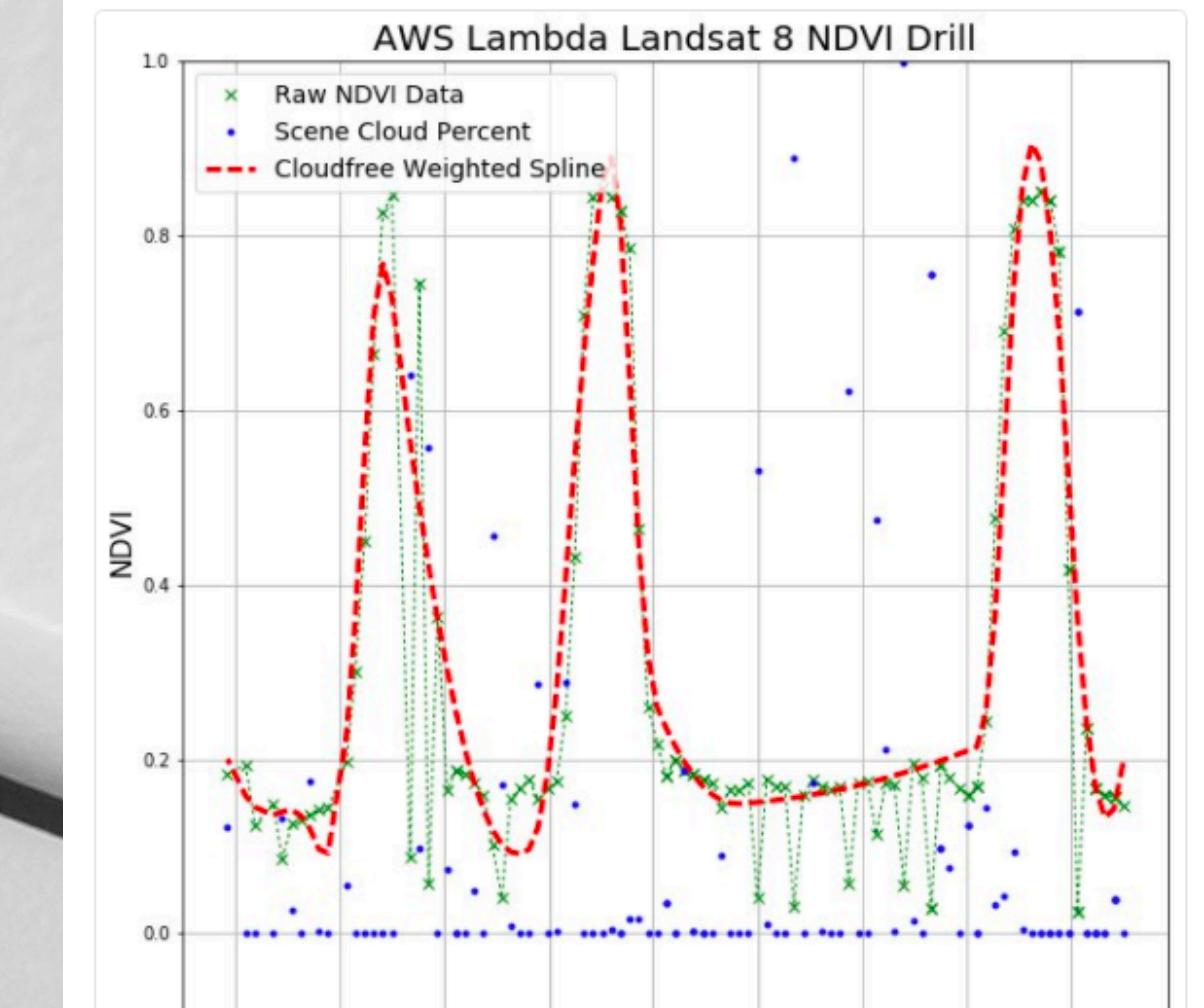


Microscopy and Optics

Peter Scarth  @petescarth

Follow

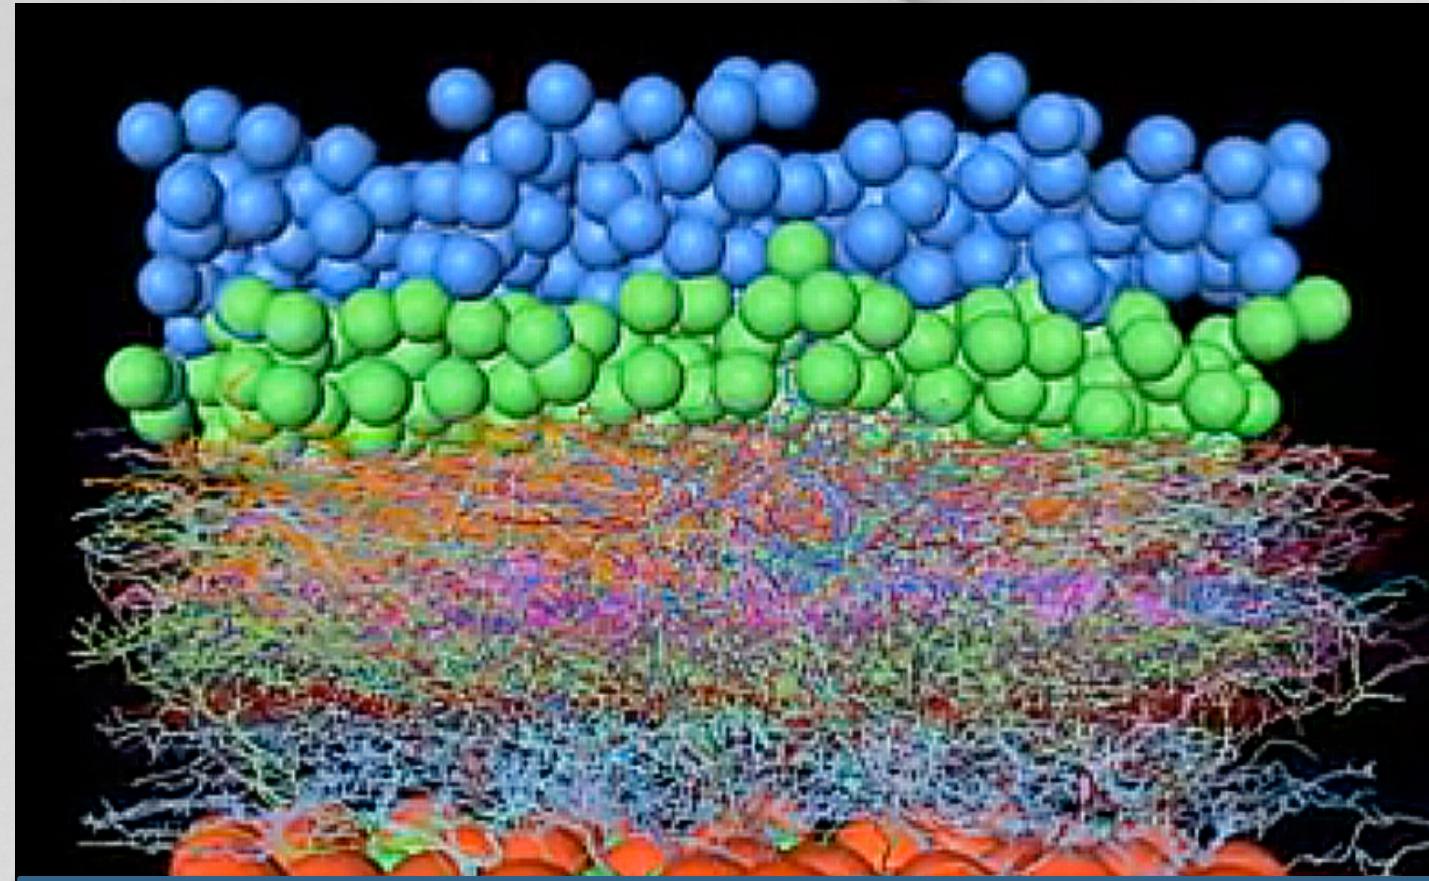
Today's little experiment - #Landsat8 time series extracted over cotton. #lambda + #pywren = #serverless query of 120 scenes in 60 seconds



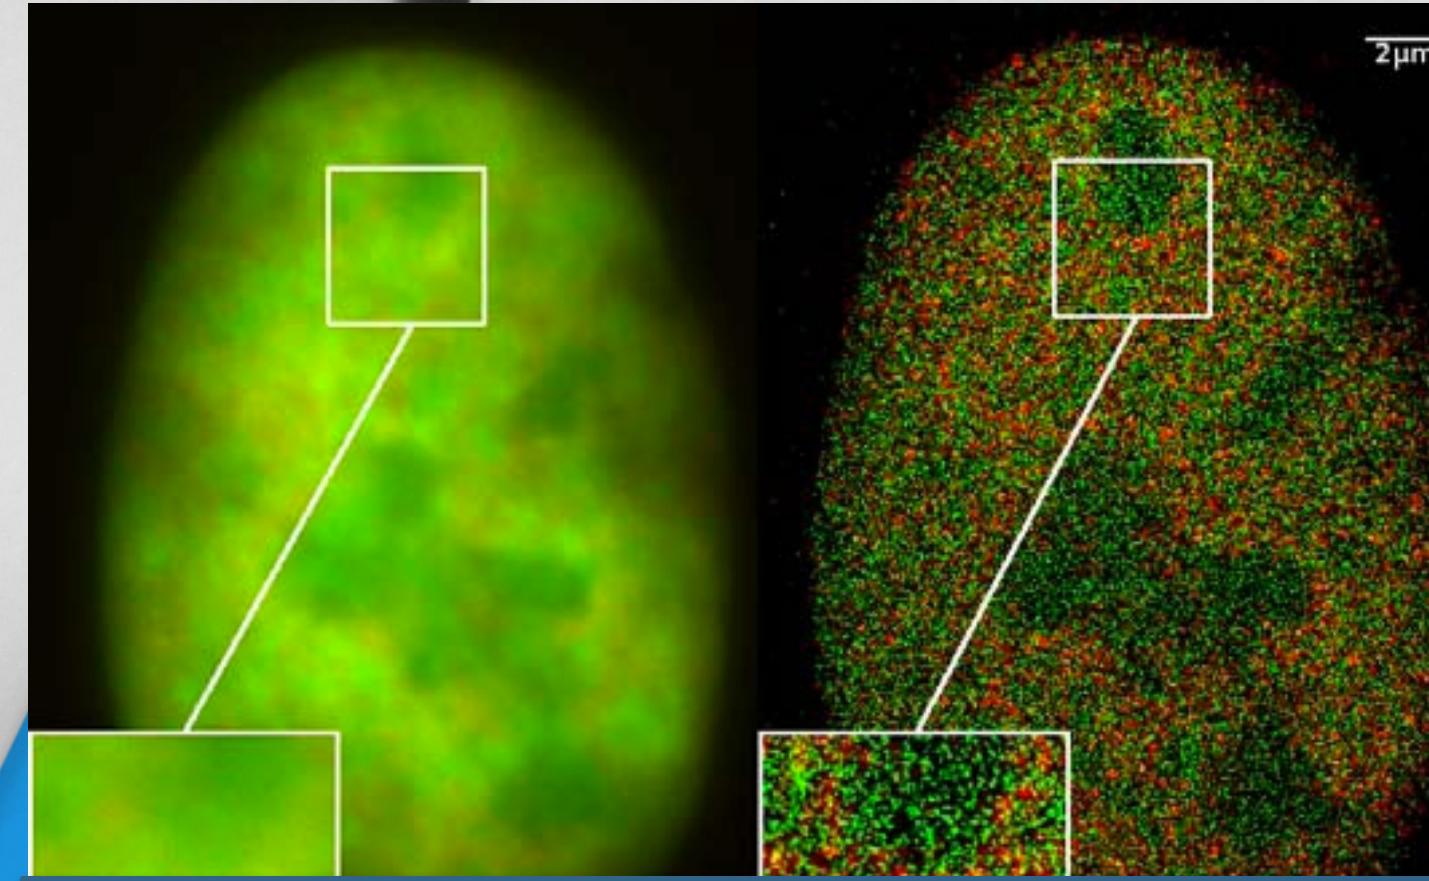
Geophysicists

PyWren: Scale For Everyone

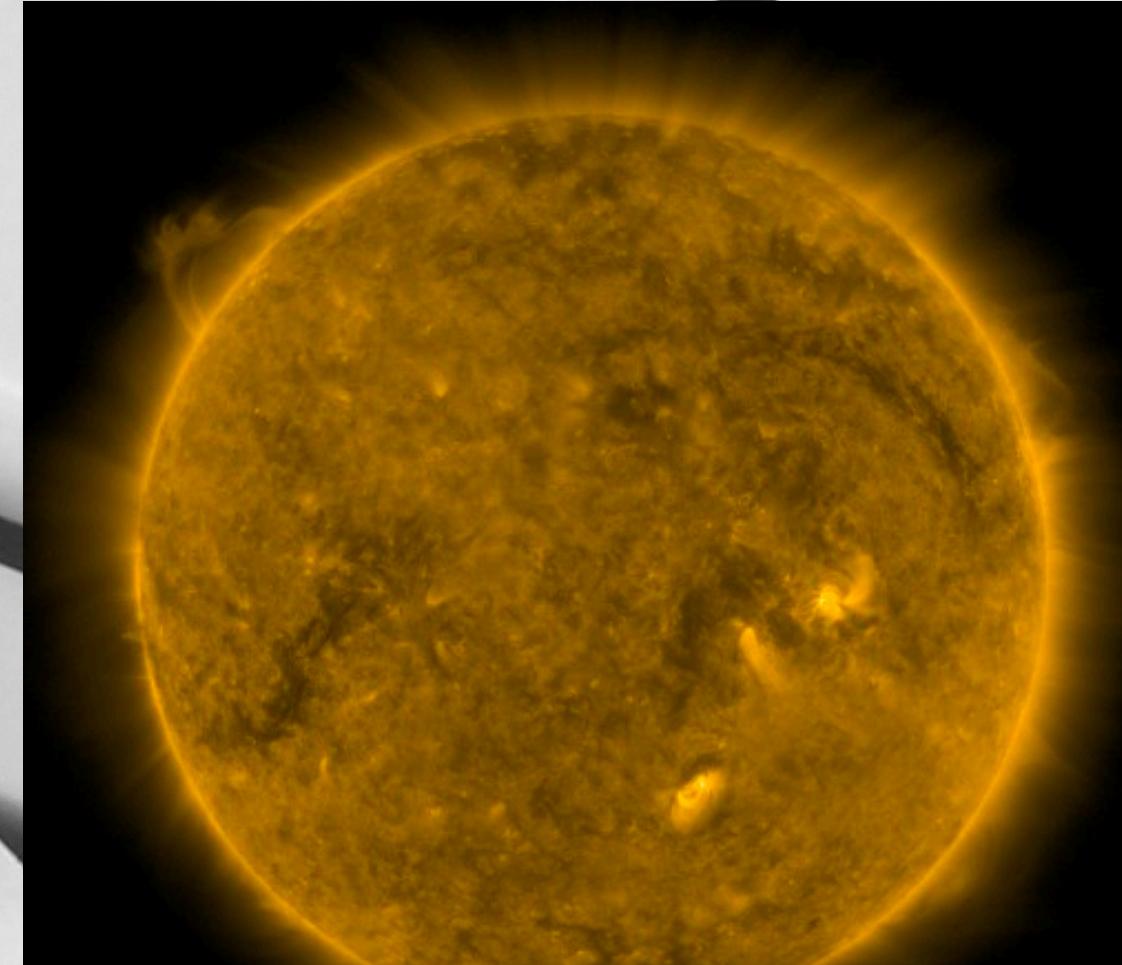
Not just computer scientists



Neuroscientists



Microscopy and Optics

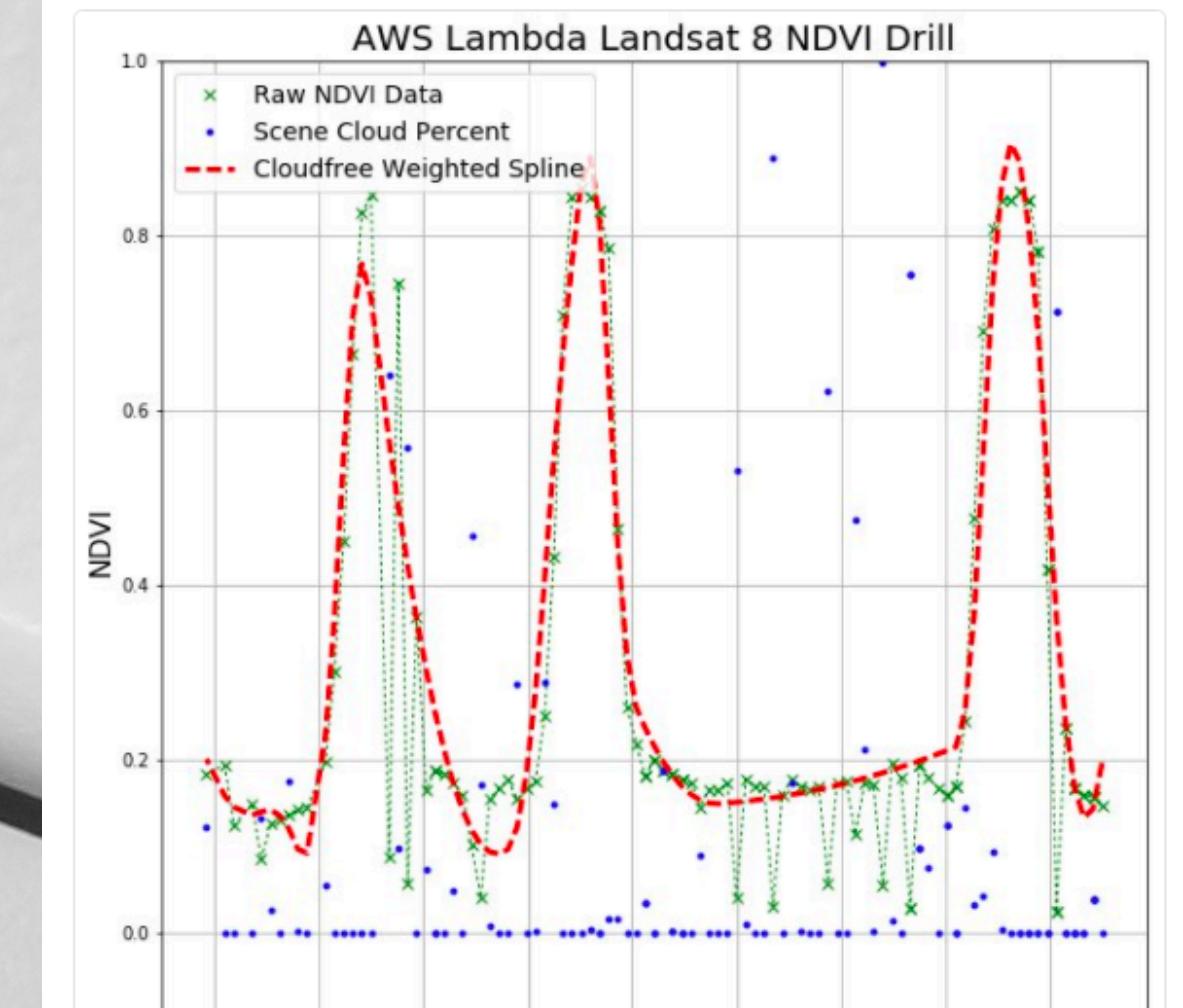


Astronomers

Peter Scarth
@petescarth

Follow

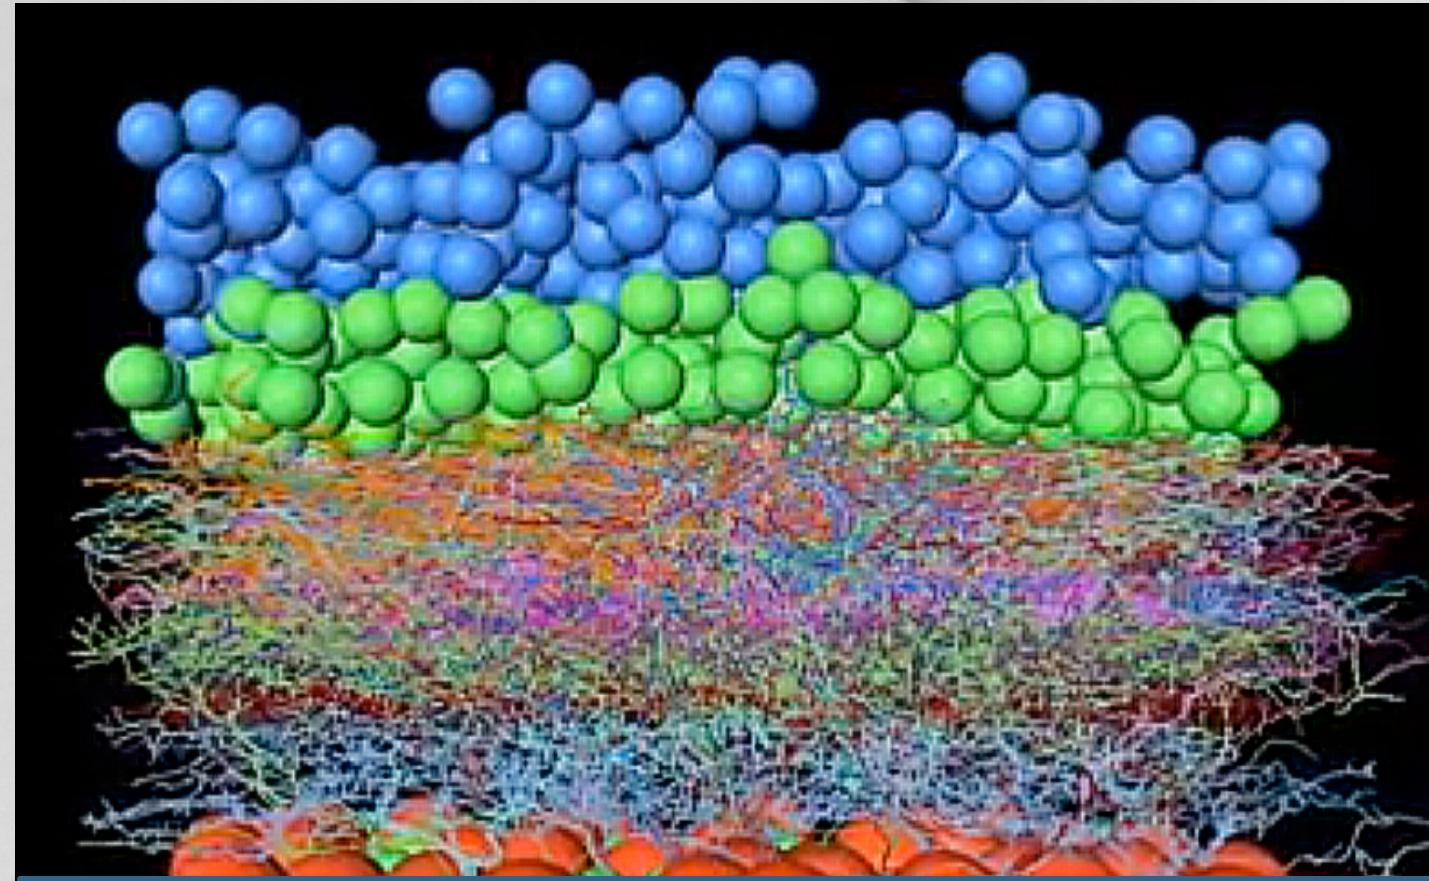
Today's little experiment - #Landsat8 time series extracted over cotton. #lambda + #pywren = #serverless query of 120 scenes in 60 seconds



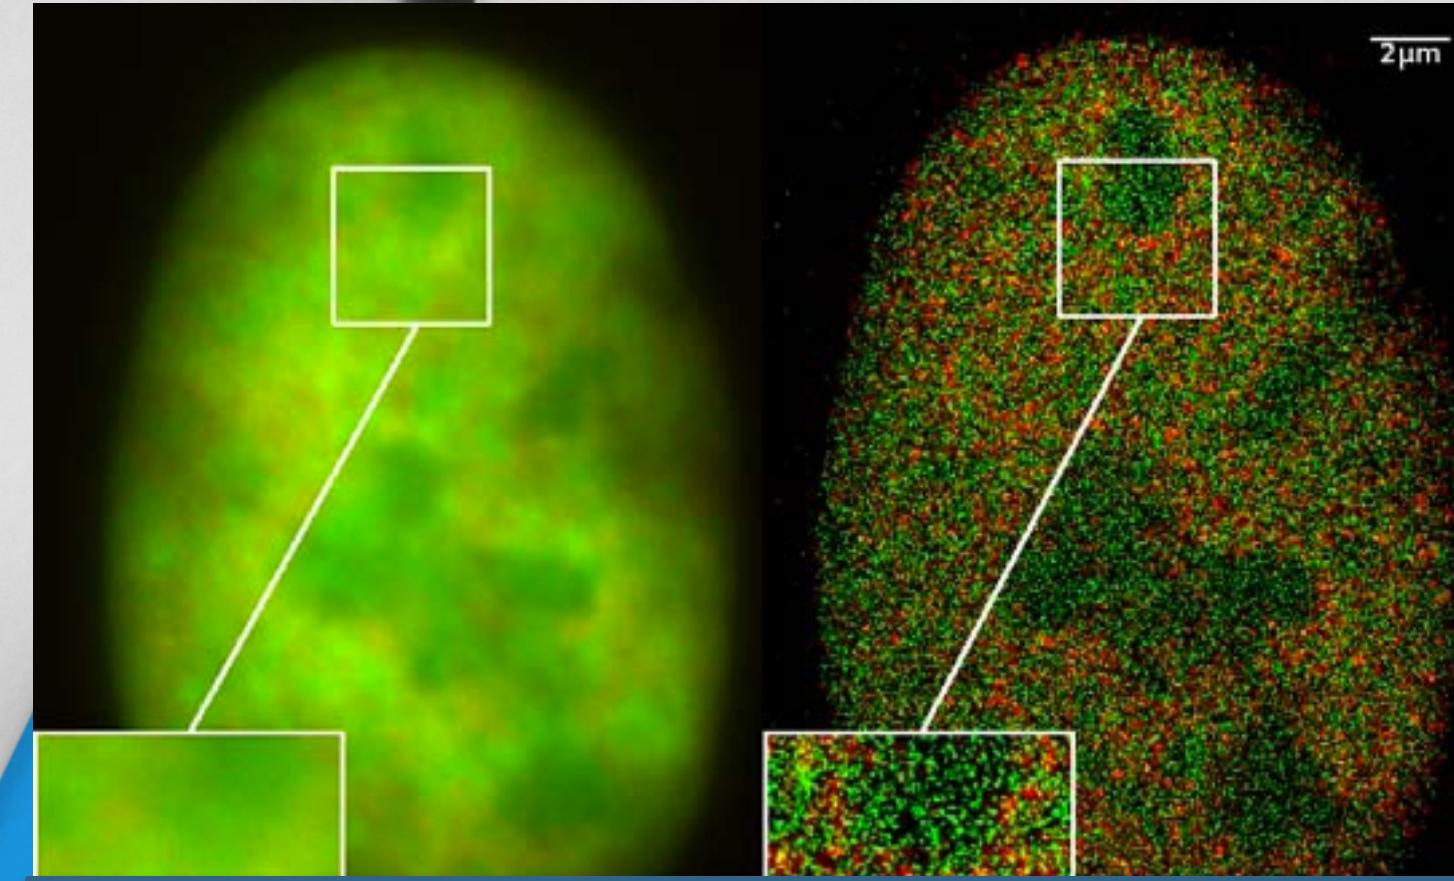
Geophysicists

PyWren: Scale For Everyone

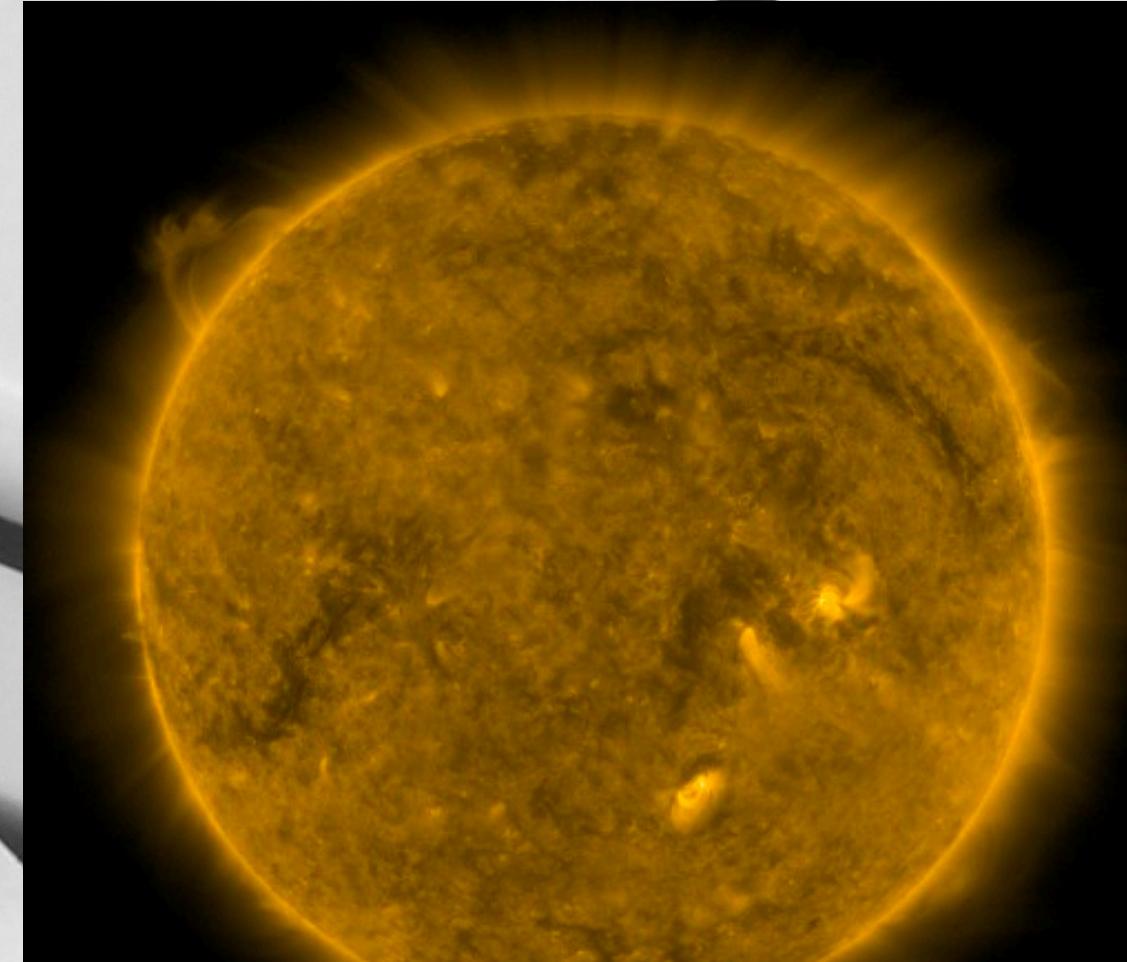
Not just computer scientists



Neuroscientists



Microscopy and Optics



Astronomers

BL ABOUT PROJECTS BLOG

305 Million Solutions to The Black-Scholes Equation in 16 Minutes with AWS Lambda

Originally Posted: May 28, 2017

The research I'm working on involves estimating a firm's probability of default over a variety of time horizons using the Merton Distance to Default model. The dataset contains daily financial information for more than 24,000 firms over the past 30 years. Given that I am calculating the probability of default over five time horizons, applying the Merton model will require solving the Black-Scholes equation roughly 305 million times. Luckily, the model is easily parallelized because the only data needed for the model, aside from the risk-free rate, is firm specific. This post shows how the Python library [Pywren](#) can

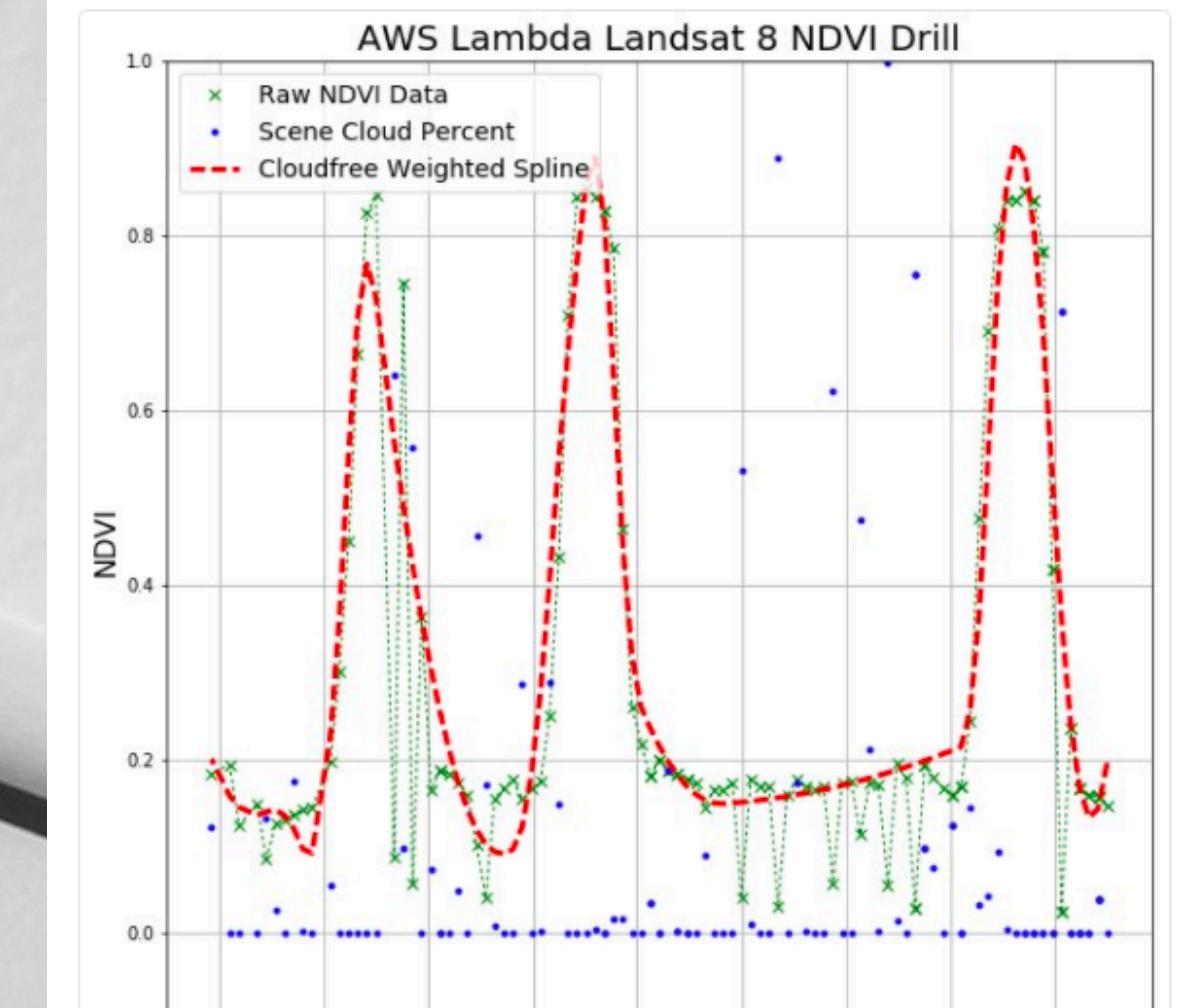
2

Finance and Credit

Peter Scarth
@petescarth

Follow

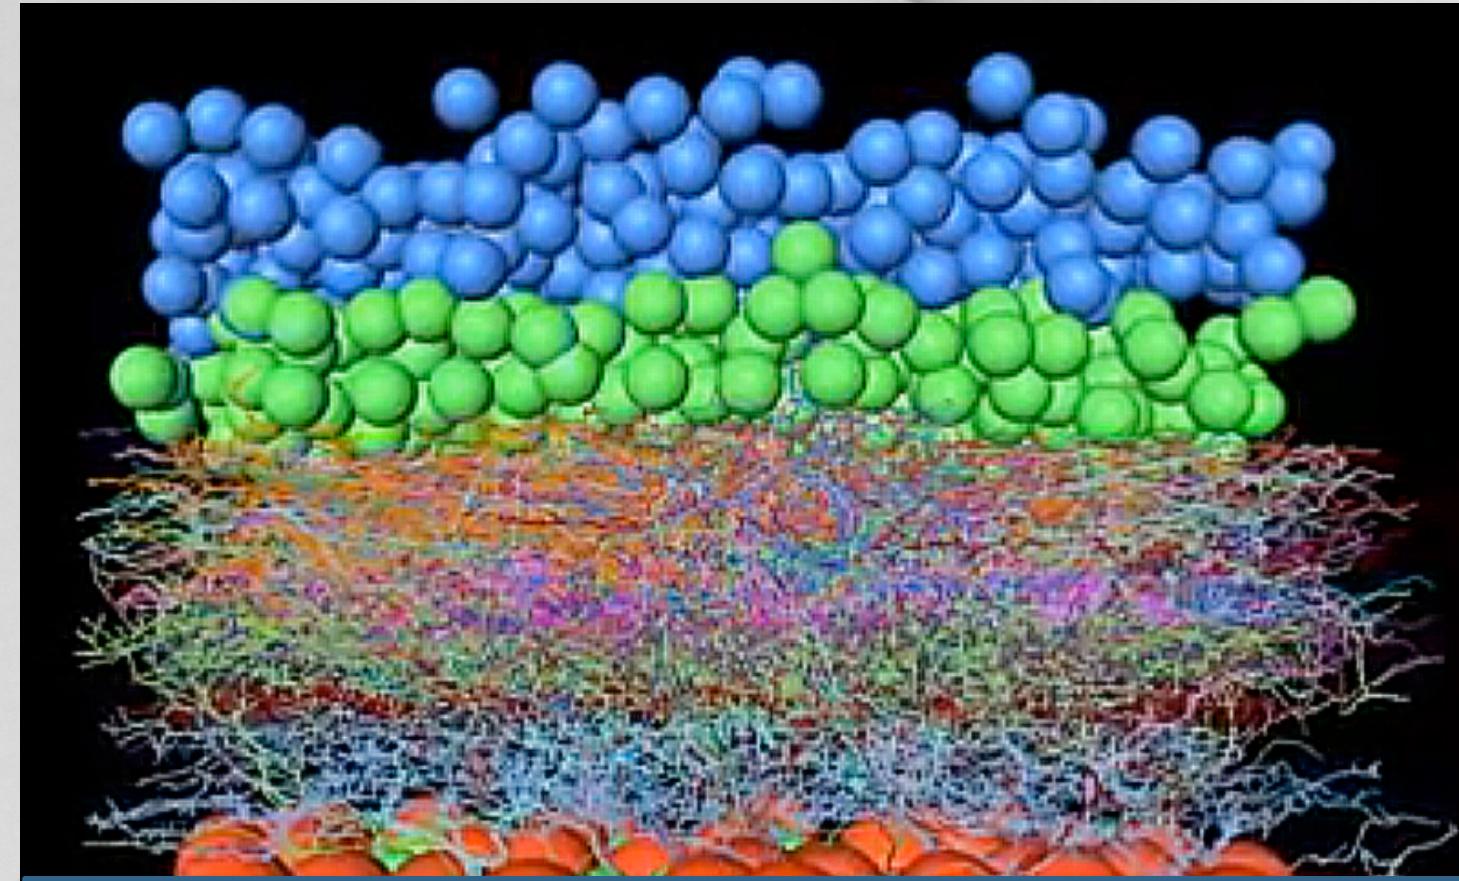
Today's little experiment - #Landsat8 time series extracted over cotton. #lambda + #pywren = #serverless query of 120 scenes in 60 seconds



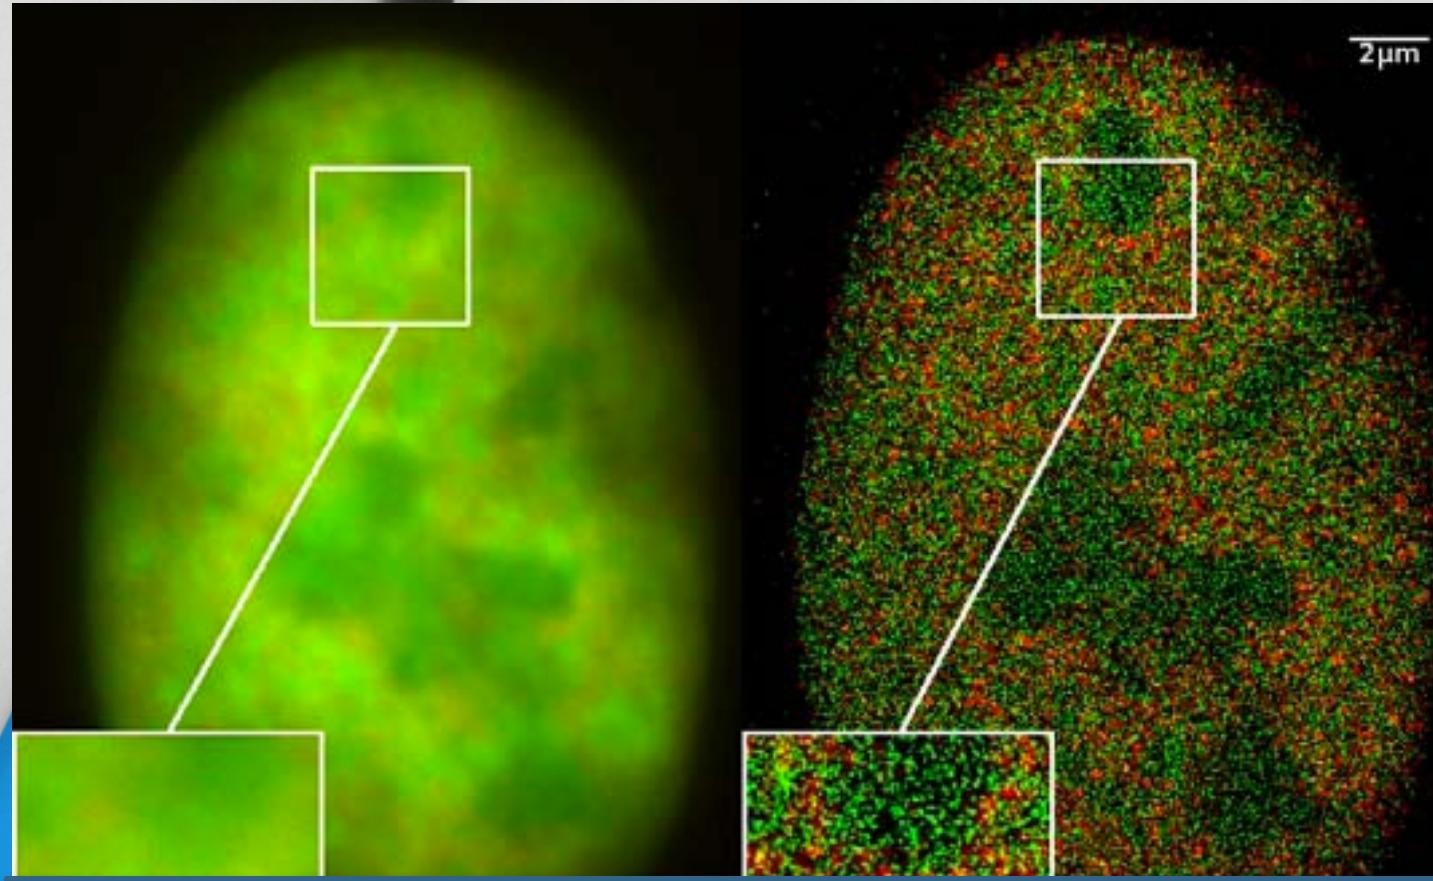
Geophysicists

PyWren: Scale For Everyone

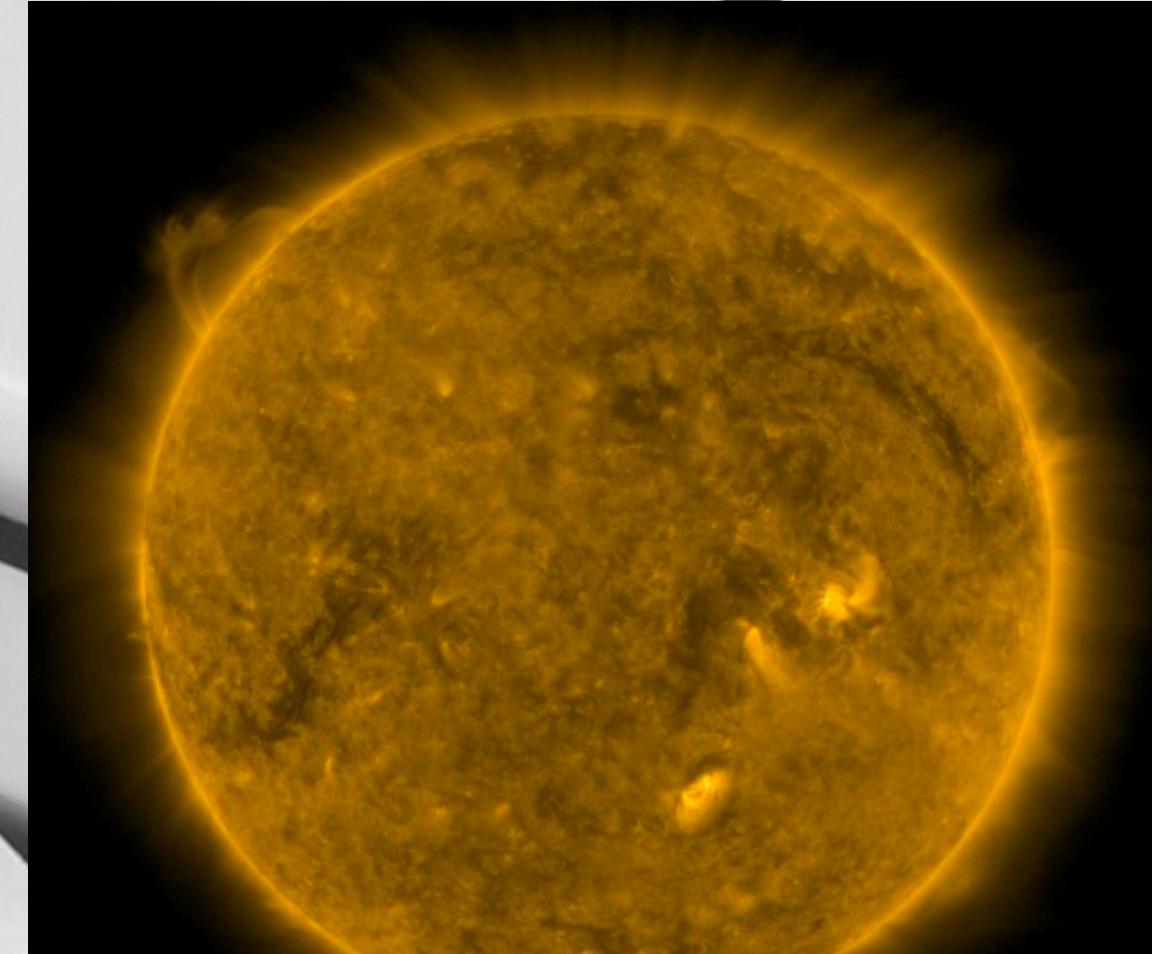
Not just computer scientists



Neuroscientists



Microscopy and Optics



Astronomers

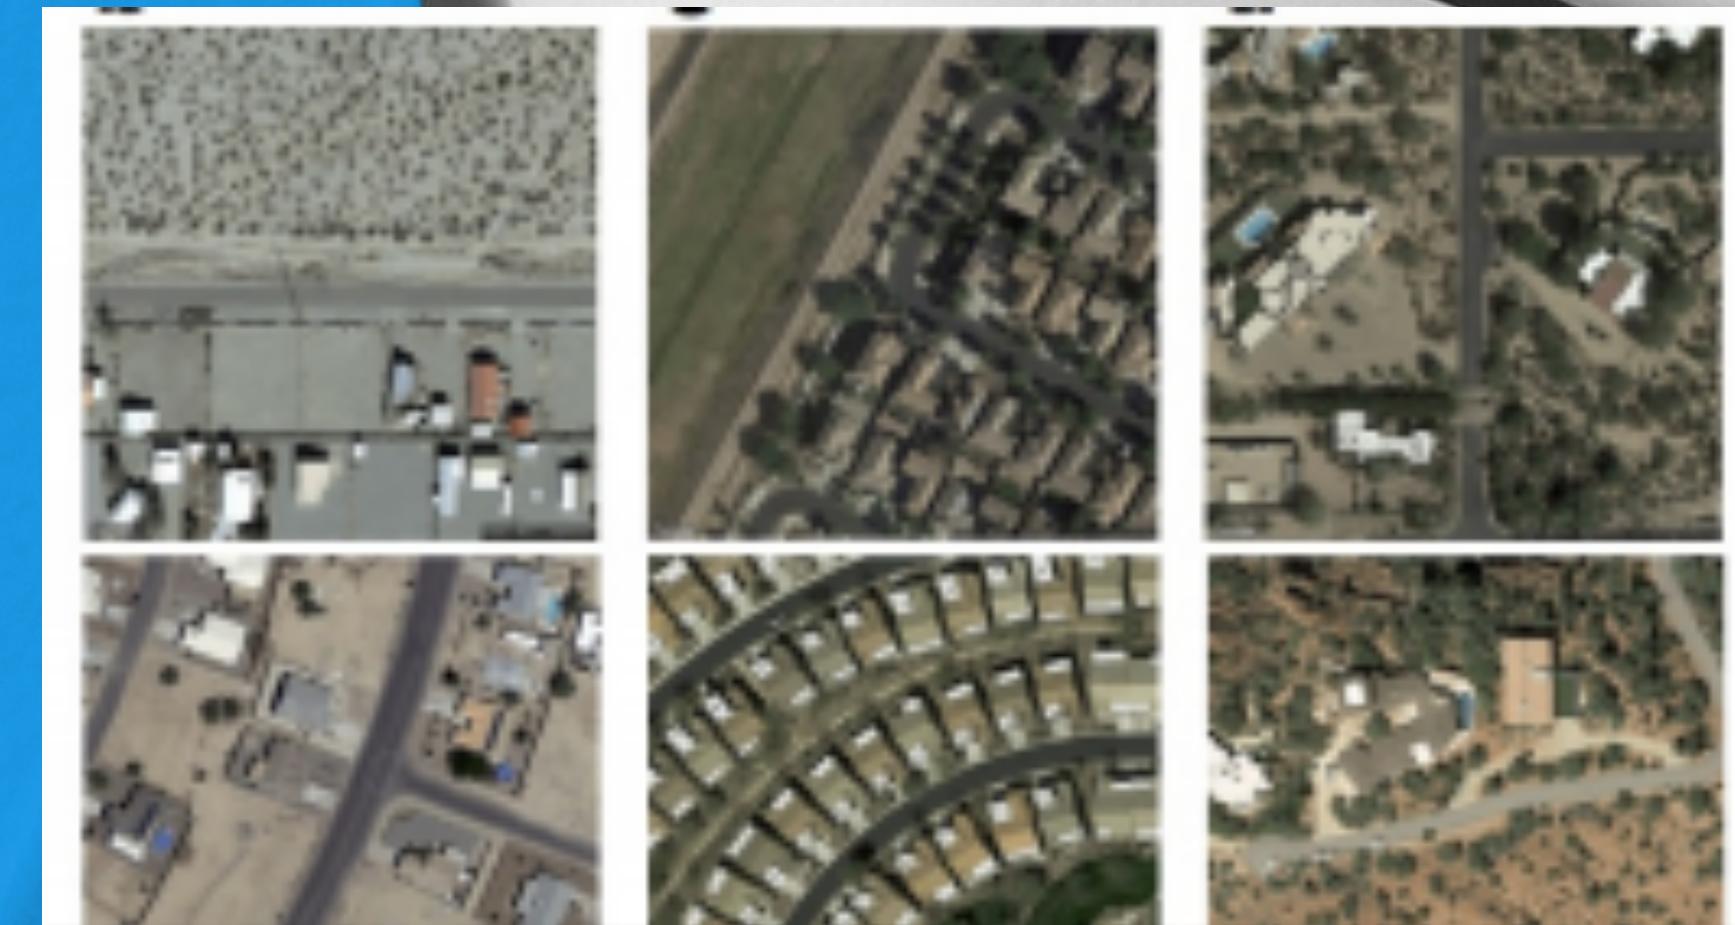
BL ABOUT PROJECTS BLOG

305 Million Solutions to The Black-Scholes Equation in 16 Minutes with AWS Lambda

Originally Posted: May 28, 2017

The research I'm working on involves estimating a firm's probability of default over a variety of time horizons using the Merton Distance to Default model. The dataset contains daily financial information for more than 24,000 firms over the past 30 years. Given that I am calculating the probability of default over five time horizons, applying the Merton model will require solving the Black-Scholes equation roughly 305 million times. Luckily, the model is easily parallelized because the only data needed for the model, aside from the risk-free rate, is firm specific. This post shows how the Python library [Pywren](#) can

Finance and Credit

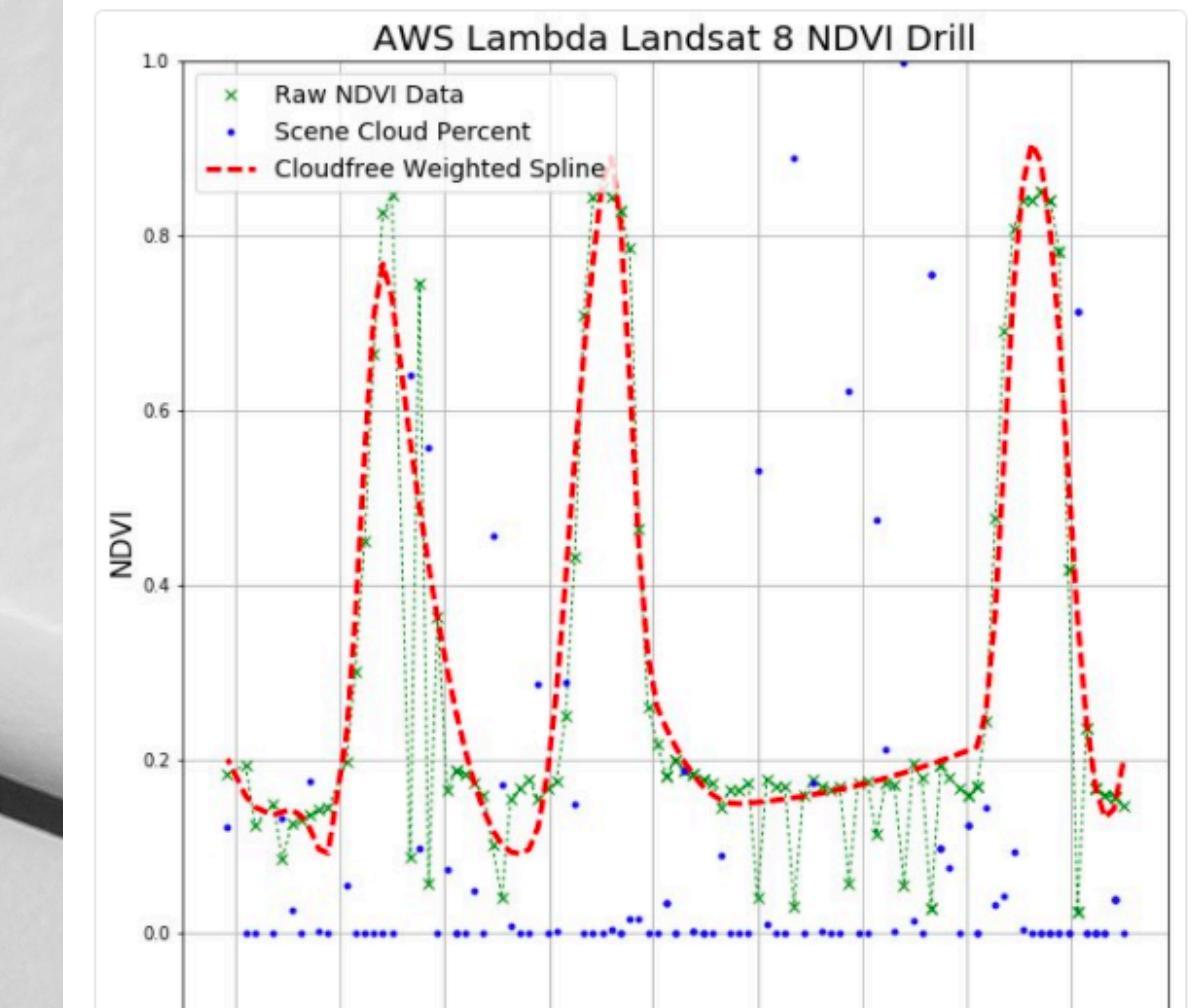


Developmental Economists

Peter Scarth @petescarth

Follow

Today's little experiment - #Landsat8 time series extracted over cotton. #lambda + #pywren = #serverless query of 120 scenes in 60 seconds





Michael H. Oshita @ijin · Apr 28

PyWren - lambda map/reduce framework. 25TFLOPS!
github.com/pywren/pywren #ServerlessConf

PyWren: a massive data framework for Lambda

- Open source MapReduce framework using Lambda
- 25 TFLOPS performance
- 60 GB/sec read and 50 GB/sec write to S3

peak GFLOPS
effective GFLOPS

time (sec)

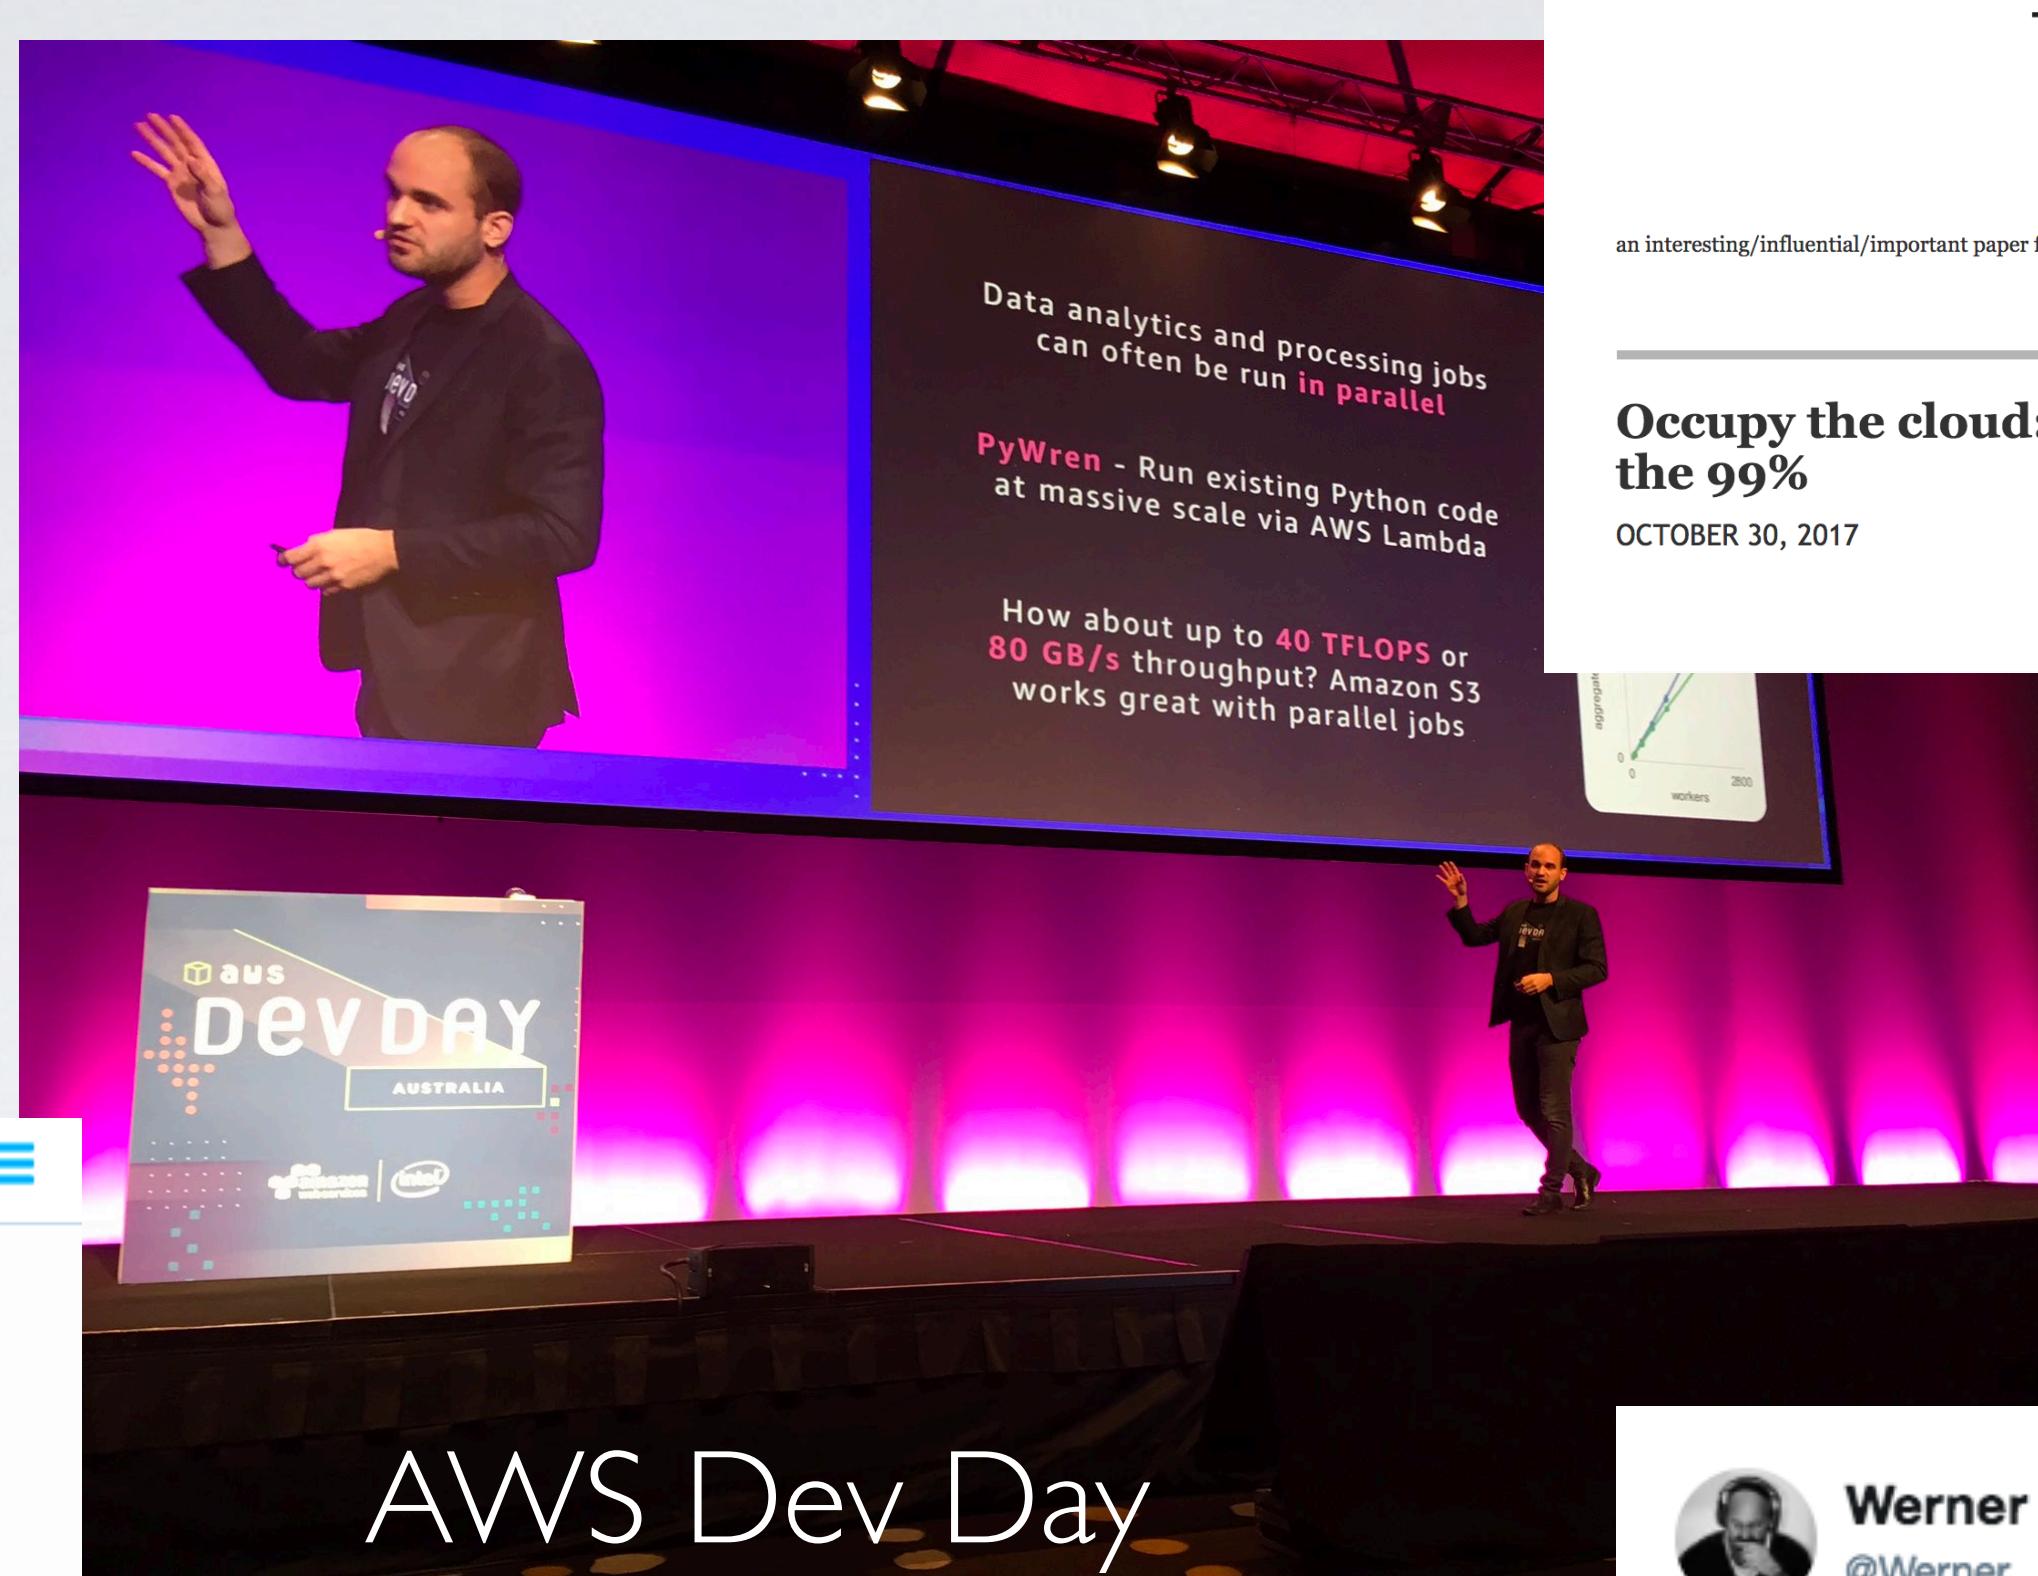
THE NEW STACK

Q ≡

EVENTS / TECHNOLOGY

With PyWren, AWS Lambda Finds an Unexpected Market in Scientific Computing

16 Feb 2017 10:26am, by Joab Jackson



AWS Dev Day

SRV424
AWS re:INVENT
Massively Parallel Data Processing
with PyWren and AWS Lambda

November 30, 2017

AWS Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

PyWren Web Scraping

I was tasked with scraping information of houses for sale in Massachusetts for my data mining class. The target site in question was [redfin.com](#), they explicitly do not tolerate web scraping and will give you a captcha if you exceeded some unknown threshold of pages per minute or had a fishy user-agent.

Not to be deterred by a captcha, I used selenium Chromedriver to write a scraper that worked pretty well and importantly was not caught by redfin's algorithm. Each page took ~2-3 seconds to scrape.



ACM Symposium
on Cloud Computing

Occupy the Cloud: Distributed Computing for the 99% [VISION]

Eric Jonas, Qifan Pu, Shivaram Venkataraman, Ion Stoica, Benjamin Recht (UC Berkeley)

the morning paper

an interesting/influential/important paper from the world of CS every weekday morning, as selected by Adrian Colyer

Home About InfoQ QR Editions Subscribe

SUBSCRIBE



never miss an issue! The

Occupy the cloud: distributed computing for the 99%

OCTOBER 30, 2017

tags: Distributed Systems



one in a million
@TearTheSky

Follow

サーバレスのトークを聞きにきてるけどFlaskだけ固有名詞で出たりPyWrenが出たり、スピーカーはPython推しなのかな？ PyWrenは科学計算フレームワークみたい。

[aws.amazon.com/jp/blogs/news/ ...](http://aws.amazon.com/jp/blogs/news/)

Translate from Japanese

9:34 PM - 30 May 2017



Werner Vogels ✅
@Werner

Follow

#Microservices and TerraFlops - Extracting 25 TFLOPS from #AWS #Lambda - @stochastician on the origin of #pywren ericjonas.com/pywren.html

map(function, data)

and... that's mostly it

map(function, data)

and... that's mostly it

```
def myfunc(x):  
    return x + 1
```

map(function, data)

and... that's mostly it

```
def myfunc(x):  
    return x + 1
```

```
futures = pwex.map(myfunc, [1, 2, 3])
```

map(function, data)

and... that's mostly it

```
def myfunc(x):  
    return x + 1
```

```
futures = pwex.map(myfunc, [1, 2, 3])
```

```
print pywren.get_all_results(futures)
```

map(function, data)

and... that's mostly it

```
def myfunc(x):  
    return x + 1
```

```
futures = pwex.map(myfunc, [1, 2, 3])
```

```
print pywren.get_all_results(futures)
```

```
[2, 3, 4]
```

map(function, data)

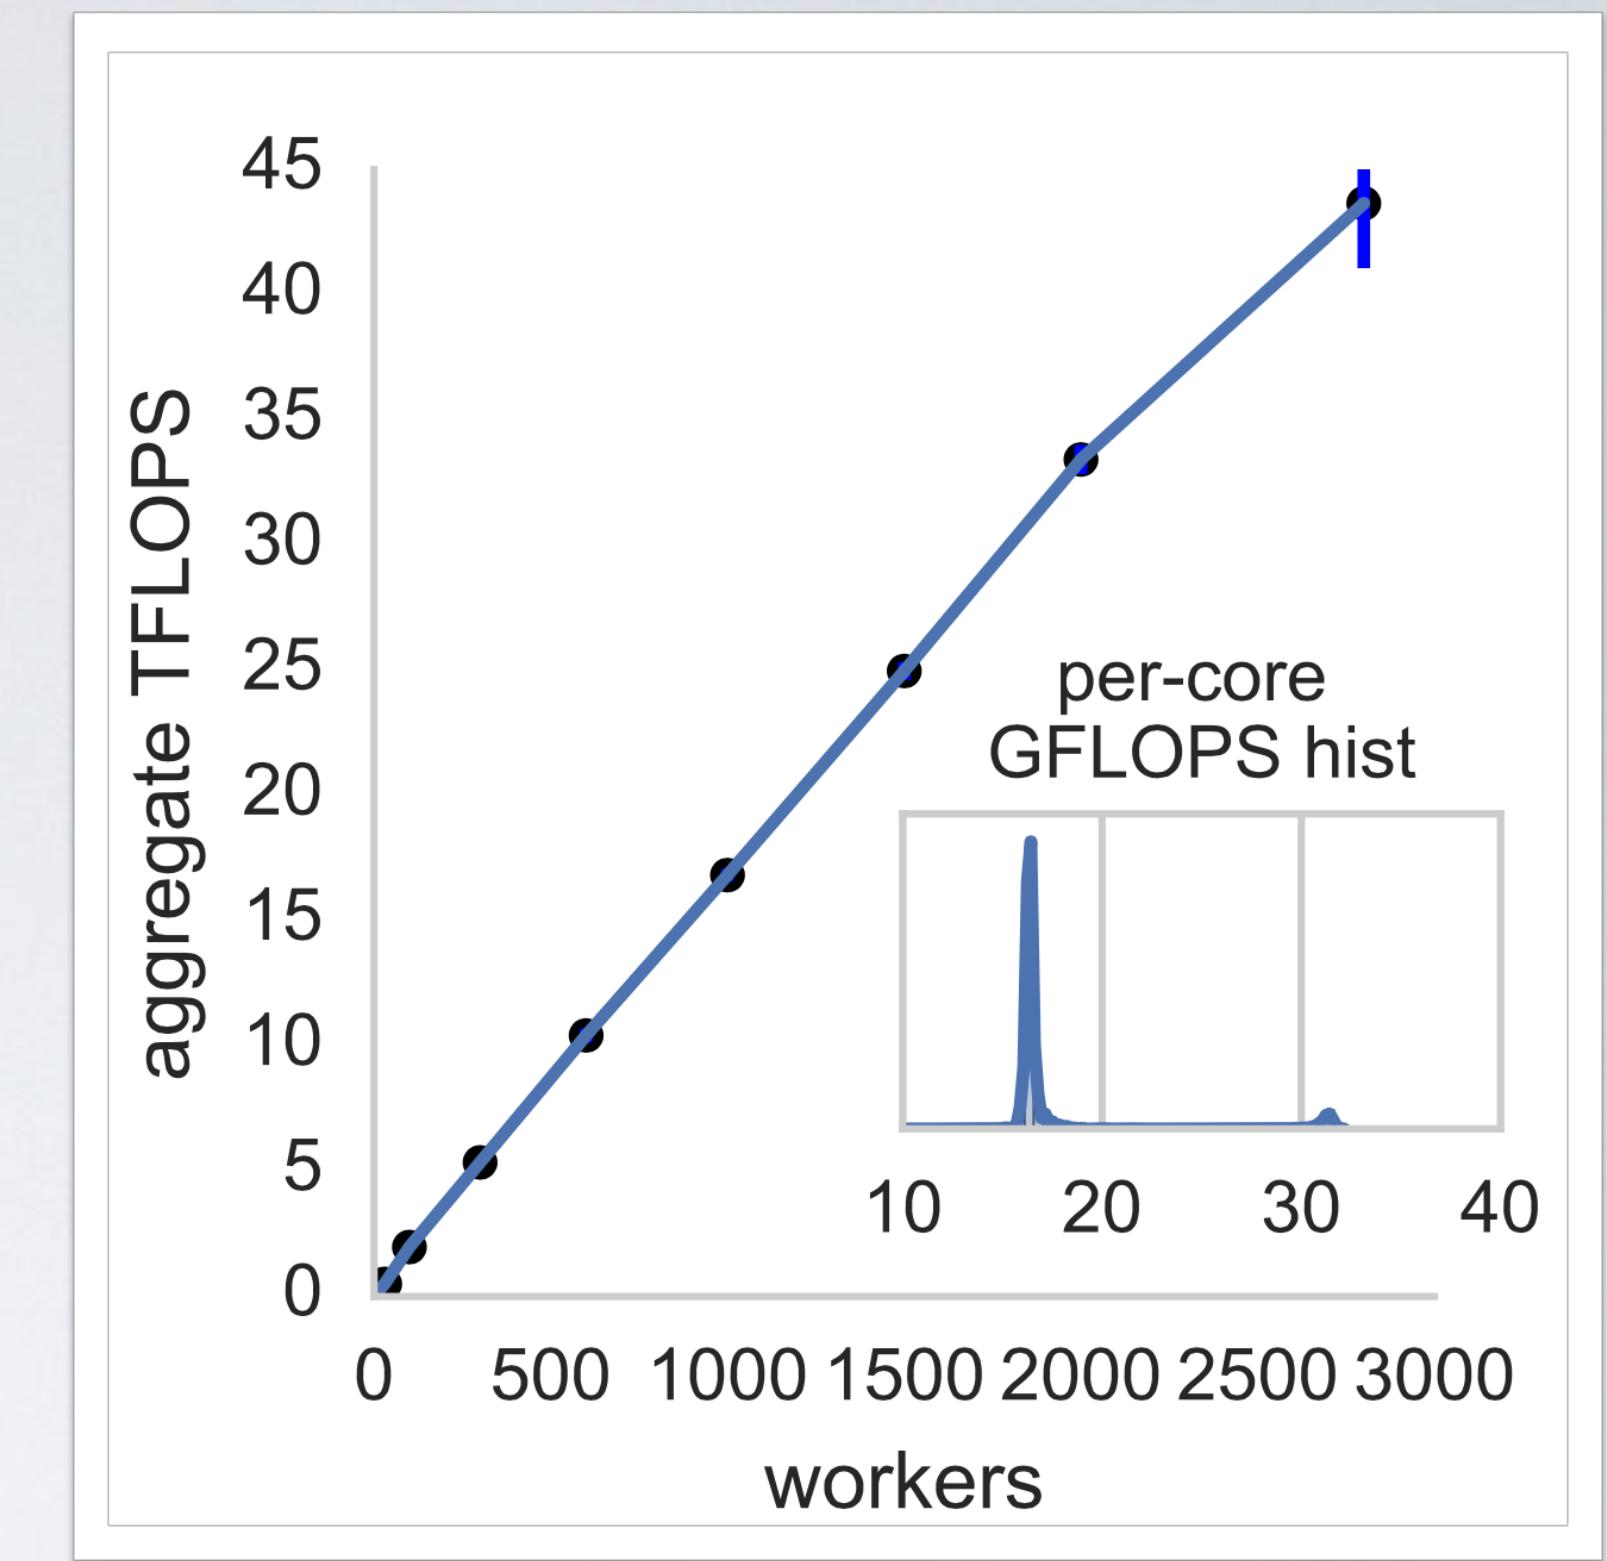
and... that's mostly it

```
def myfunc(x):  
    return x + 1
```

```
futures = pwex.map(myfunc, [1, 2, 3])
```

```
print pywren.get_all_results(futures)
```

```
[2, 3, 4]
```



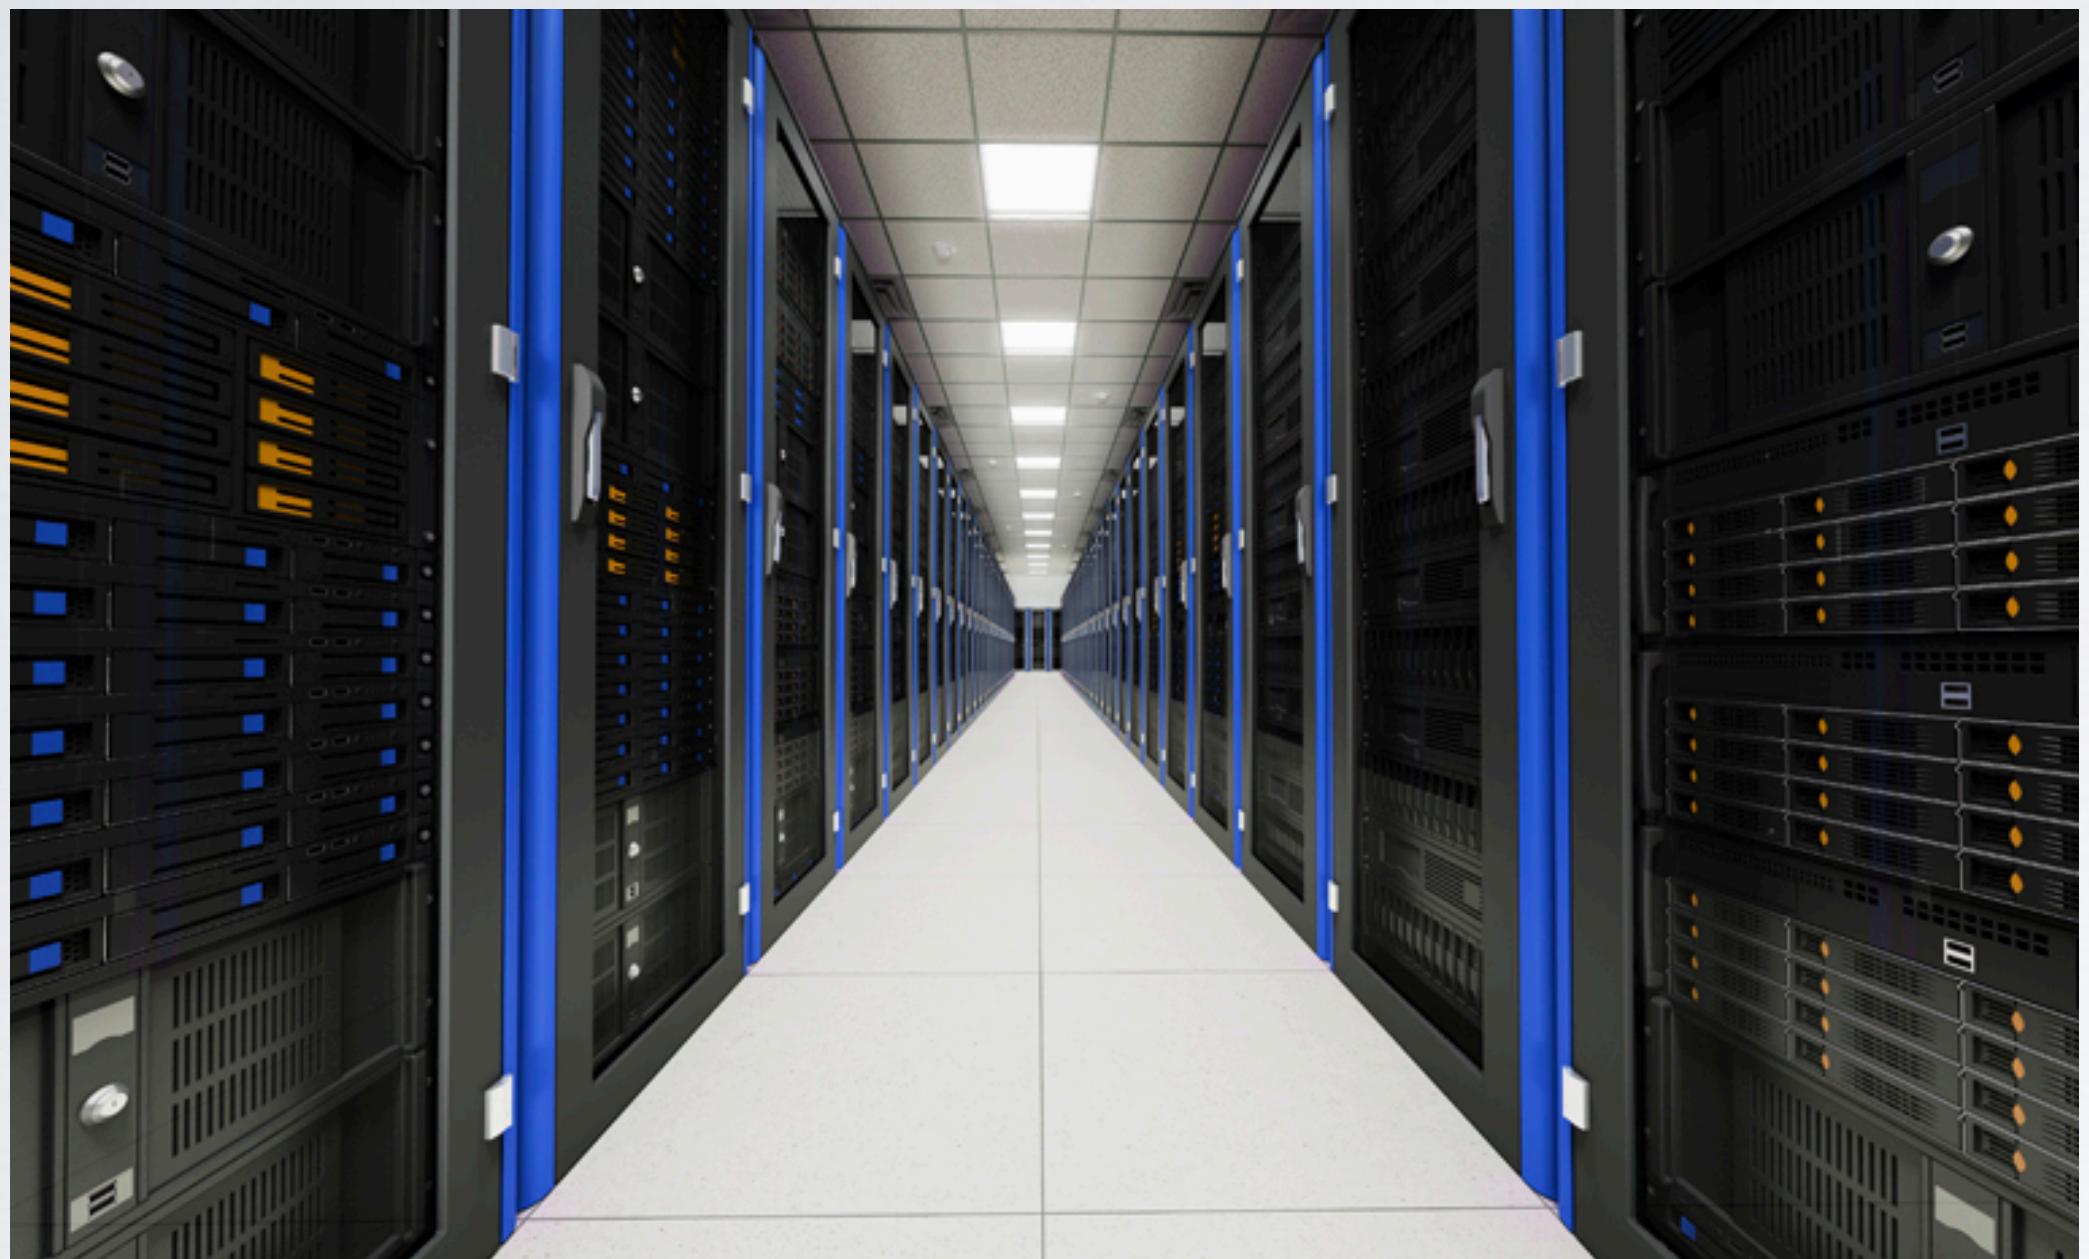
Beyond map?



MOTIVATING LINEAR ALGEBRA

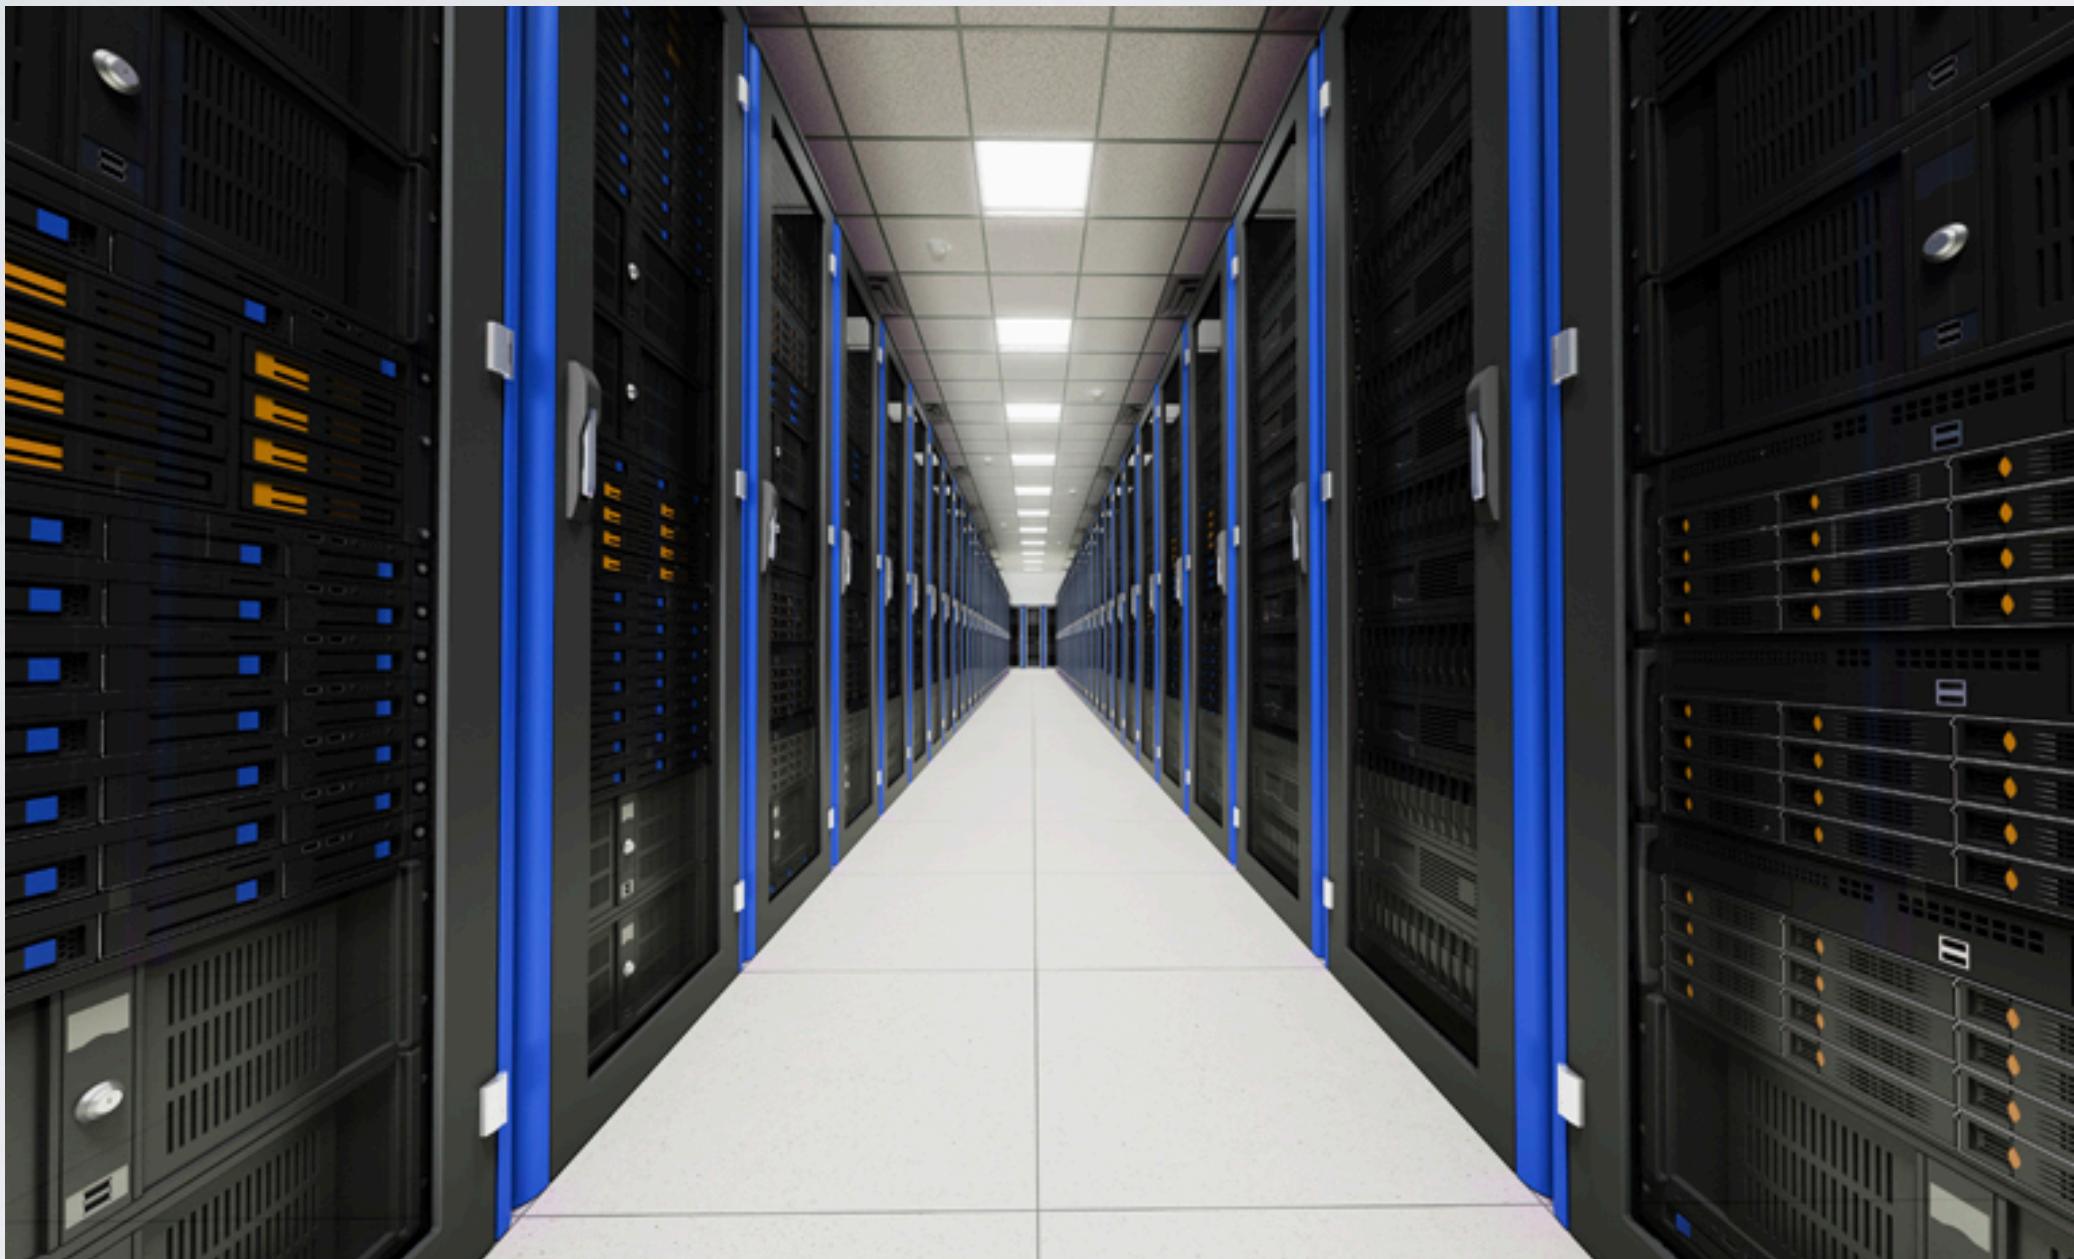
MOTIVATING LINEAR ALGEBRA

High Performance Computing (HPC)



MOTIVATING LINEAR ALGEBRA

High Performance Computing (HPC)



Expensive capital outlay

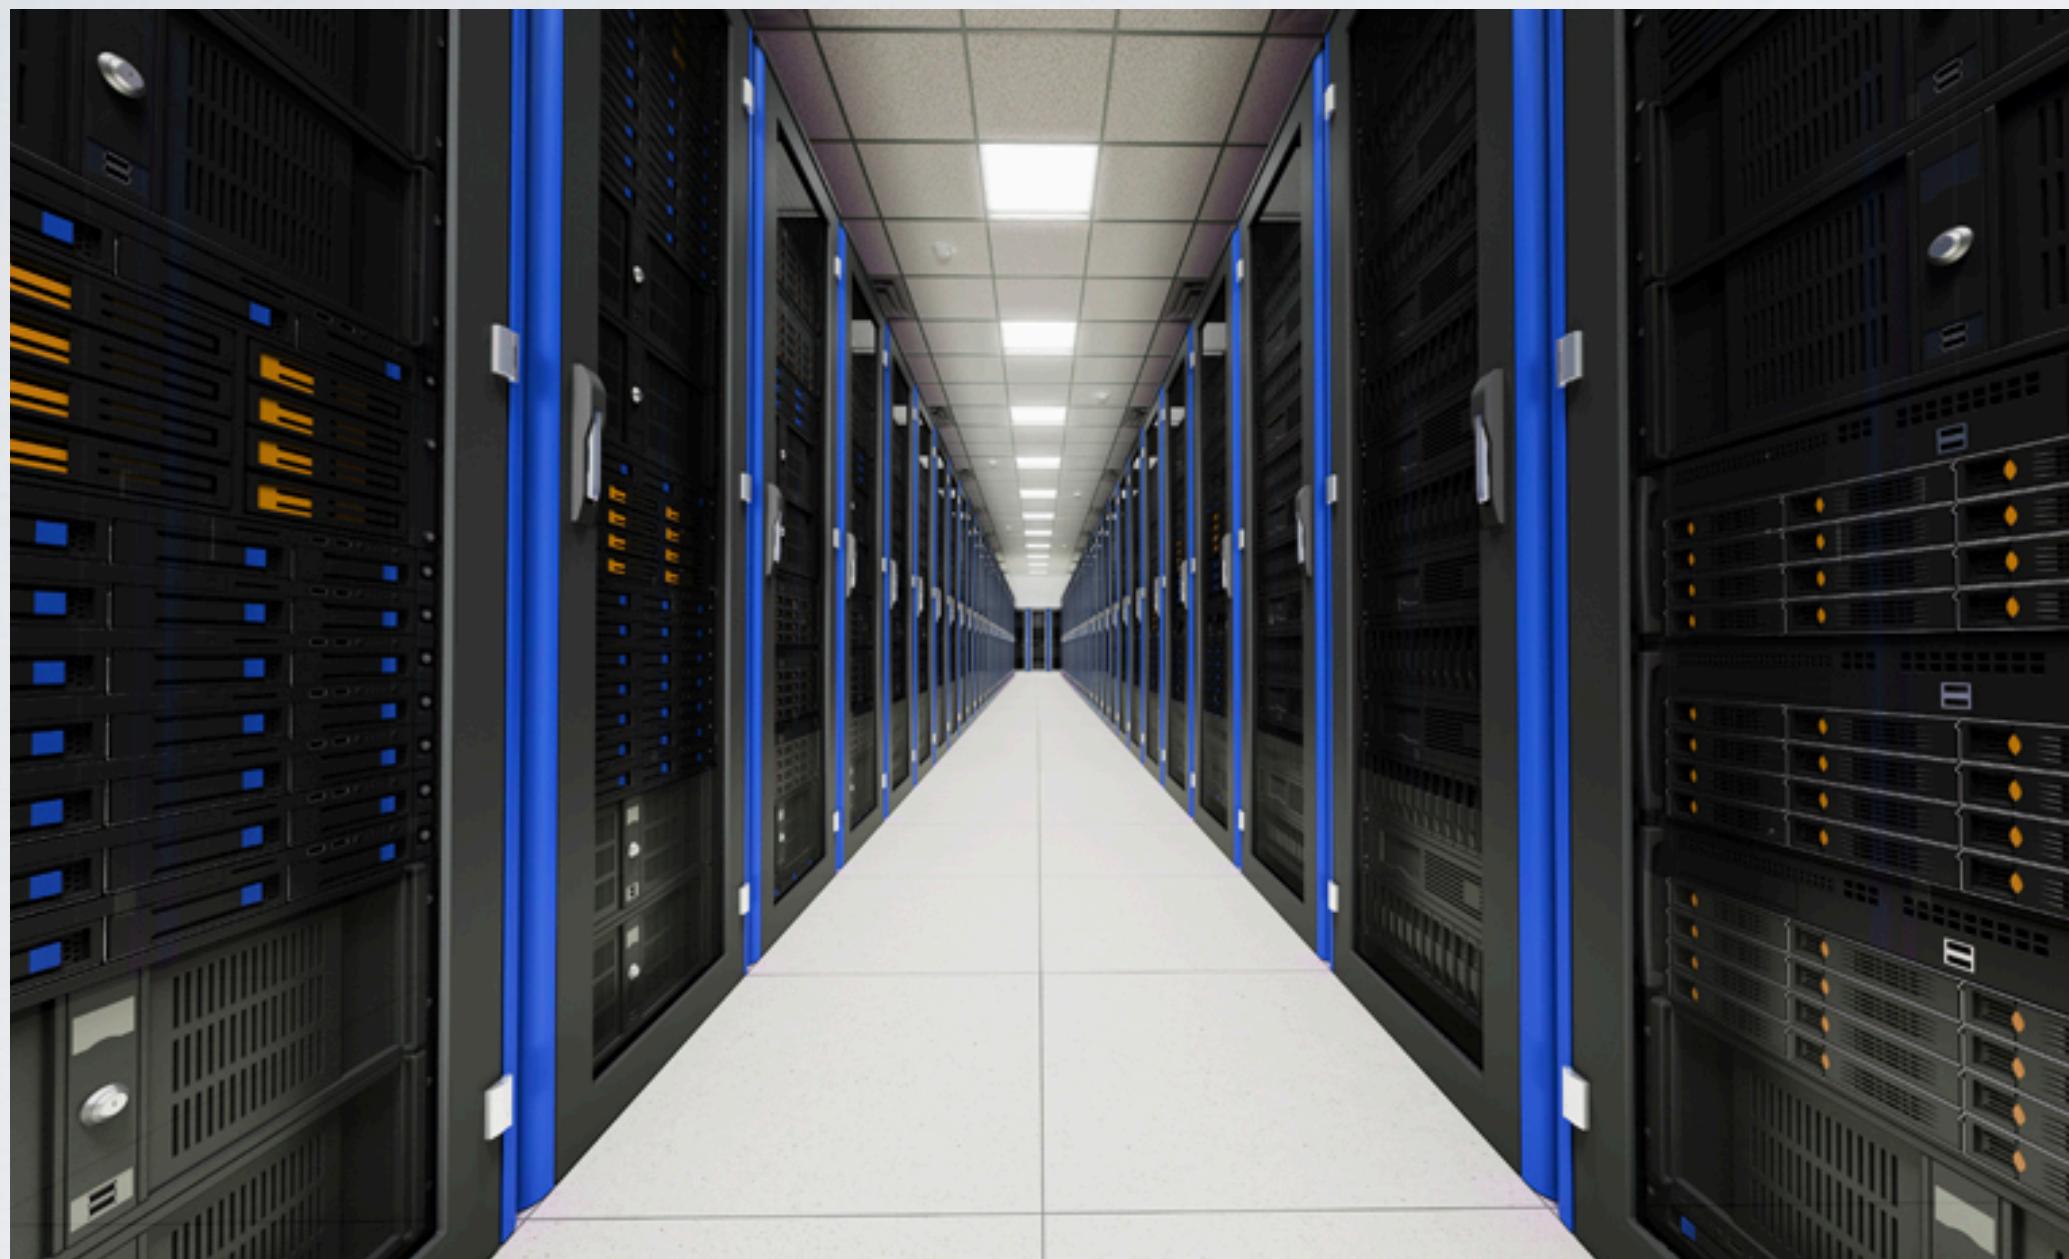
High speed interconnect

Speed is #1 job

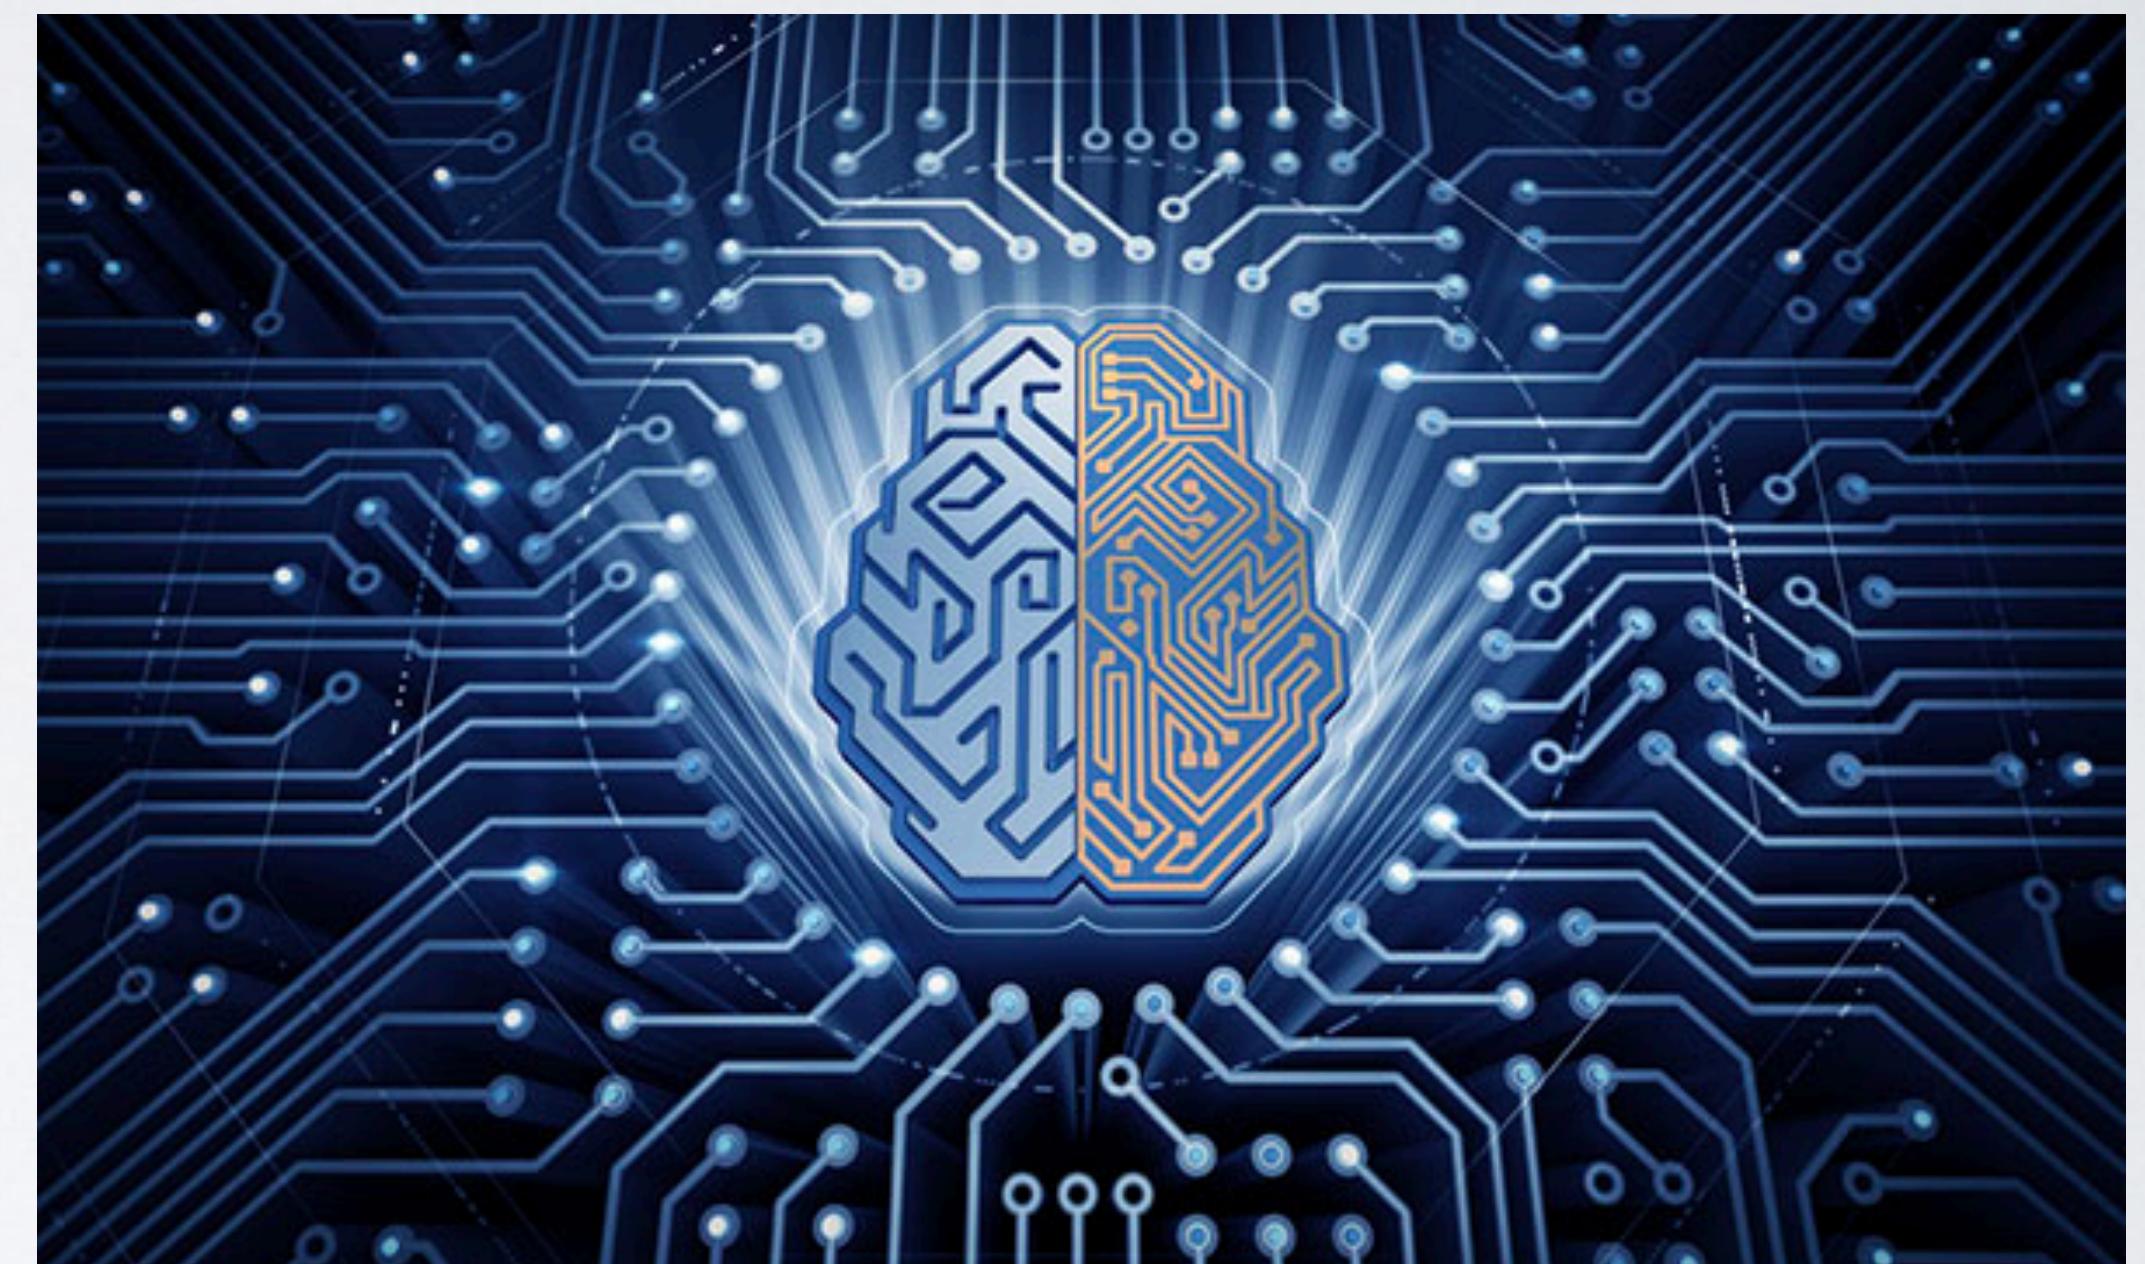
Older technology stack

MOTIVATING LINEAR ALGEBRA

High Performance Computing (HPC)



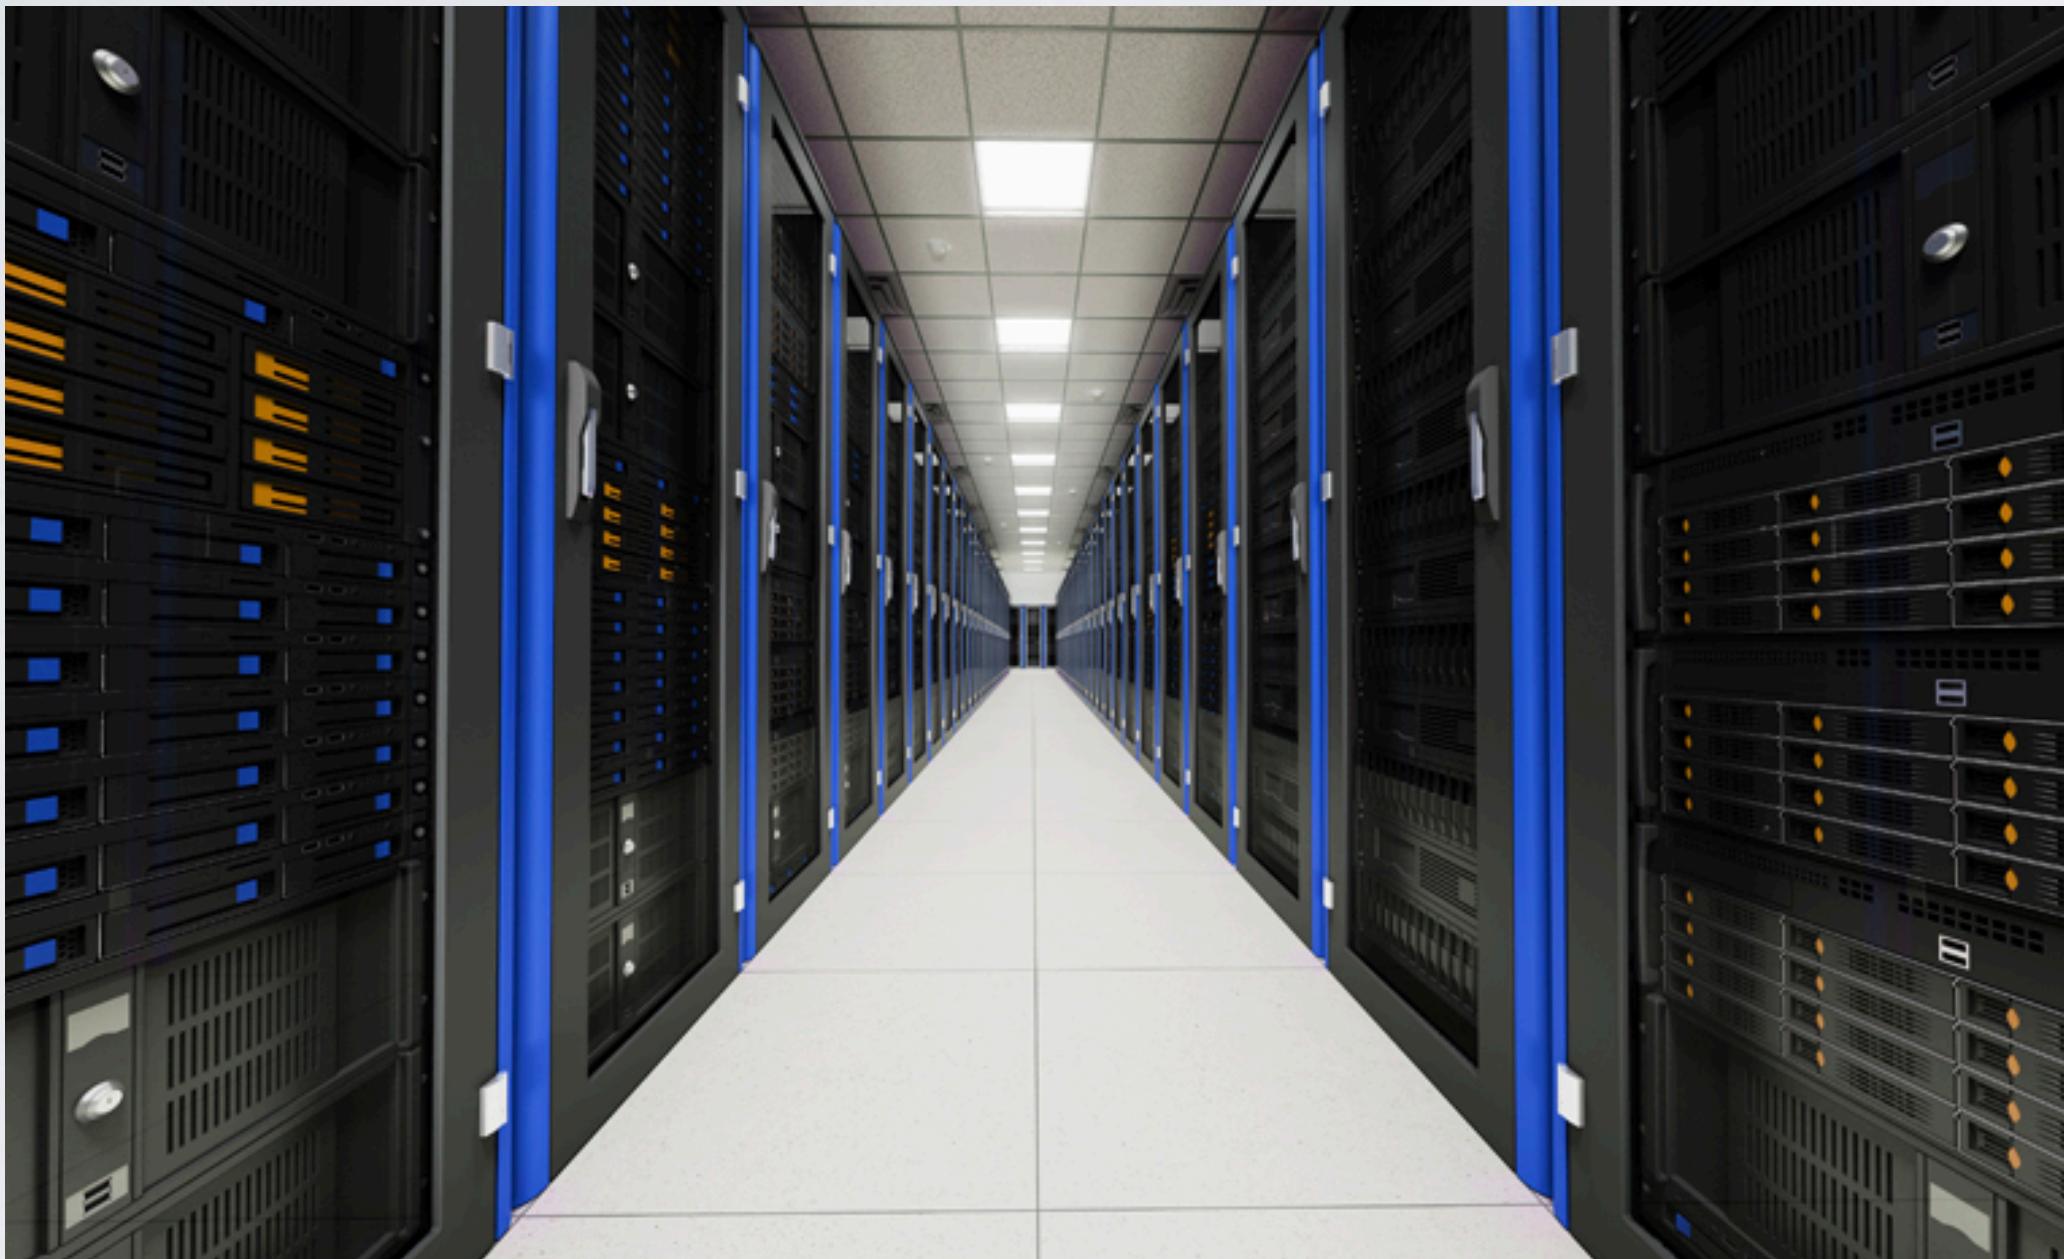
Machine Learning



Expensive capital outlay
High speed interconnect
Speed is #1 job
Older technology stack

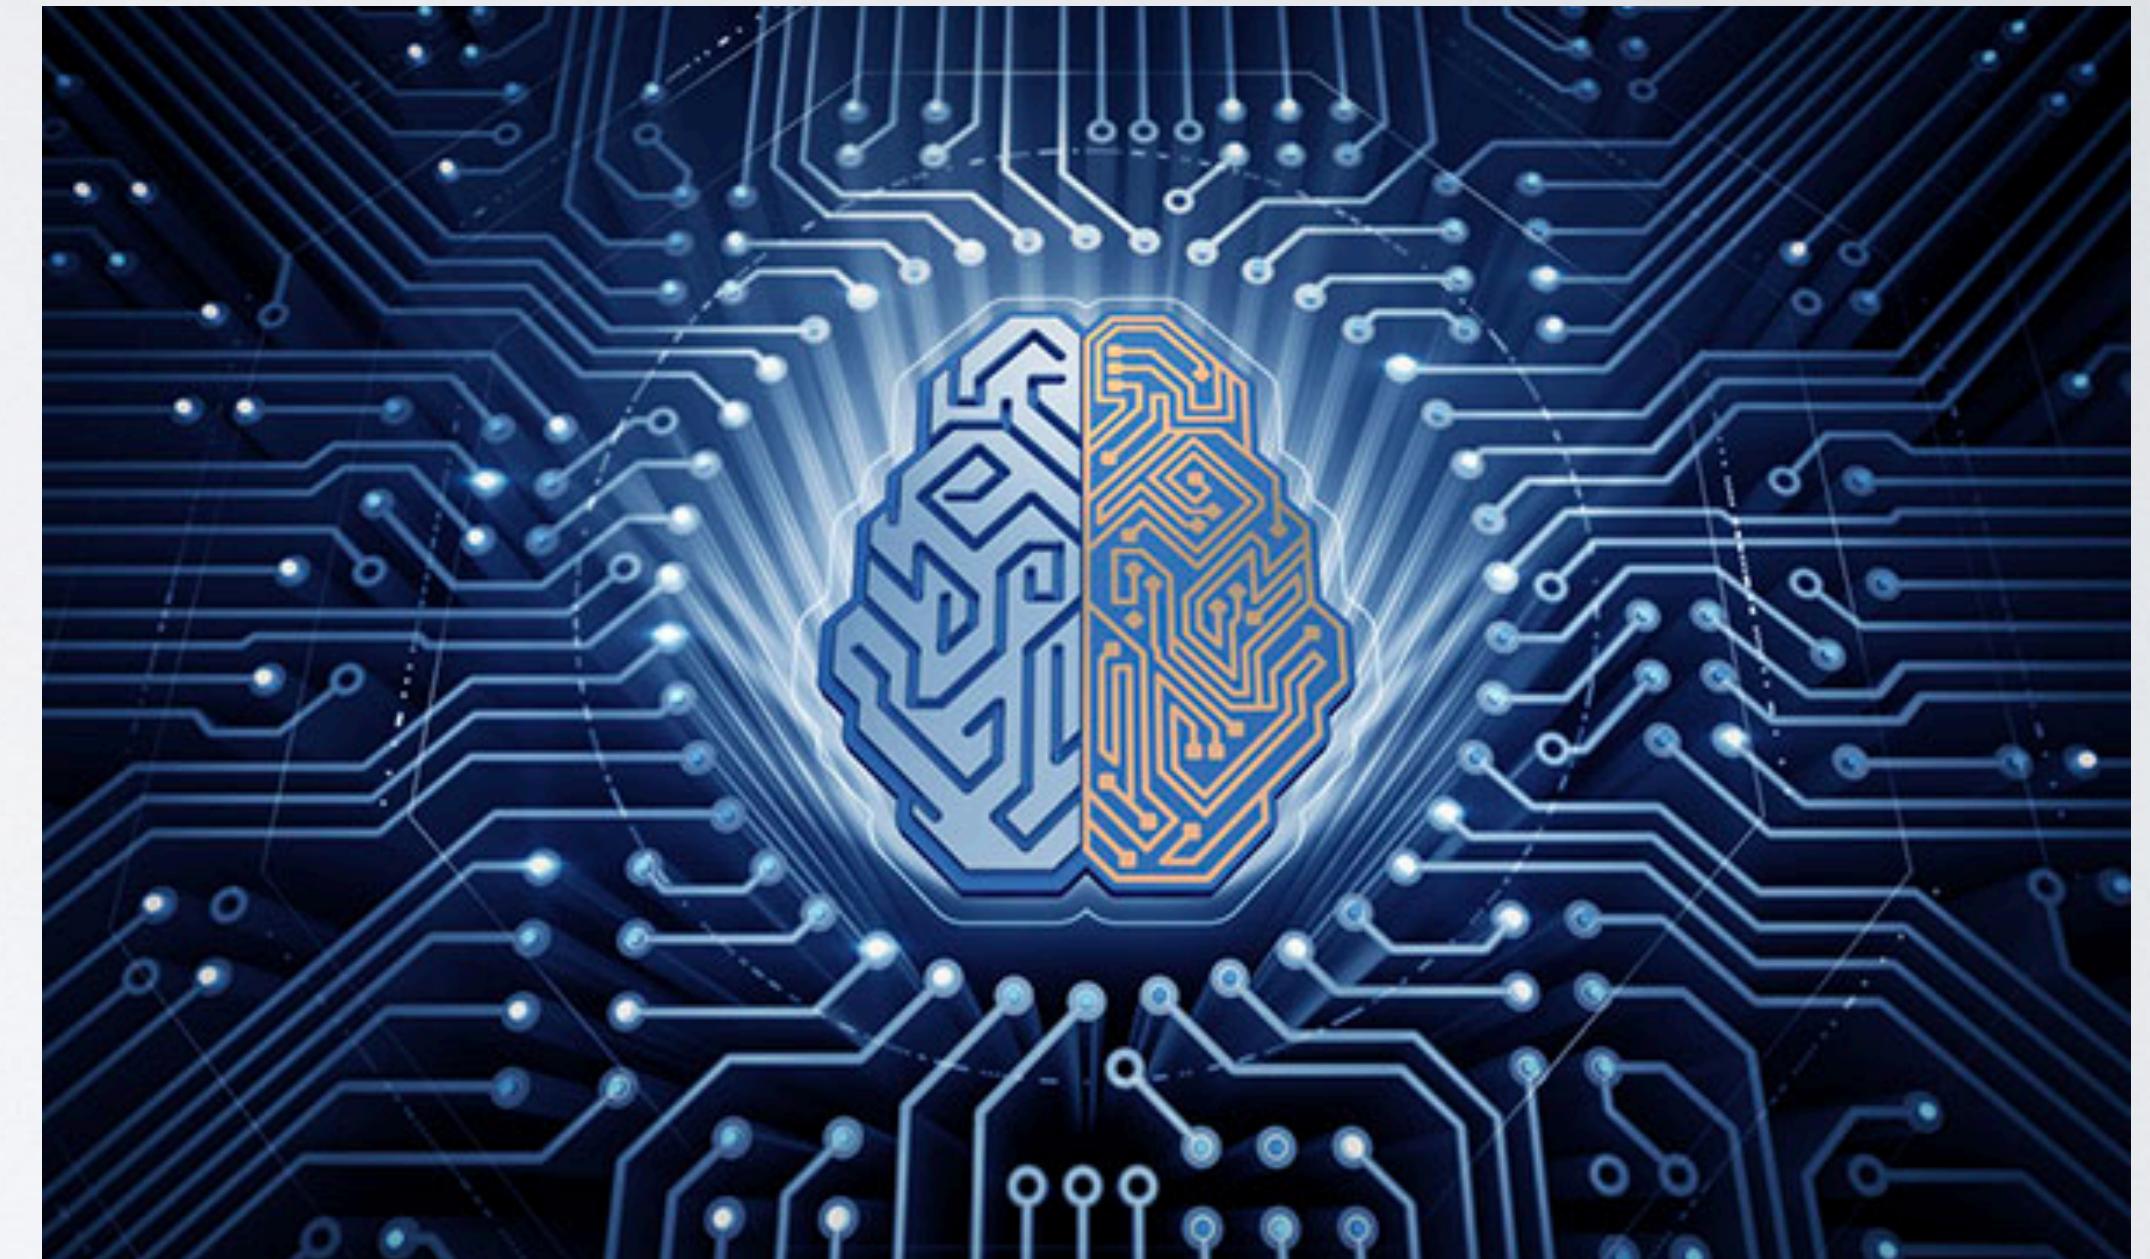
MOTIVATING LINEAR ALGEBRA

High Performance Computing (HPC)



Expensive capital outlay
High speed interconnect
Speed is #1 job
Older technology stack

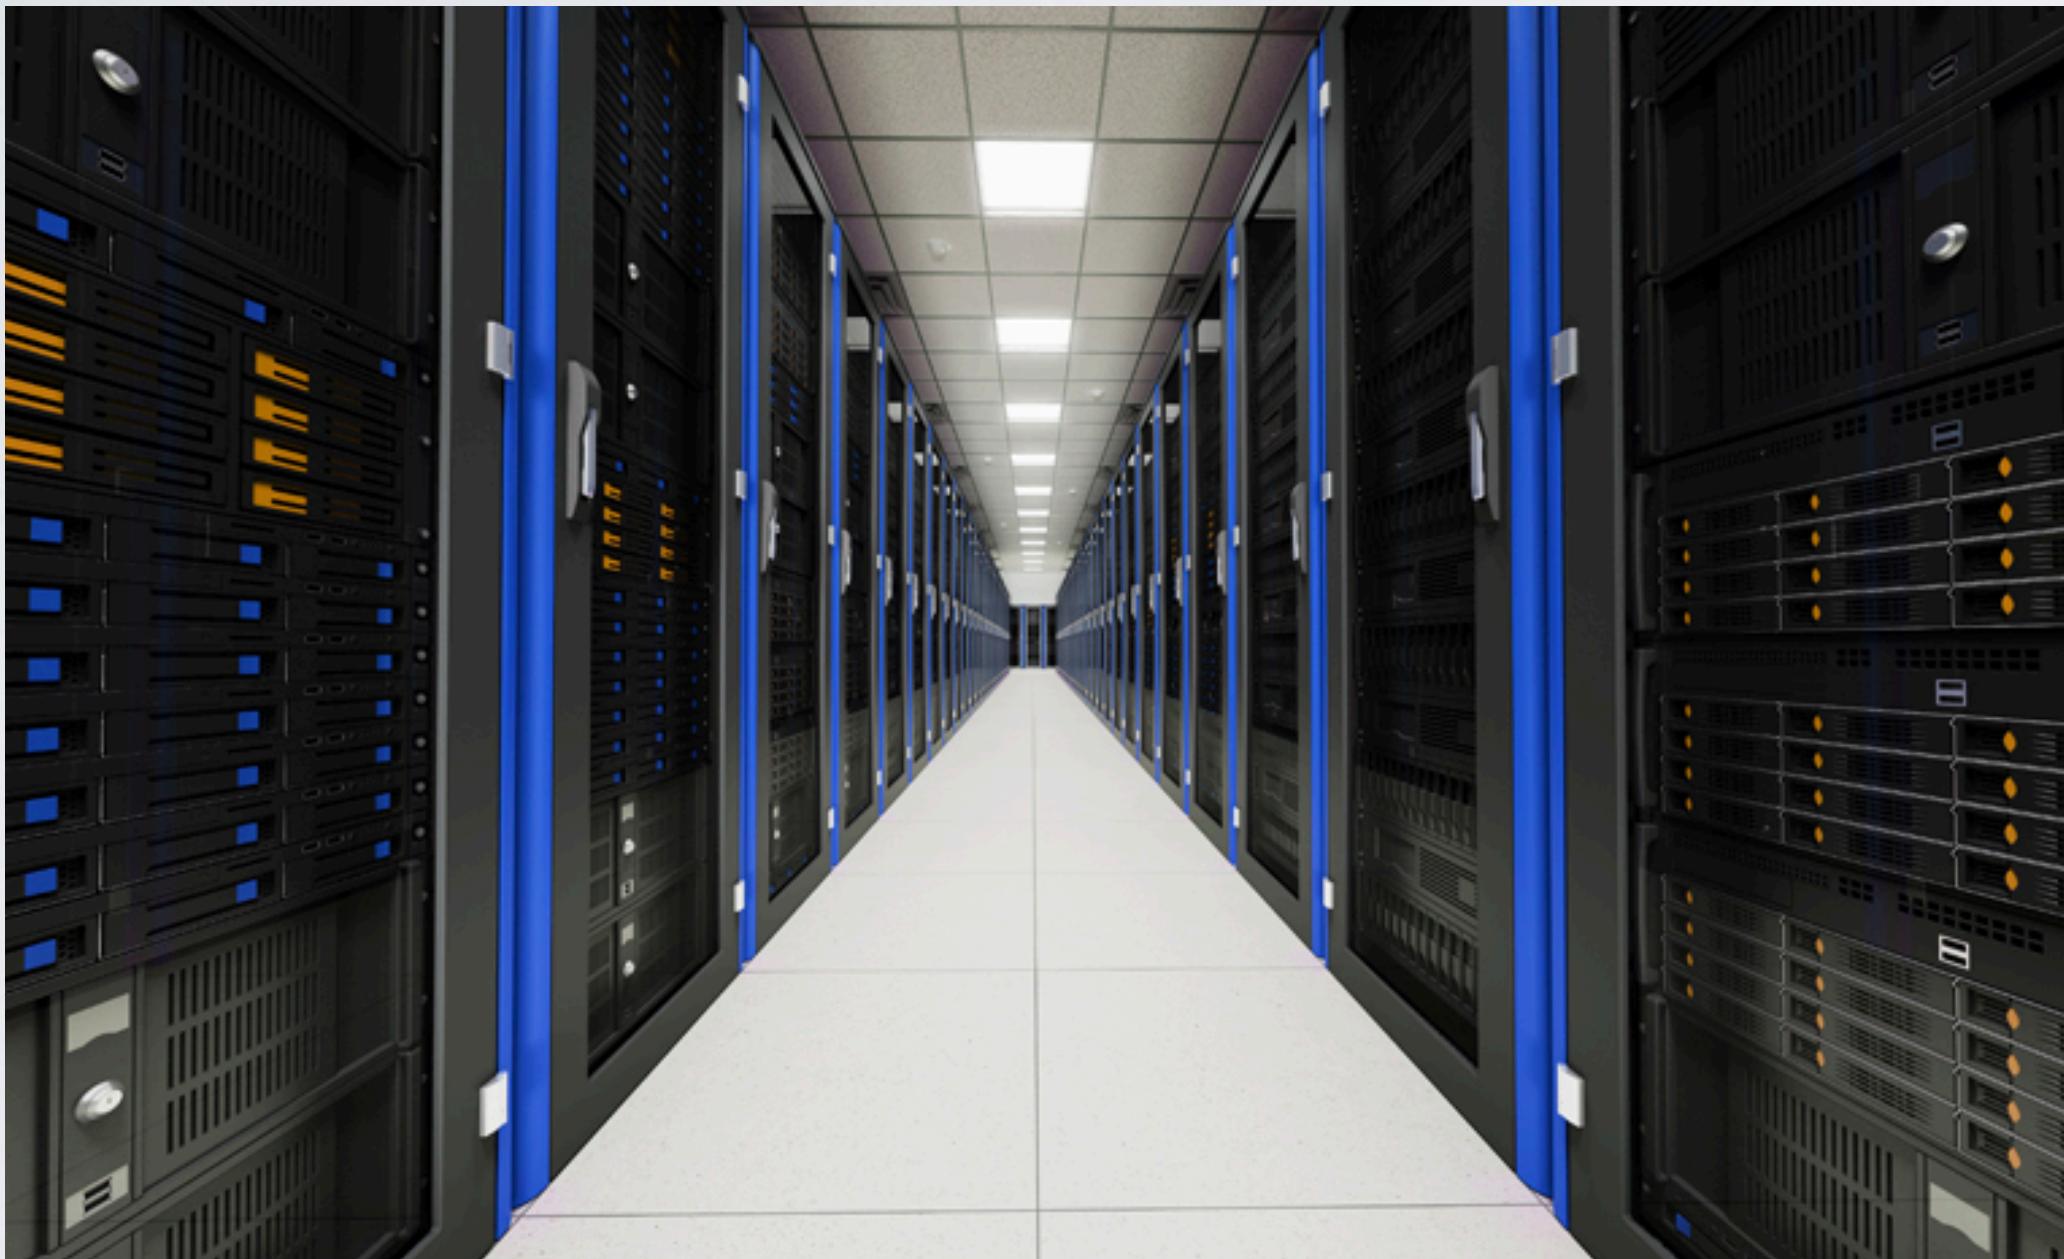
Machine Learning



Focus on deep method
Everything is streaming
Does this really work?

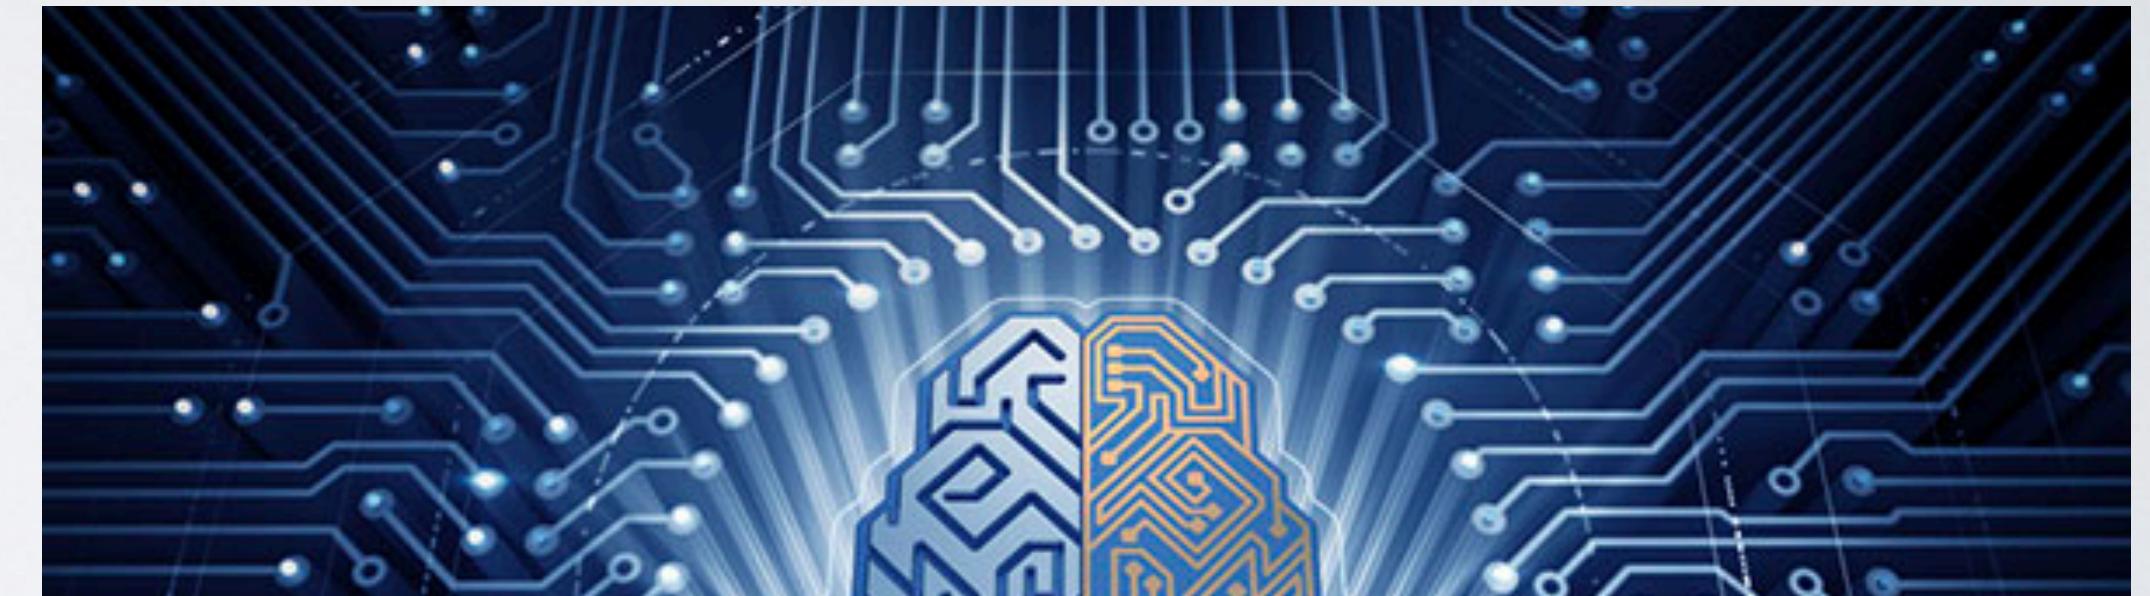
MOTIVATING LINEAR ALGEBRA

High Performance Computing (HPC)



Expensive capital outlay
High speed interconnect
Speed is #1 job
Older technology stack

Machine Learning



"It's easier to train a deep neural bidirectional LSTM with attention than it is to compute the SVD of a large matrix" - Chris Re

Focus on deep method
Everything is streaming
Does this really work?

TRENDS AND OBSERVATIONS

Compute more precious



Fast cheap disaggregated state



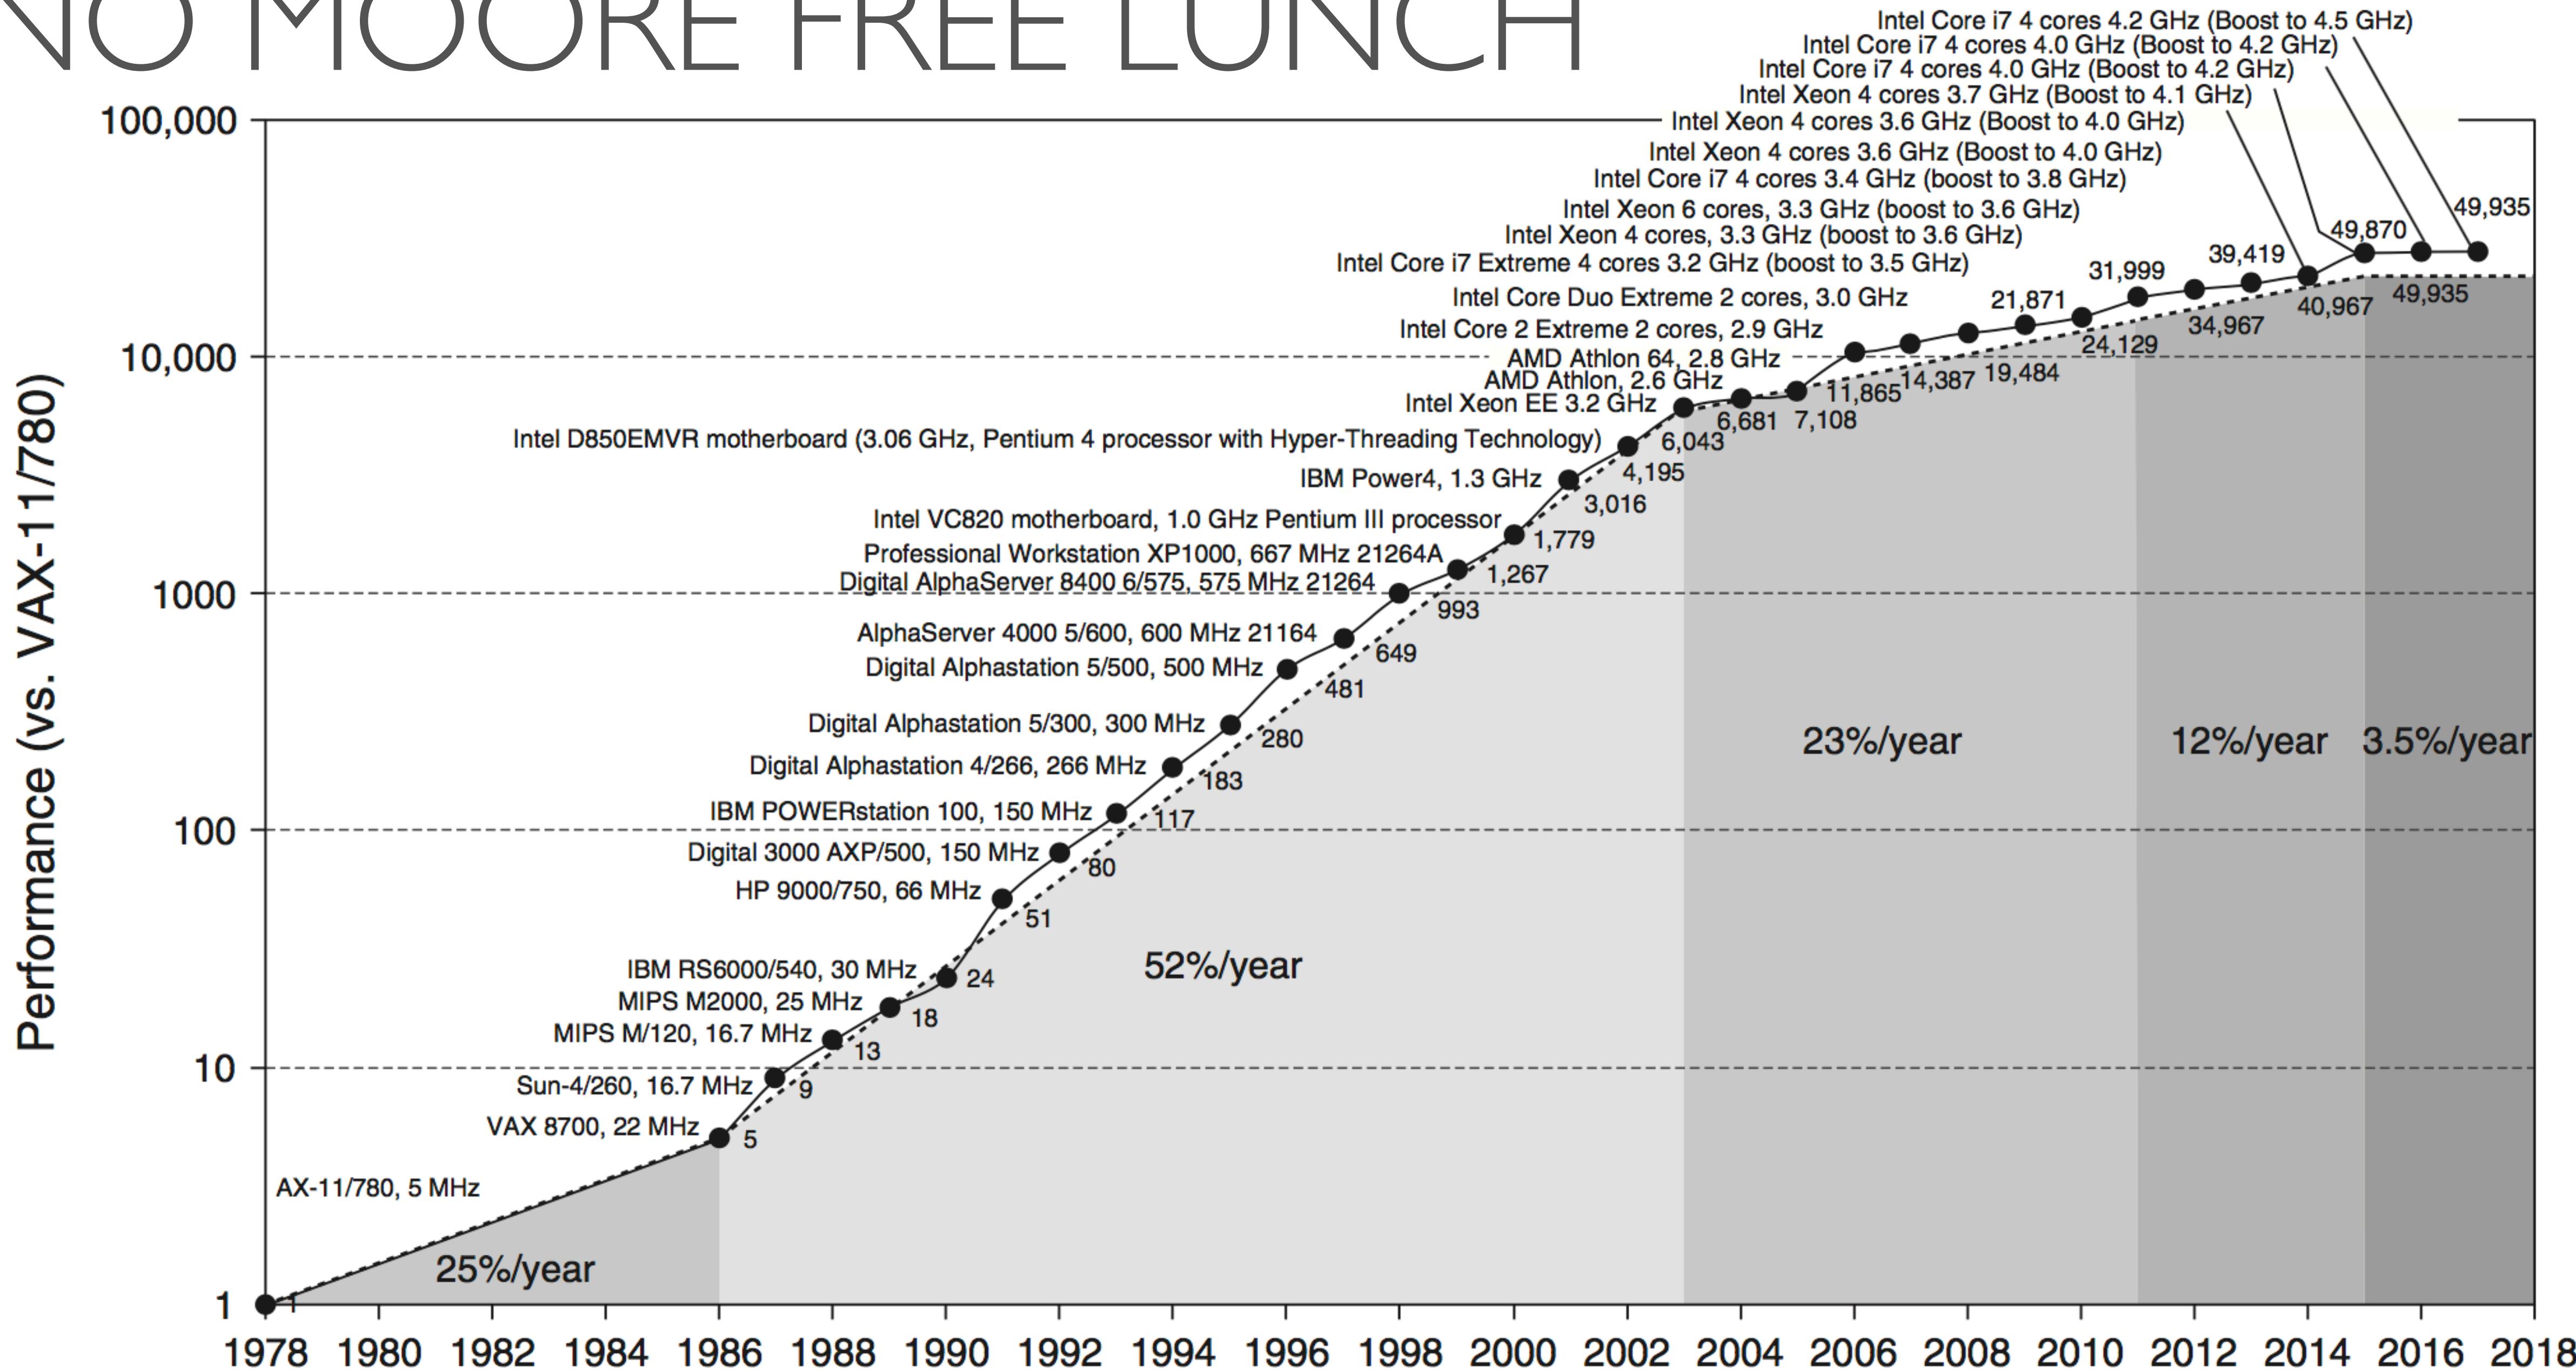
Algorithms with dynamic parallelism



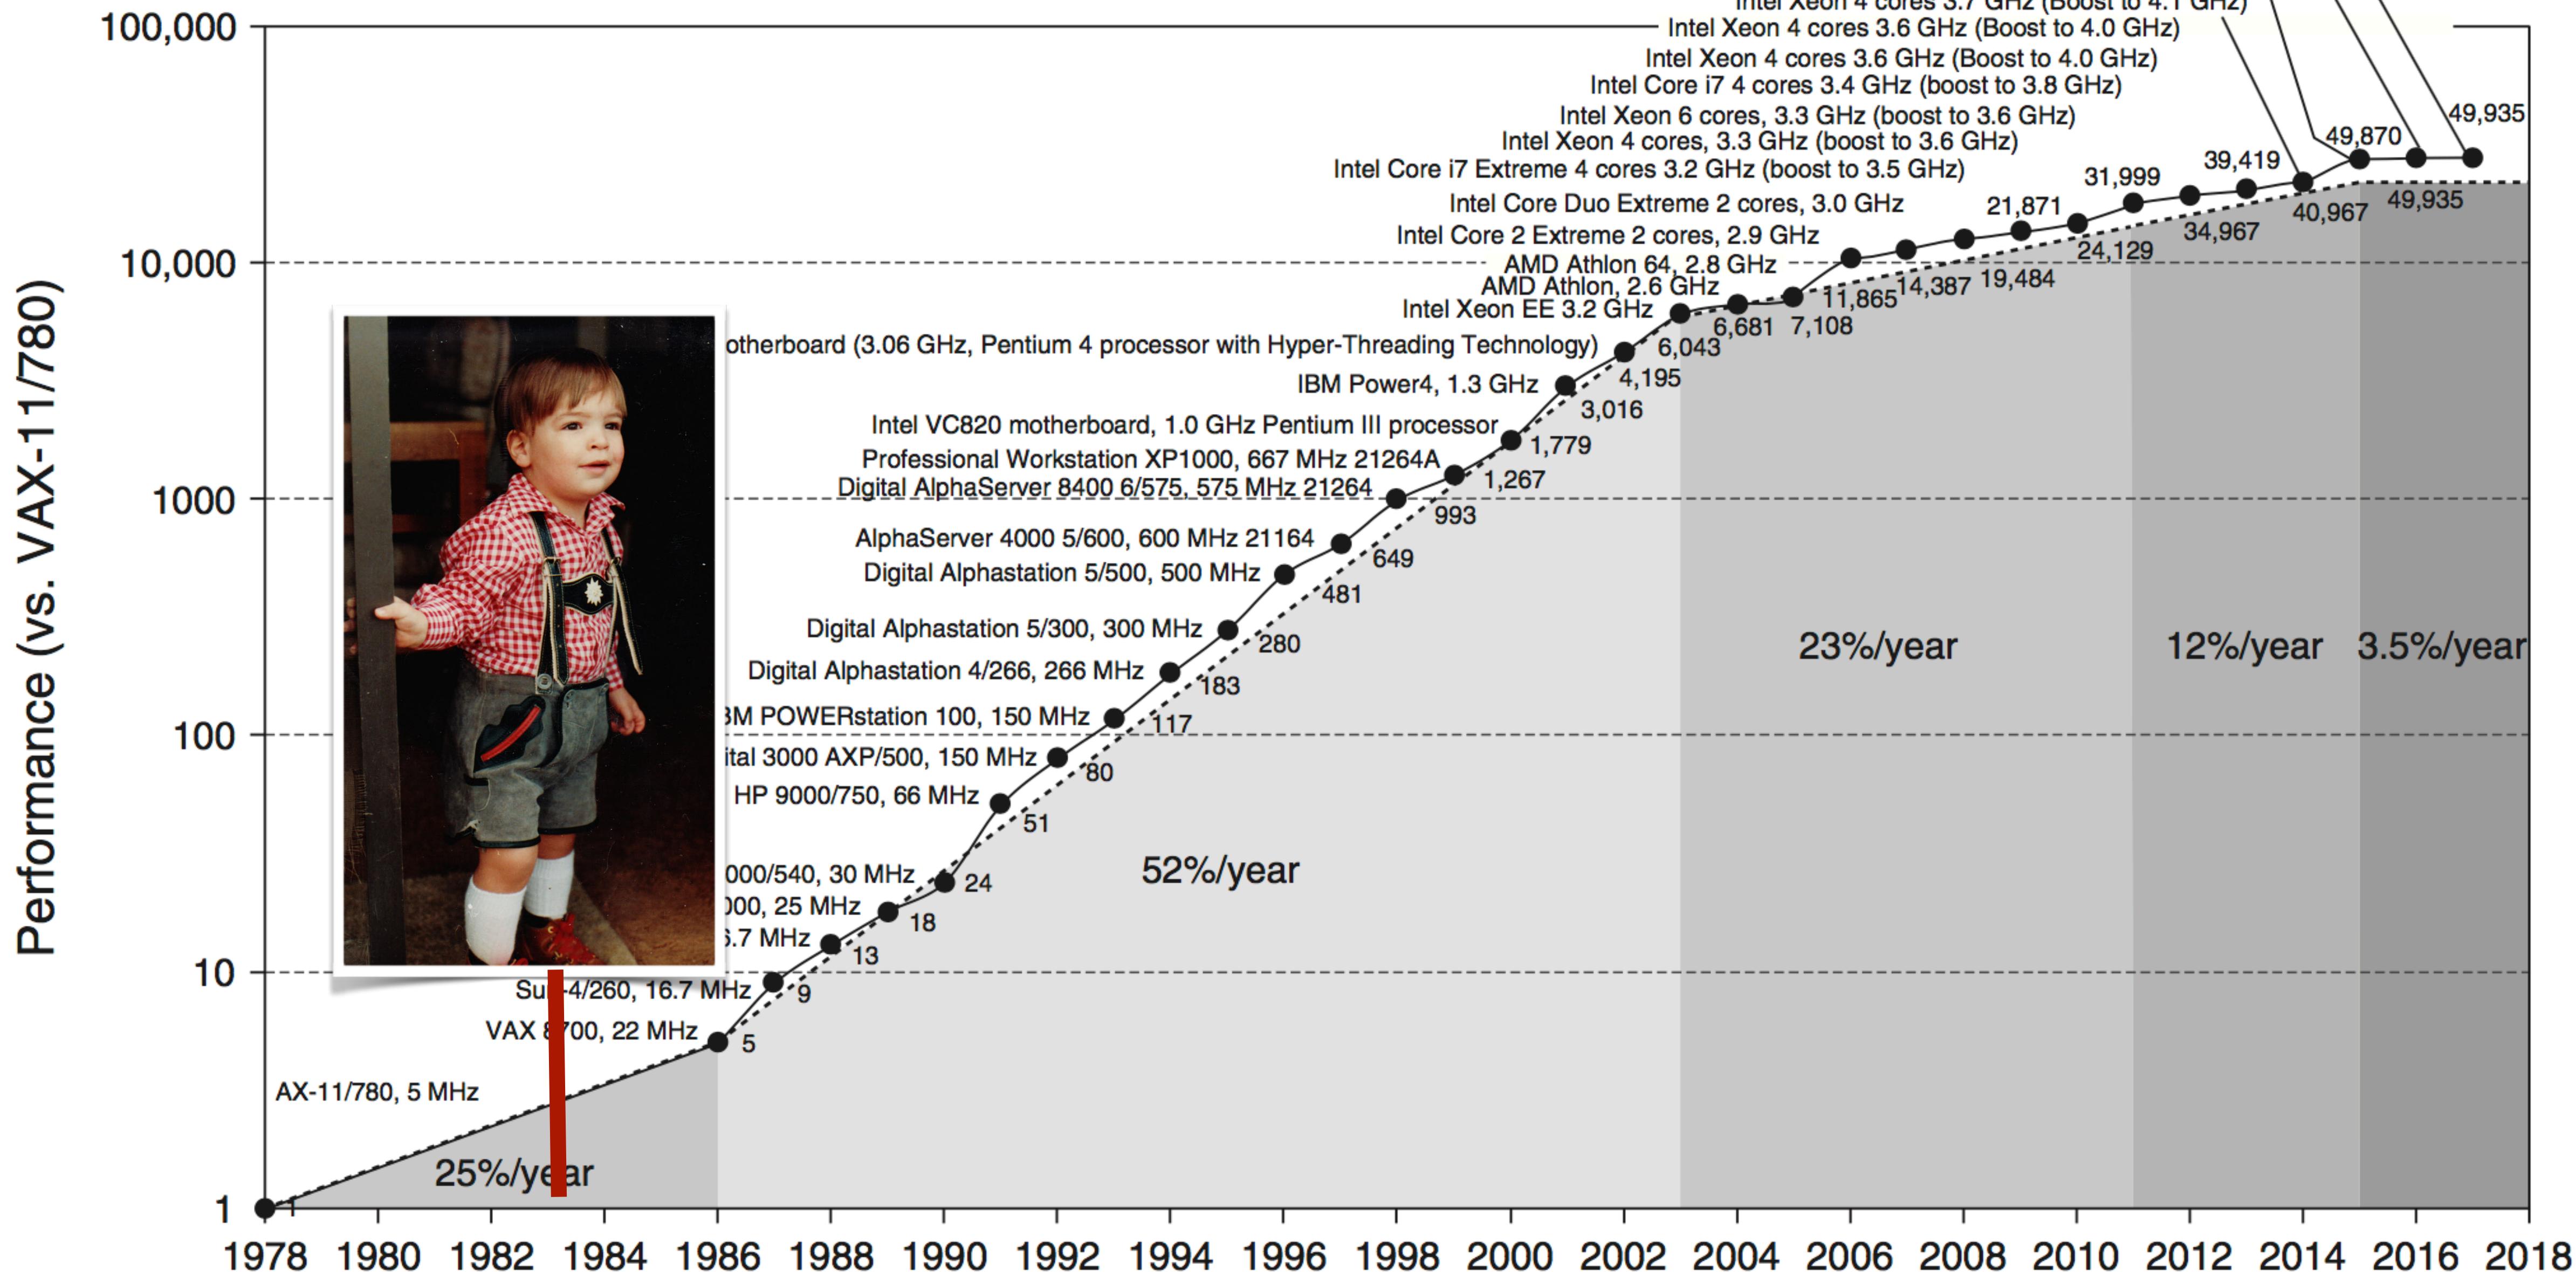
Operations where compute dominates IO

$$O(n^3) > O(n^2)$$

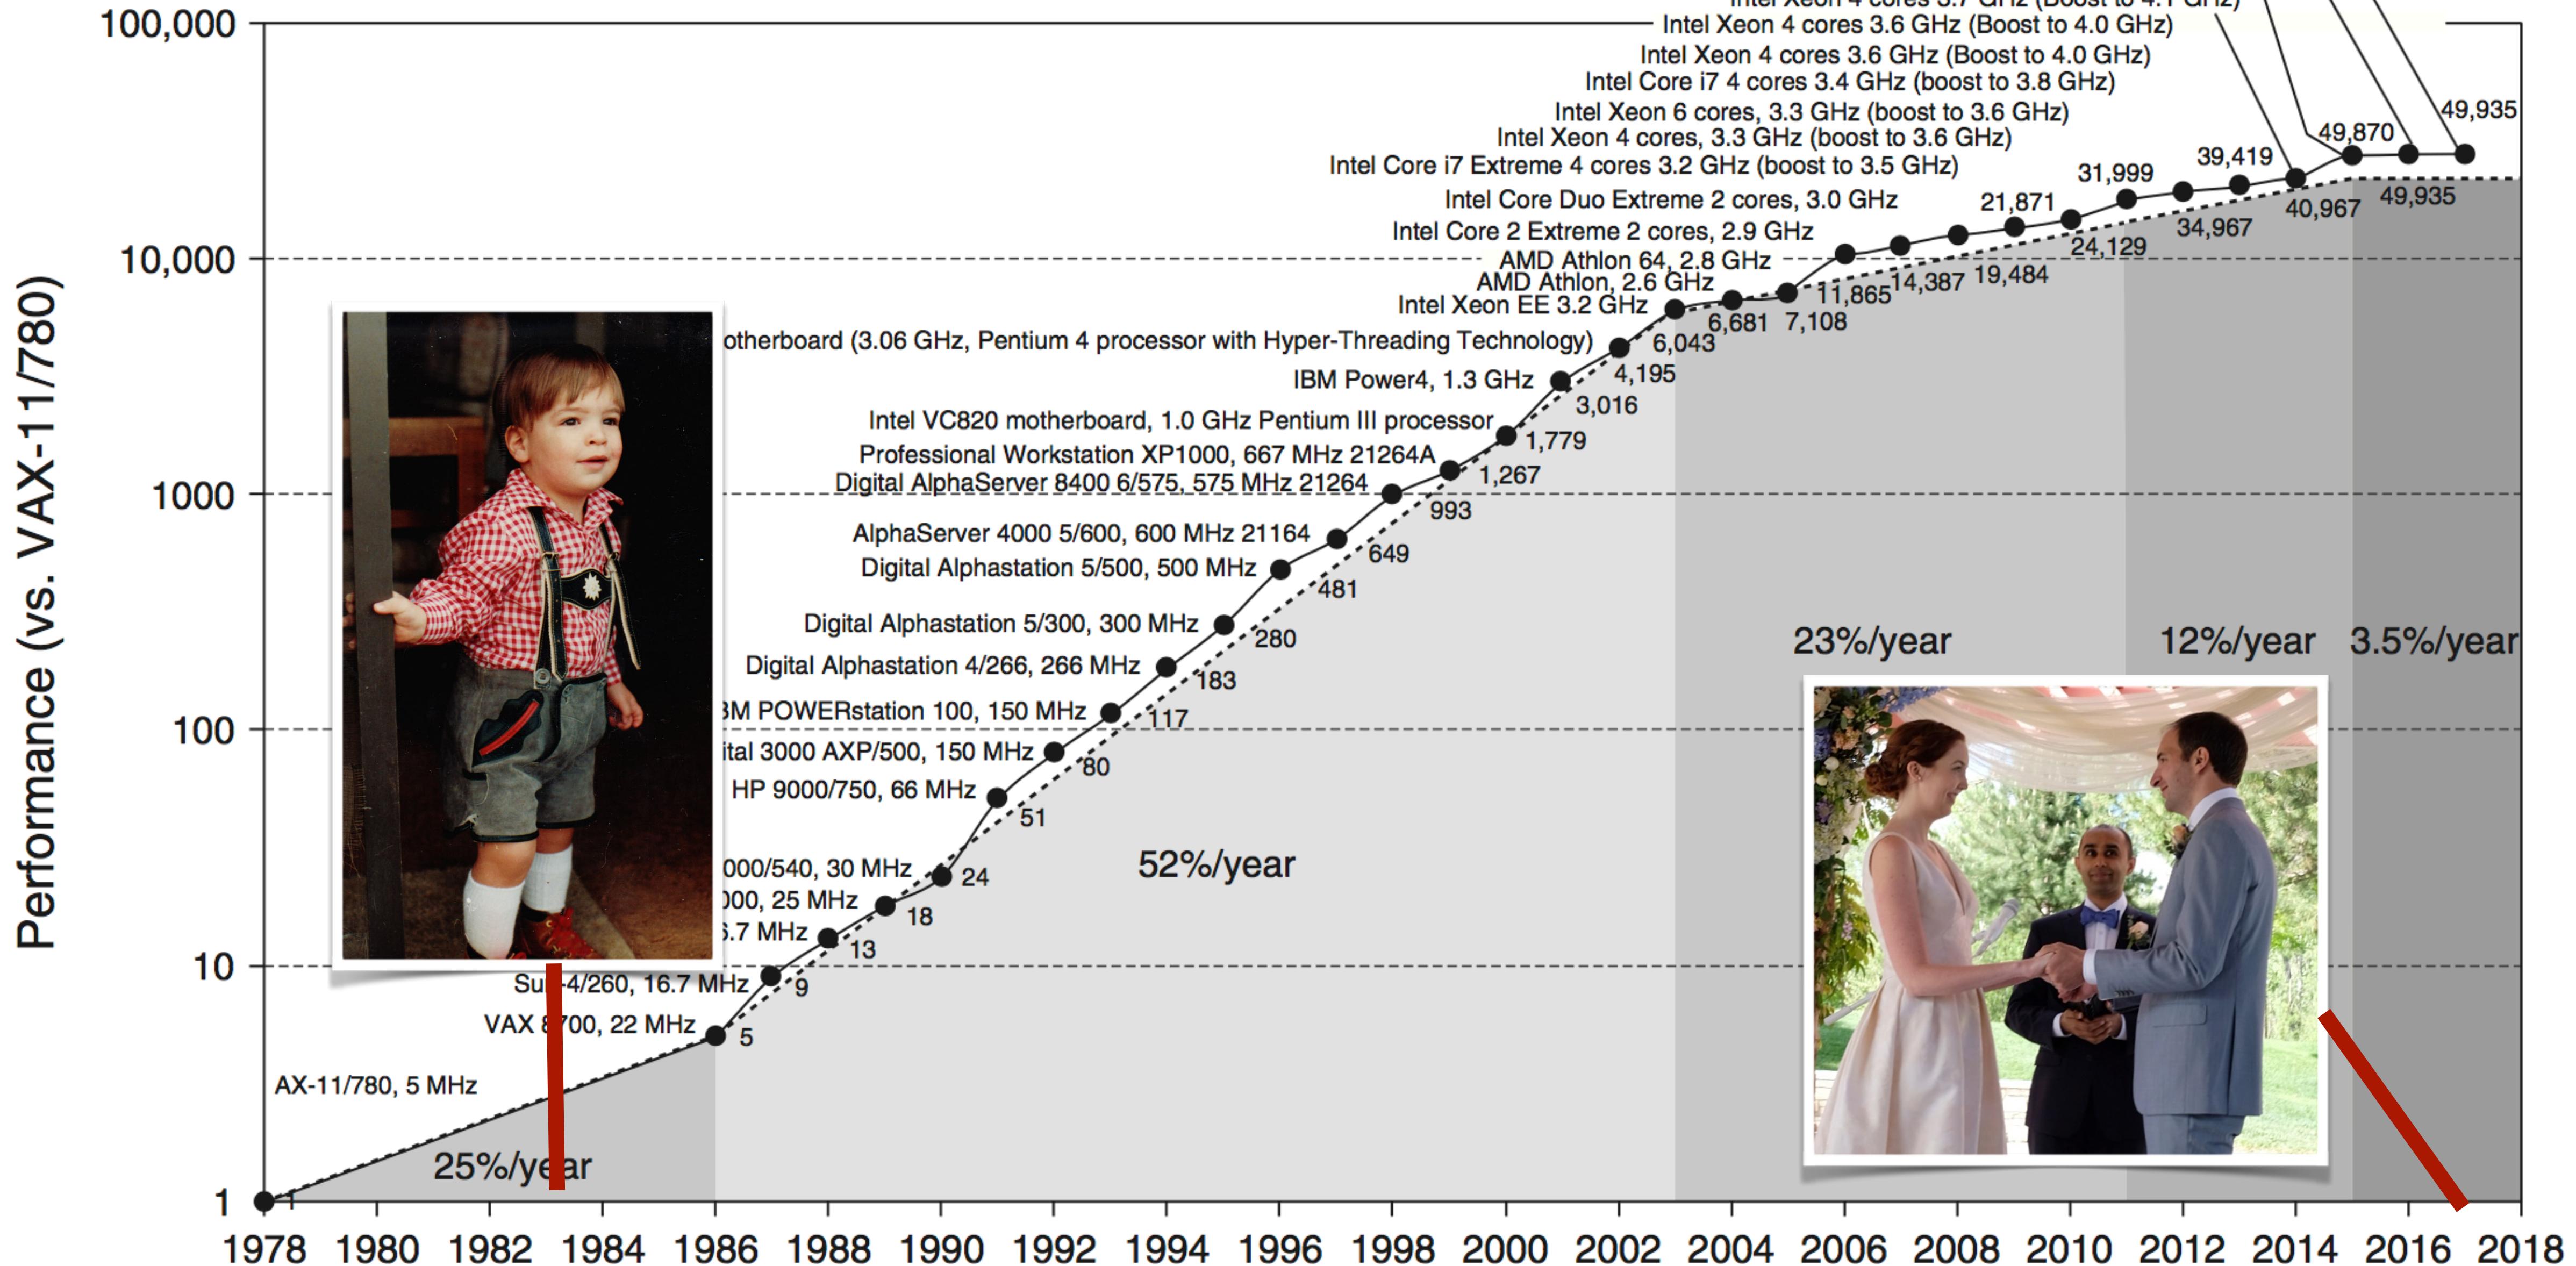
NO MOORE FREE LUNCH



NO MOORE FREE LUNCH



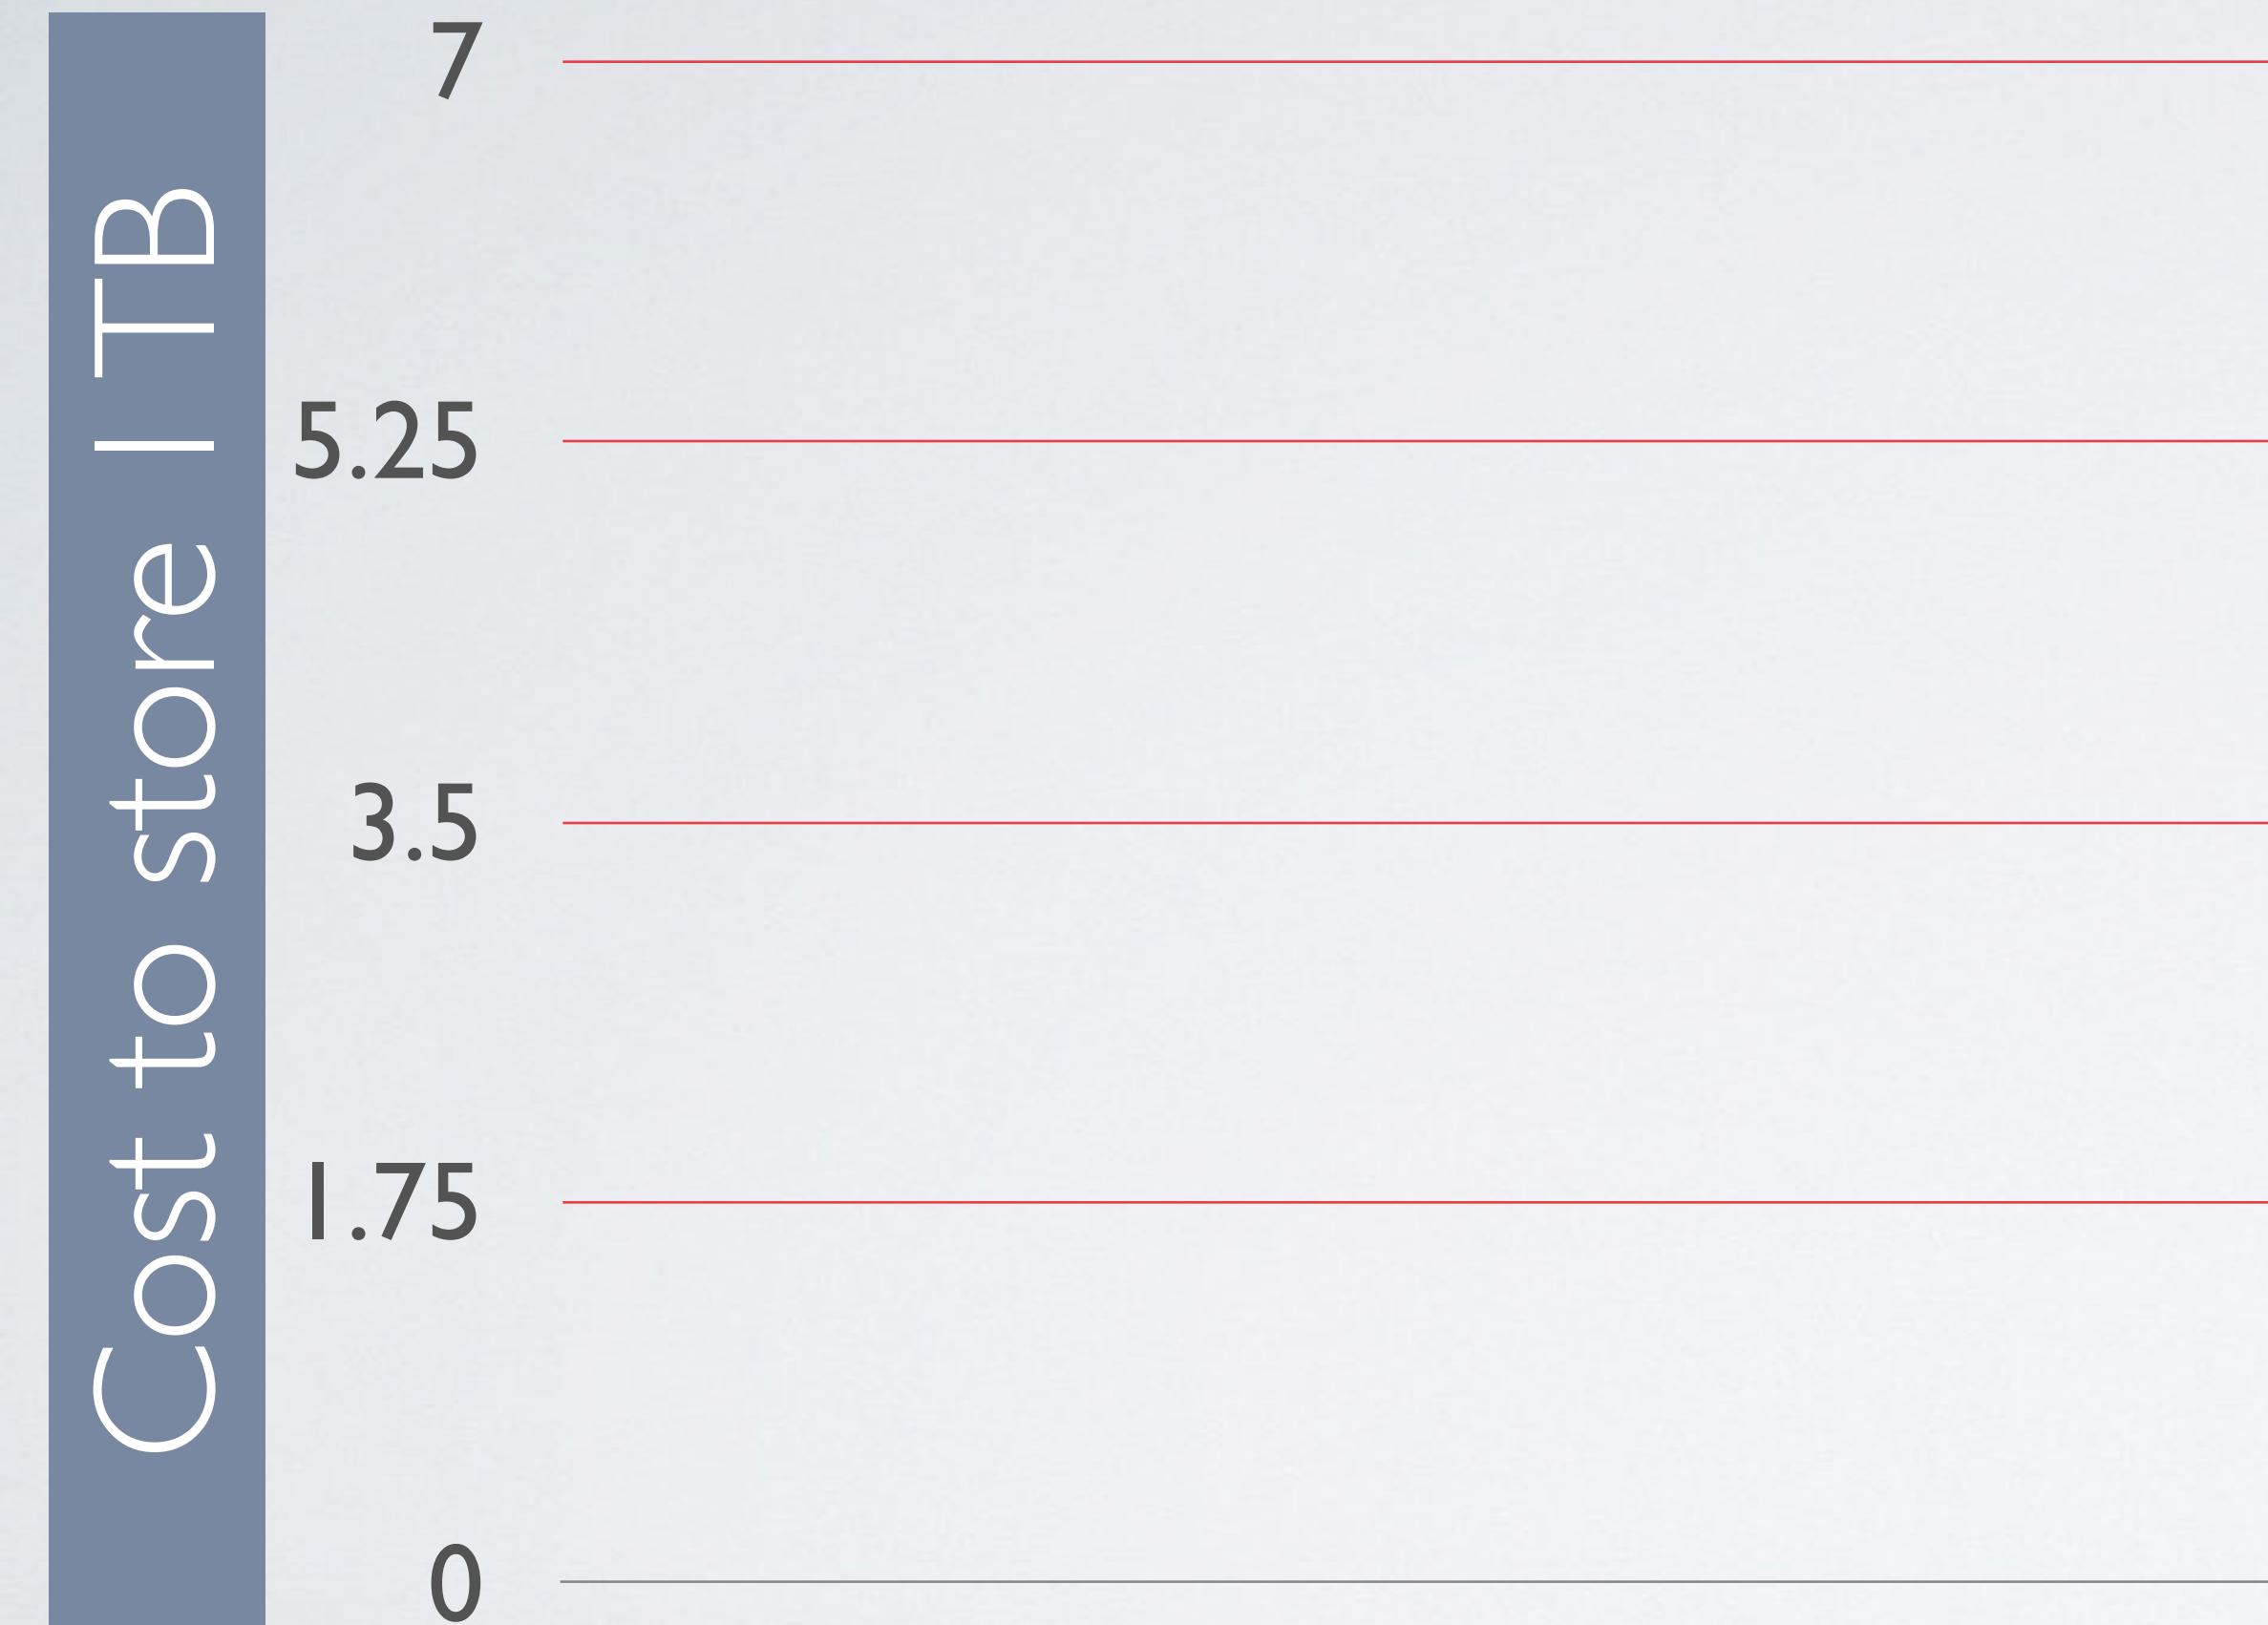
NO MOORE FREE LUNCH



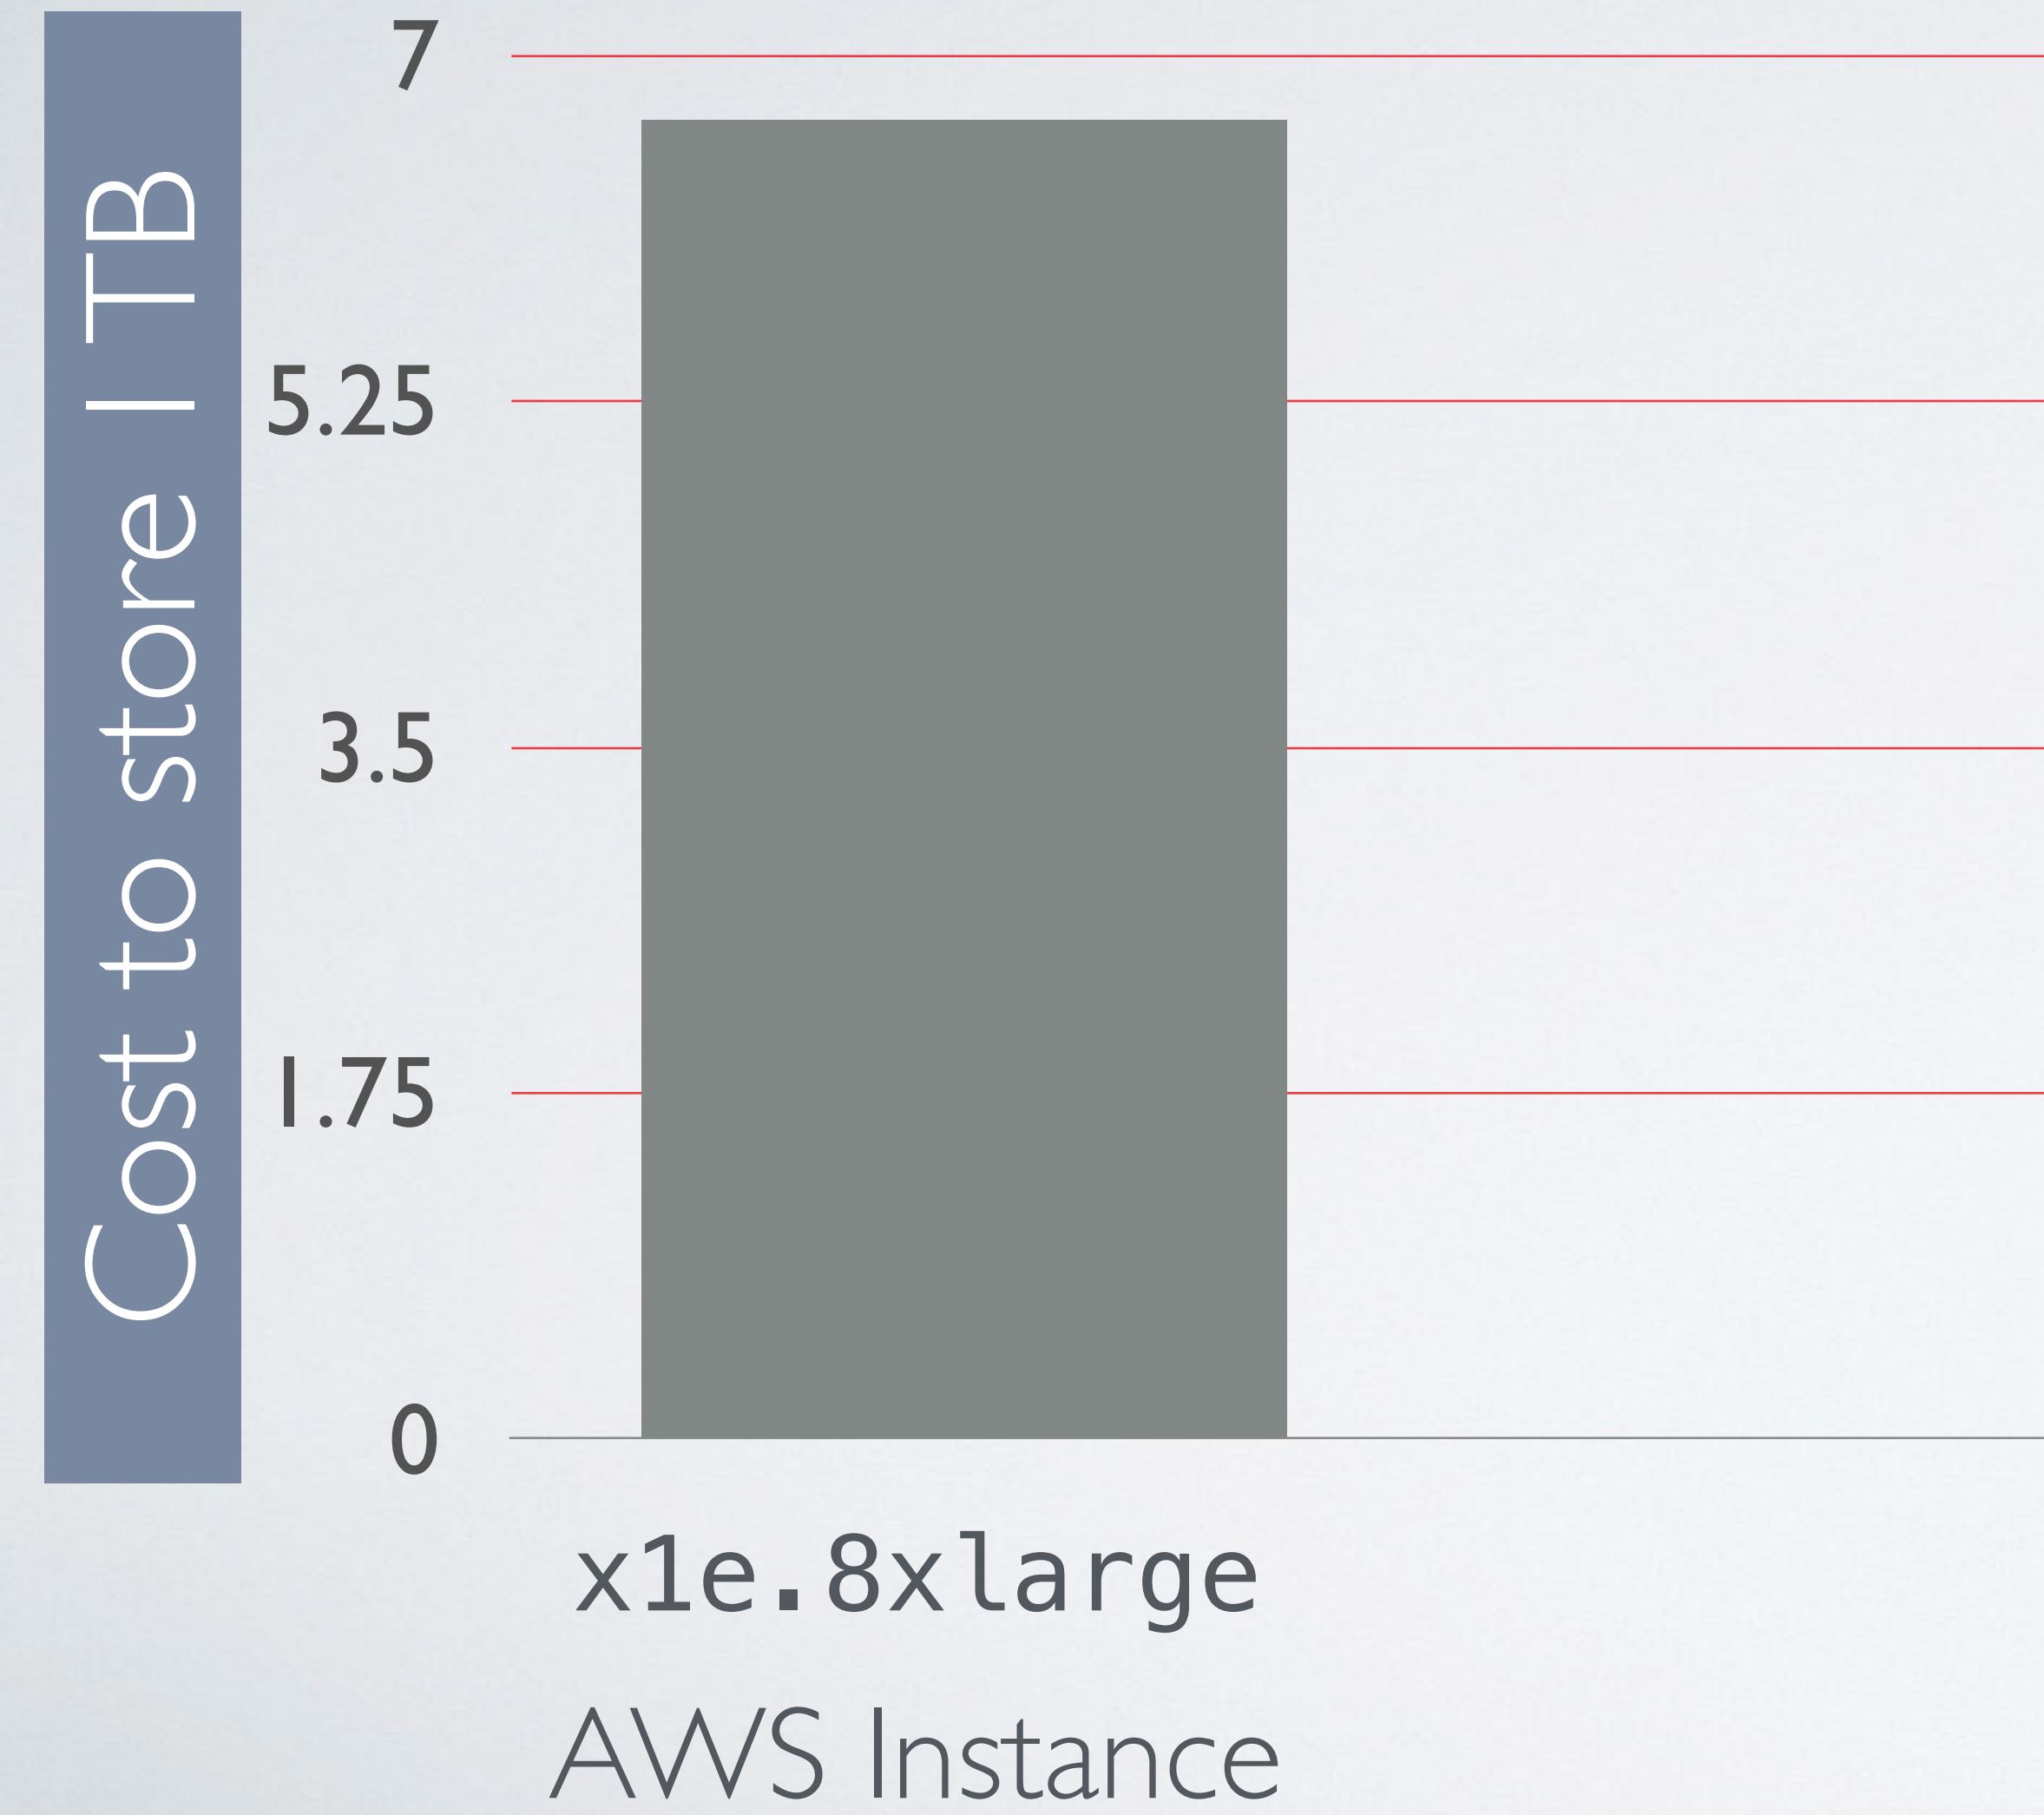
DATA CENTER DISAGGREGATION

Cost to store | TB

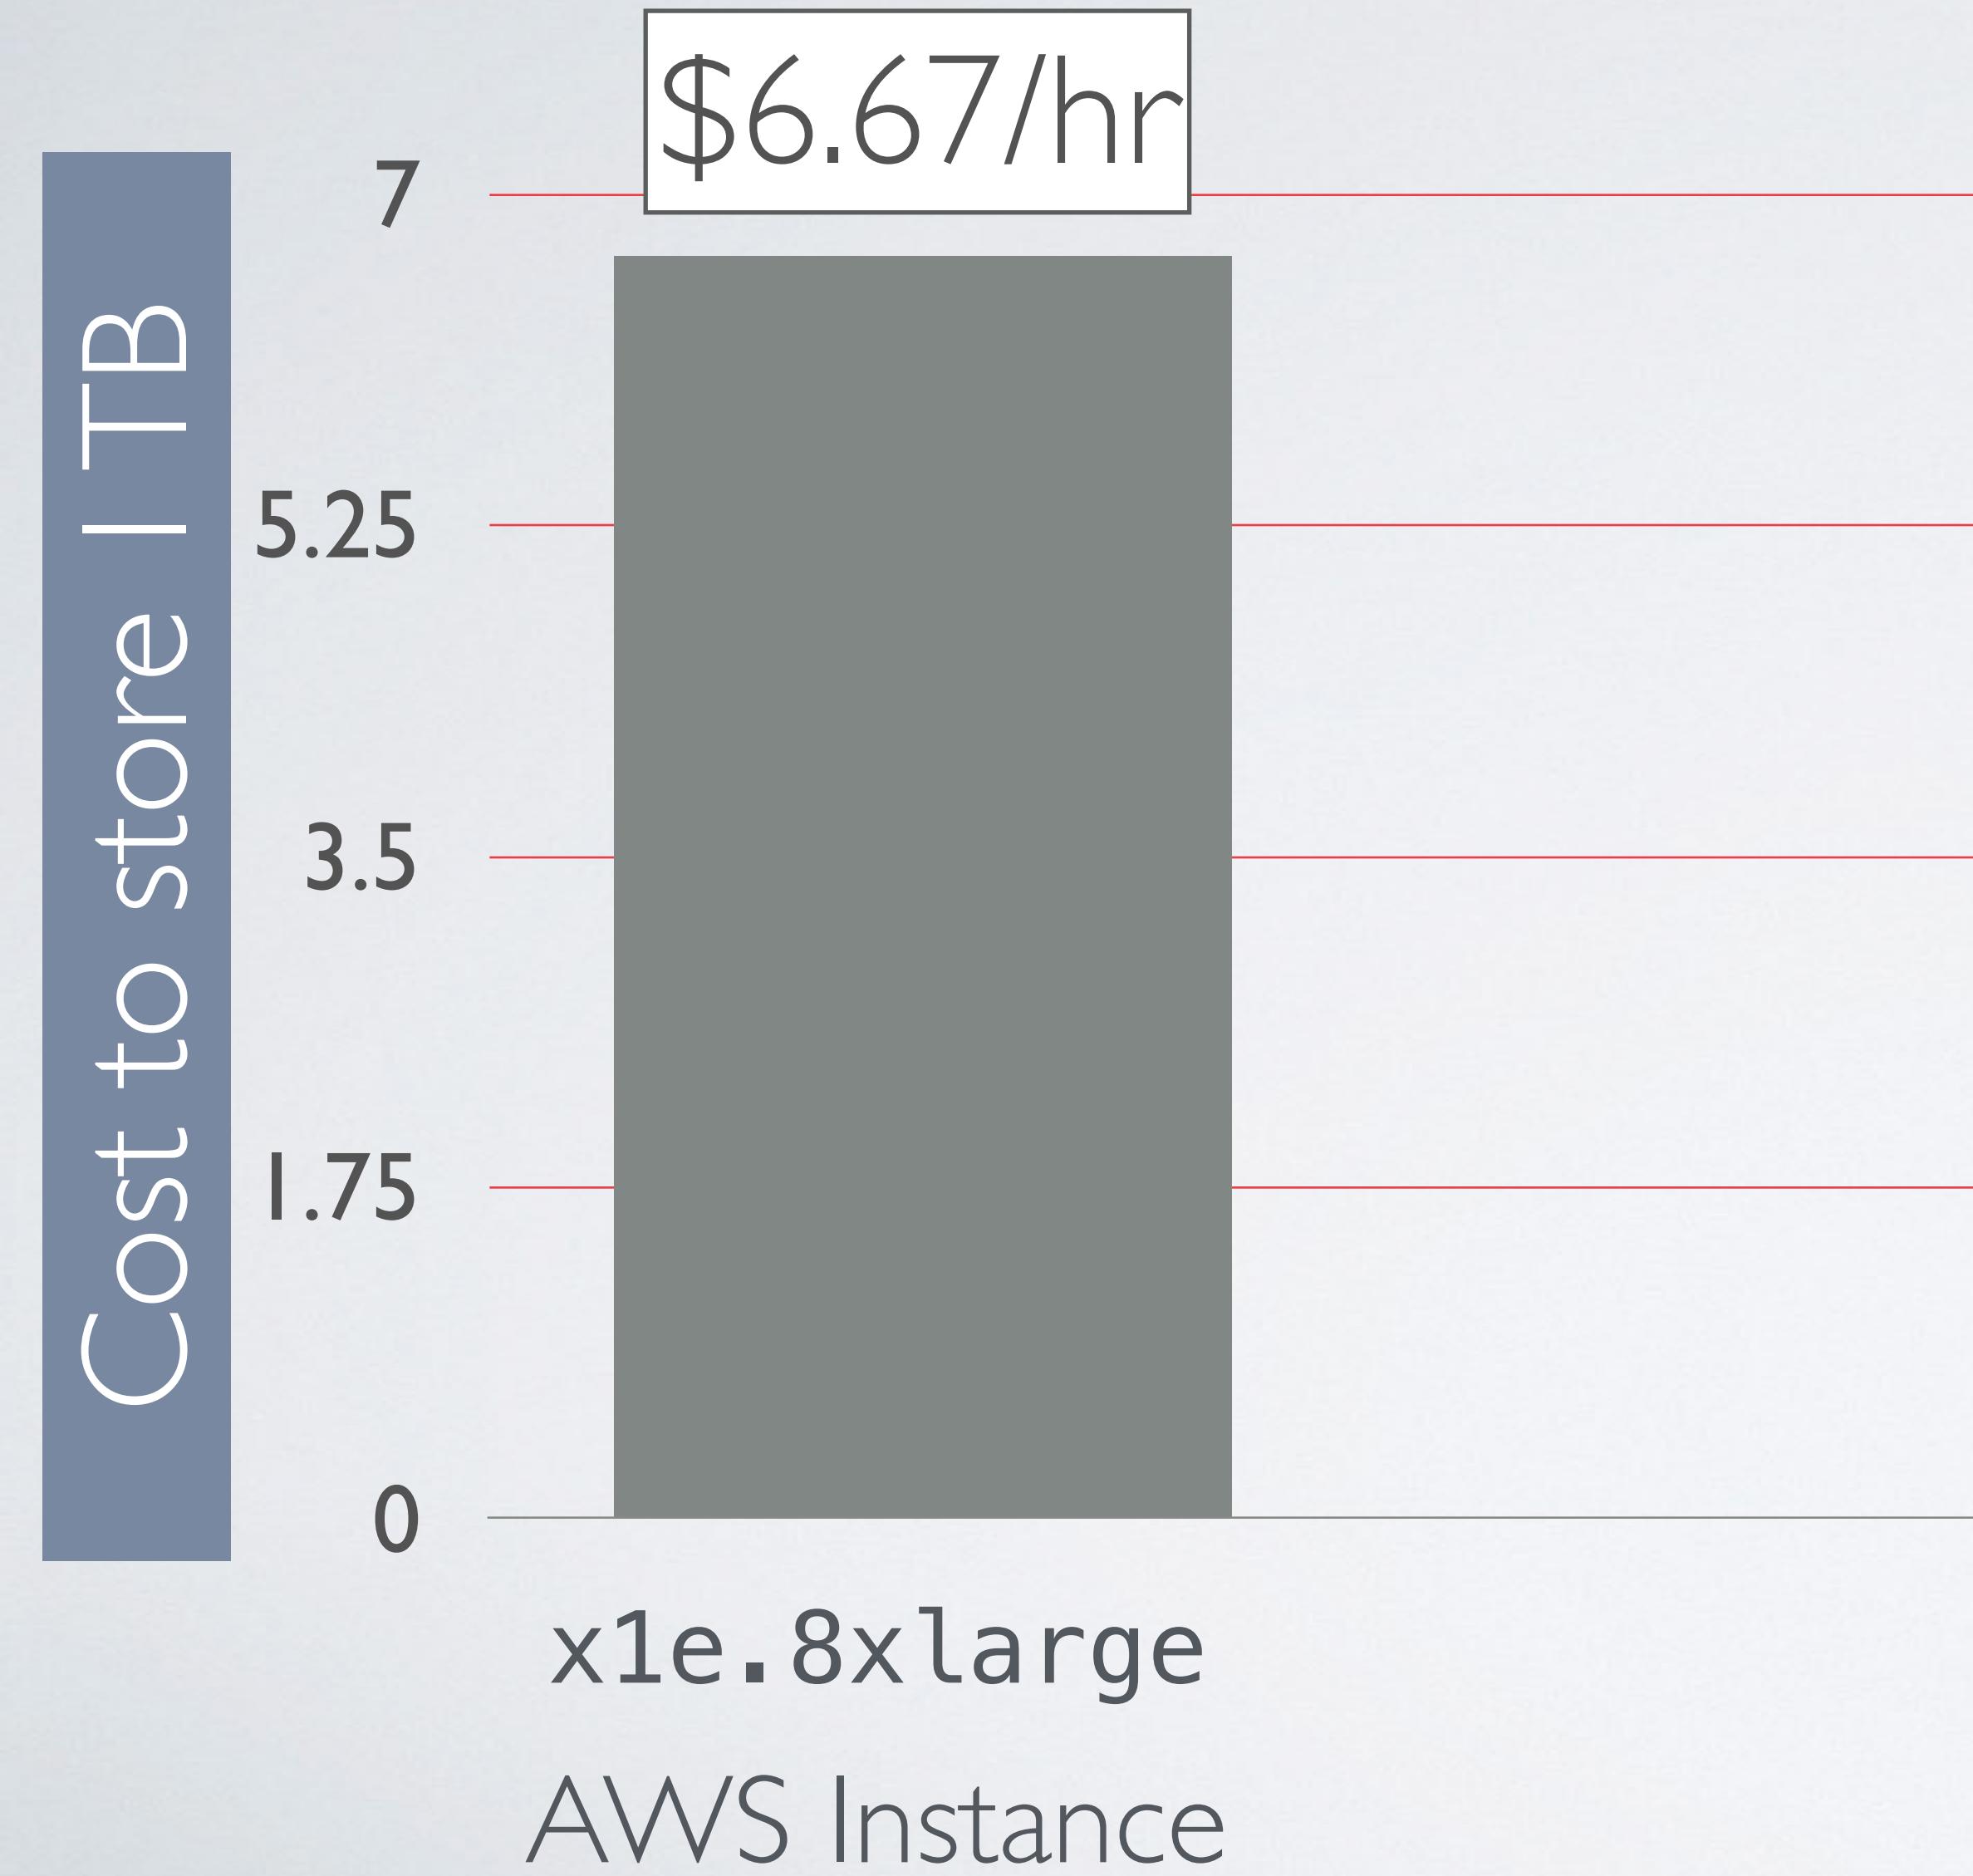
DATA CENTER DISAGGREGATION



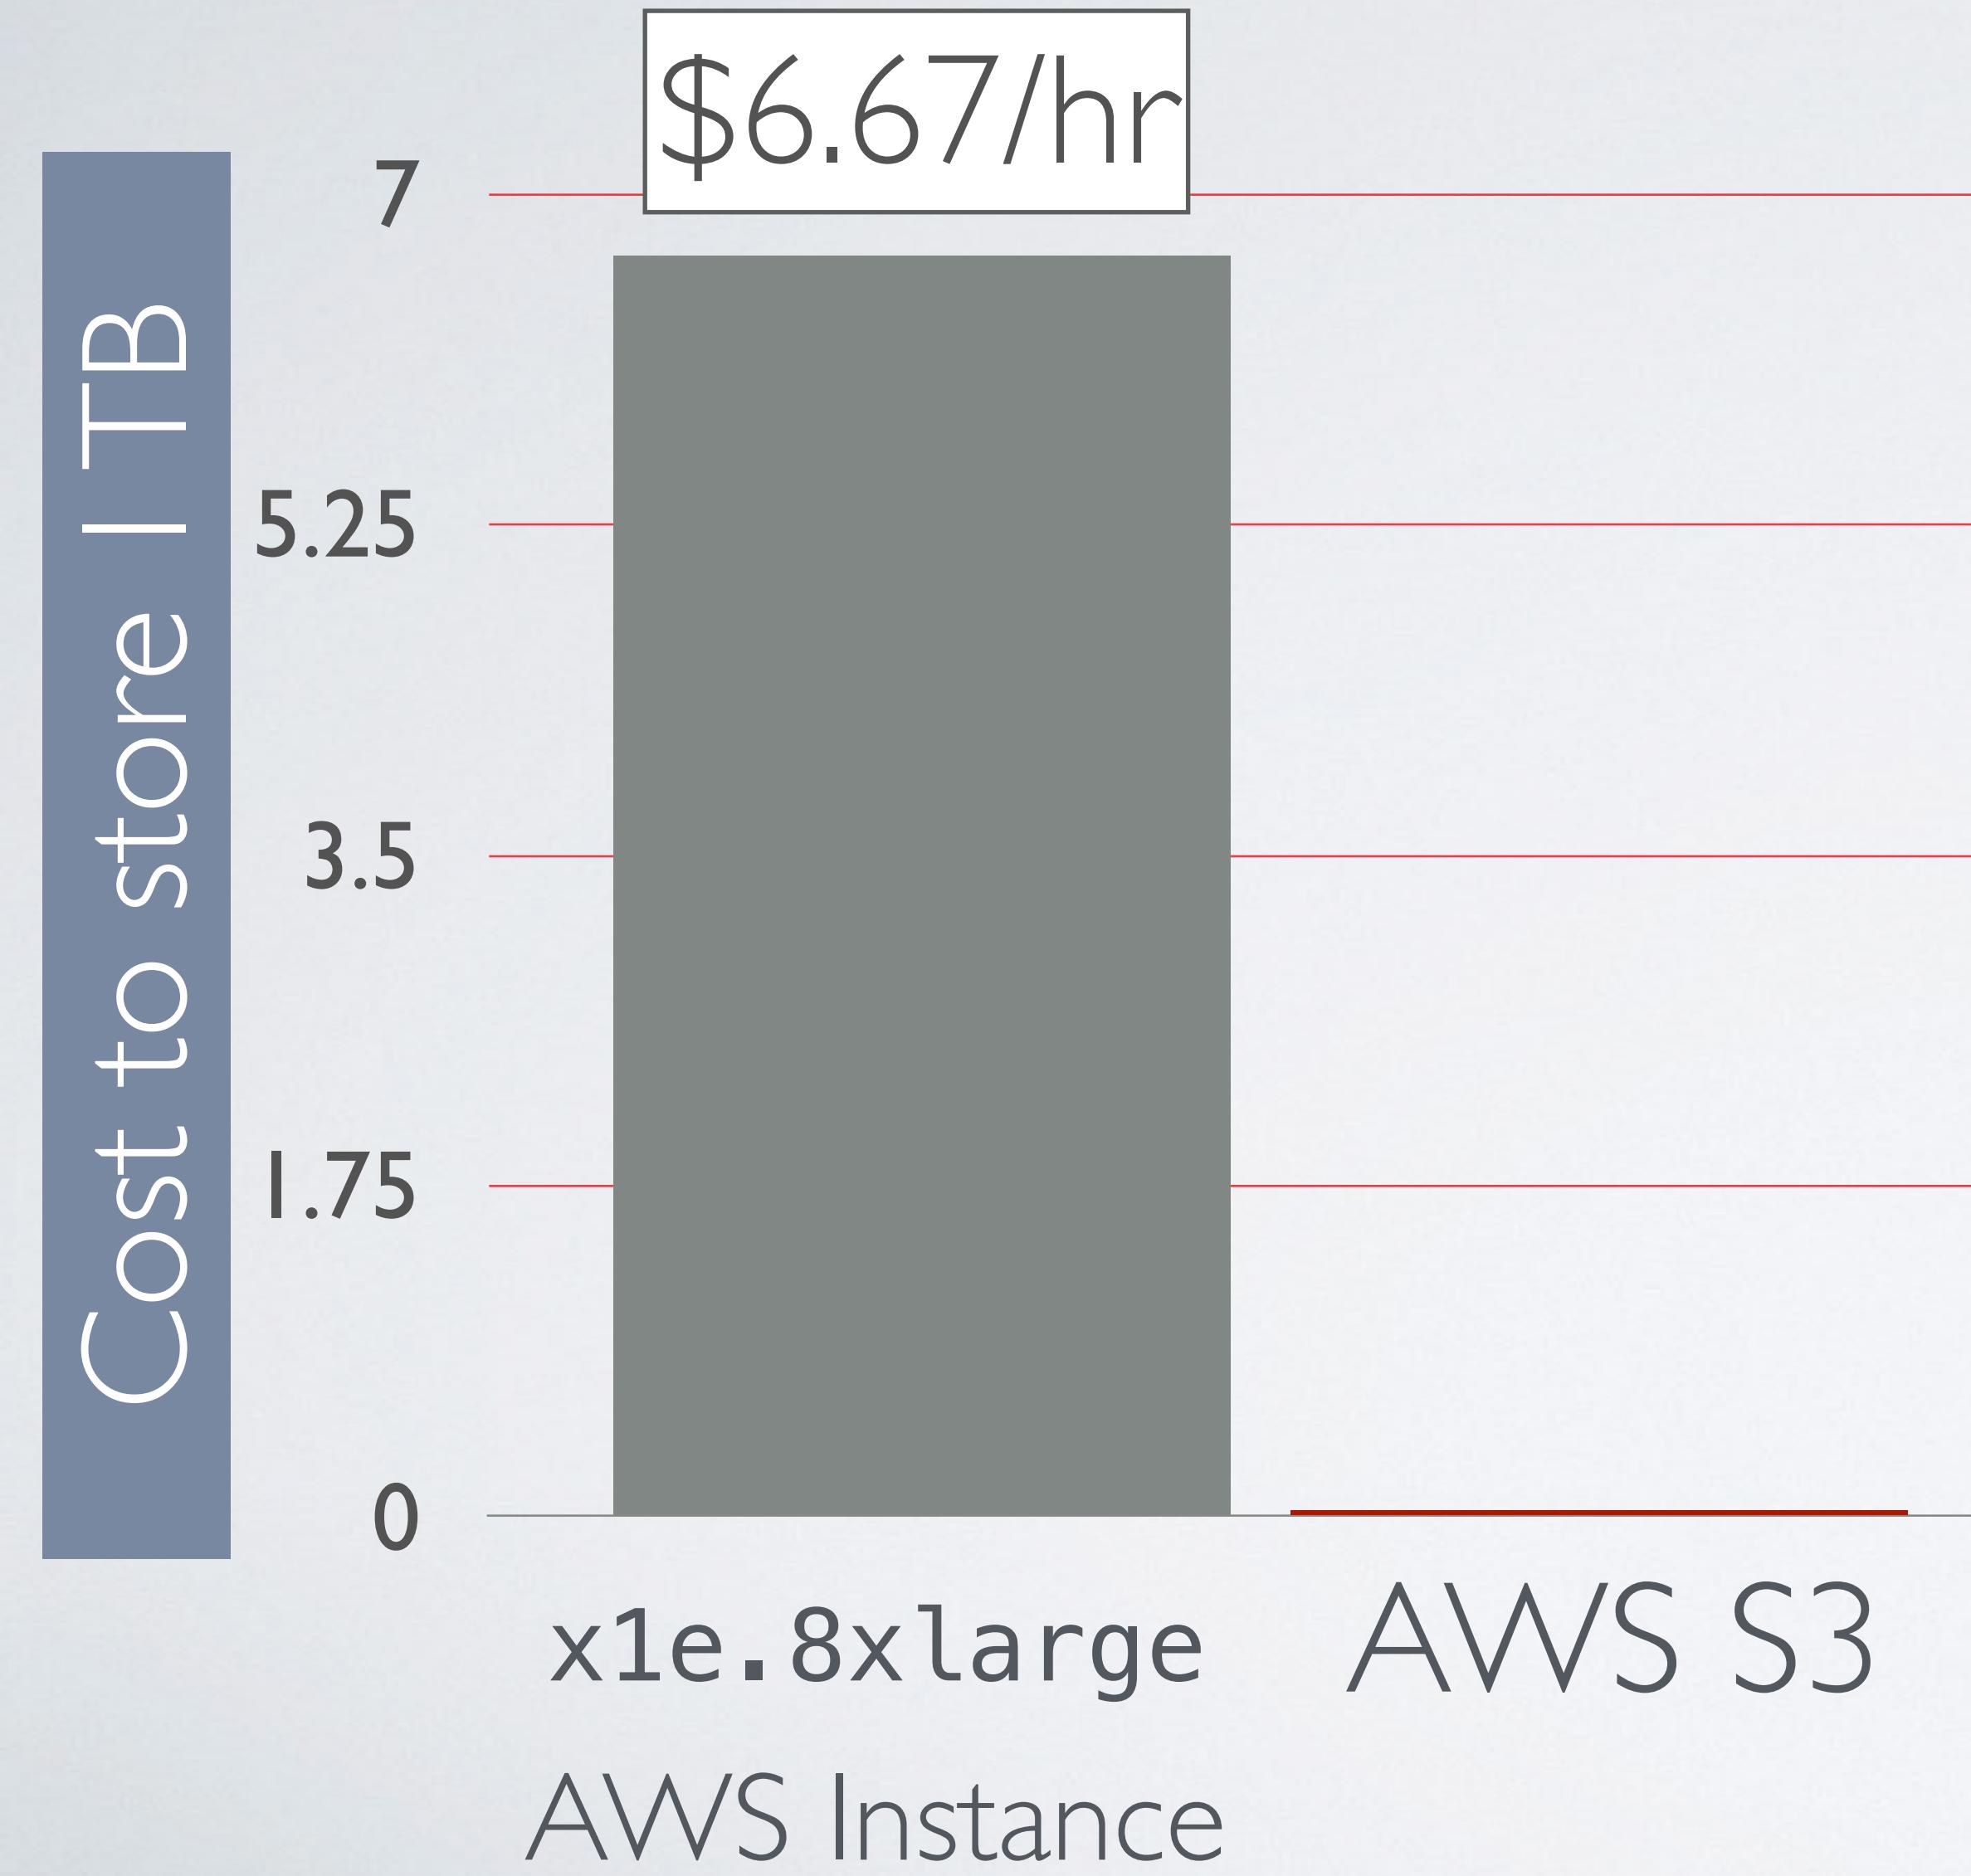
DATA CENTER DISAGGREGATION



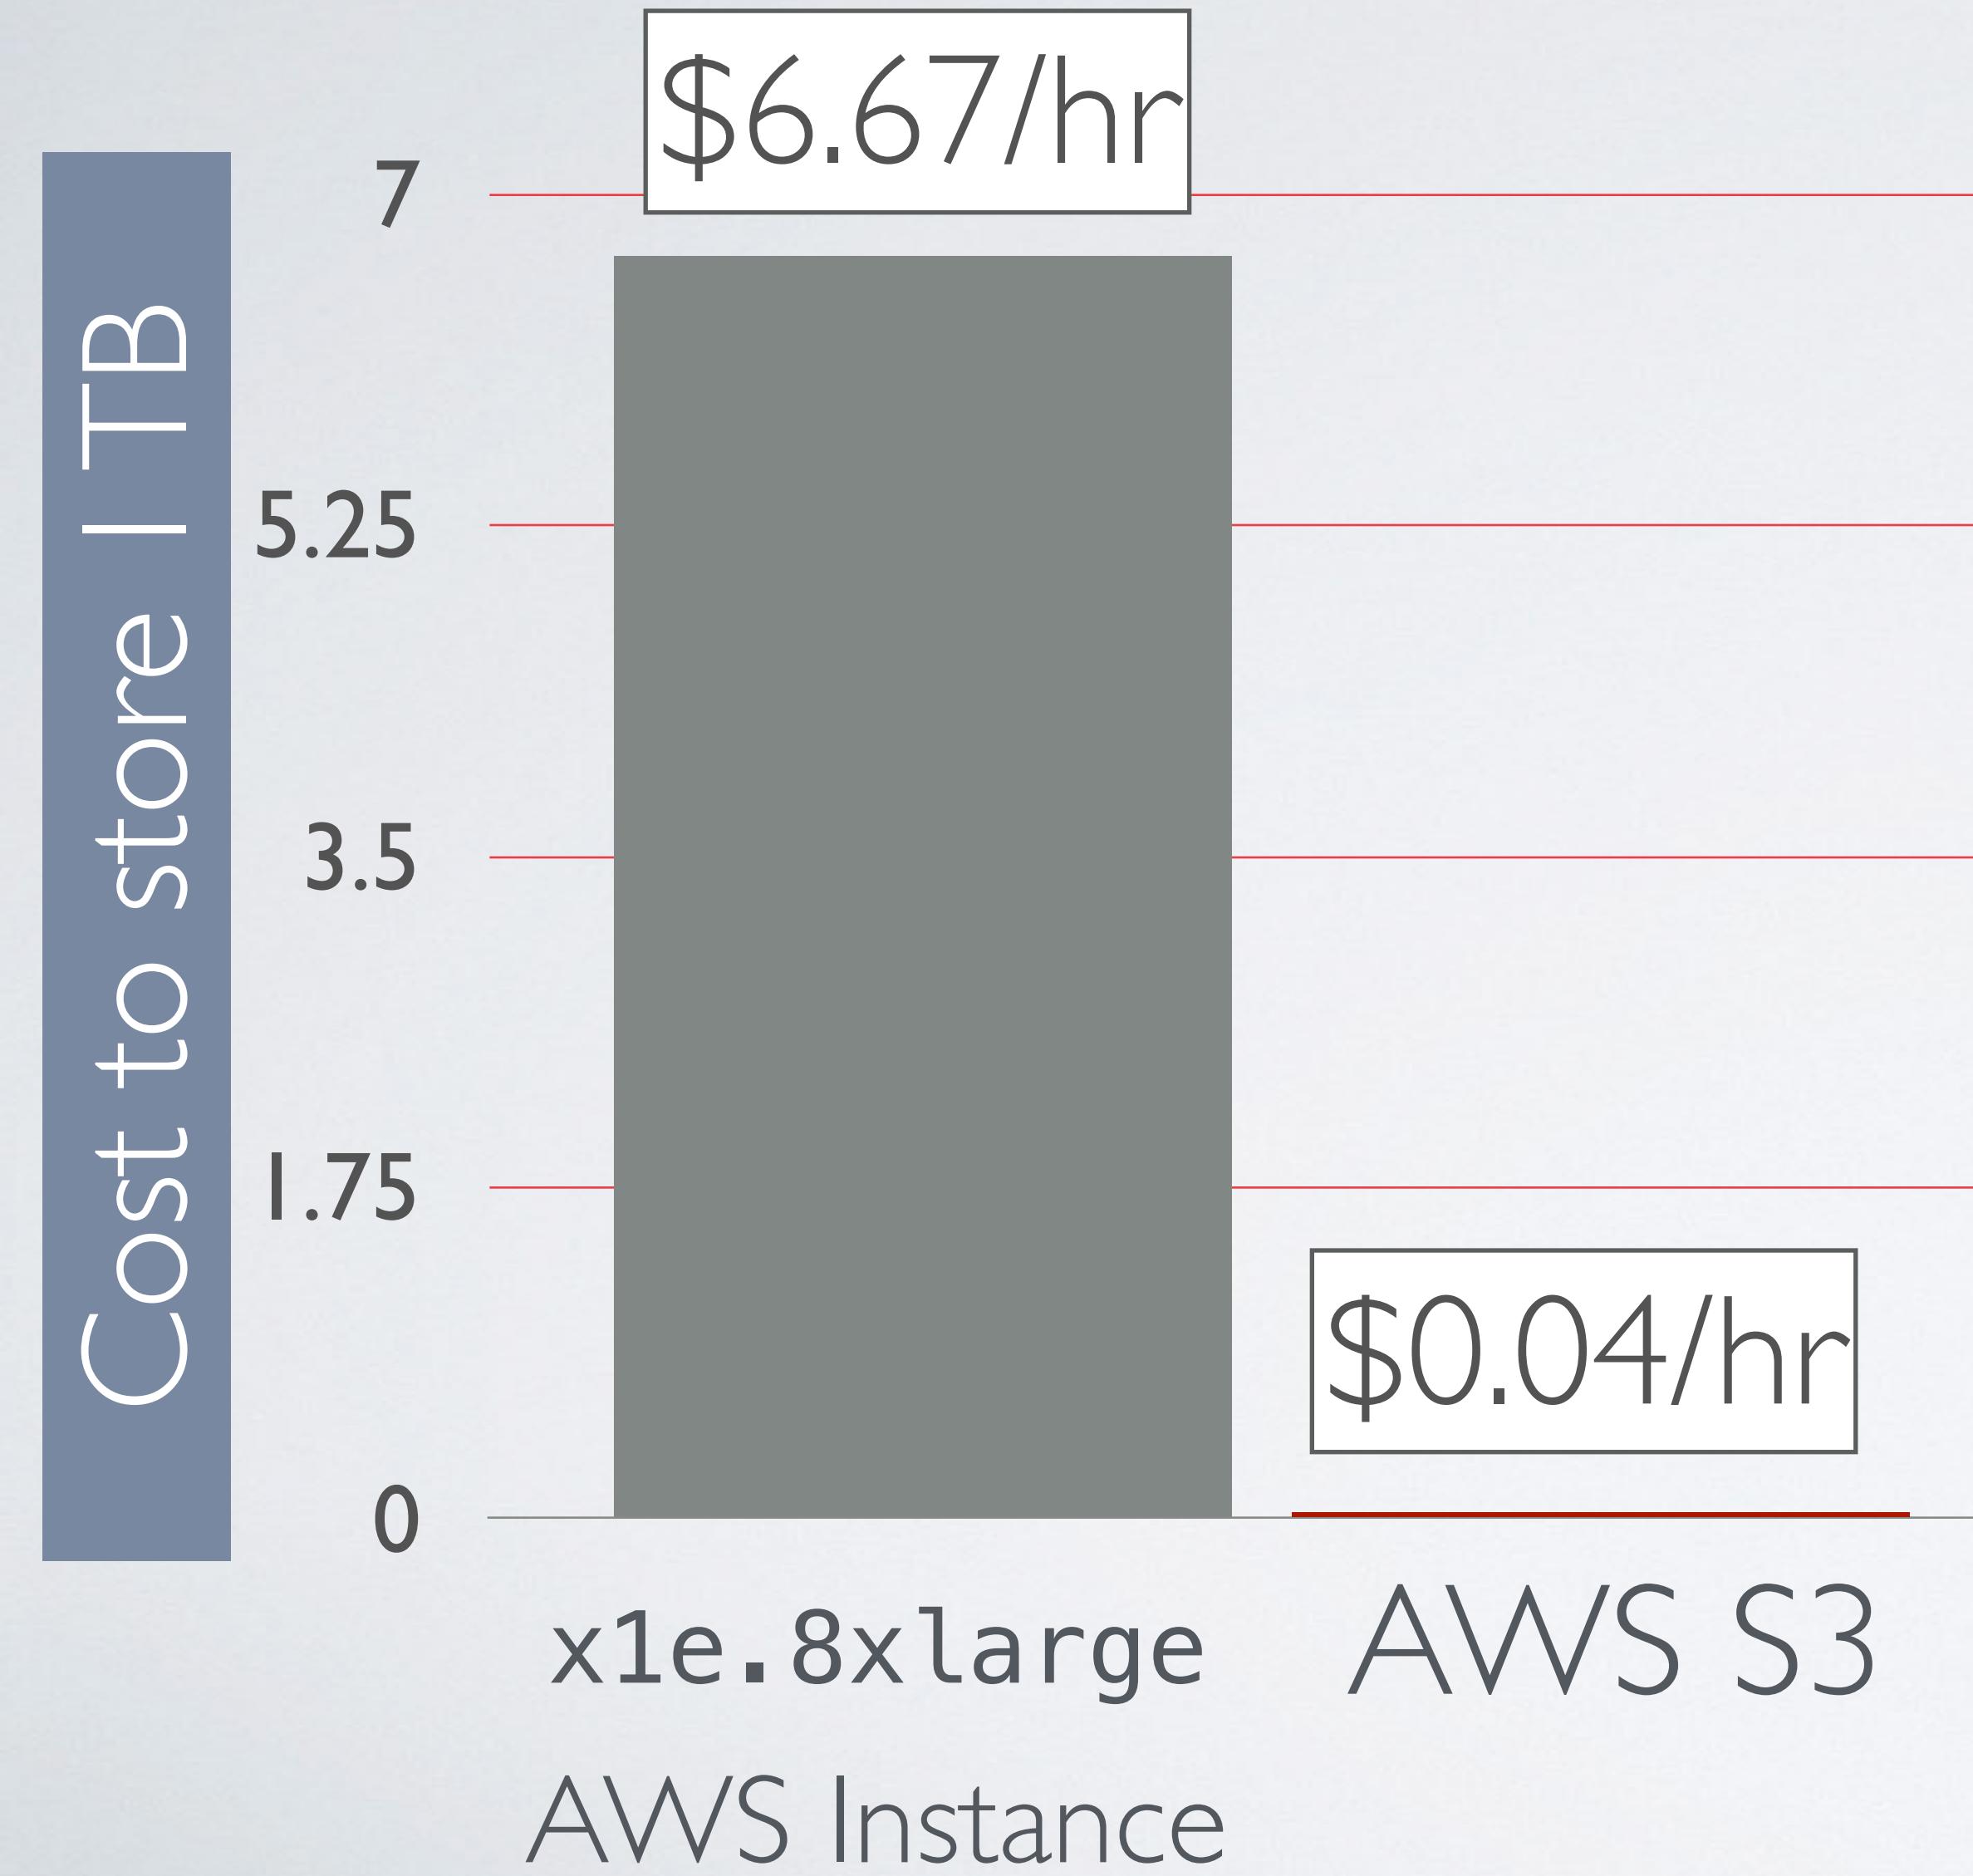
DATA CENTER DISAGGREGATION



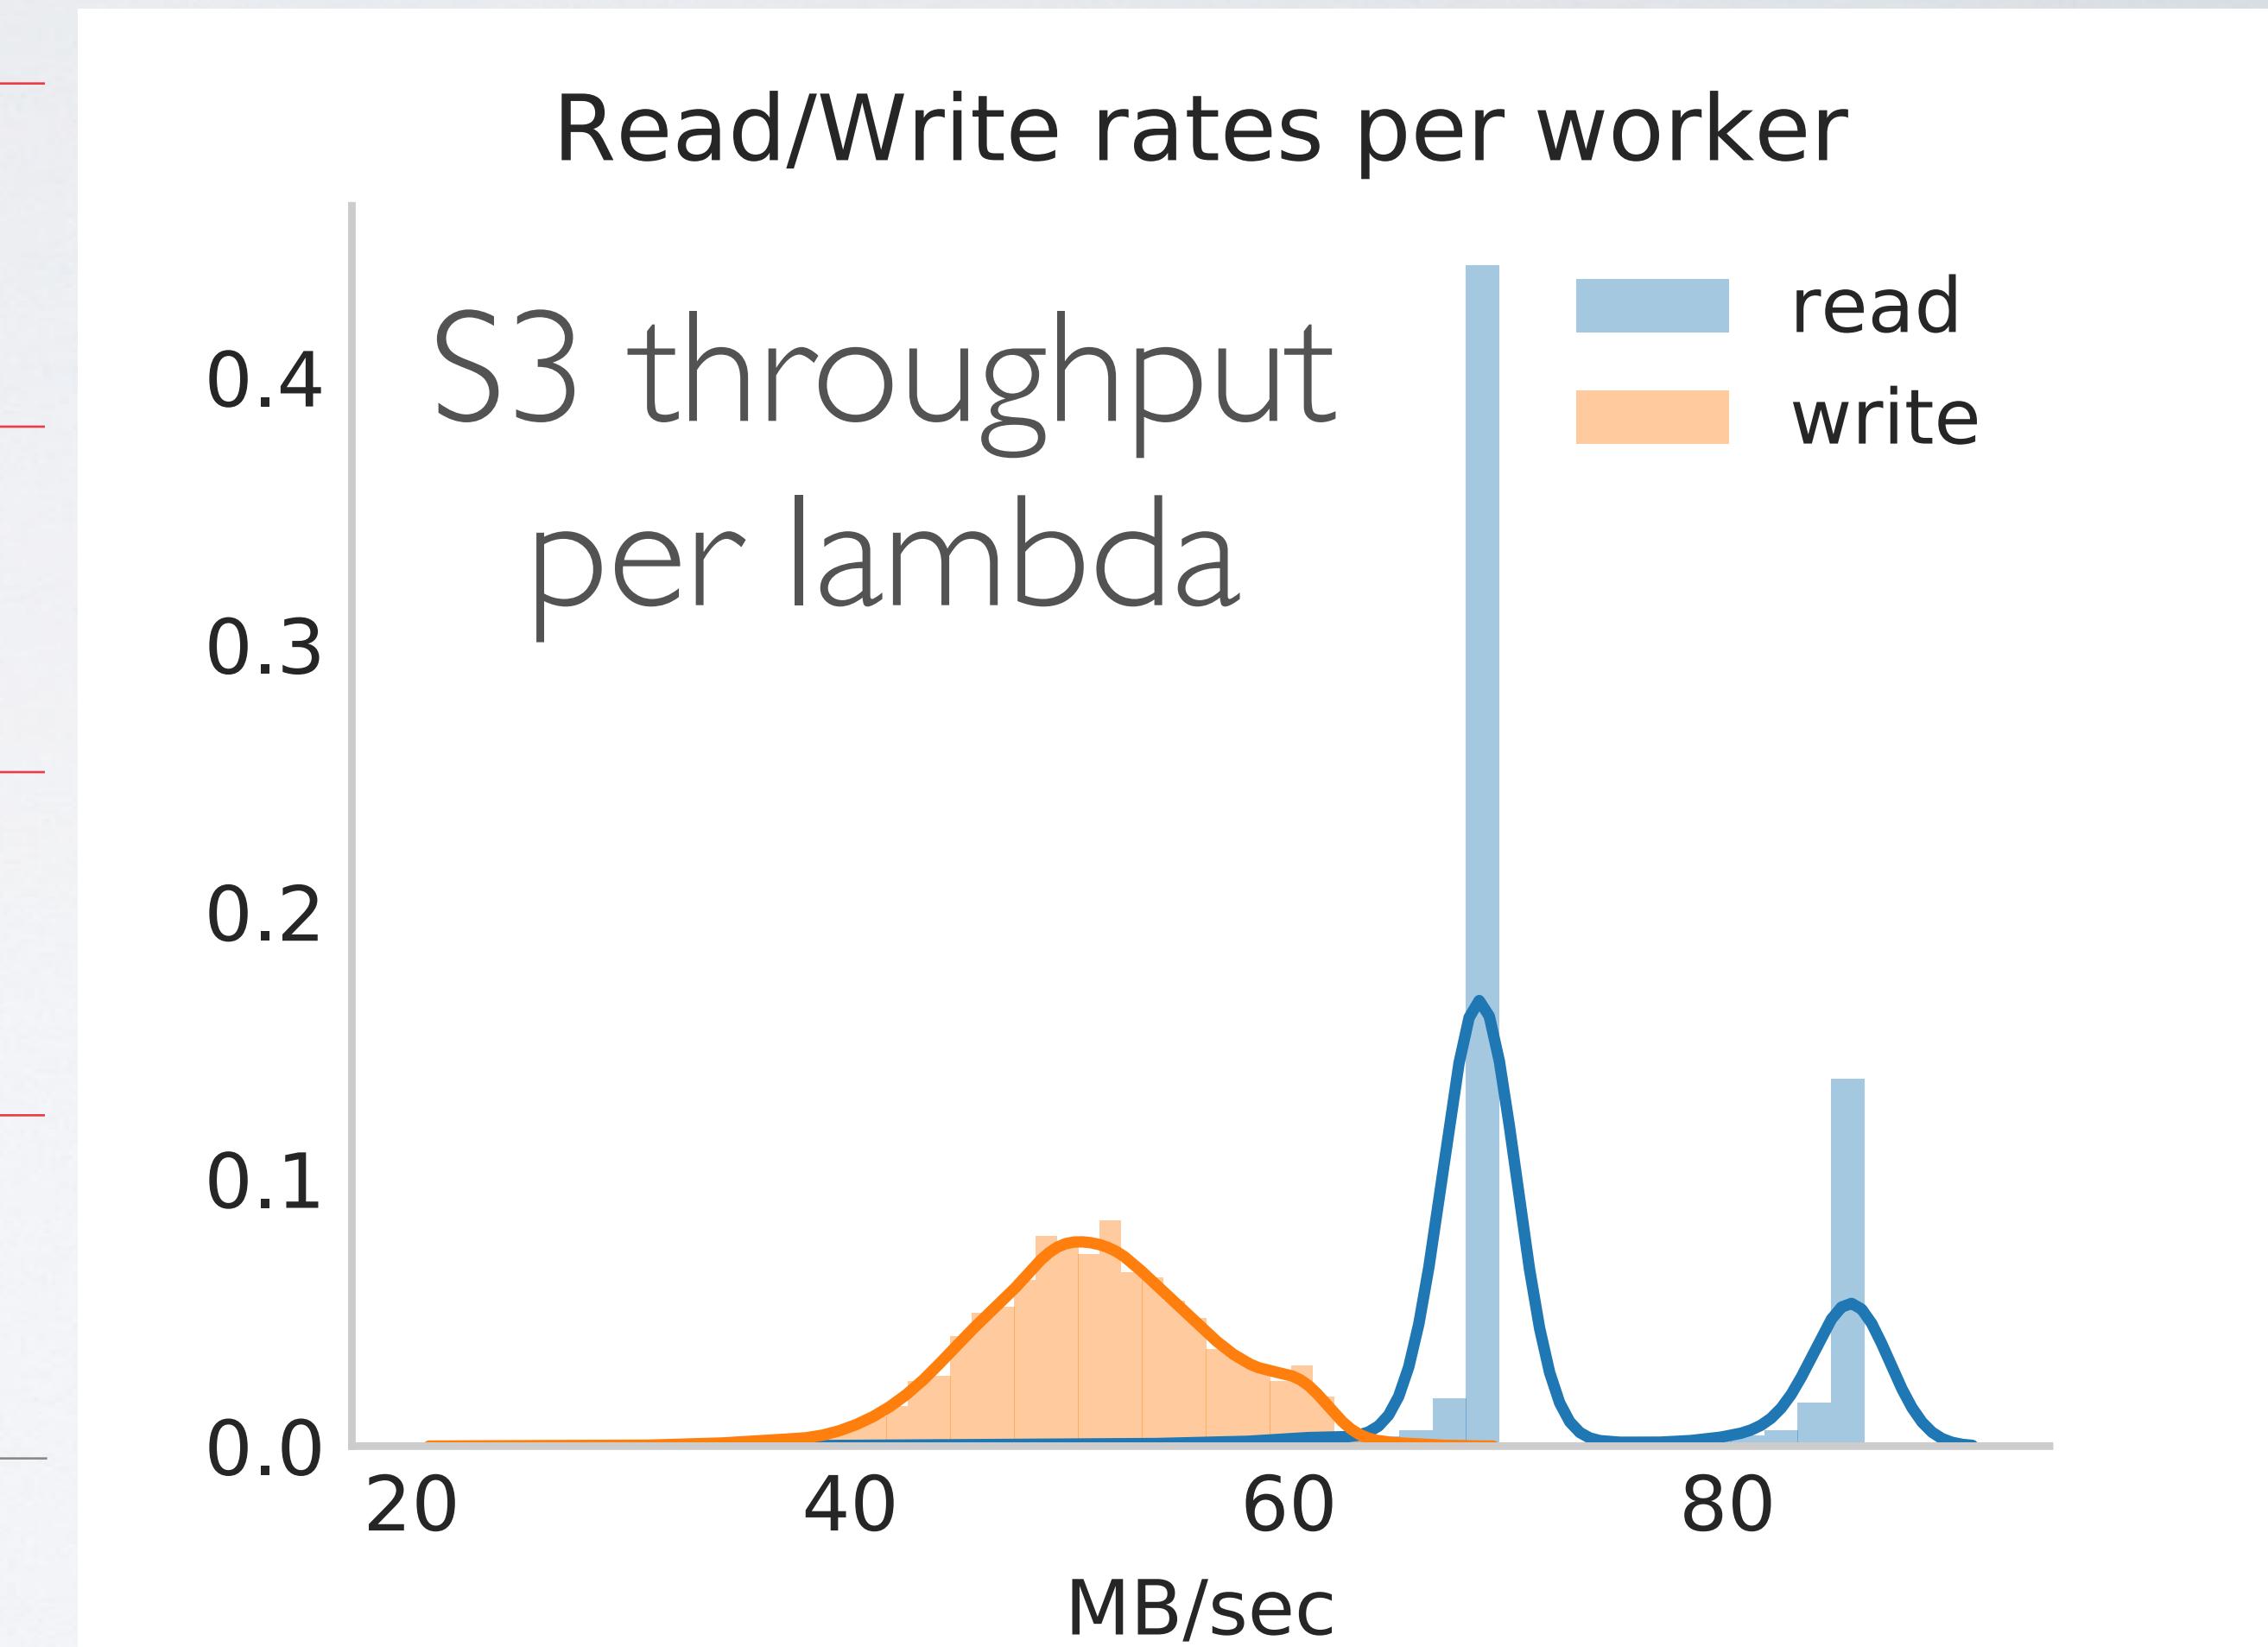
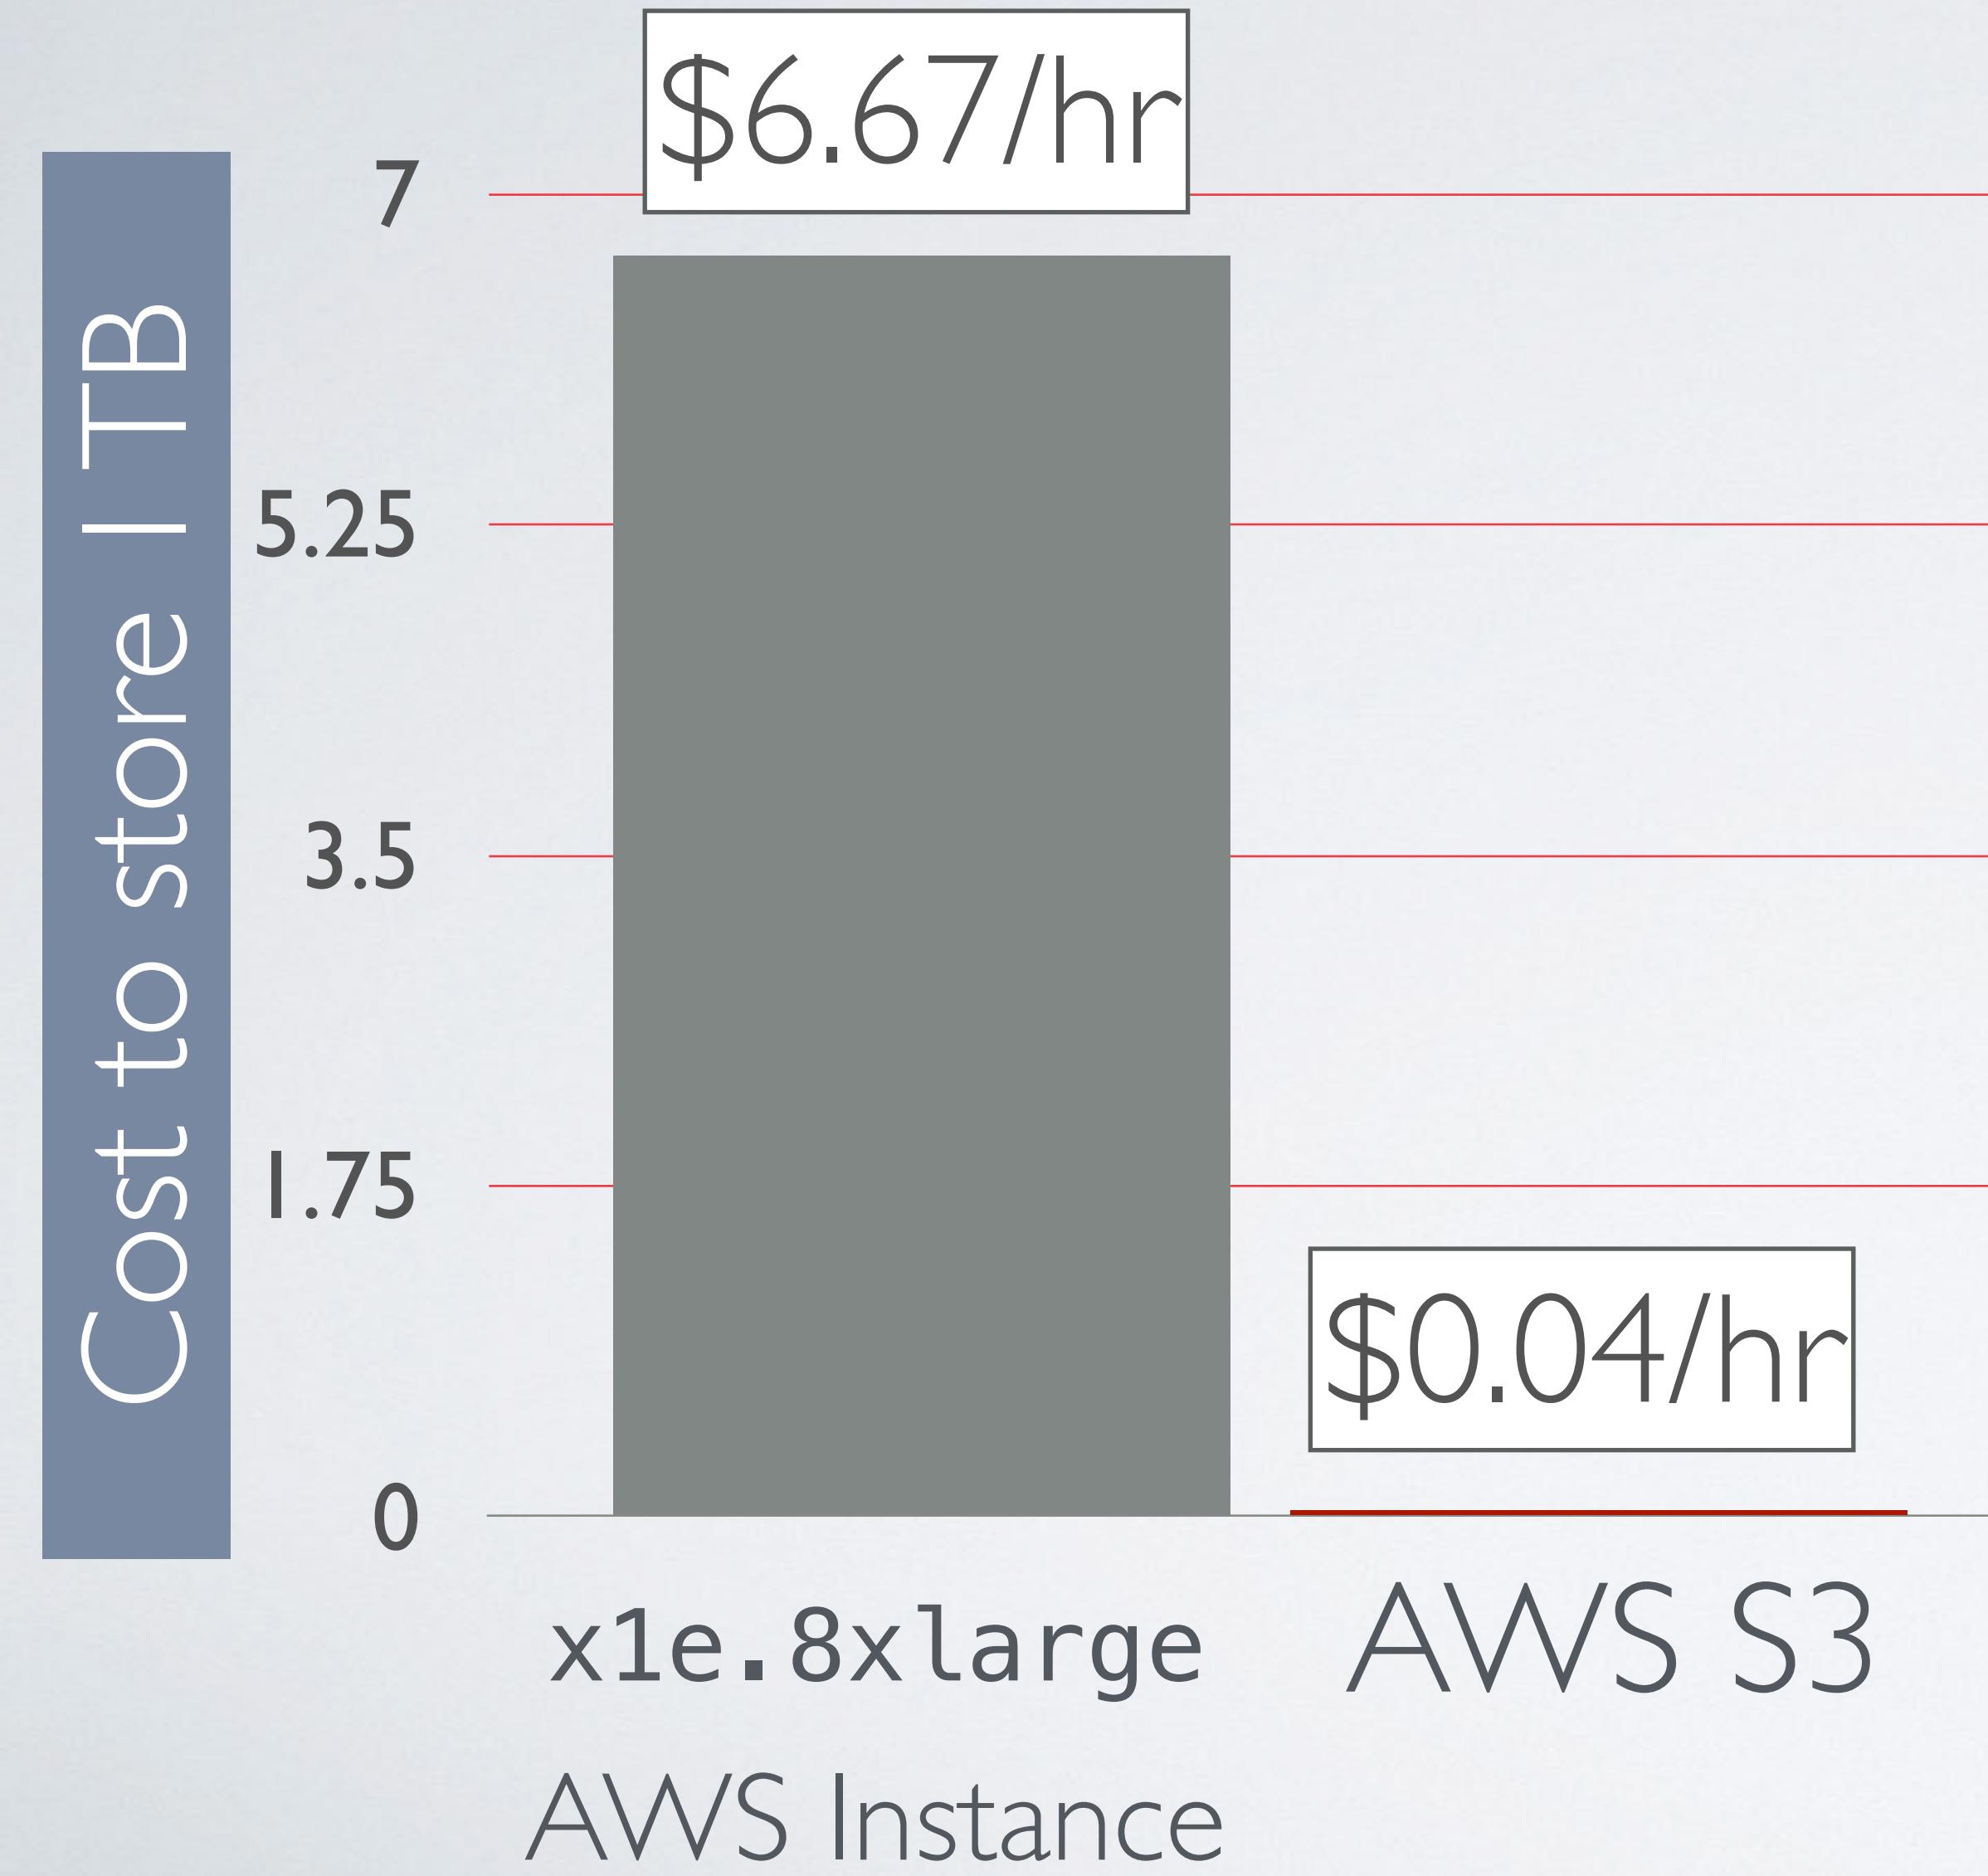
DATA CENTER DISAGGREGATION



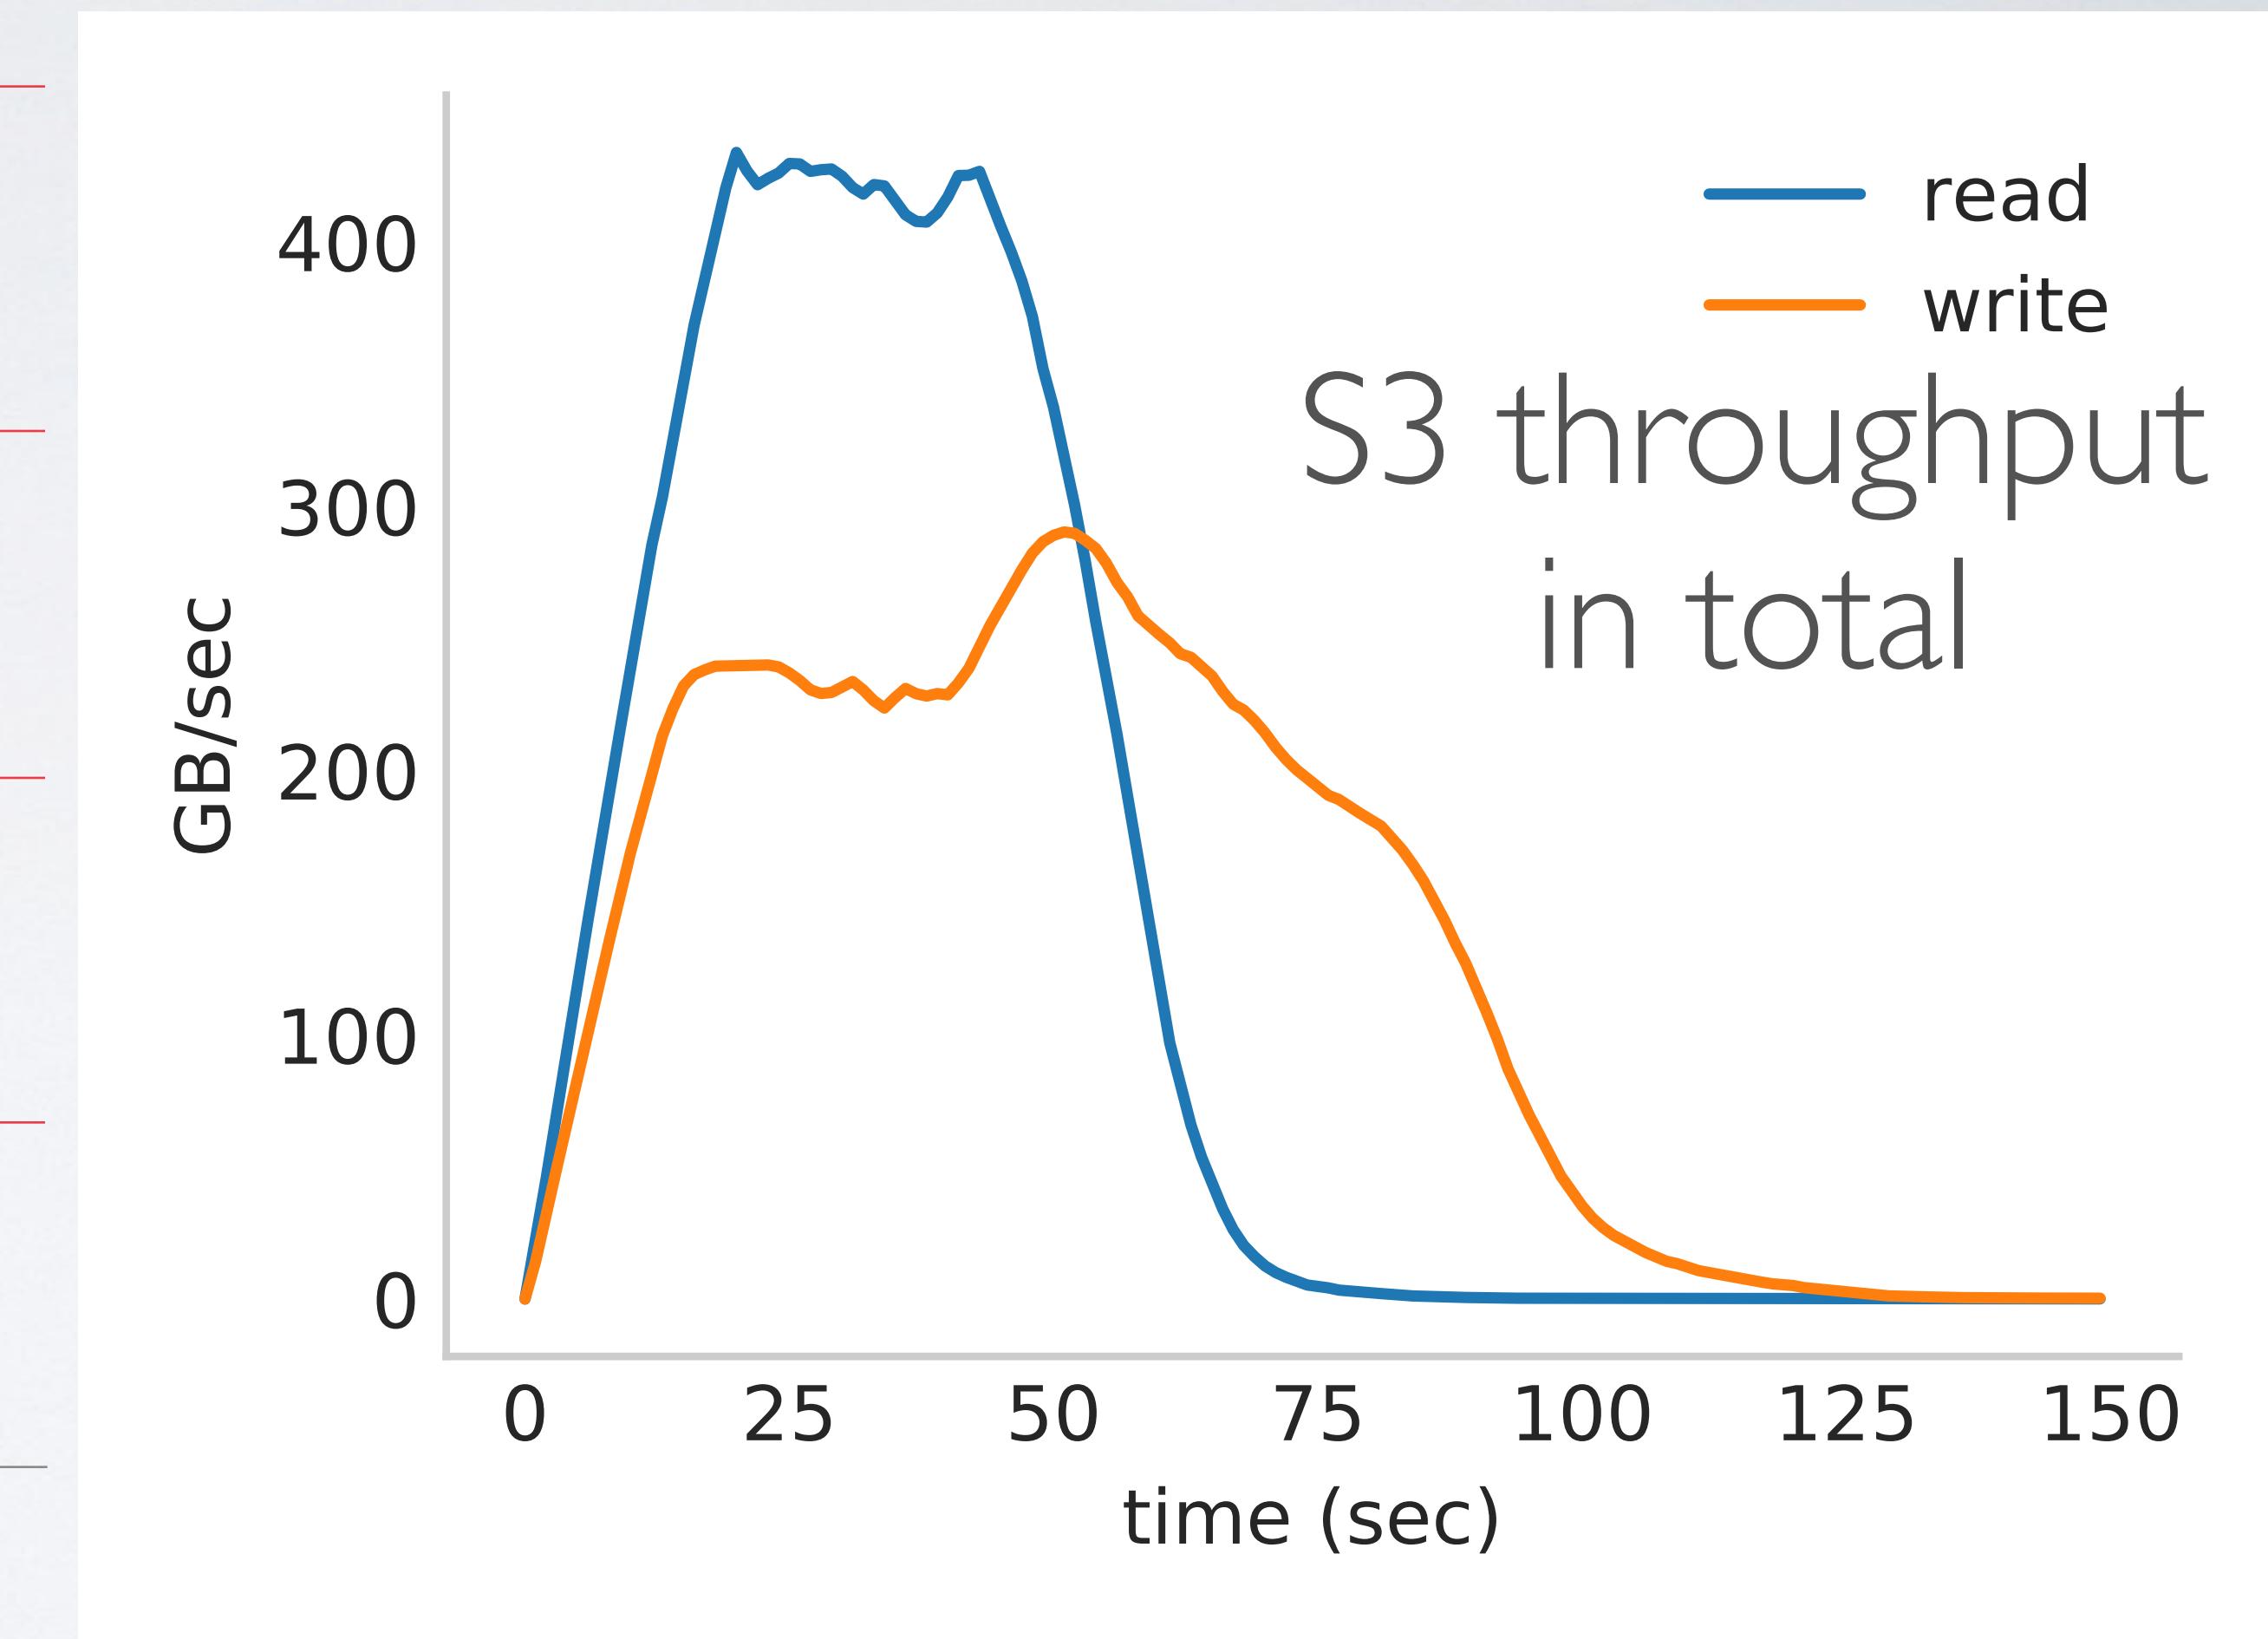
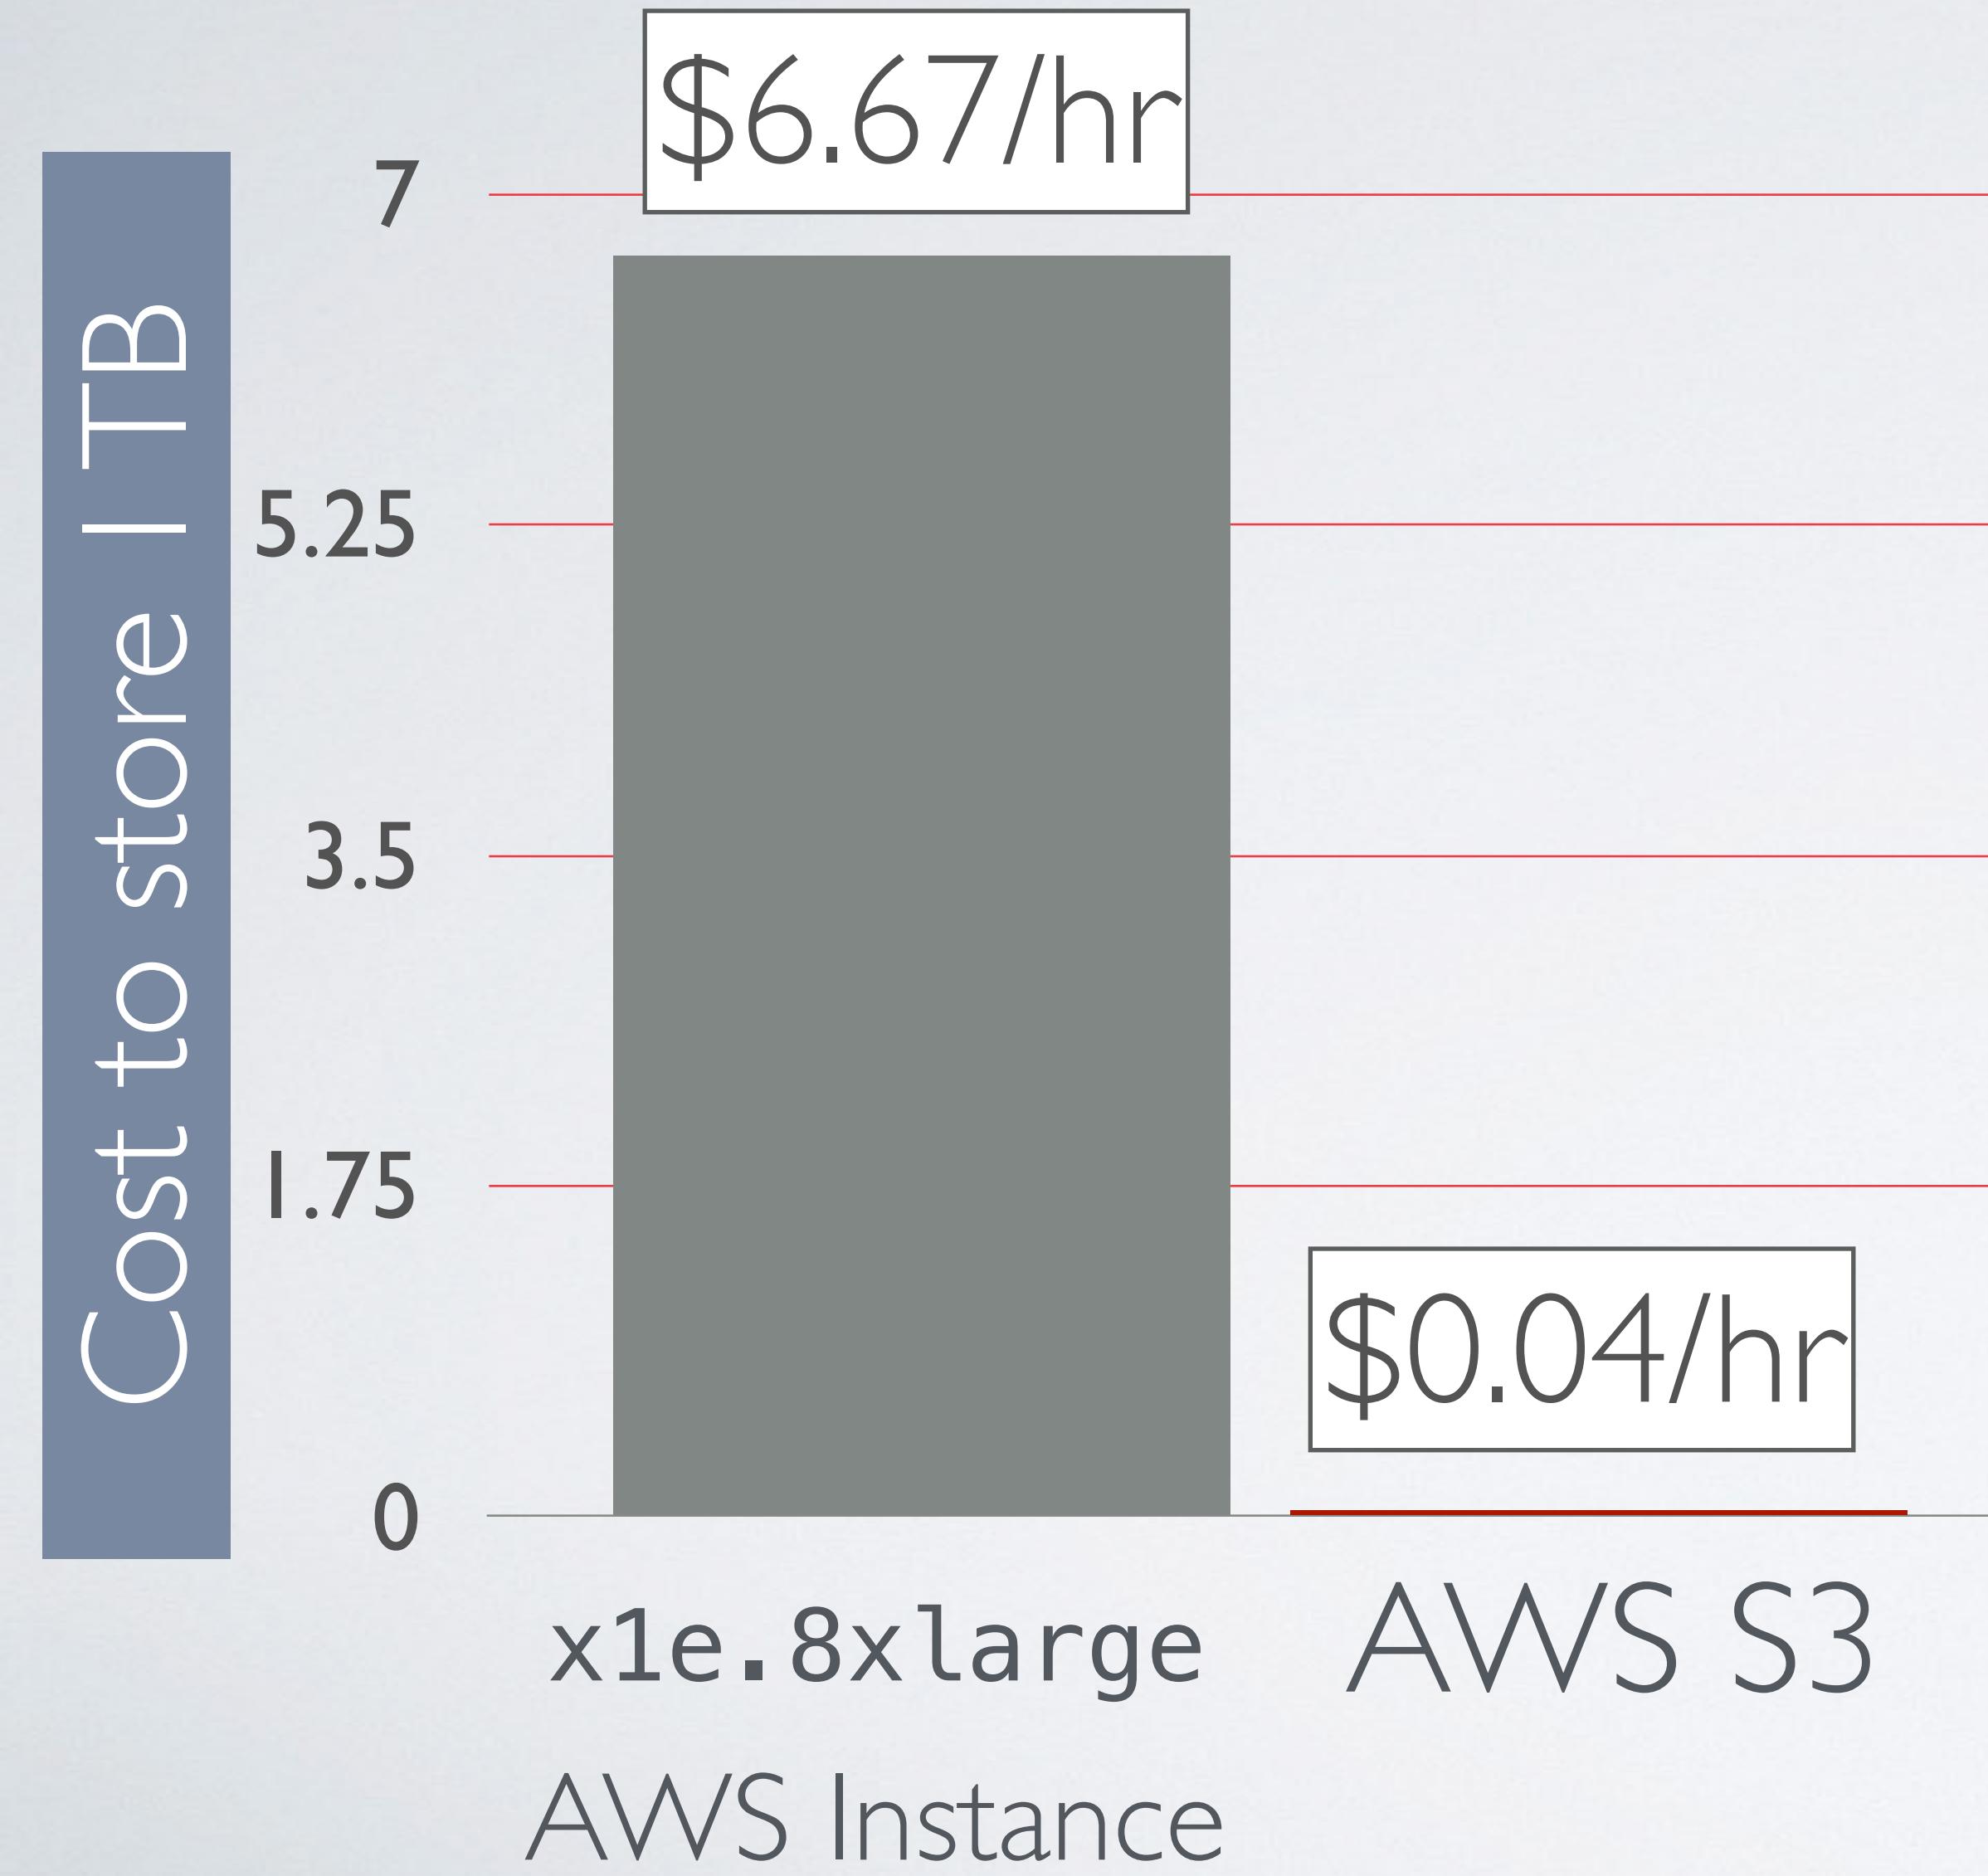
DATA CENTER DISAGGREGATION



DATA CENTER DISAGGREGATION

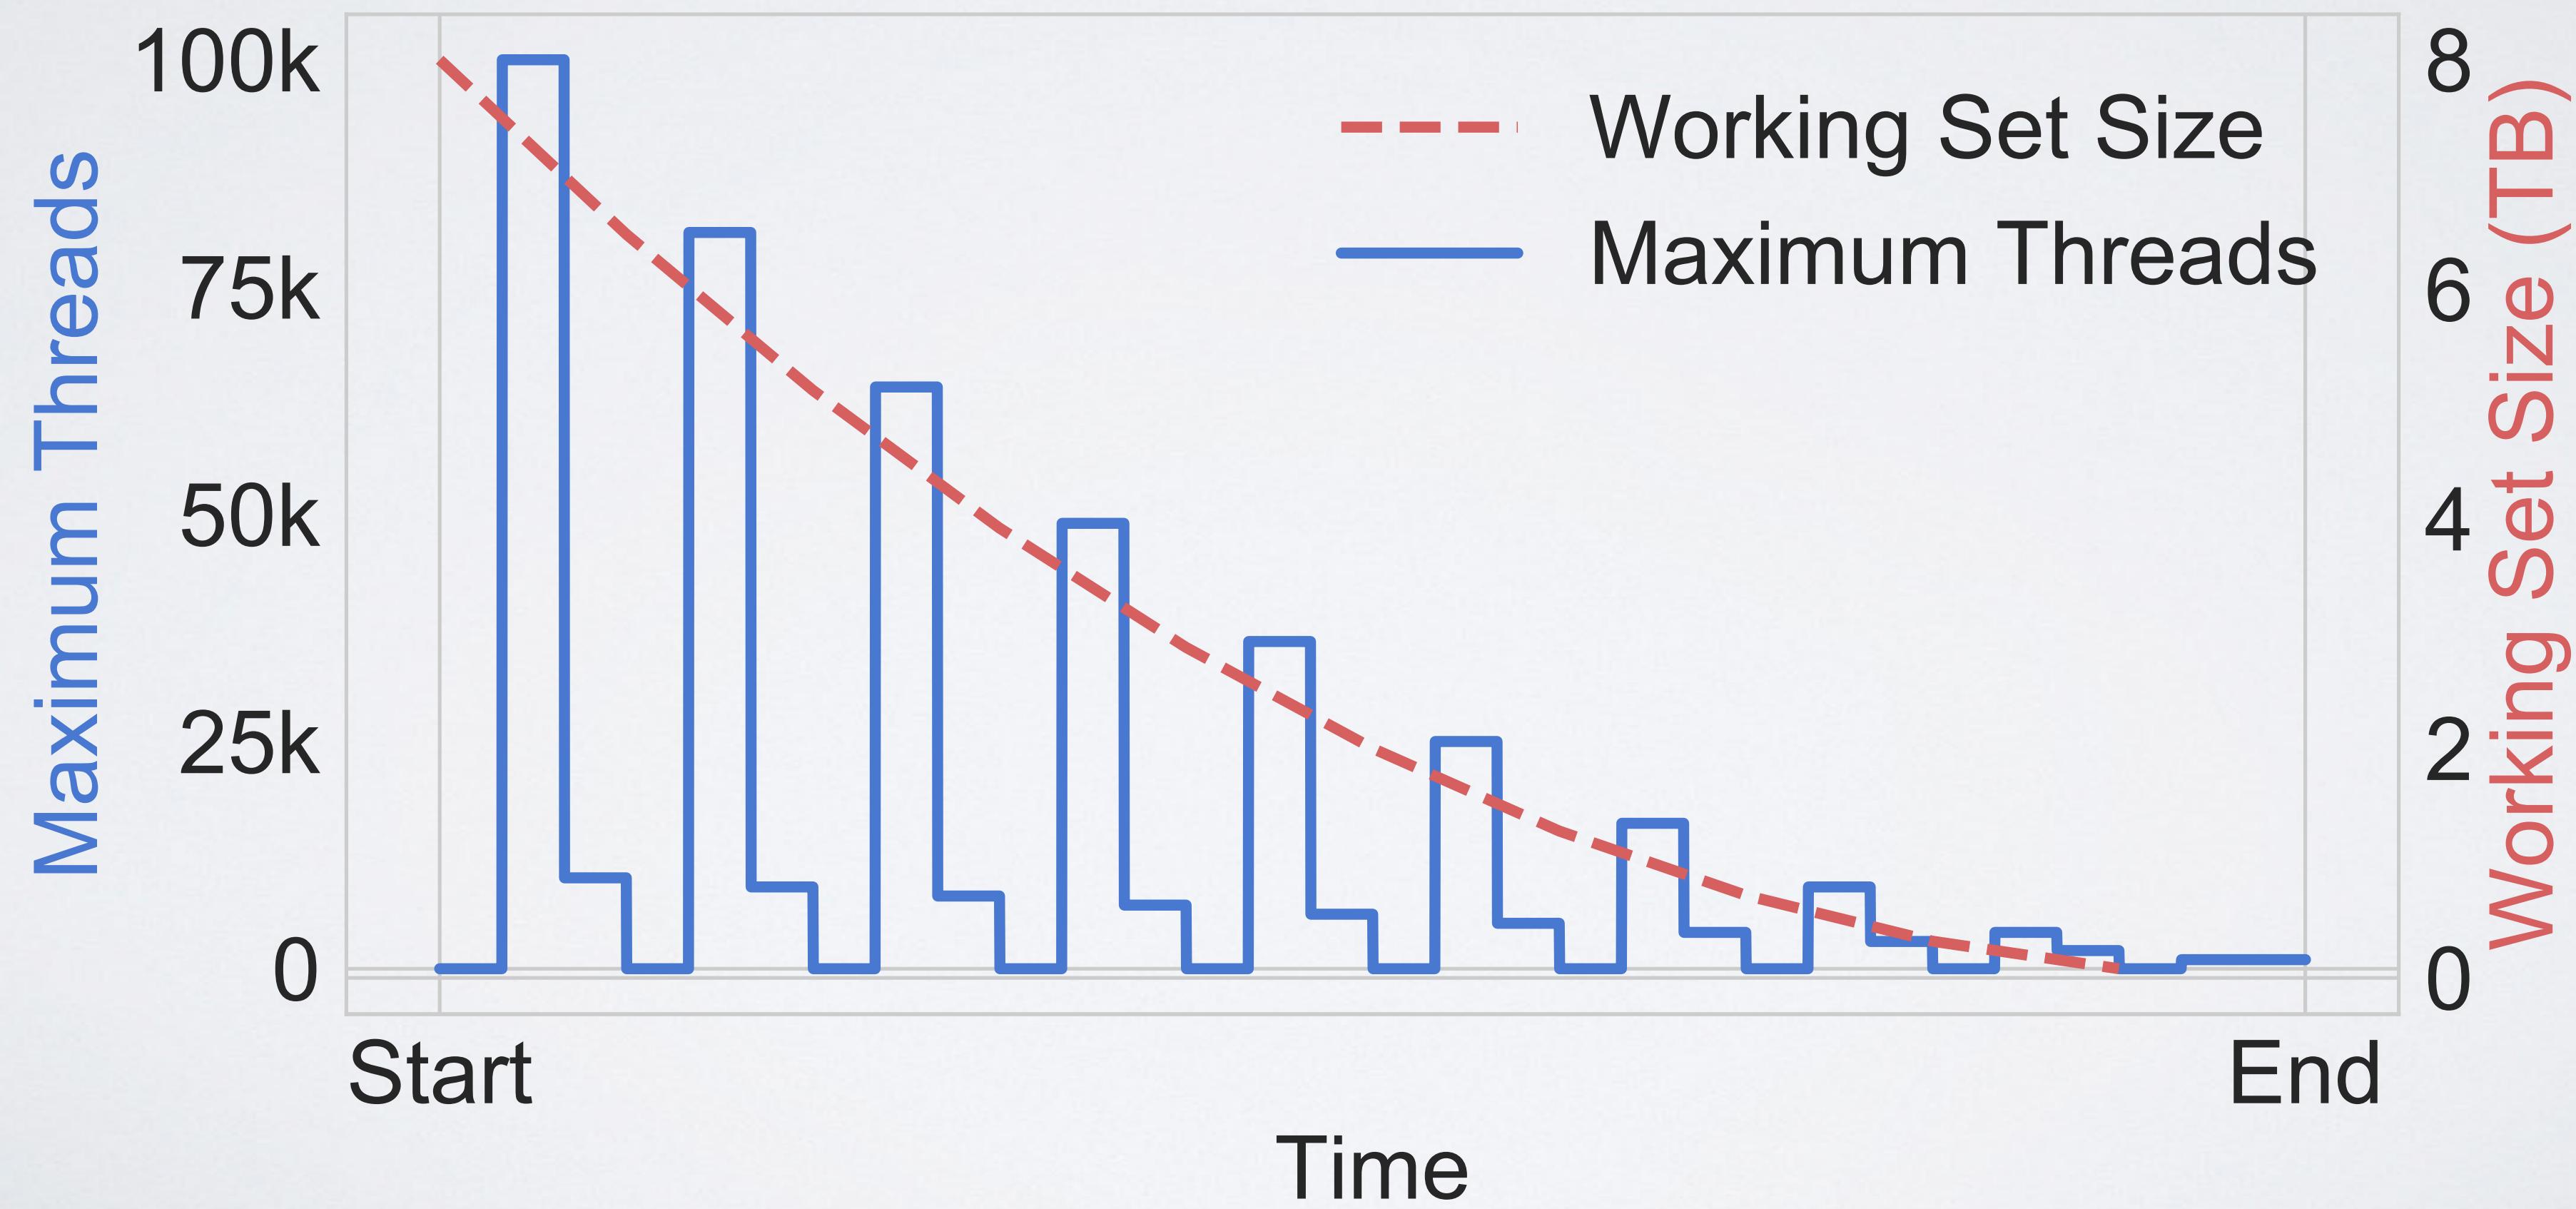


DATA CENTER DISAGGREGATION



Linear algebra operations have

DYNAMIC PARALLELISM AND WORKING SET SIZE



Linear algebra operations have

Linear algebra operations have

Compute

$O(n^3)$

Linear algebra operations have

Compute

$$O(n^3) >$$

Communication

$$O(n^2)$$

Linear algebra operations have

Compute

$O(n^3)$

Communication

$O(n^2)$

Matrix-Matrix product

Singular Value Decomposition

Least Squares Solve

Cholesky Factorization

TRENDS AND OBSERVATIONS

Compute more
precious



Fast cheap
disaggregated
state



Algorithms with
dynamic
parallelism



Operations
where compute
dominates IO

$$O(n^3) > O(n^2)$$

NUMPYWREN GOALS

NUMPYWREN GOALS

- No expensive setup (ala PyWren)

NUMPYWREN GOALS

- No expensive setup (ala PyWren)
- Decouple computation and storage

NUMPYWREN GOALS

- No expensive setup (ala PyWren)
- Decouple computation and storage
 - More cores->faster

NUMPYWREN GOALS

- No expensive setup (ala PyWren)
- Decouple computation and storage
 - More cores->faster
 - More storage -> bigger

NUMPYWREN GOALS

- No expensive setup (ala PyWren)
- Decouple computation and storage
 - More cores->faster
 - More storage -> bigger
- Elastic parallelism — be careful with compute

NUMPYWREN GOALS

user facing numpy/matlab-like interface
numpywren

Low Level IR aimed at LA primitives
lambdapack

Execution Framework
pywren

NUMPYWREN GOALS

user facing numpy/matlab-like interface
numpywren

Low Level IR aimed at LA primitives
lambdapack

Execution Framework
pywren

- Usable by anyone who knows Numpy

NUMPYWREN GOALS

user facing numpy/matlab-like interface
numpywren

Low Level IR aimed at LA primitives
lambdapack

Execution Framework
pywren

- Usable by anyone who knows Numpy
- All big matrices live transparently in S3

NUMPYWREN GOALS

user facing numpy/matlab-like interface
numpywren

Low Level IR aimed at LA primitives
lambdapack

Execution Framework
pywren

- Usable by anyone who knows Numpy
- All big matrices live transparently in S3
- All intermediate state is retained

NEAREST NEIGHBOR

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(x_train, x_test, y_train, y_test):
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(x_train, x_test, y_train, y_test):  
    npwex = numpywren.default_executor()
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(x_train, x_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    x_train_sharded = npwex.matrix_init(x_train)
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(x_train, x_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    x_train_sharded = npwex.matrix_init(x_train)  
    x_test_sharded = npwex.matrix_init(x_test)
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)  
    XYT *= -2
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)  
    XYT *= -2  
    norms_train = npwex.linalg.norm(X_train, axis=1)
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)  
    XYT *= -2  
    norms_train = npwex.linalg.norm(X_train, axis=1)  
    norms_test = npwex.linalg.norm(X_test, axis=1)
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)  
    XYT *= -2  
    norms_train = npwex.linalg.norm(X_train, axis=1)  
    norms_test = npwex.linalg.norm(X_test, axis=1)  
    distances = norms_train + XYT + norms_test.T
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)  
    XYT *= -2  
    norms_train = npwex.linalg.norm(X_train, axis=1)  
    norms_test = npwex.linalg.norm(X_test, axis=1)  
    distances = norms_train + XYT + norms_test.T  
    argmins = npwex.argmin(distances, axis=0).numpy()
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)  
    XYT *= -2  
    norms_train = npwex.linalg.norm(X_train, axis=1)  
    norms_test = npwex.linalg.norm(X_test, axis=1)  
    distances = norms_train + XYT + norms_test.T  
    argmins = npwex.argmin(distances, axis=0).numpy()
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):
    npwex = npywren.default_executor()

    X_train_sharded = npwex.matrix_init(X_train)
    X_test_sharded = npwex.matrix_init(X_test)

    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)
    XYT *= -2
    norms_train = npwex.linalg.norm(X_train, axis=1)
    norms_test = npwex.linalg.norm(X_test, axis=1)
    distances = norms_train + XYT + norms_test.T
    argmins = npwex.argmin(distances, axis=0).numpy()

    return metrics.accuracy_score(y_train[argmins], y_test)
```

NEAREST NEIGHBOR

```
def nearest_neighbor_numpywren(X_train, X_test, y_train, y_test):  
    npwex = npywren.default_executor()  
  
    X_train_sharded = npwex.matrix_init(X_train)  
    X_test_sharded = npwex.matrix_init(X_test)  
  
    XYT = npwex.dot(X_train_sharded, X_test_sharded.T)  
    XYT *= -2  
    norms_train = npwex.linalg.norm(X_train, axis=1)  
    norms_test = npwex.linalg.norm(X_test, axis=1)  
    distances = norms_train + XYT + norms_test.T  
    argmins = npwex.argmin(distances, axis=0).numpy()  
  
    return metrics.accuracy_score(y_train[argmins], y_test)
```

SOLVING A LINEAR SYSTEM $Ax=B$

SOLVING A LINEAR SYSTEM $Ax=B$

I. Compute Cholesky Factorization $LL^T=A$

SOLVING A LINEAR SYSTEM $Ax=B$

1. Compute Cholesky Factorization $LL^T=A$
2. Forward substitution to solve $Lz = B$

SOLVING A LINEAR SYSTEM $Ax=B$

1. Compute Cholesky Factorization $LL^T=A$
2. Forward substitution to solve $Lz = B$
3. Backward substitution $L^Tx = z$

SOLVING A LINEAR SYSTEM $Ax=B$

1. Compute Cholesky Factorization $LL^T=A$ $O(n^3)$
2. Forward substitution to solve $Lz = B$
3. Backward substitution $L^Tx = z$

SOLVING A LINEAR SYSTEM $Ax=B$

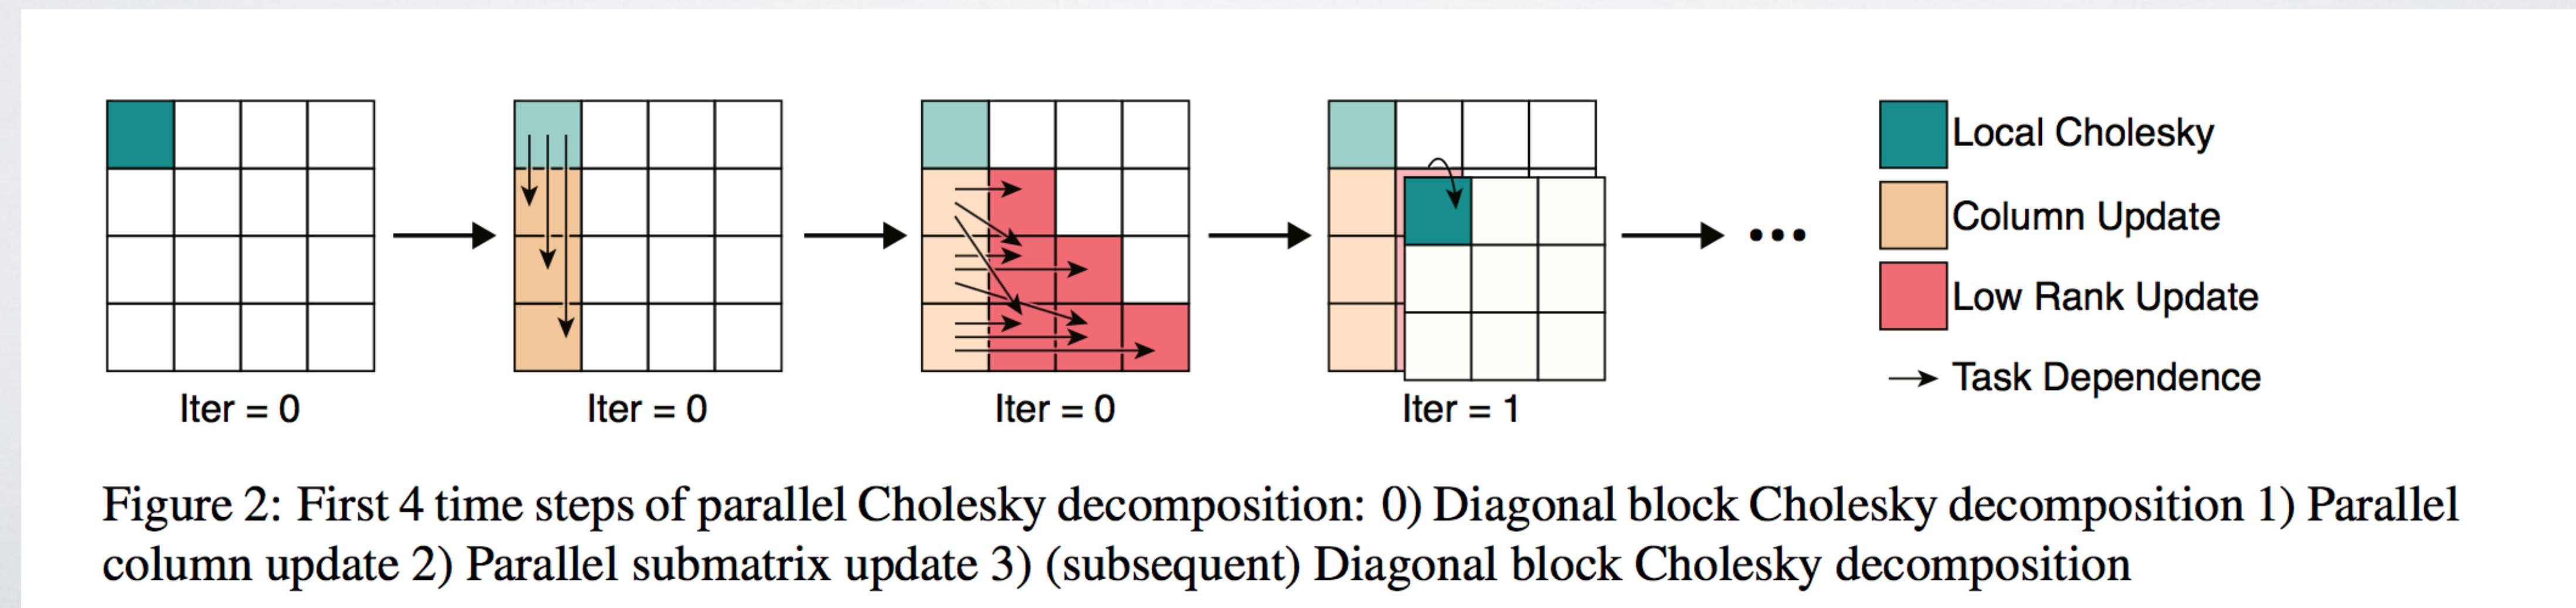
1. Compute Cholesky Factorization $LL^T=A$ $O(n^3)$
2. Forward substitution to solve $Lz = B$ $O(n^2)$
3. Backward substitution $L^Tx = z$

SOLVING A LINEAR SYSTEM $Ax=B$

1. Compute Cholesky Factorization $LL^T=A$ $O(n^3)$
2. Forward substitution to solve $Lz = B$ $O(n^2)$
3. Backward substitution $L^Tx = z$ $O(n^2)$

SOLVING A LINEAR SYSTEM $Ax=B$

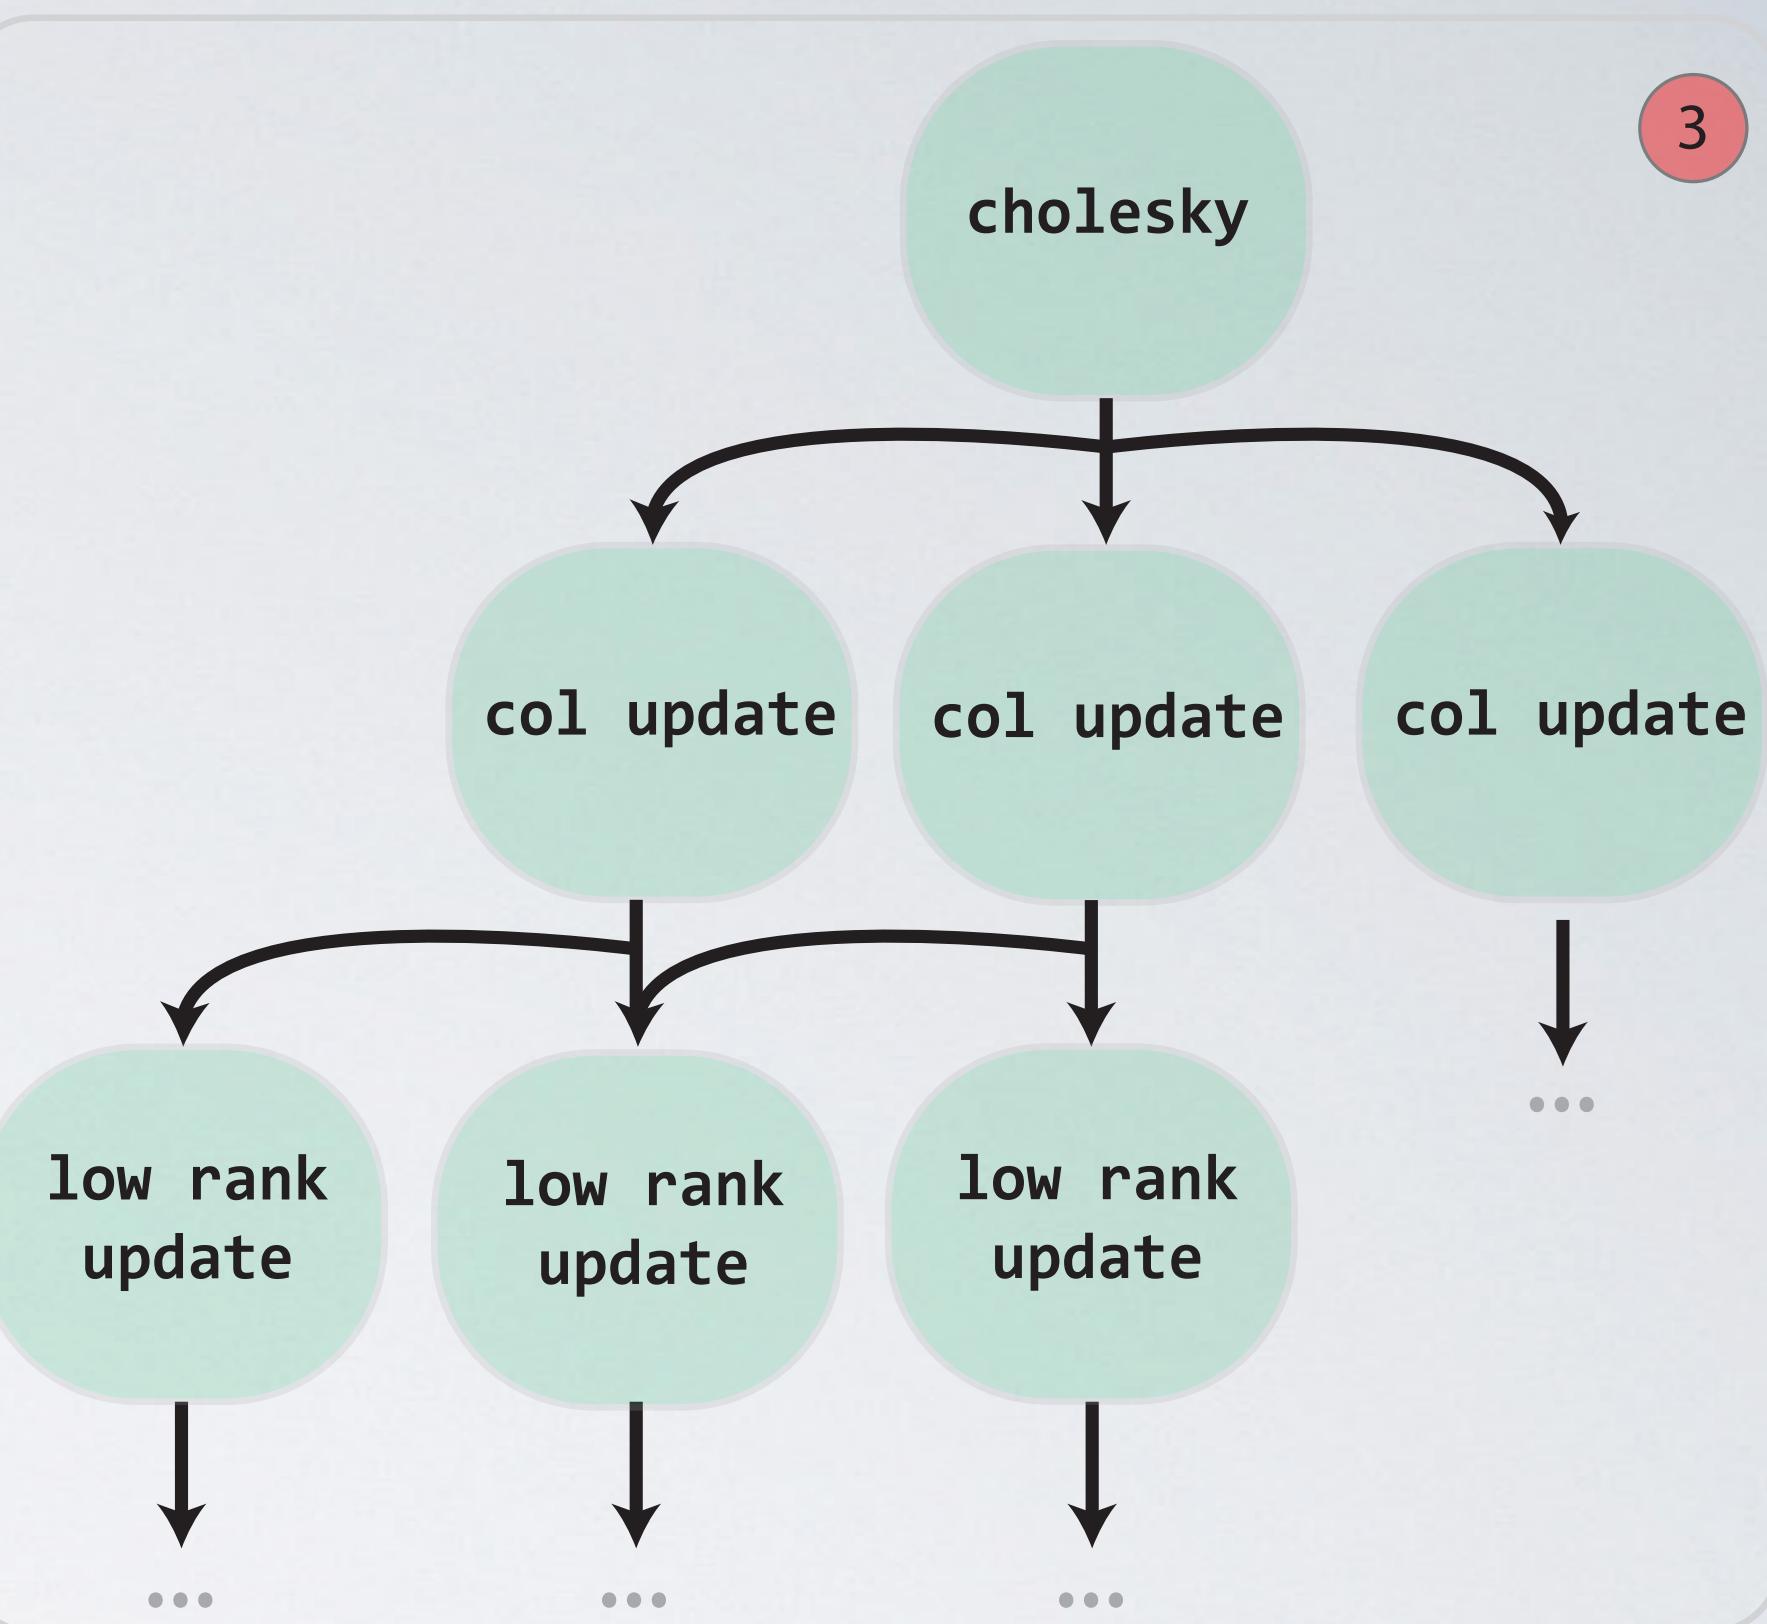
1. Compute Cholesky Factorization $LL^T=A$ $O(n^3)$
2. Forward substitution to solve $Lz = B$ $O(n^2)$
3. Backward substitution $L^T x = z$ $O(n^2)$



1 numpywren.cholesky

2

```
cholesky(iter=0)
  0 = LOAD BigMatrix(X)[0, 0]
  1 = CHOL 0
  2 = WRITE chol(BigMatrix(X))[0, 0]
col_update(row=1, col=0)
  3 = LOAD chol(BigMatrix(X))[0, 0]
  4 = LOAD BigMatrix(X)[1, 0]
  5 = TRSM 3 4
  6 = WRITE chol(BigMatrix(X))[1, 0]
col_update(row=2, col=0)
  ...
low_rank_update(iter=0, row=1, col=2)
  15 = LOAD chol(BigMatrix(X))[0, 0]
  16 = LOAD chol(BigMatrix(X))[1, 0]
  17 = LOAD chol(BigMatrix(X))[2, 0]
  18 = SYRK 15 16 17
  19 = WRITE temp(BigMatrix(X))[0, 1, 2]
low_rank_update(iter=0, row=1, col=3)
  ...
```

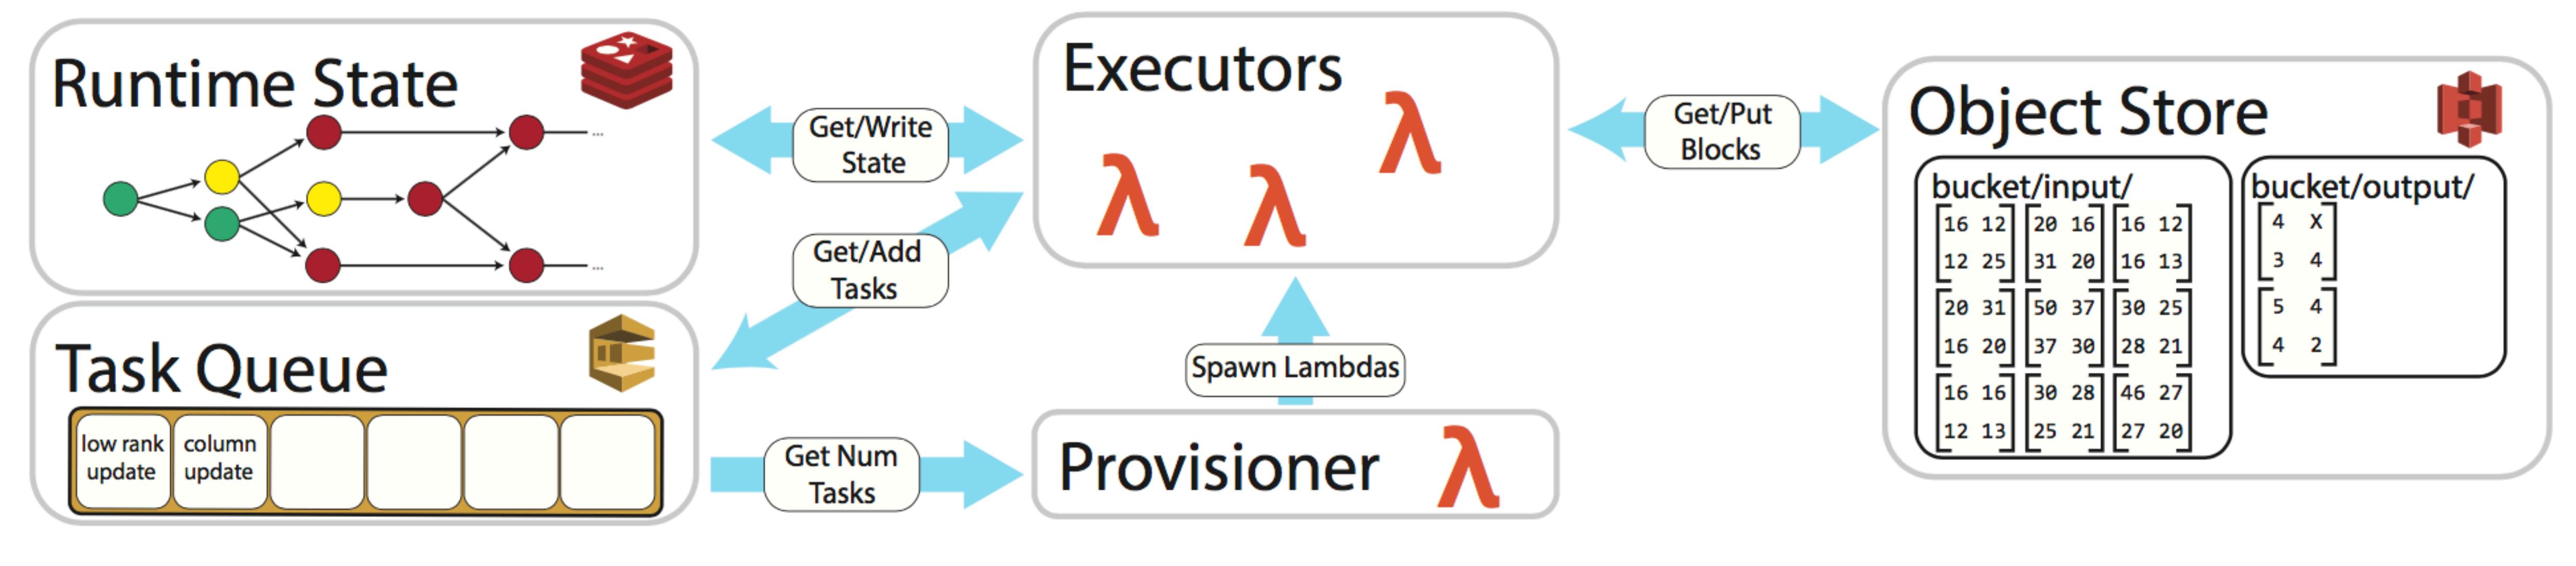


cholesky(iterator=0)
 0 = LOAD BigMatrix(X)[0, 0]
 1 = CHOL 0
 2 = WRITE chol(BigMatrix(X))[0, 0]
 col_update(row=1, col=0)
 3 = LOAD chol(BigMatrix(X))[0, 0]

2

cholesky

3

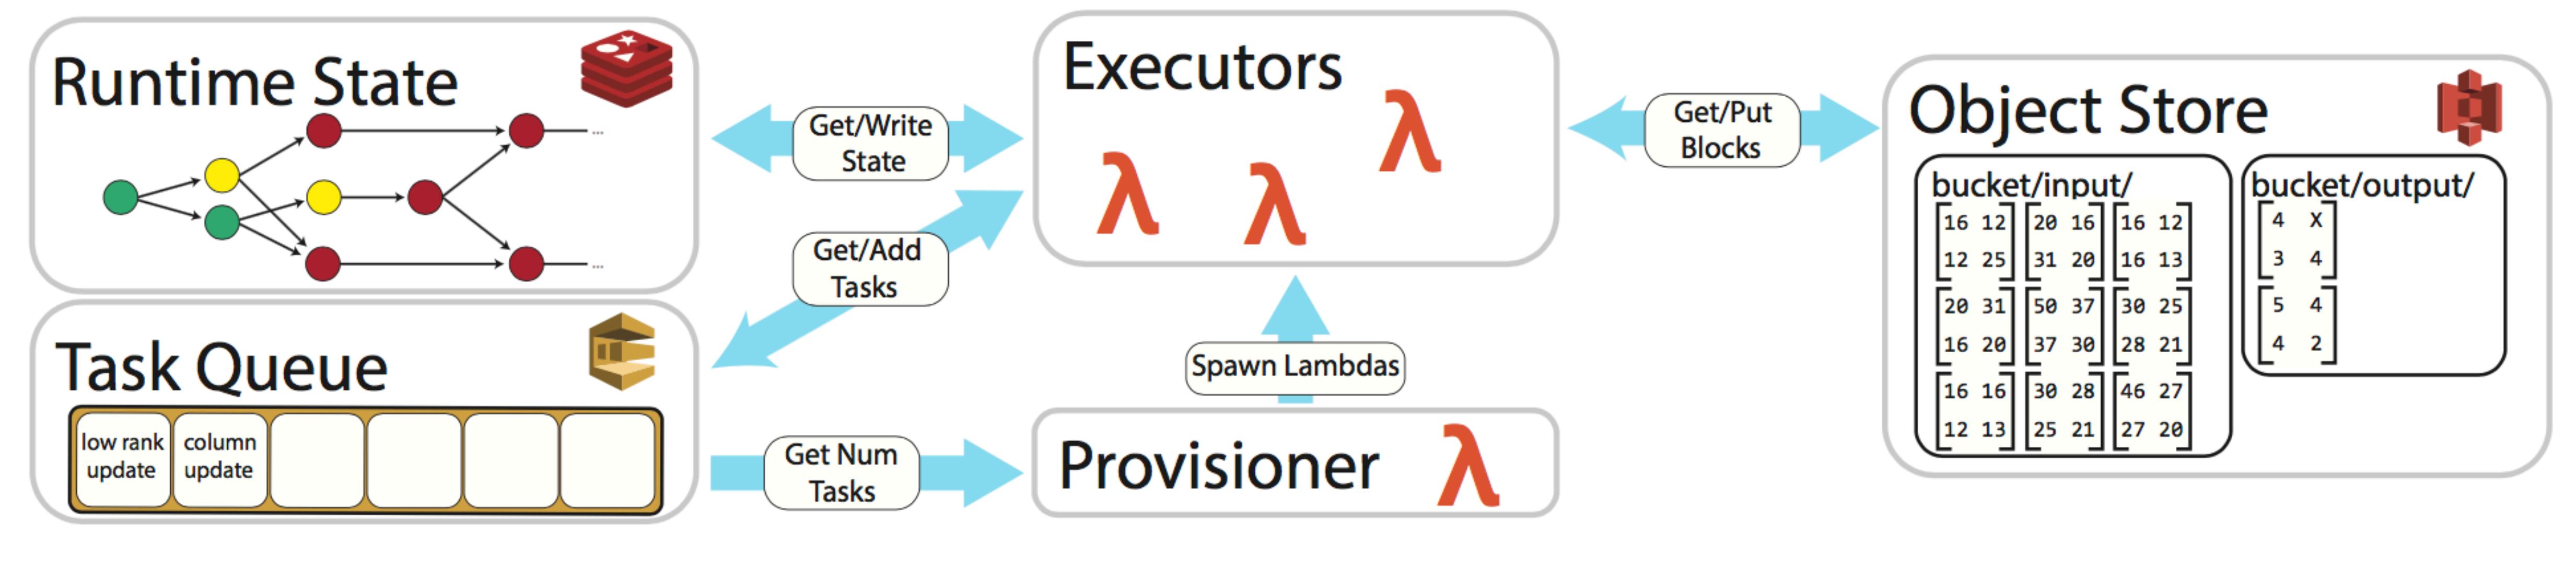


cholesky(iterator=0)
 0 = LOAD BigMatrix(X)[0, 0]
 1 = CHOL 0
 2 = WRITE chol(BigMatrix(X))[0, 0]
 col_update(row=1, col=0)
 3 = LOAD chol(BigMatrix(X))[0, 0]

2

cholesky

3

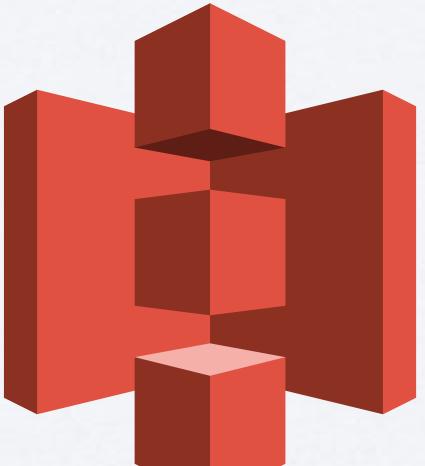


Compute



Lambda

Storage



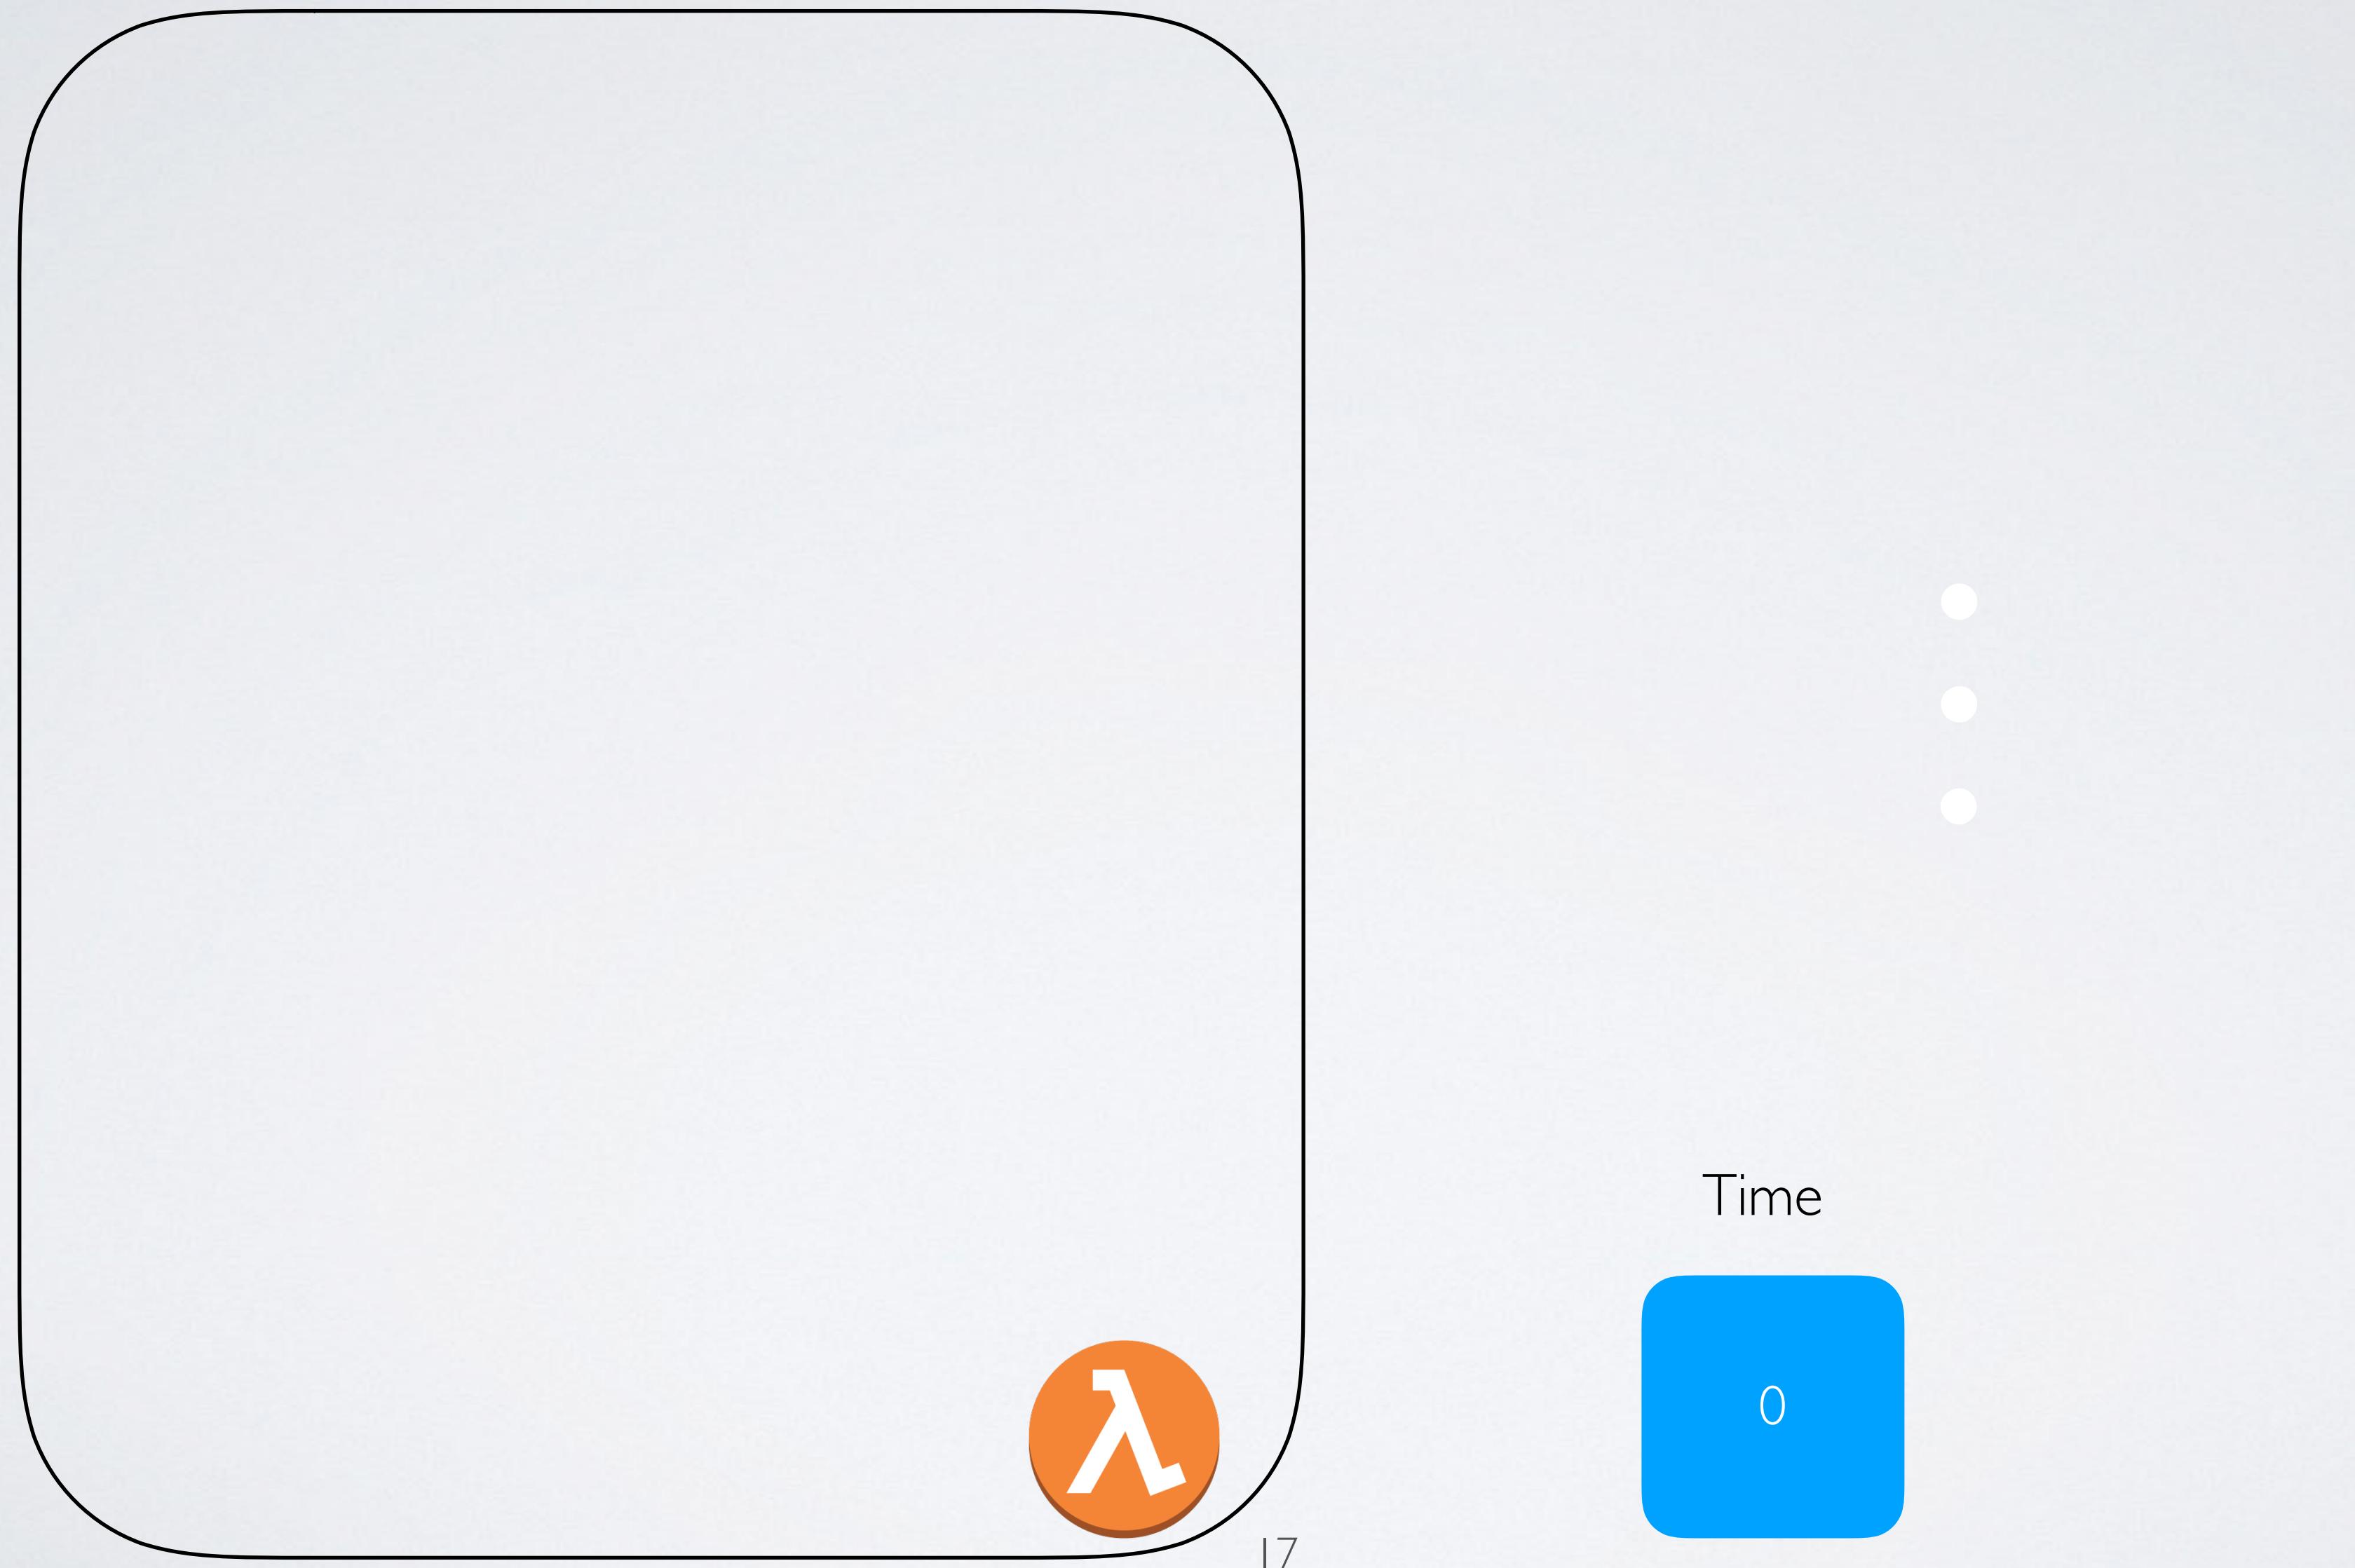
S3

Control Flow

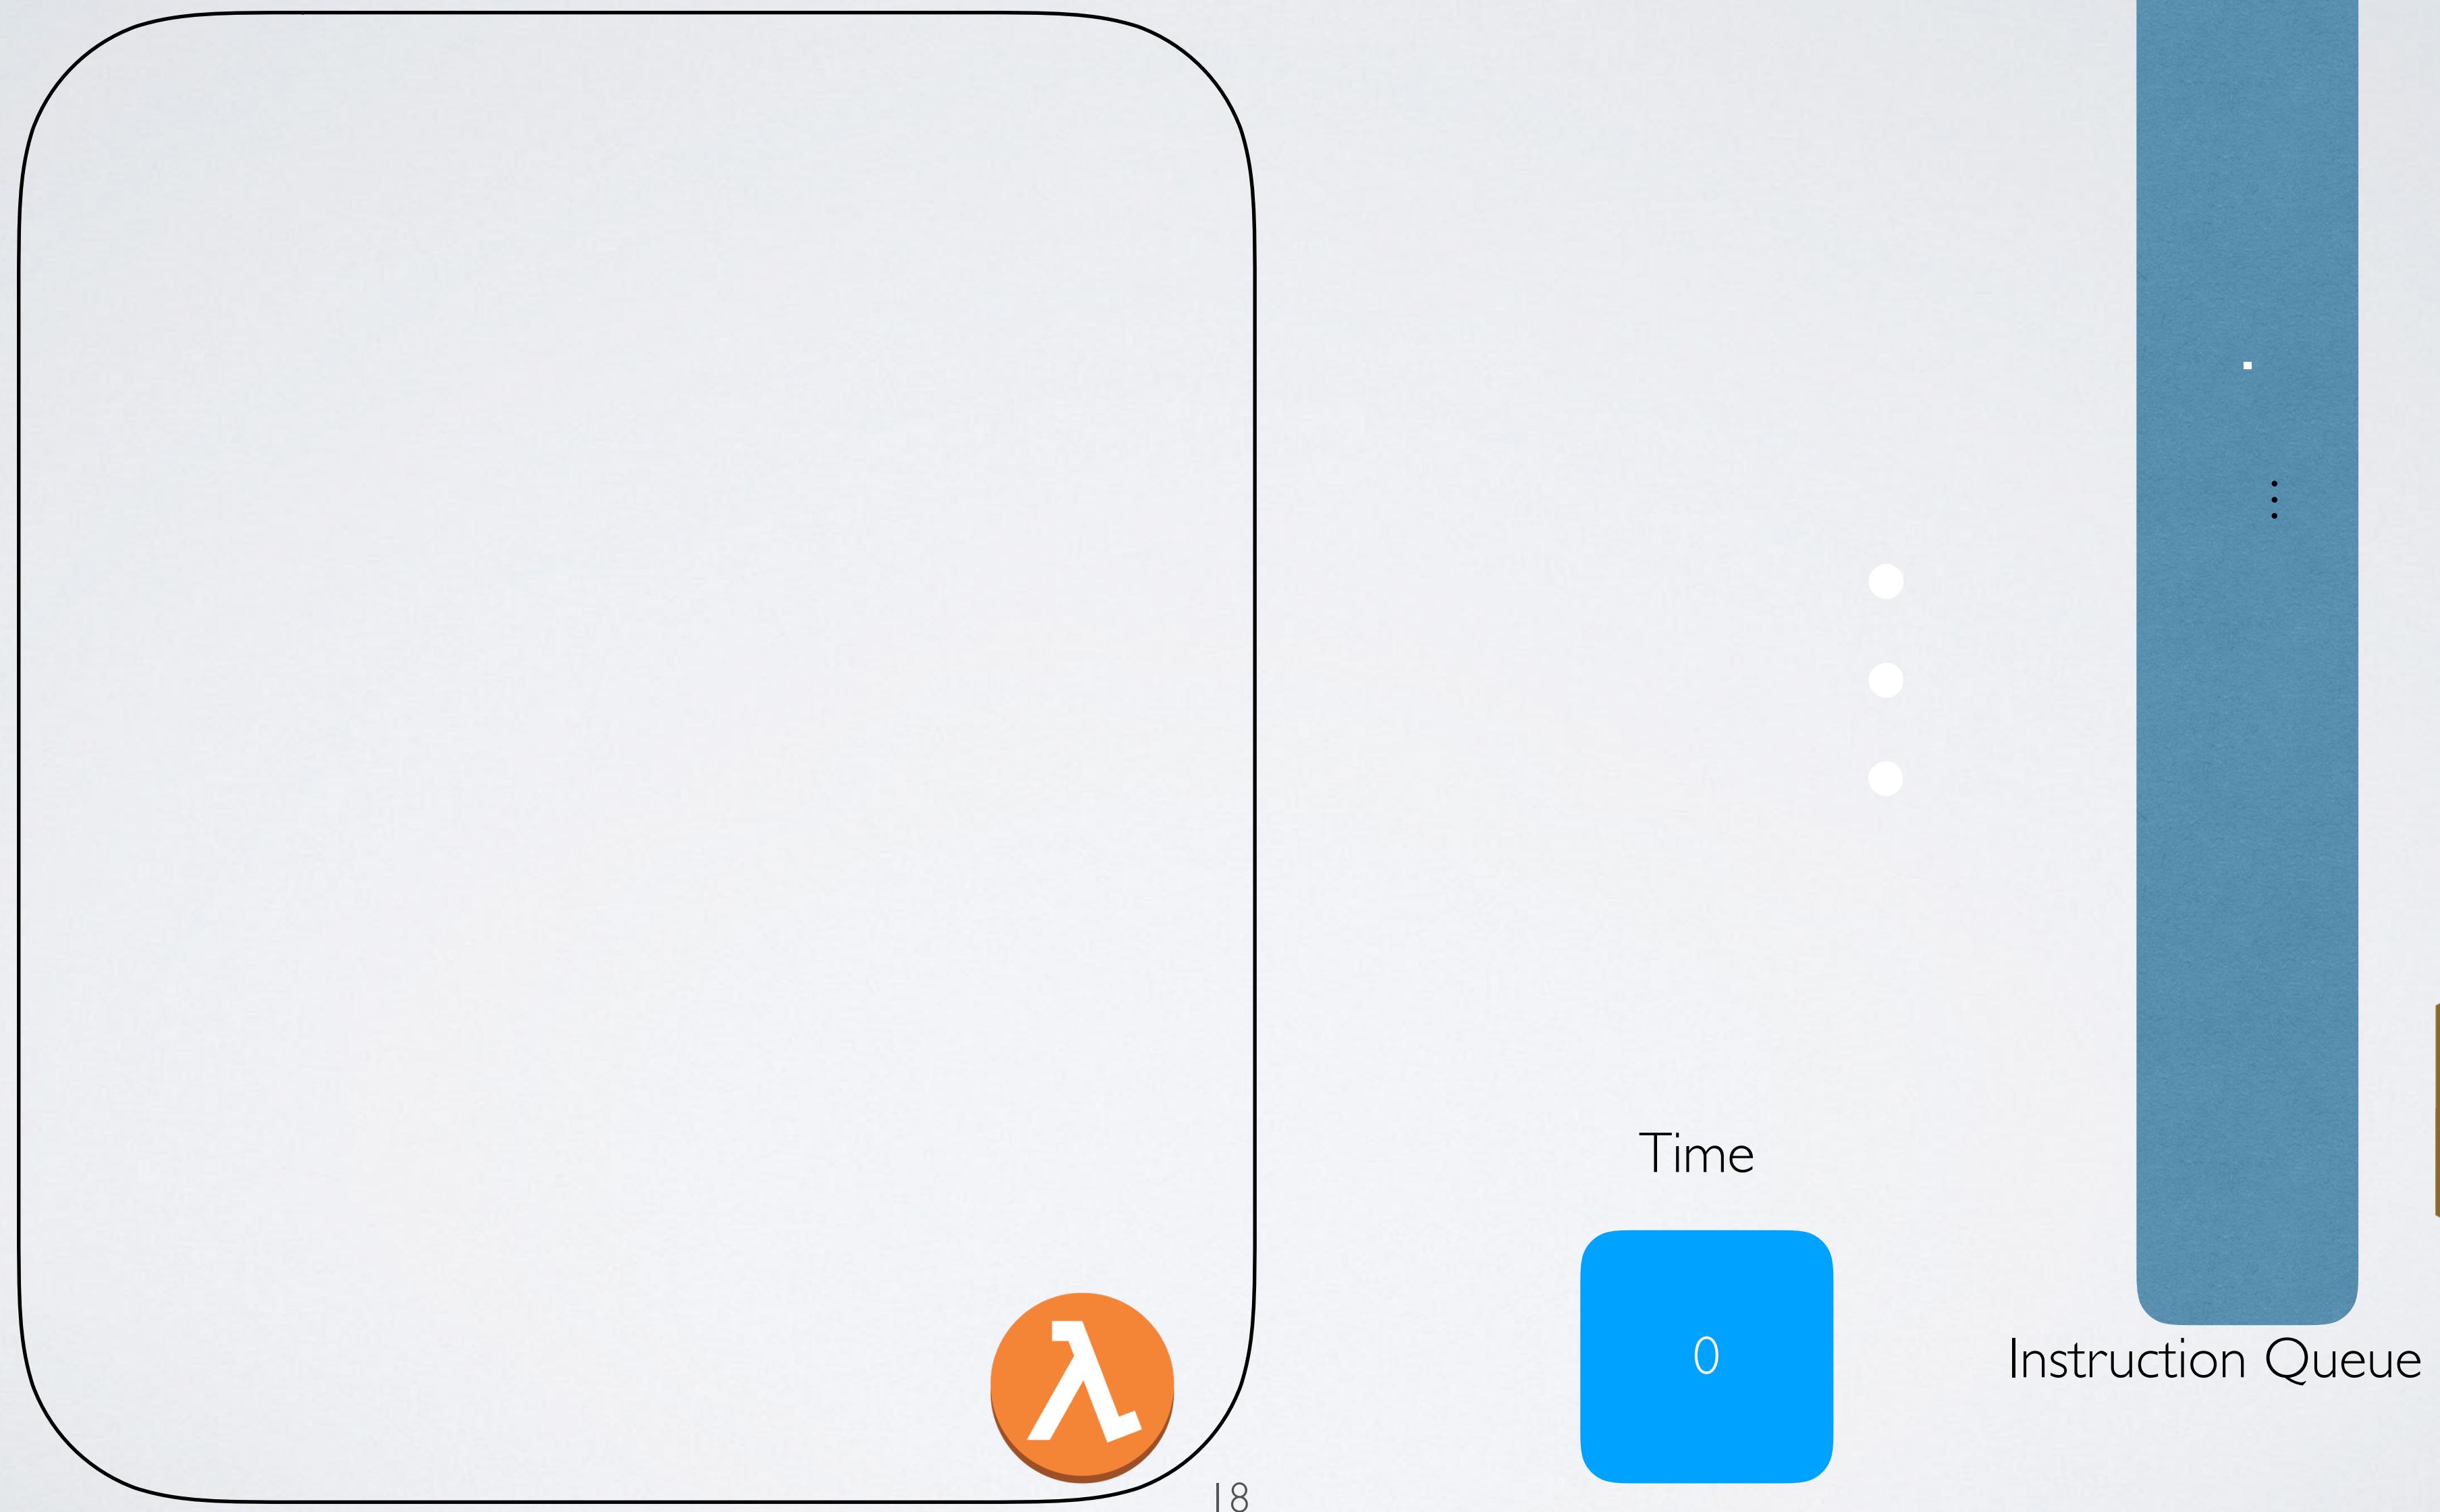


SQS

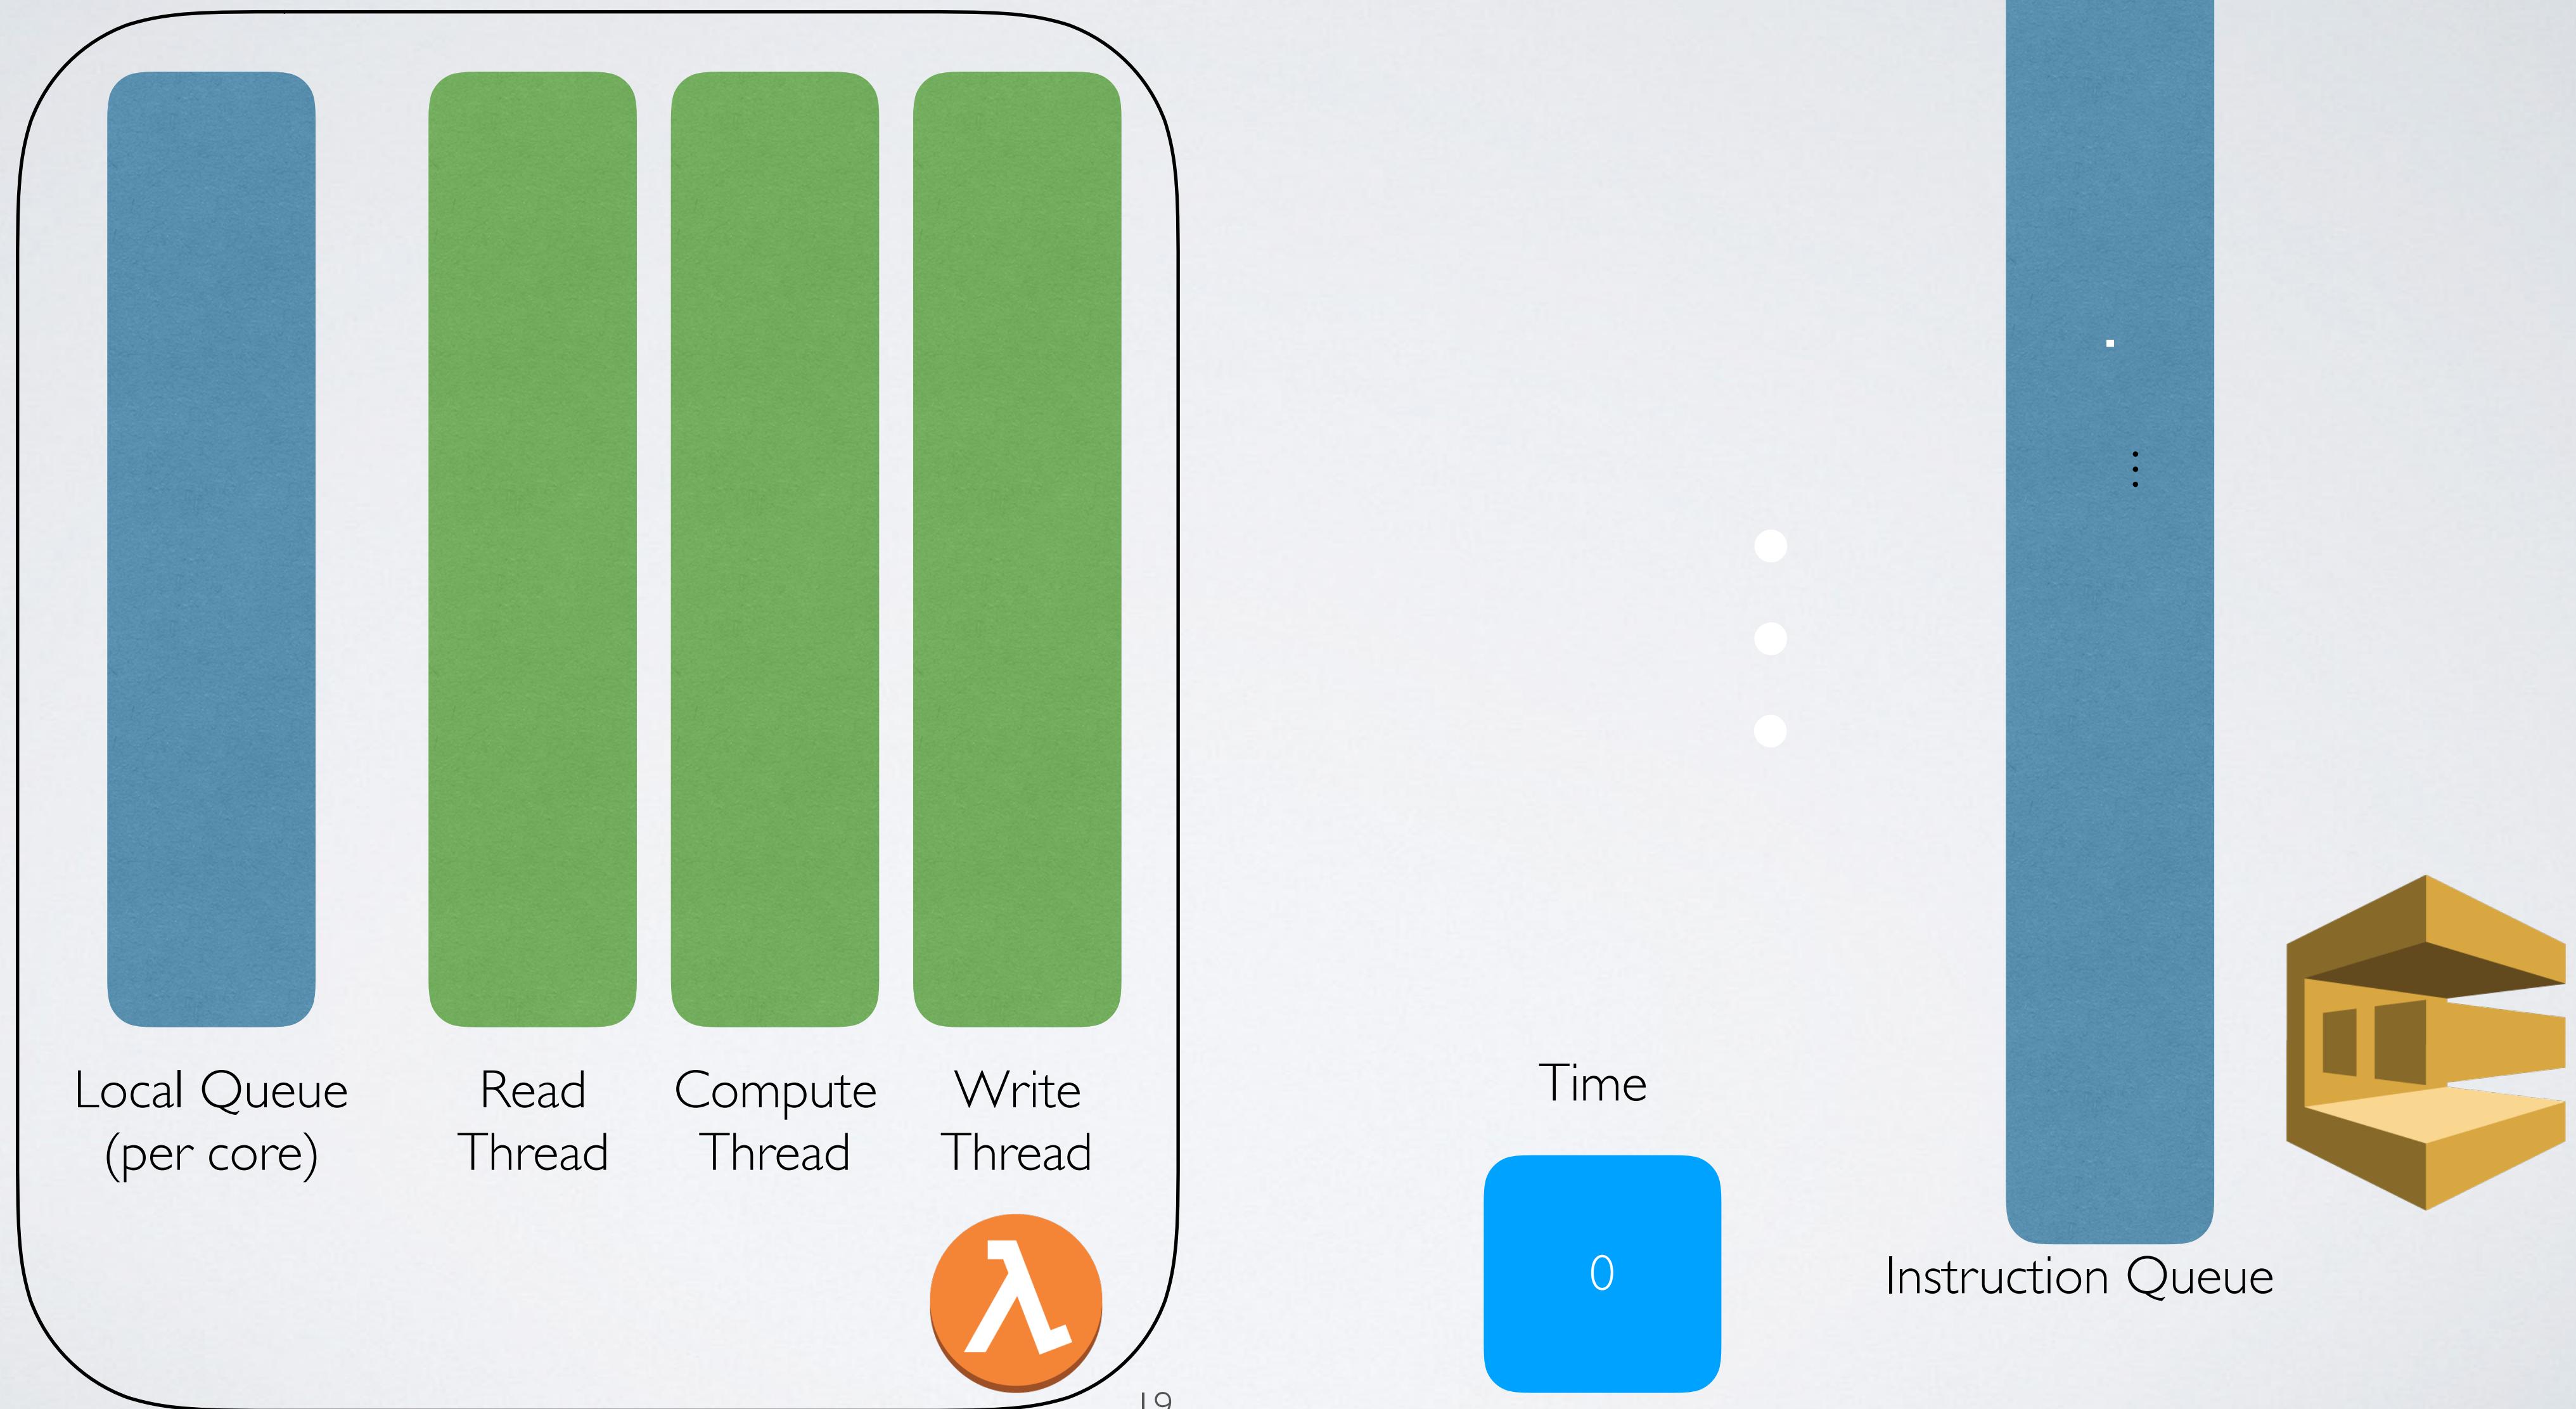
EXECUTION



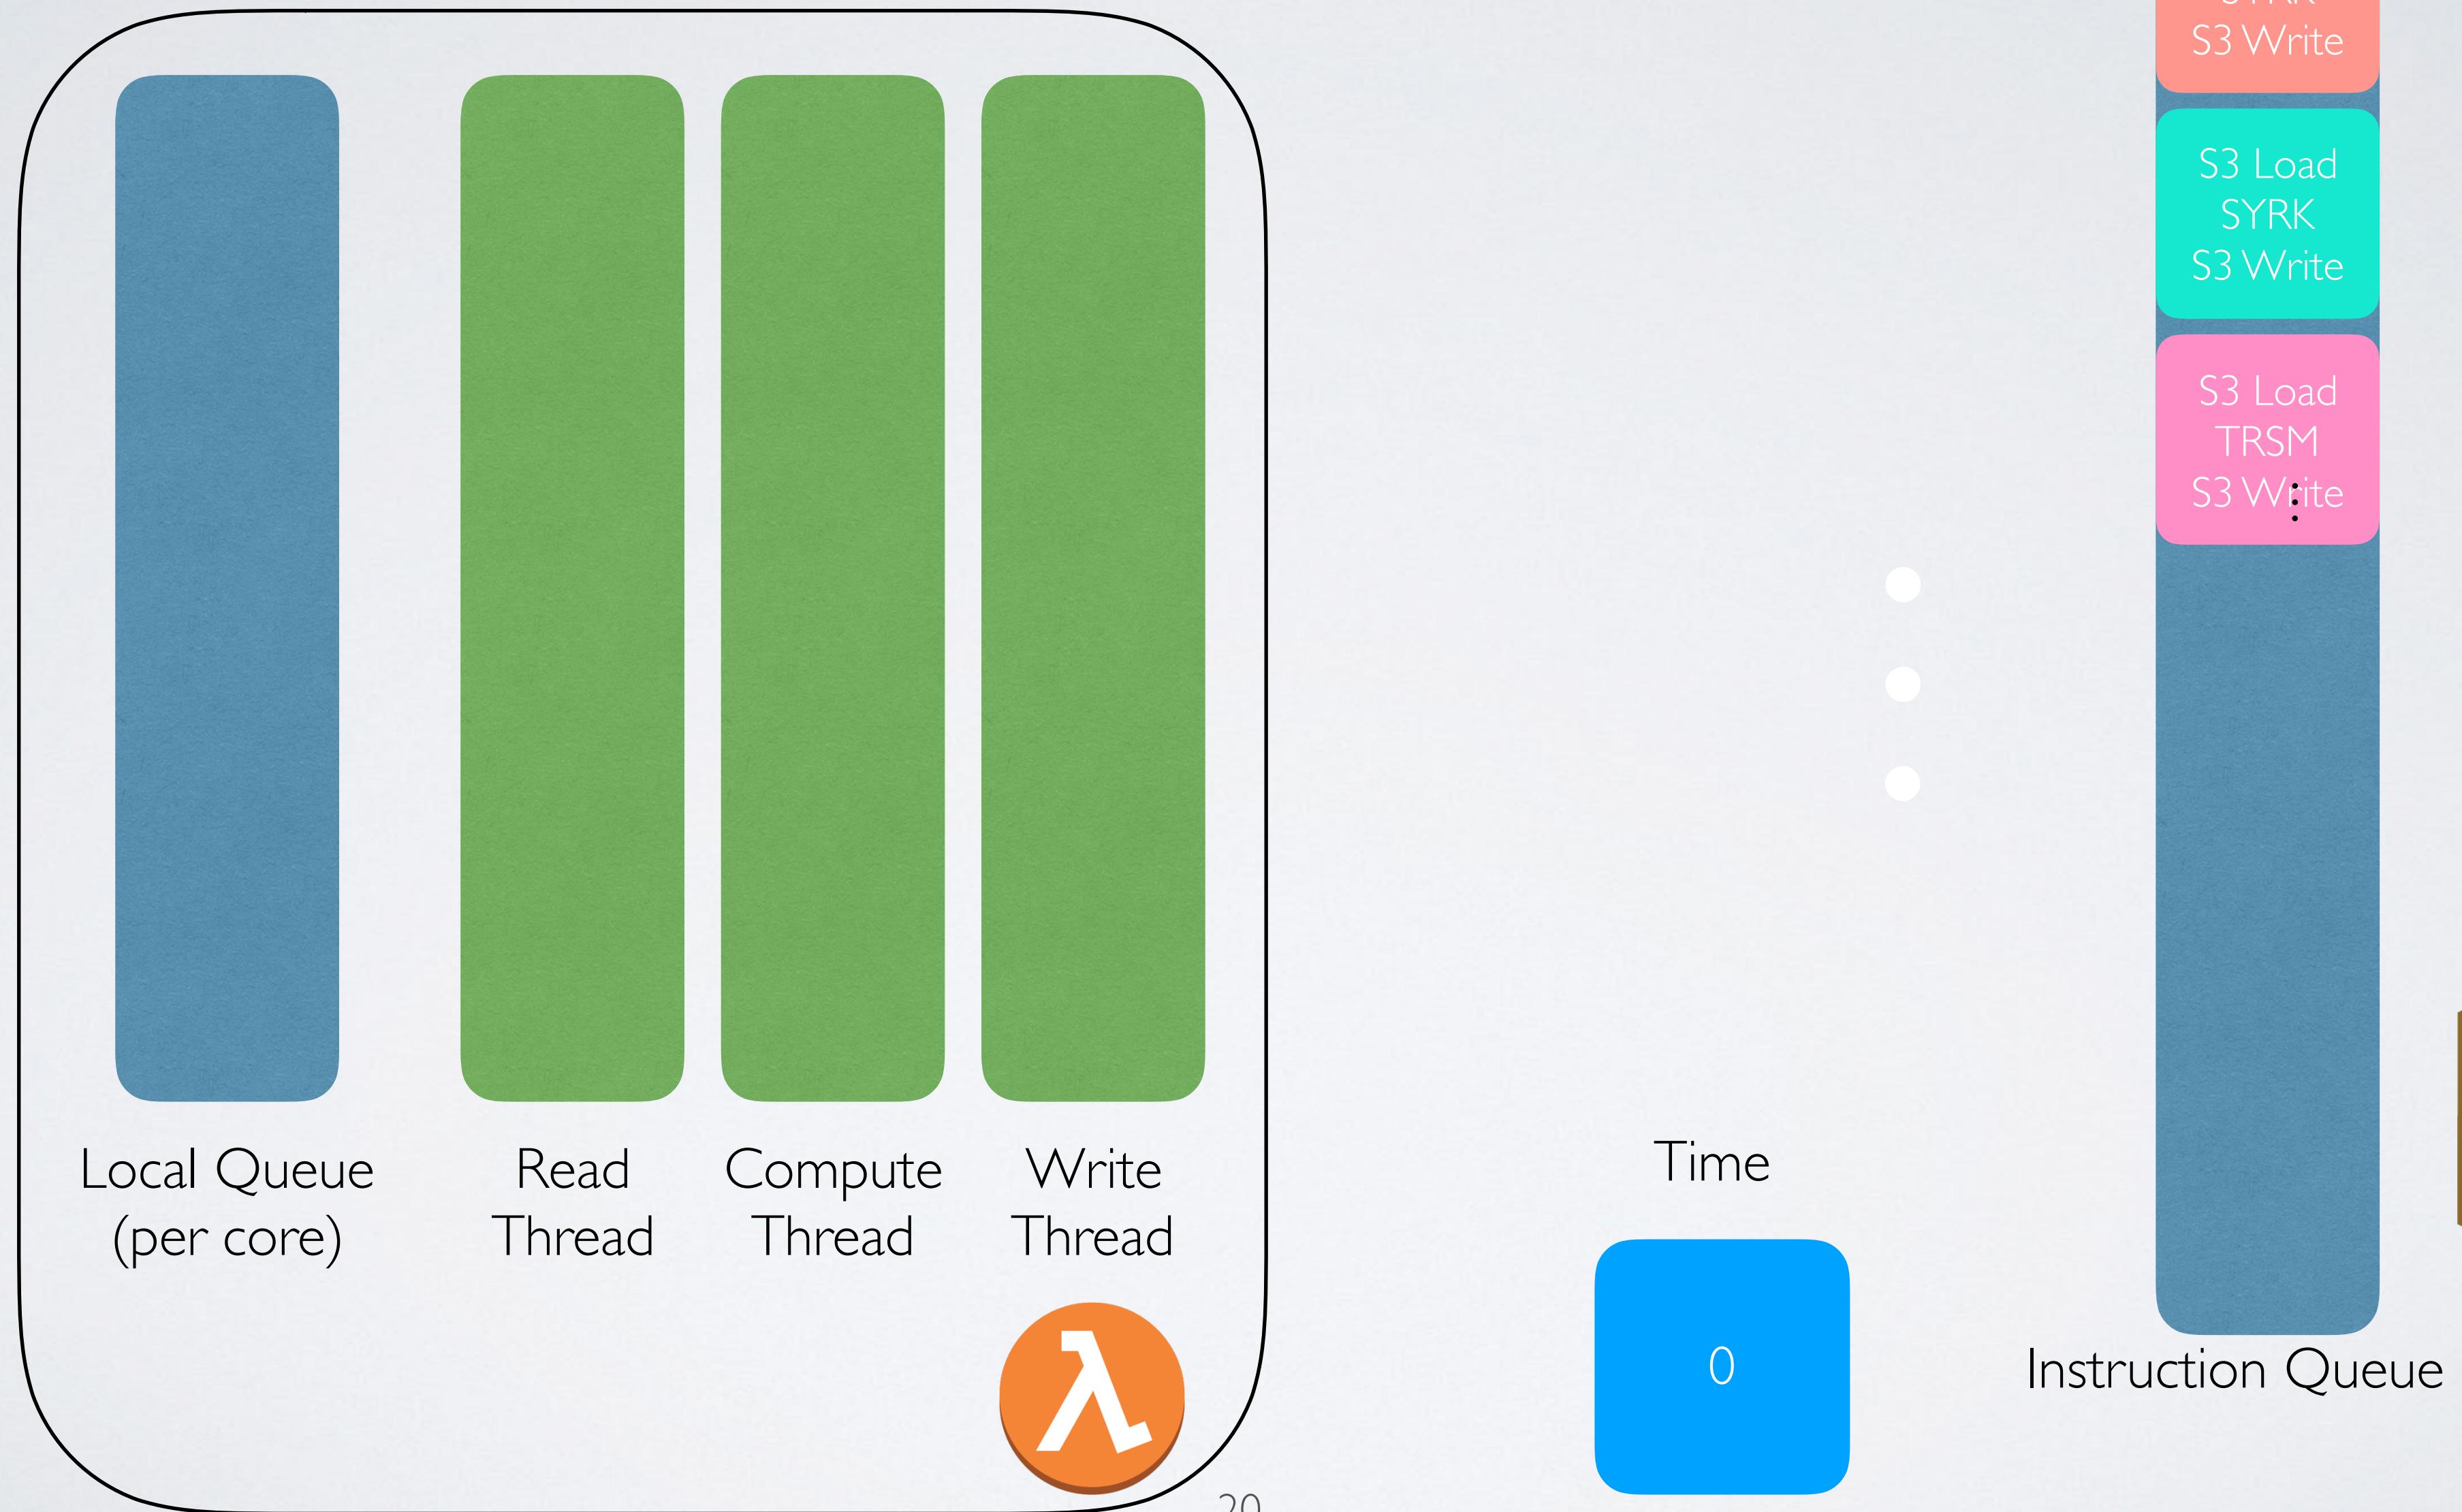
EXECUTION



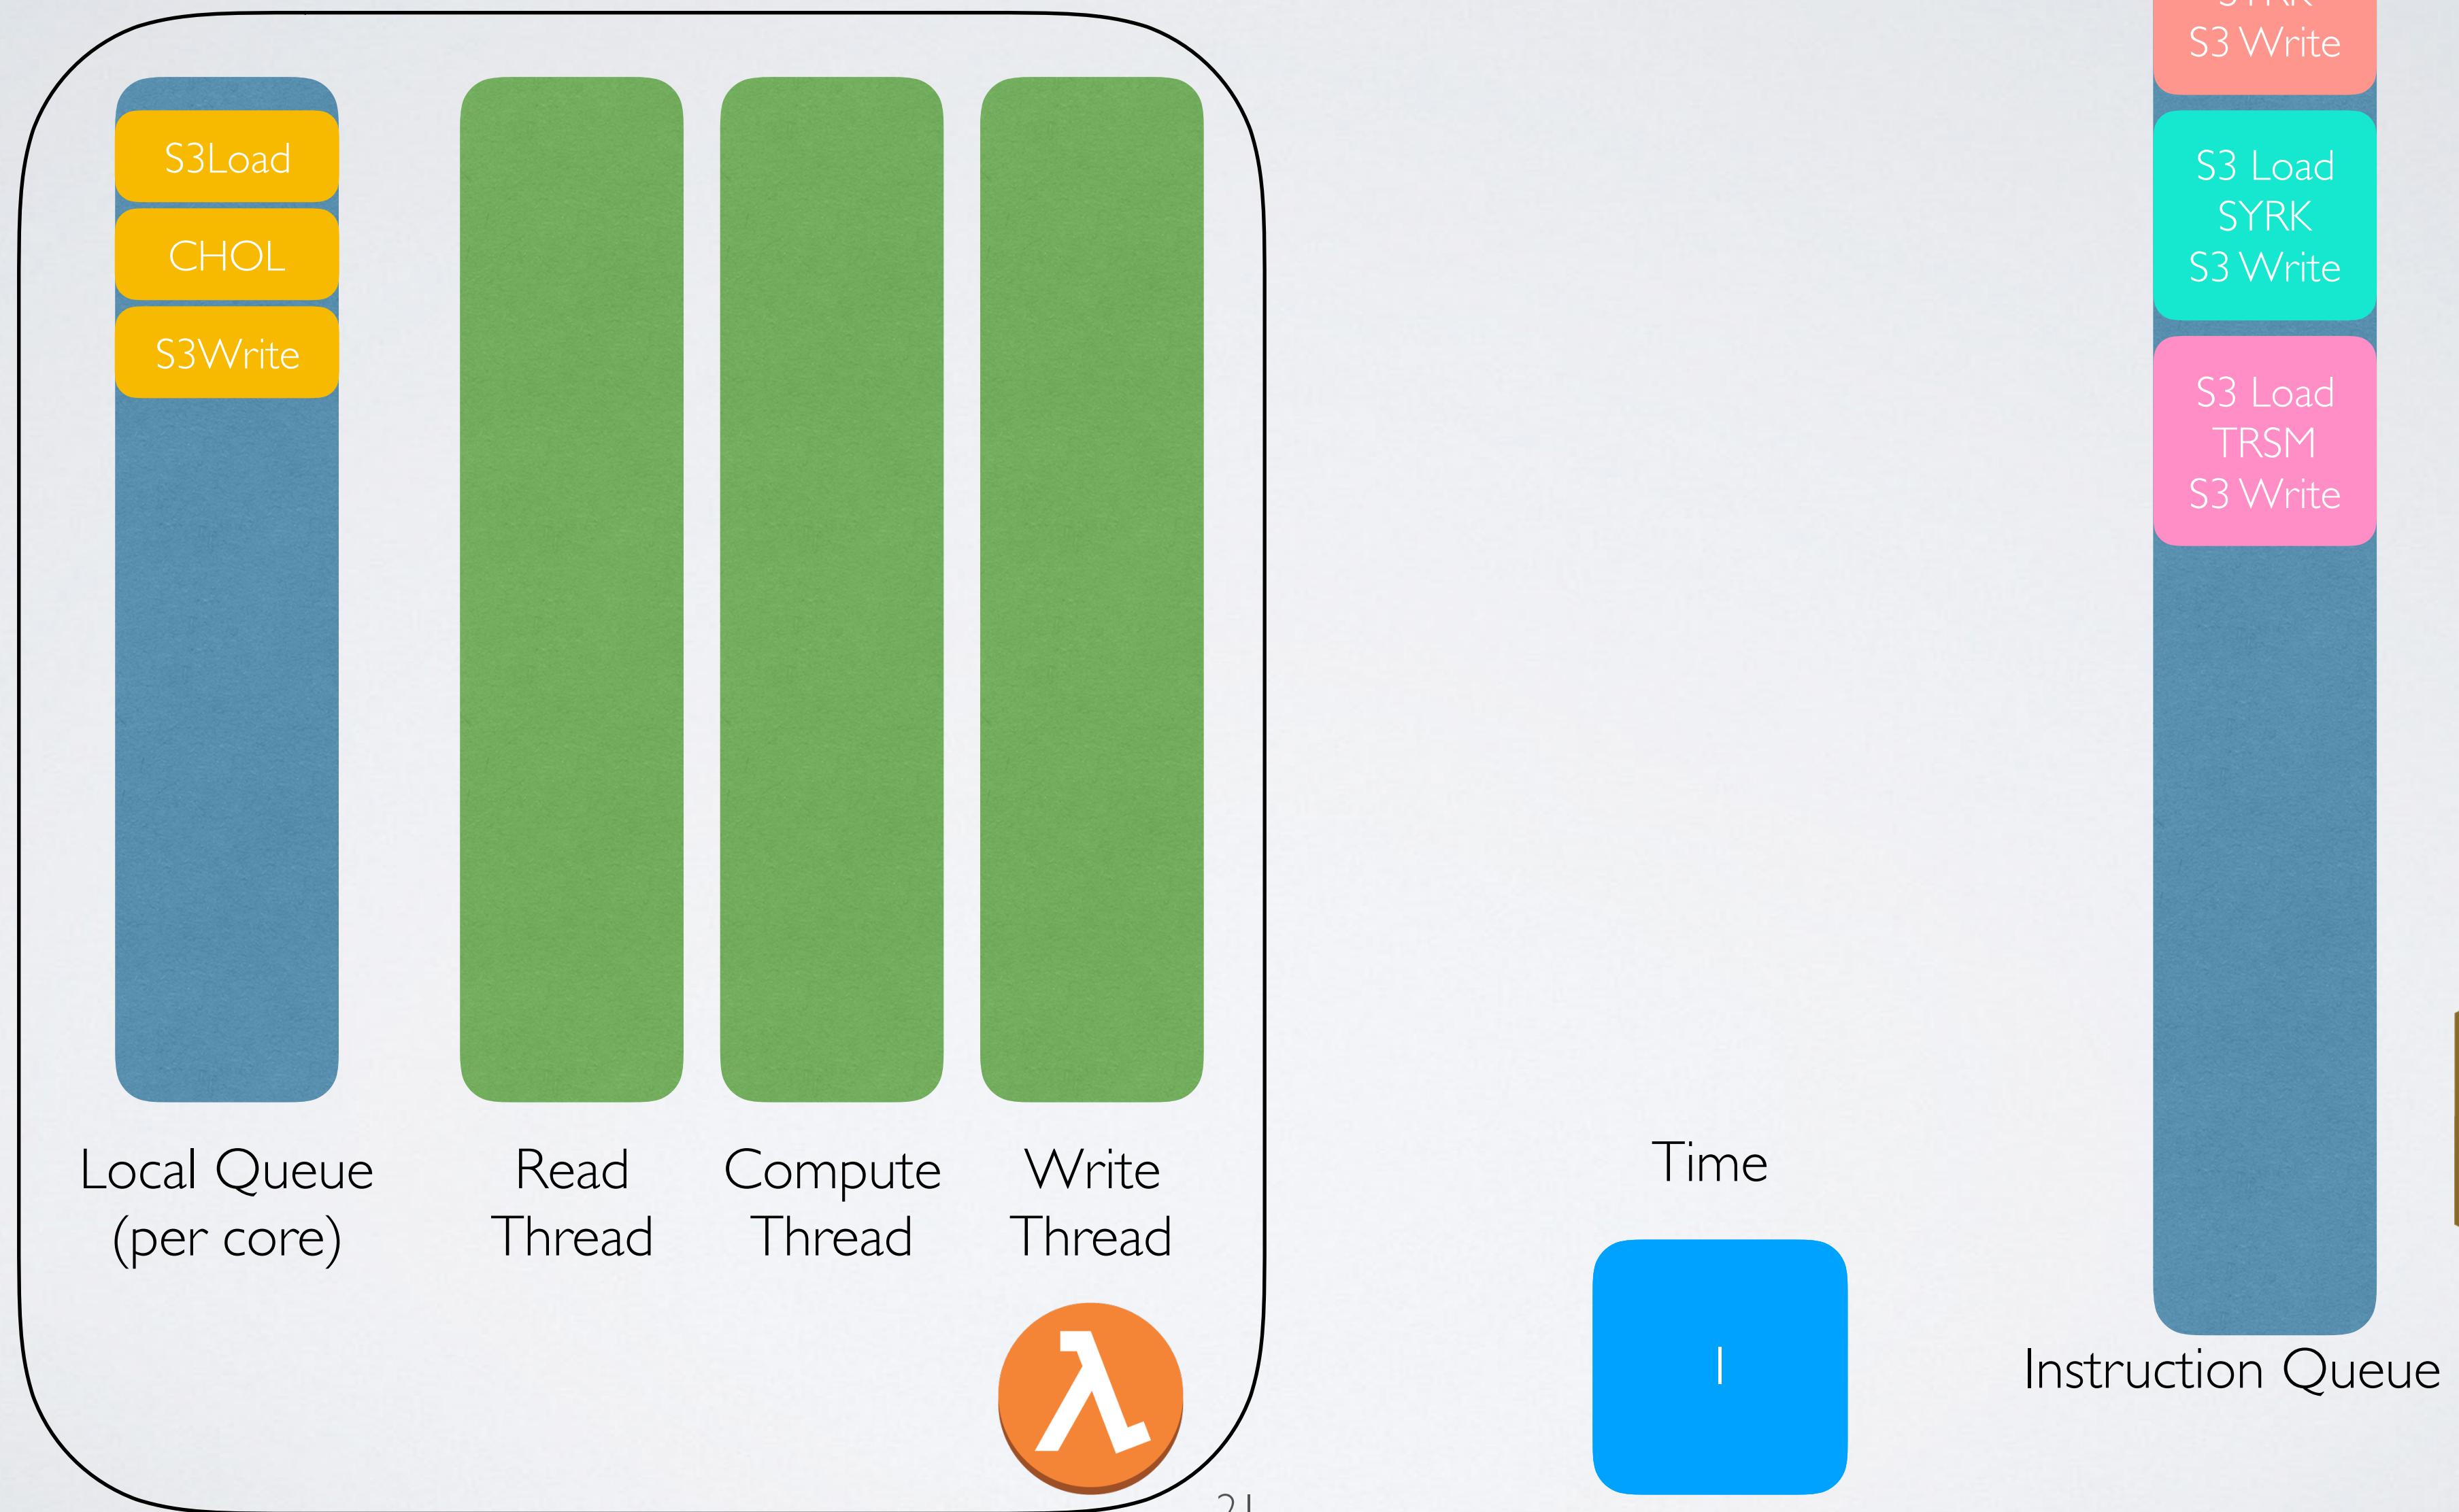
EXECUTION



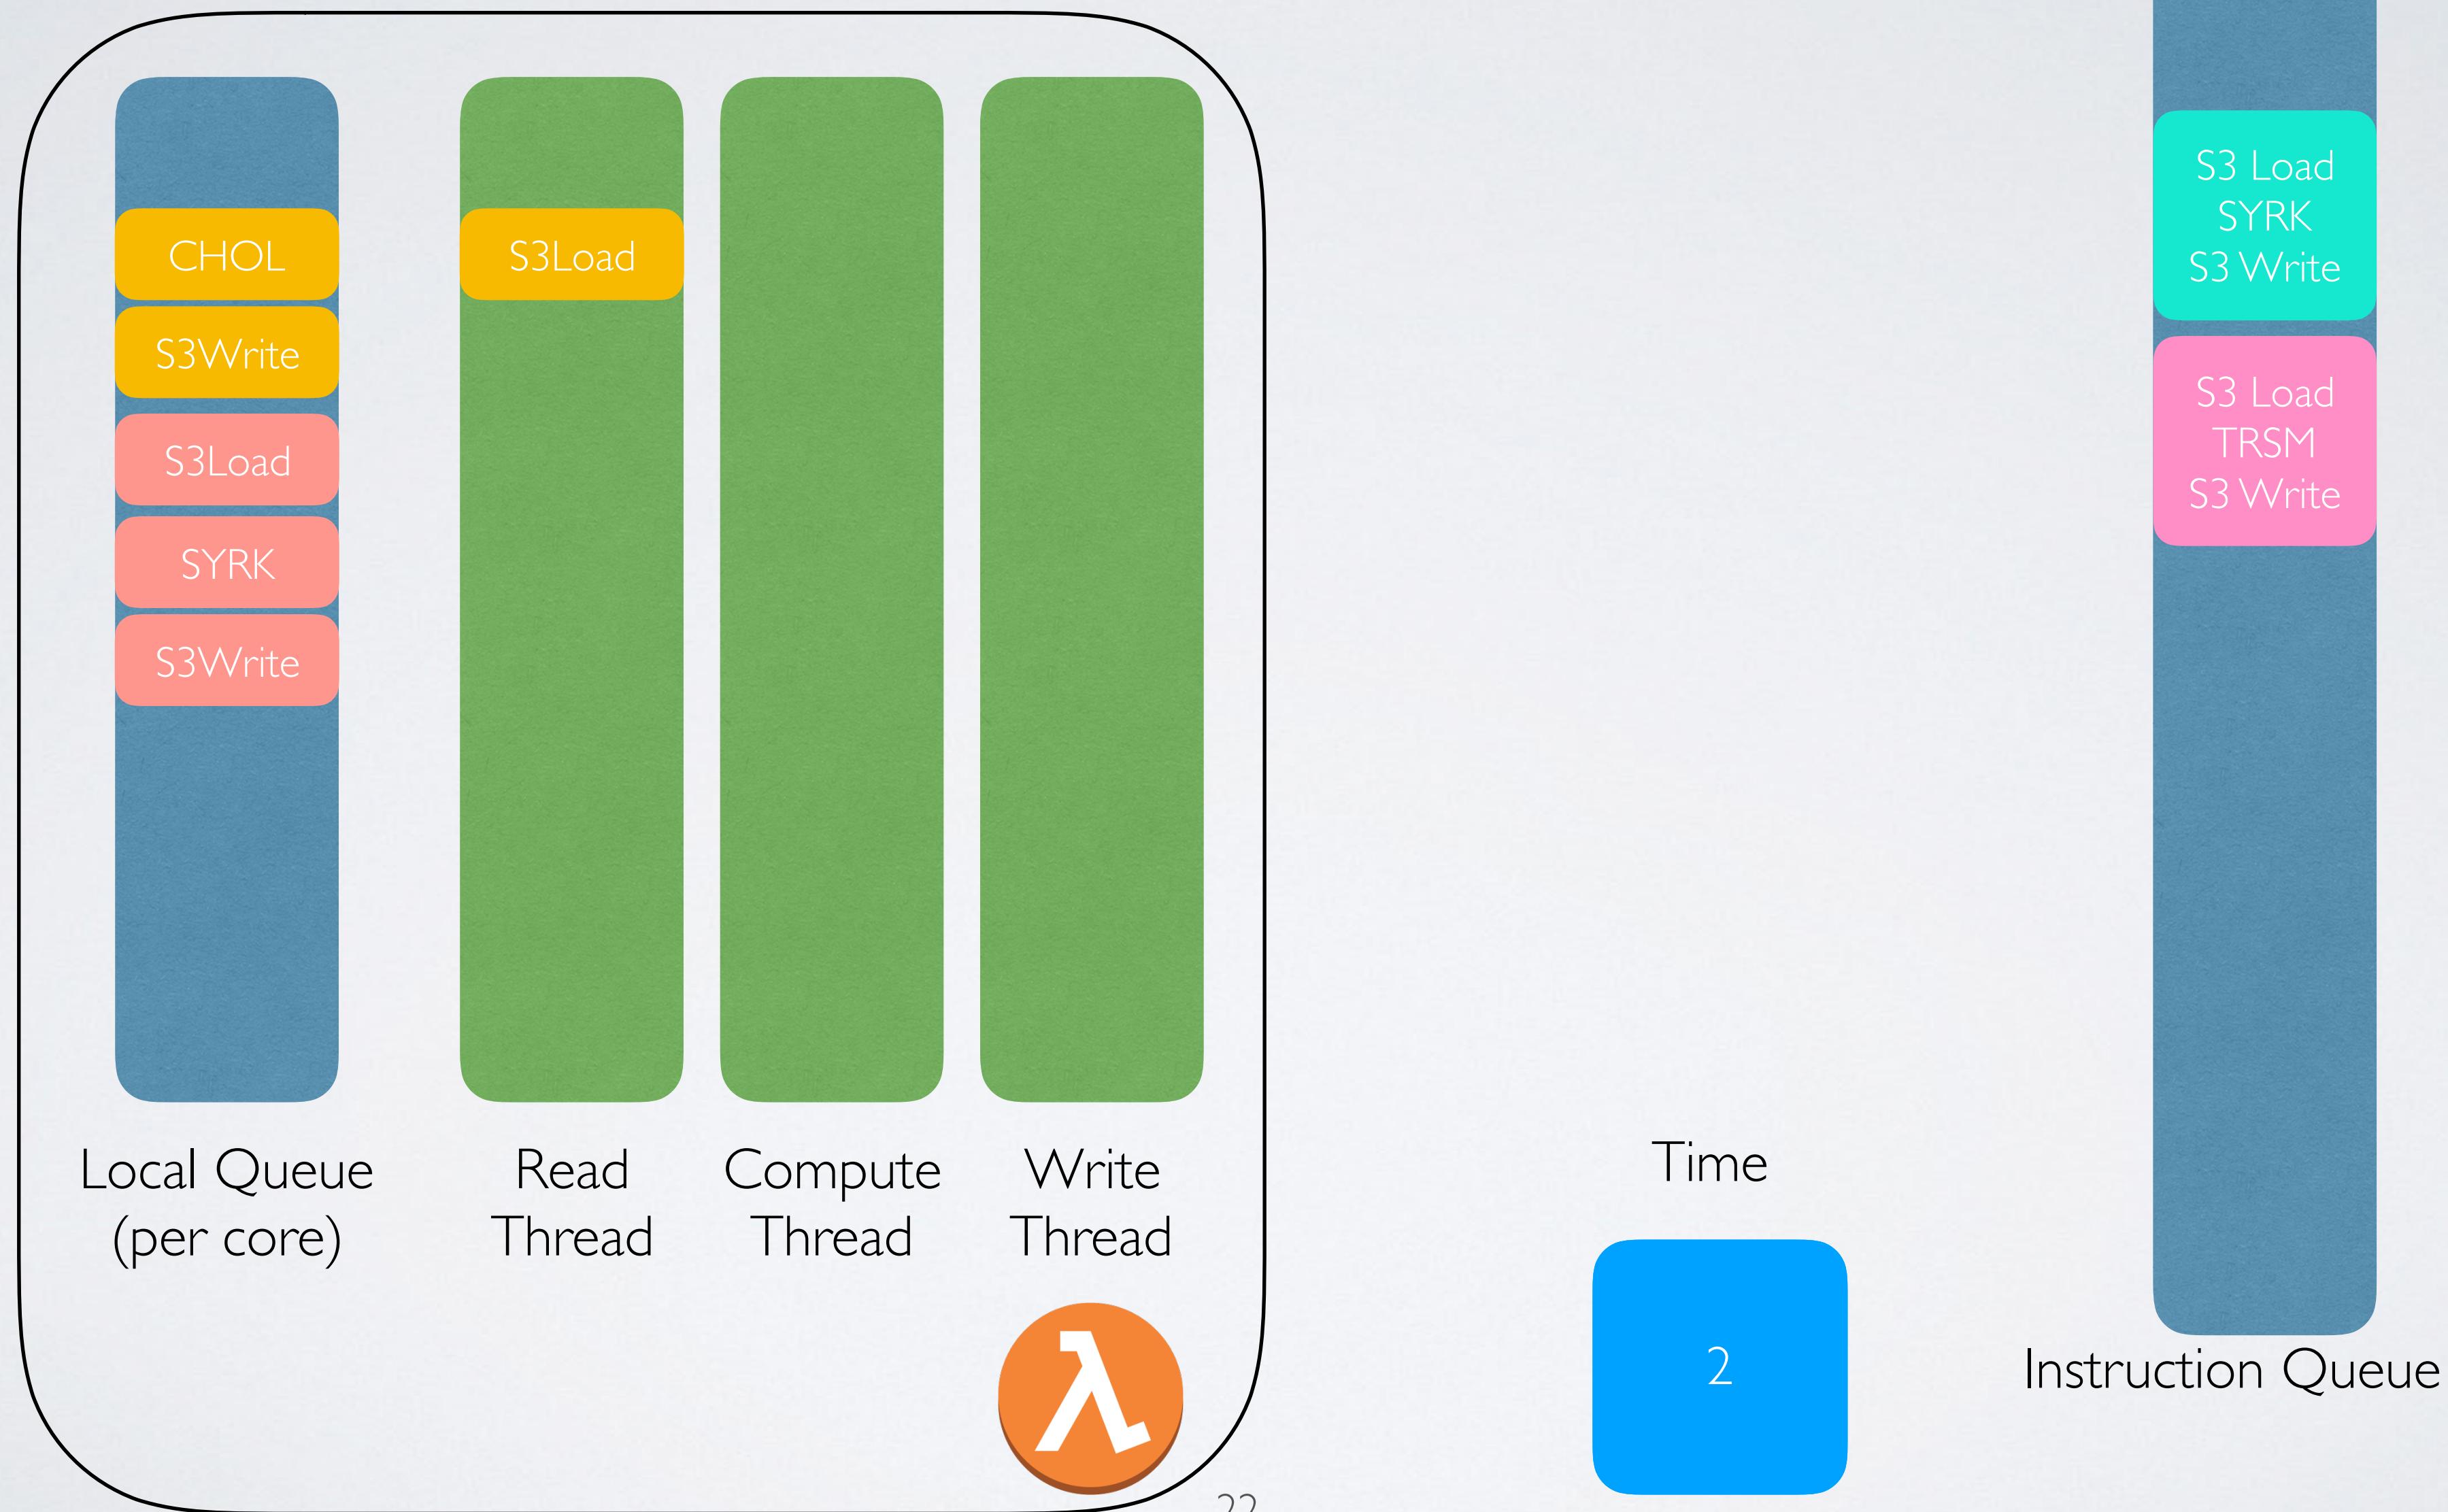
EXECUTION



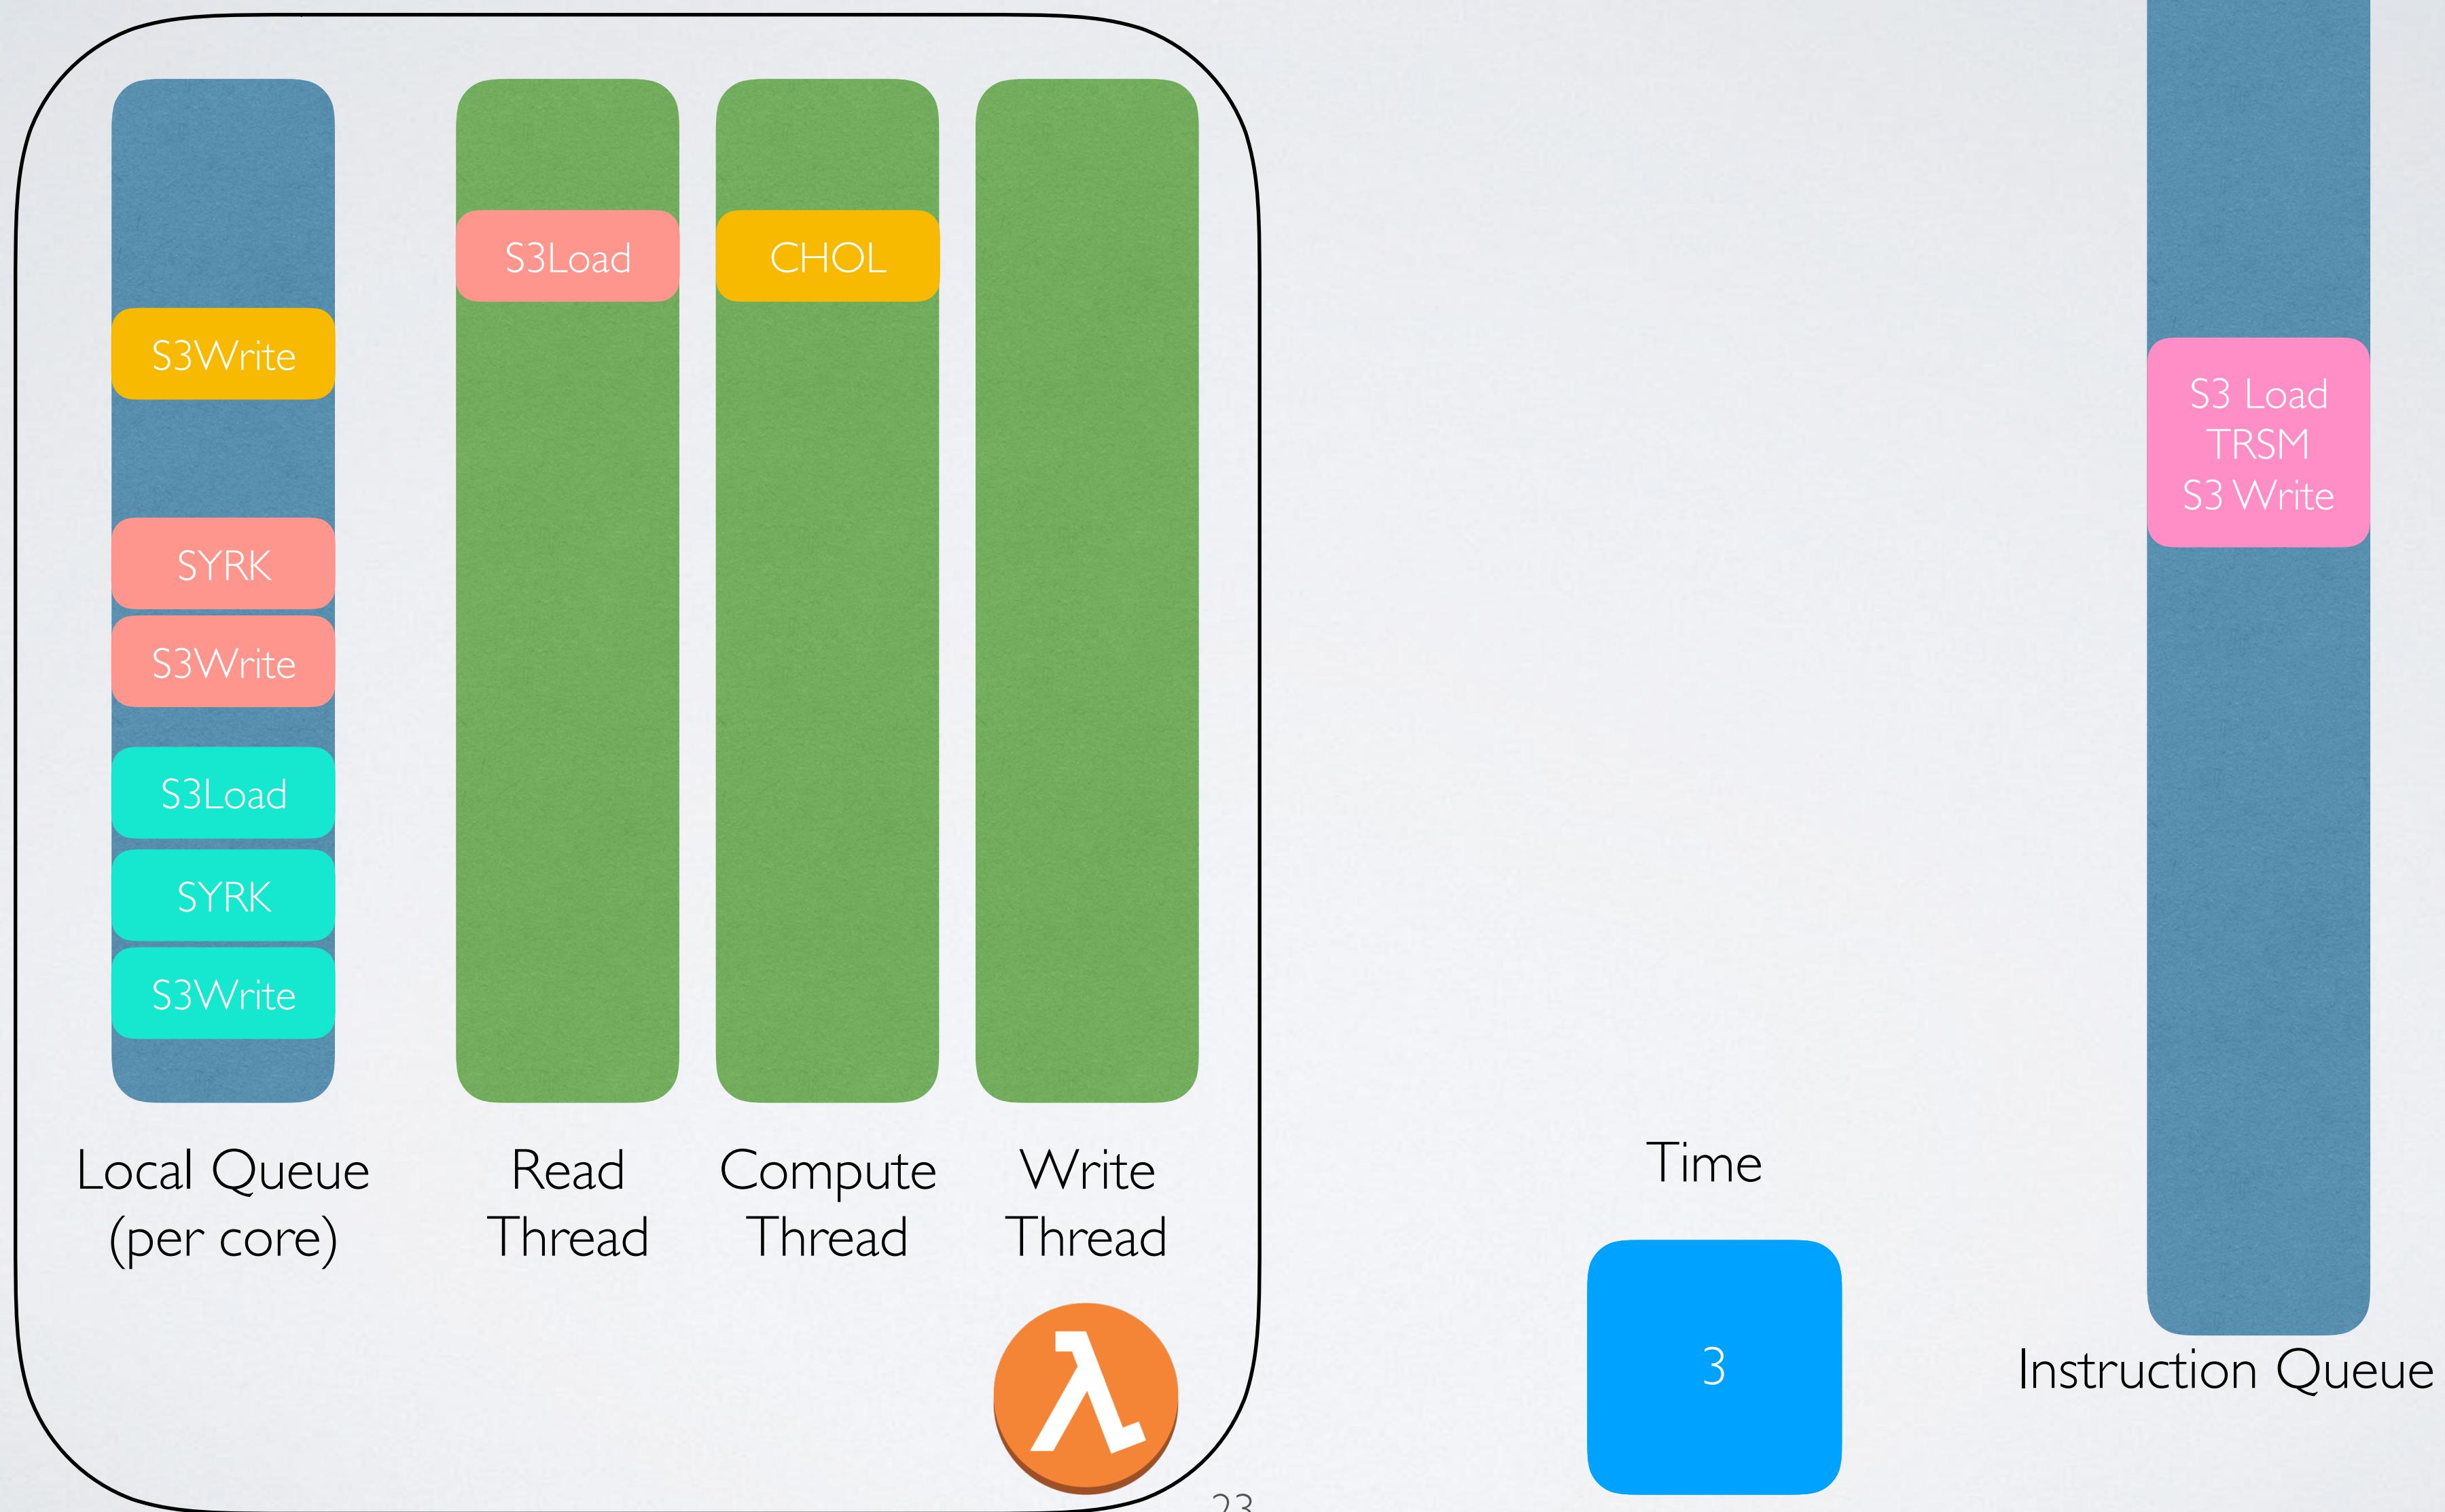
EXECUTION



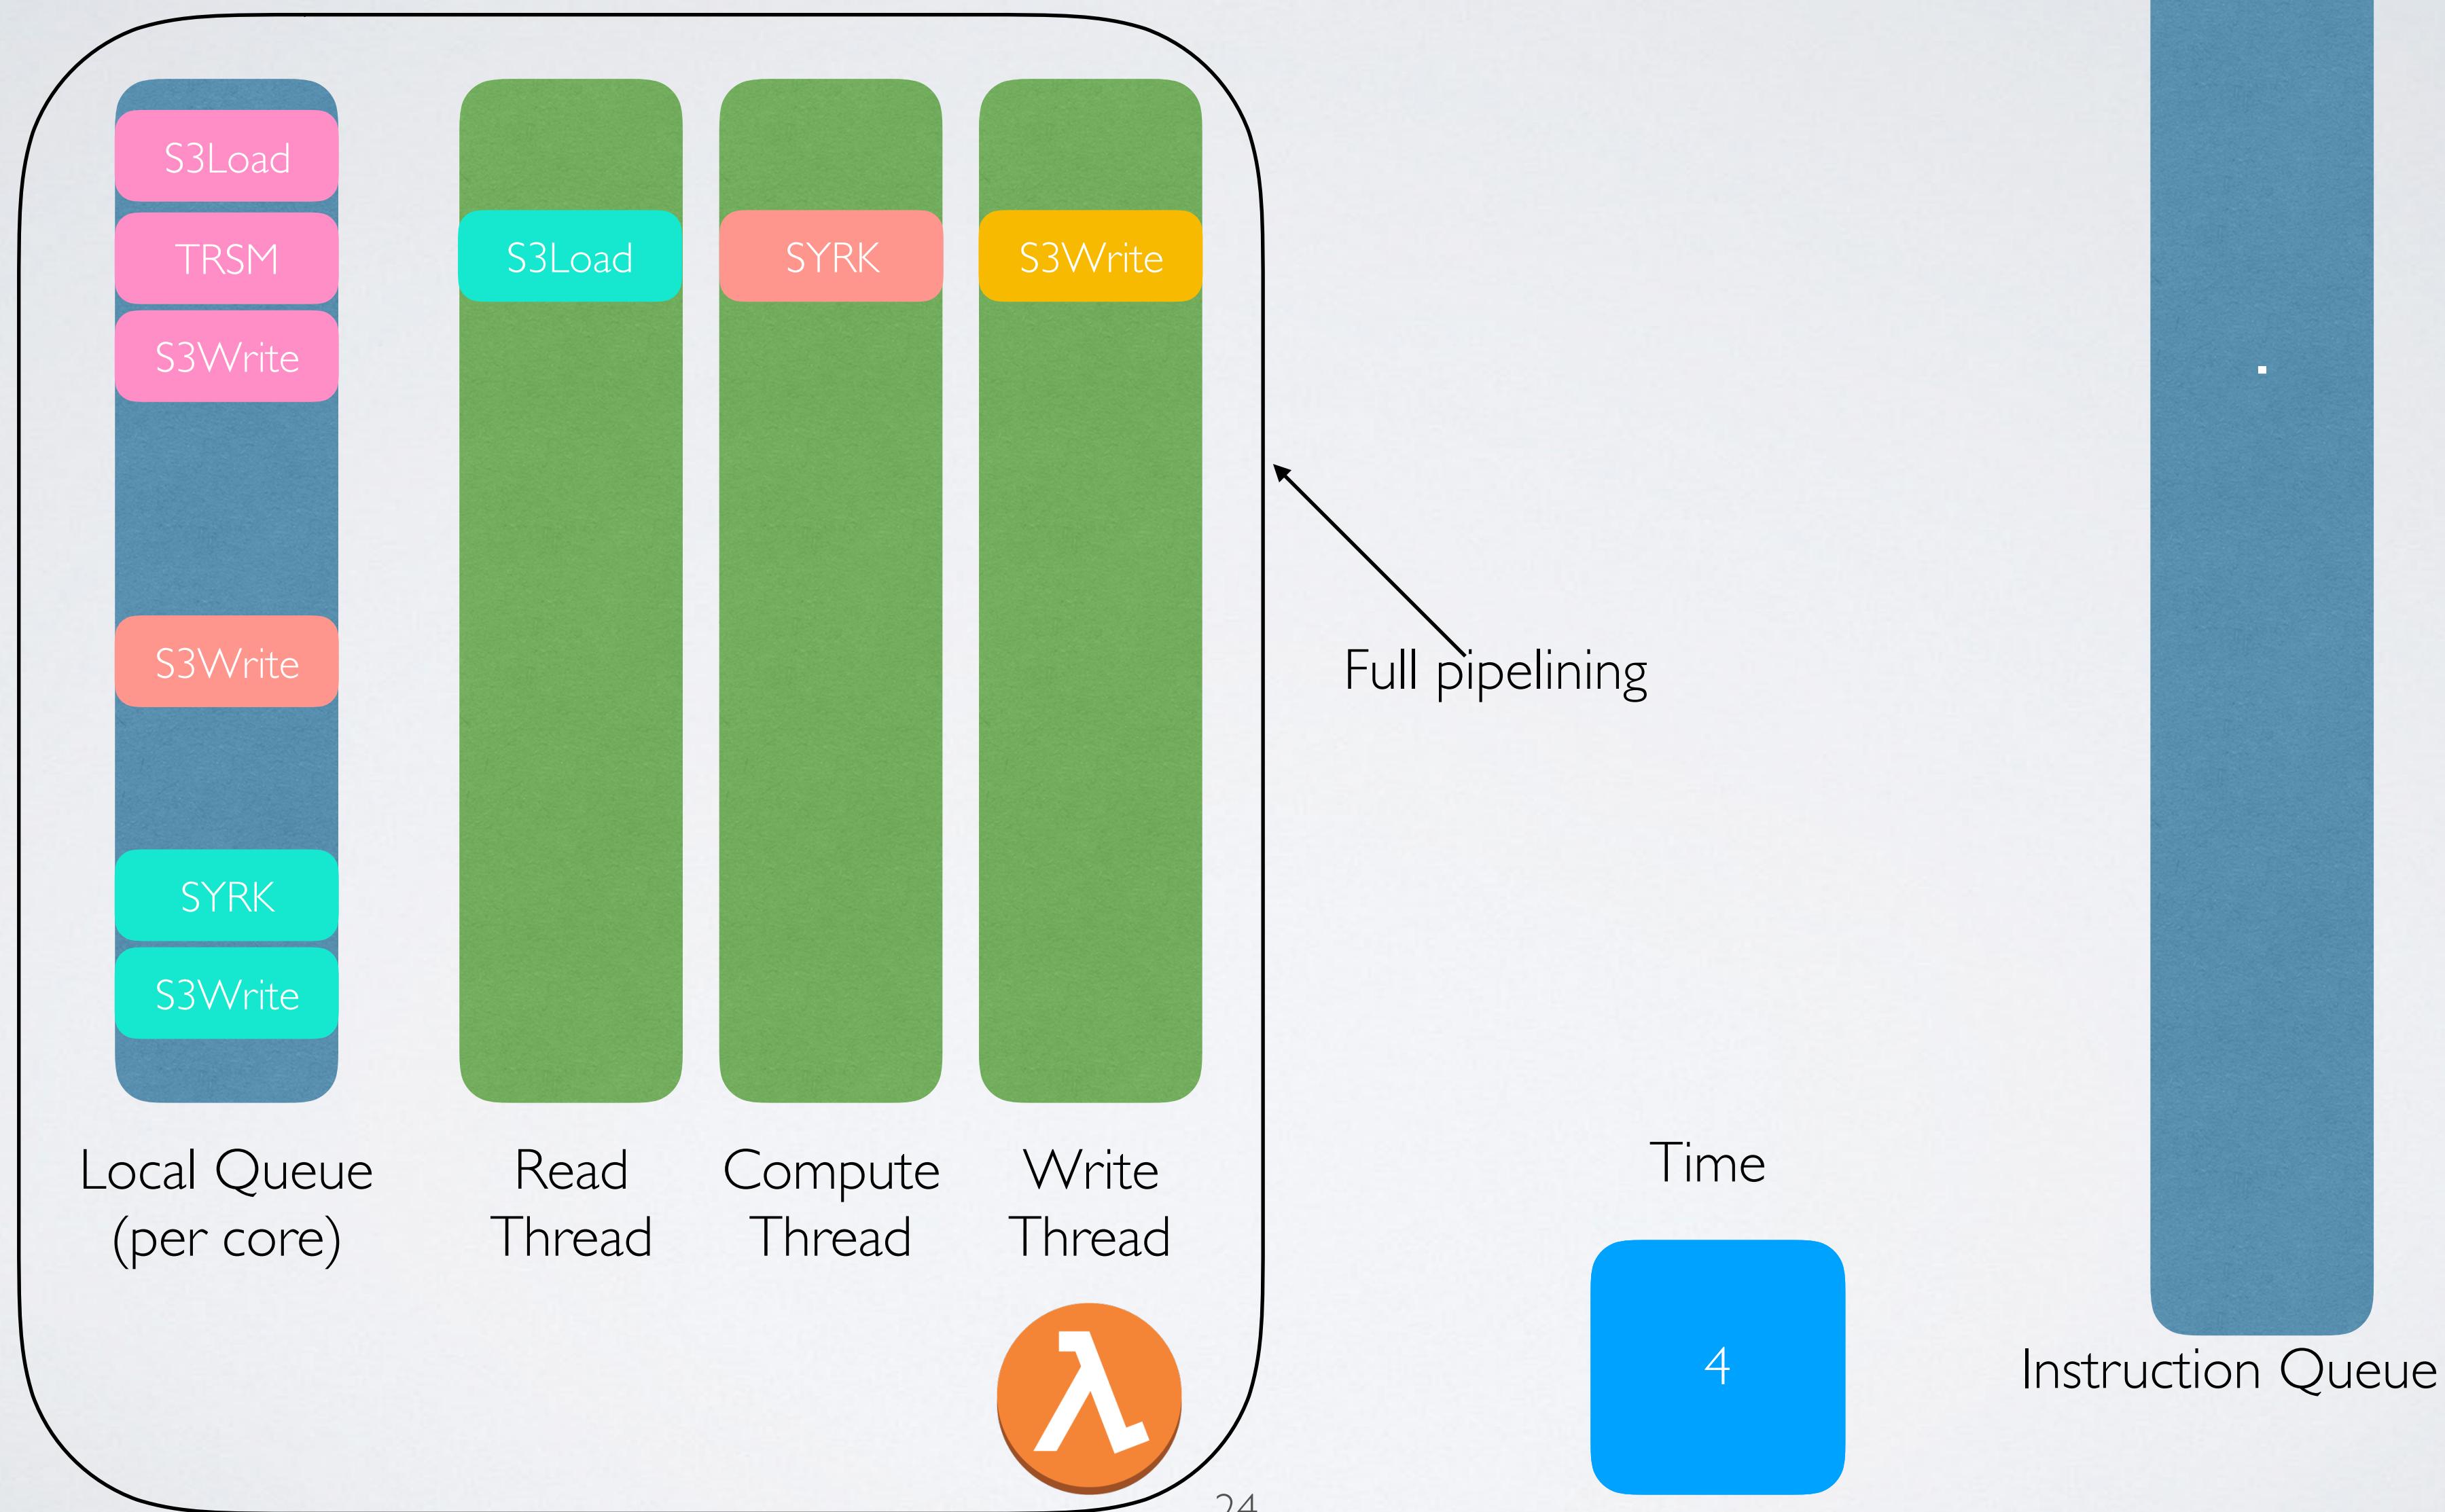
EXECUTION



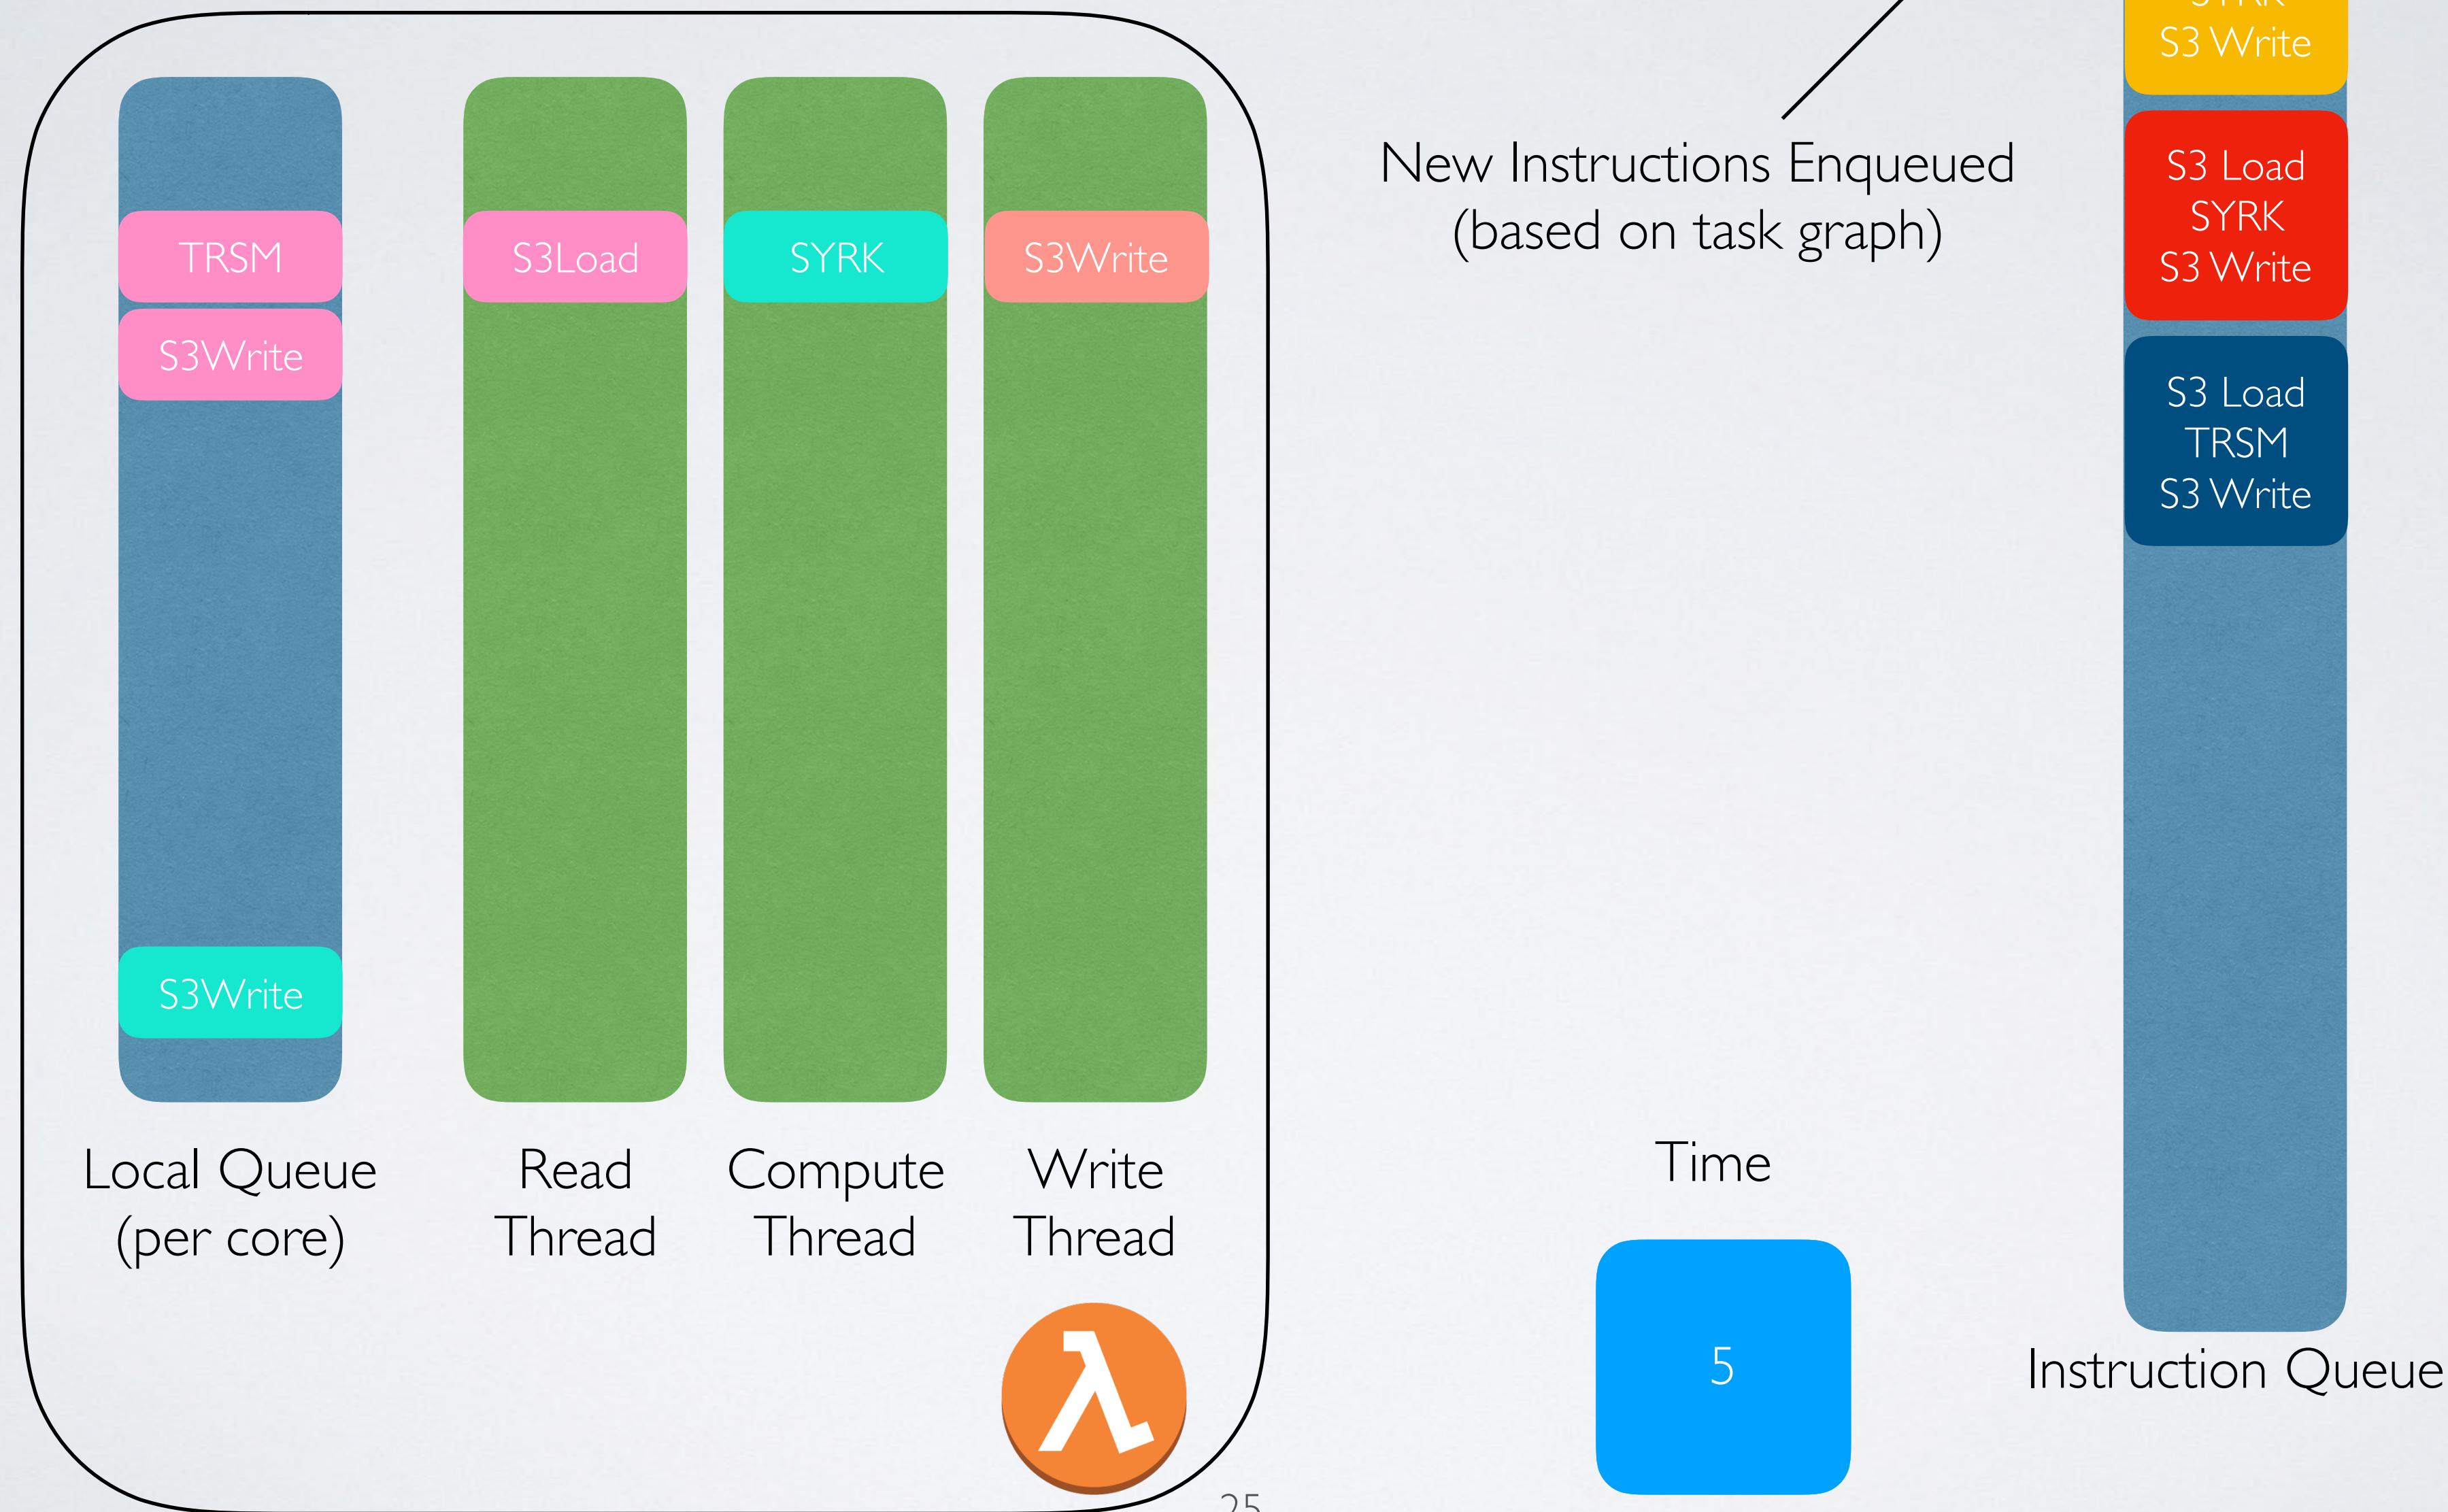
EXECUTION



EXECUTION



EXECUTION



PERFORMANCE

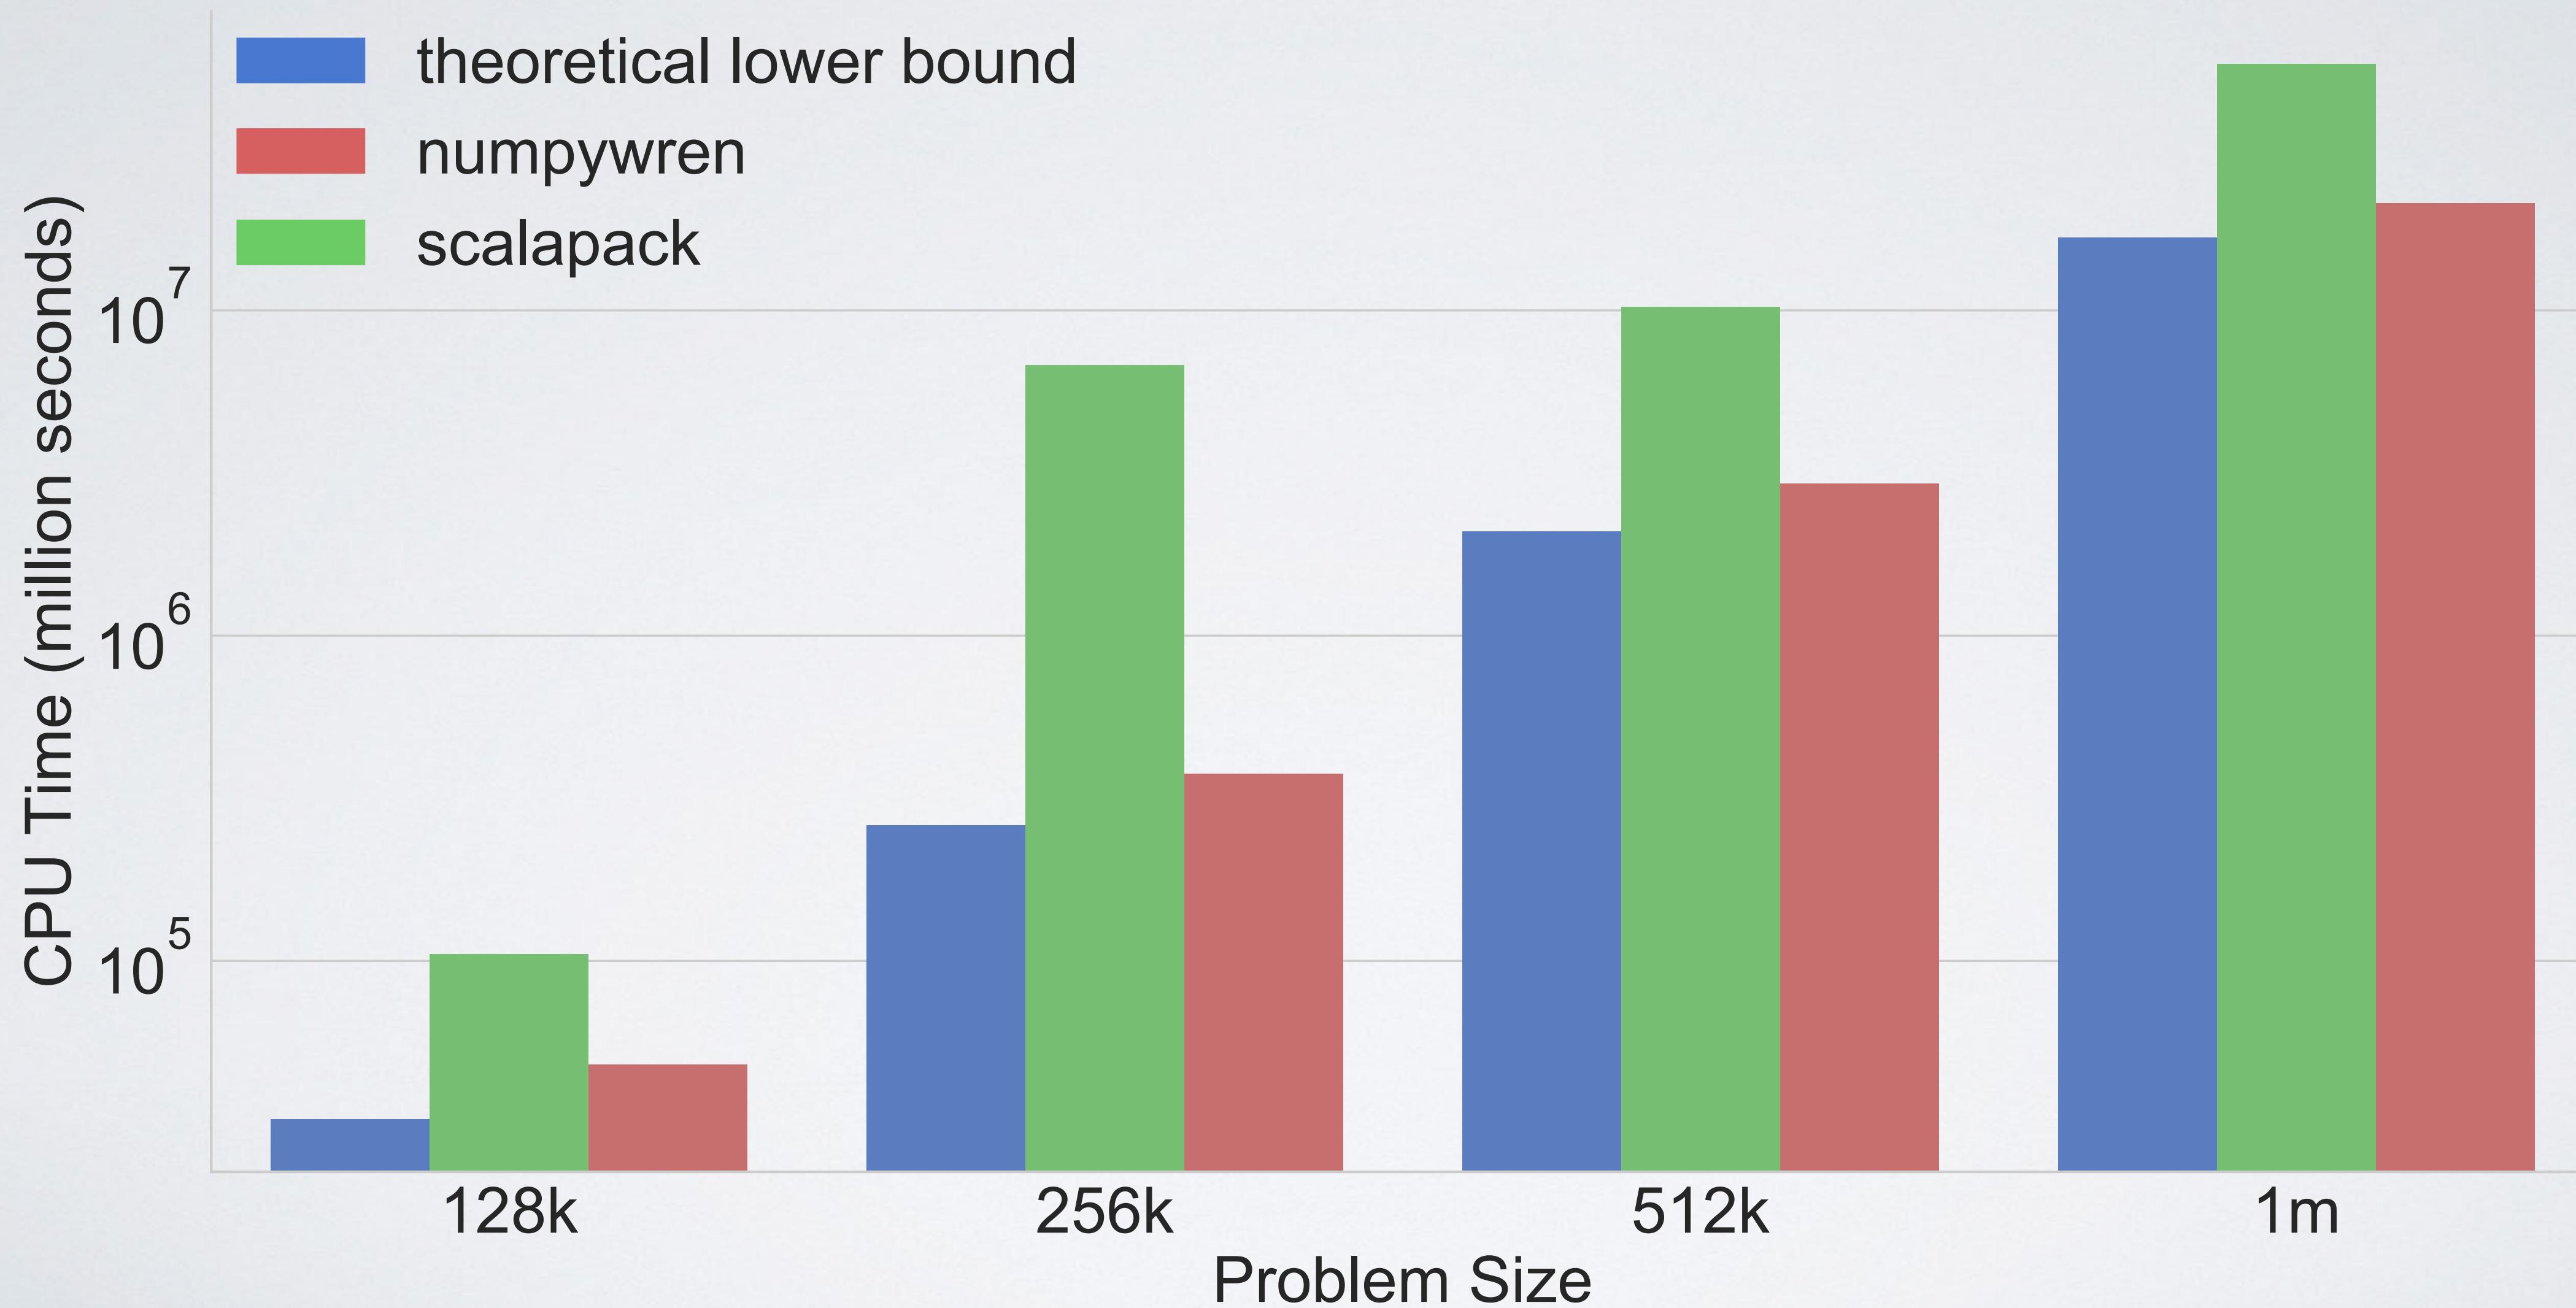
Efficiency

How efficiently did I use
my resources

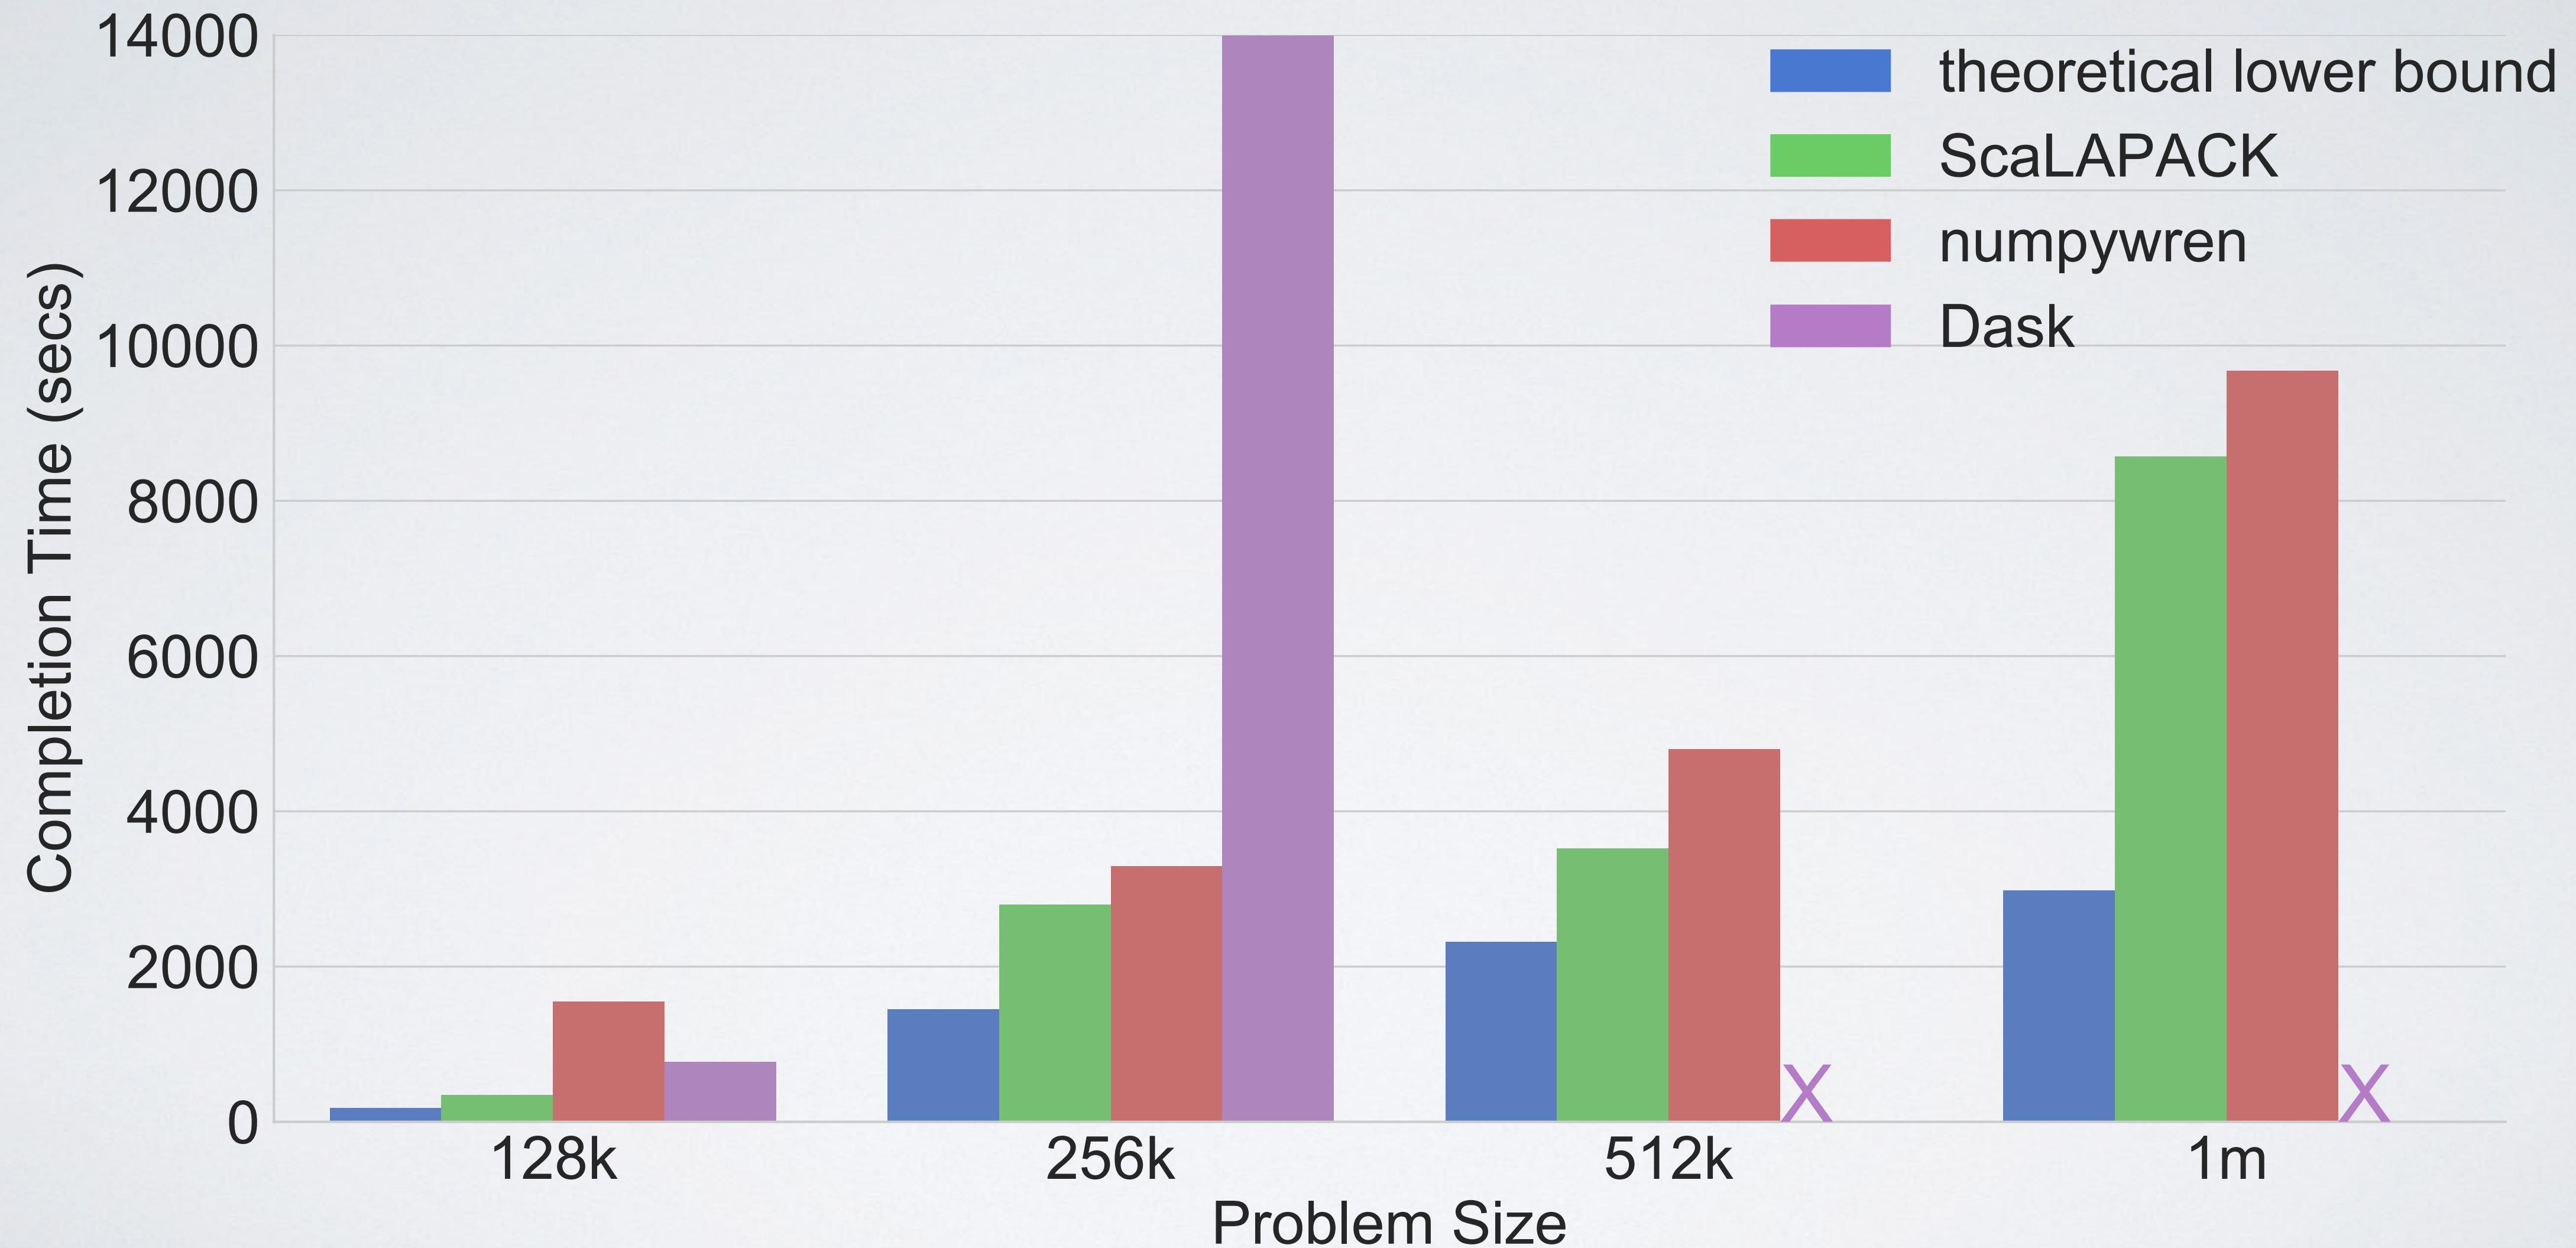
End to end
runtime

How long did it take
to get an answer

TOTAL CORE SECONDS USED



END TO END COMPLETION TIME





NumPyWren

Eric Jonas

jonas@eecs.berkeley.edu
@stochastician



Vaishaal Shankar



Karl Krauth



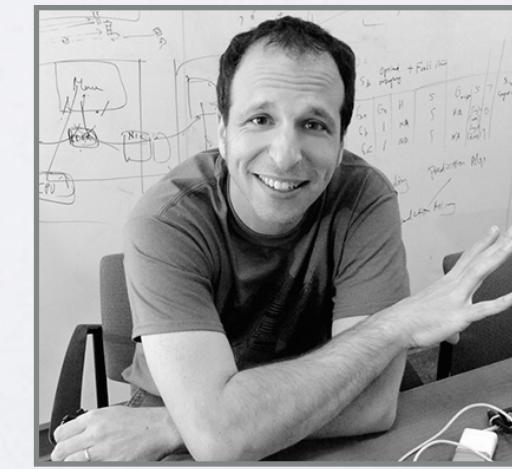
Qifan Pu



Shivaram
Venkataraman



Ion
Stoica



Ben
Recht



Jonathan
Ragan-Kelly



NumPyWren

- Serverless linear algebra is possible, performant, elastic, and easy

Eric Jonas

jonas@eecs.berkeley.edu
@stochastician



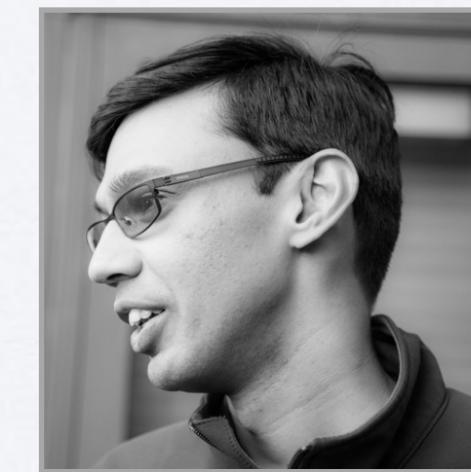
Vaishaal Shankar



Karl Krauth



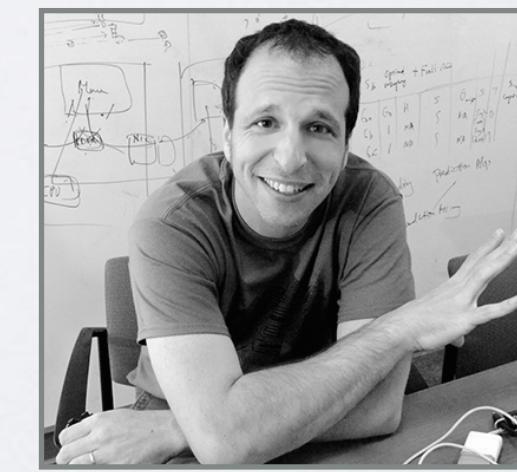
Qifan Pu



Shivaram
Venkataraman



Ion
Stoica



Ben
Recht



Jonathan
Ragan-Kelly



NumPyWren

- Serverless linear algebra is possible, performant, elastic, and easy
- Releasing code this month

Eric Jonas

jonas@eecs.berkeley.edu
@stochastician



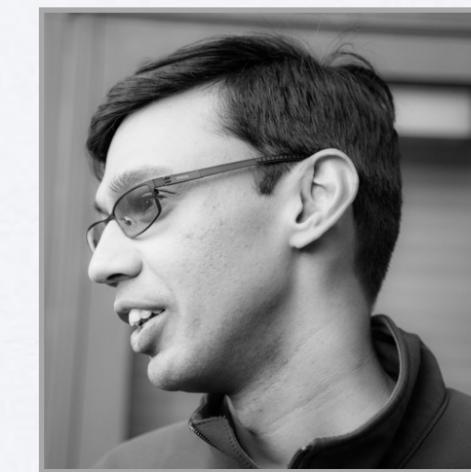
Vaishaal Shankar



Karl Krauth



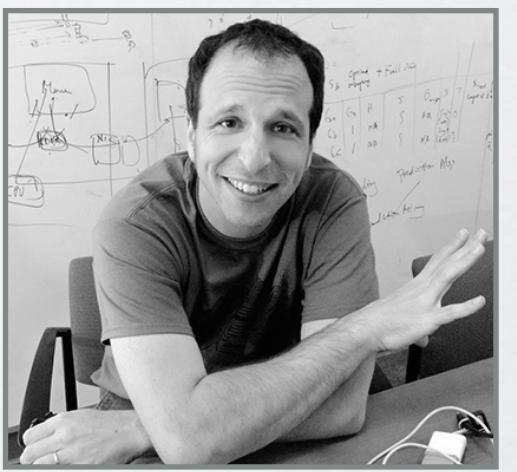
Qifan Pu



Shivaram
Venkataraman



Ion
Stoica



Ben
Recht



Jonathan
Ragan-Kelly

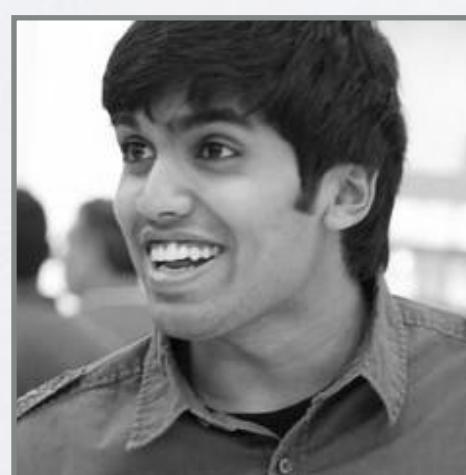


Eric Jonas

jonas@eecs.berkeley.edu
@stochastician

NumPyWren

- Serverless linear algebra is possible, performant, elastic, and easy
- Releasing code this month
- Next steps: Op fusion, straggler mitigation, even higher-level interfaces



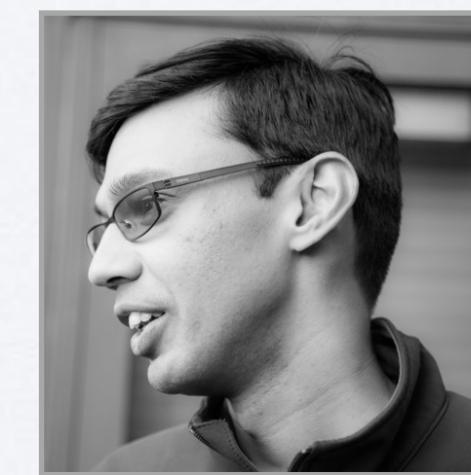
Vaishaal Shankar



Karl Krauth



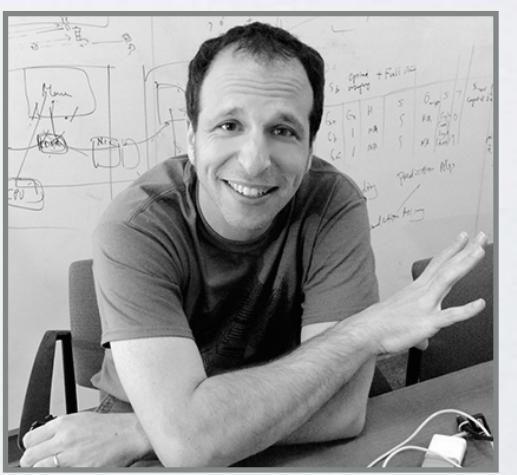
Qifan Pu



Shivaram
Venkataraman



Ion
Stoica



Ben
Recht



Jonathan
Ragan-Kelly

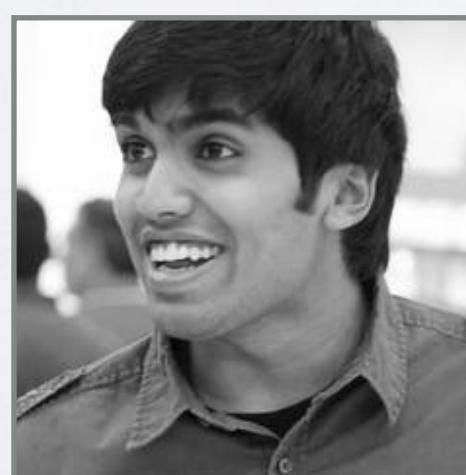


Eric Jonas

jonas@eecs.berkeley.edu
@stochastician

NumPyWren

- Serverless linear algebra is possible, performant, elastic, and easy
- Releasing code this month
- Next steps: Op fusion, straggler mitigation, even higher-level interfaces
- Questions?



Vaishaal Shankar



Karl Krauth



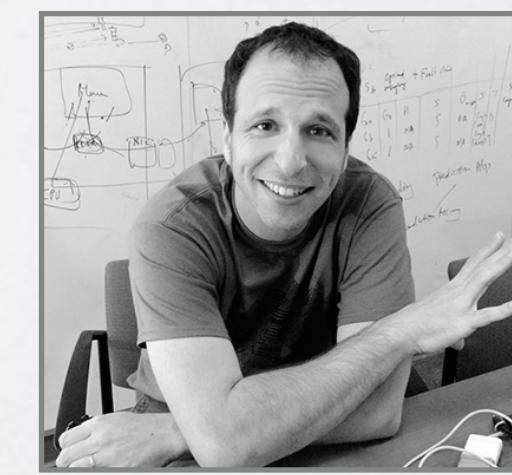
Qifan Pu



Shivaram
Venkataraman



Ion
Stoica



Ben
Recht



Jonathan
Ragan-Kelly

DISCUSSION SLIDE

- What additional services need to be truly elastic to make these sorts of applications possible?
- How much control do we want/need over queues, timing, latency, etc?
- What is the equilibrium price for serverless architectures?
- How can we expand this as a development platform for others