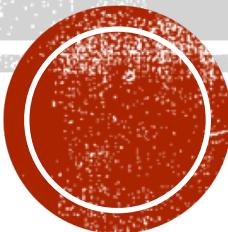


# FAAS ORCHESTRATION OF PARALLEL WORKLOADS

Daniel Barcelona Pons, Pedro García López, Álvaro Ruiz  
Ollobarren, Amanda Gómez-Gómez, Gerard París, Marc  
Sánchez Artigas



# H2020 CLOUDBUTTON: SERVERLESS DATA ANALYTICS

- Large Serverless European Research Project (4.3M)
- The main goal is to create CloudButton: a Serverless Data Analytics Platform. CloudButton will “democratize big data” by overly simplifying the overall life cycle and programming model thanks to serverless technologies
- To demonstrate the impact of the project, we target two settings with large data volumes: bioinformatics (genomics, metabolomics) and geospatial data (LiDAR, satellital)



European Molecular  
Biology Laboratory



UNIVERSITAT ROVIRA i VIRGILI

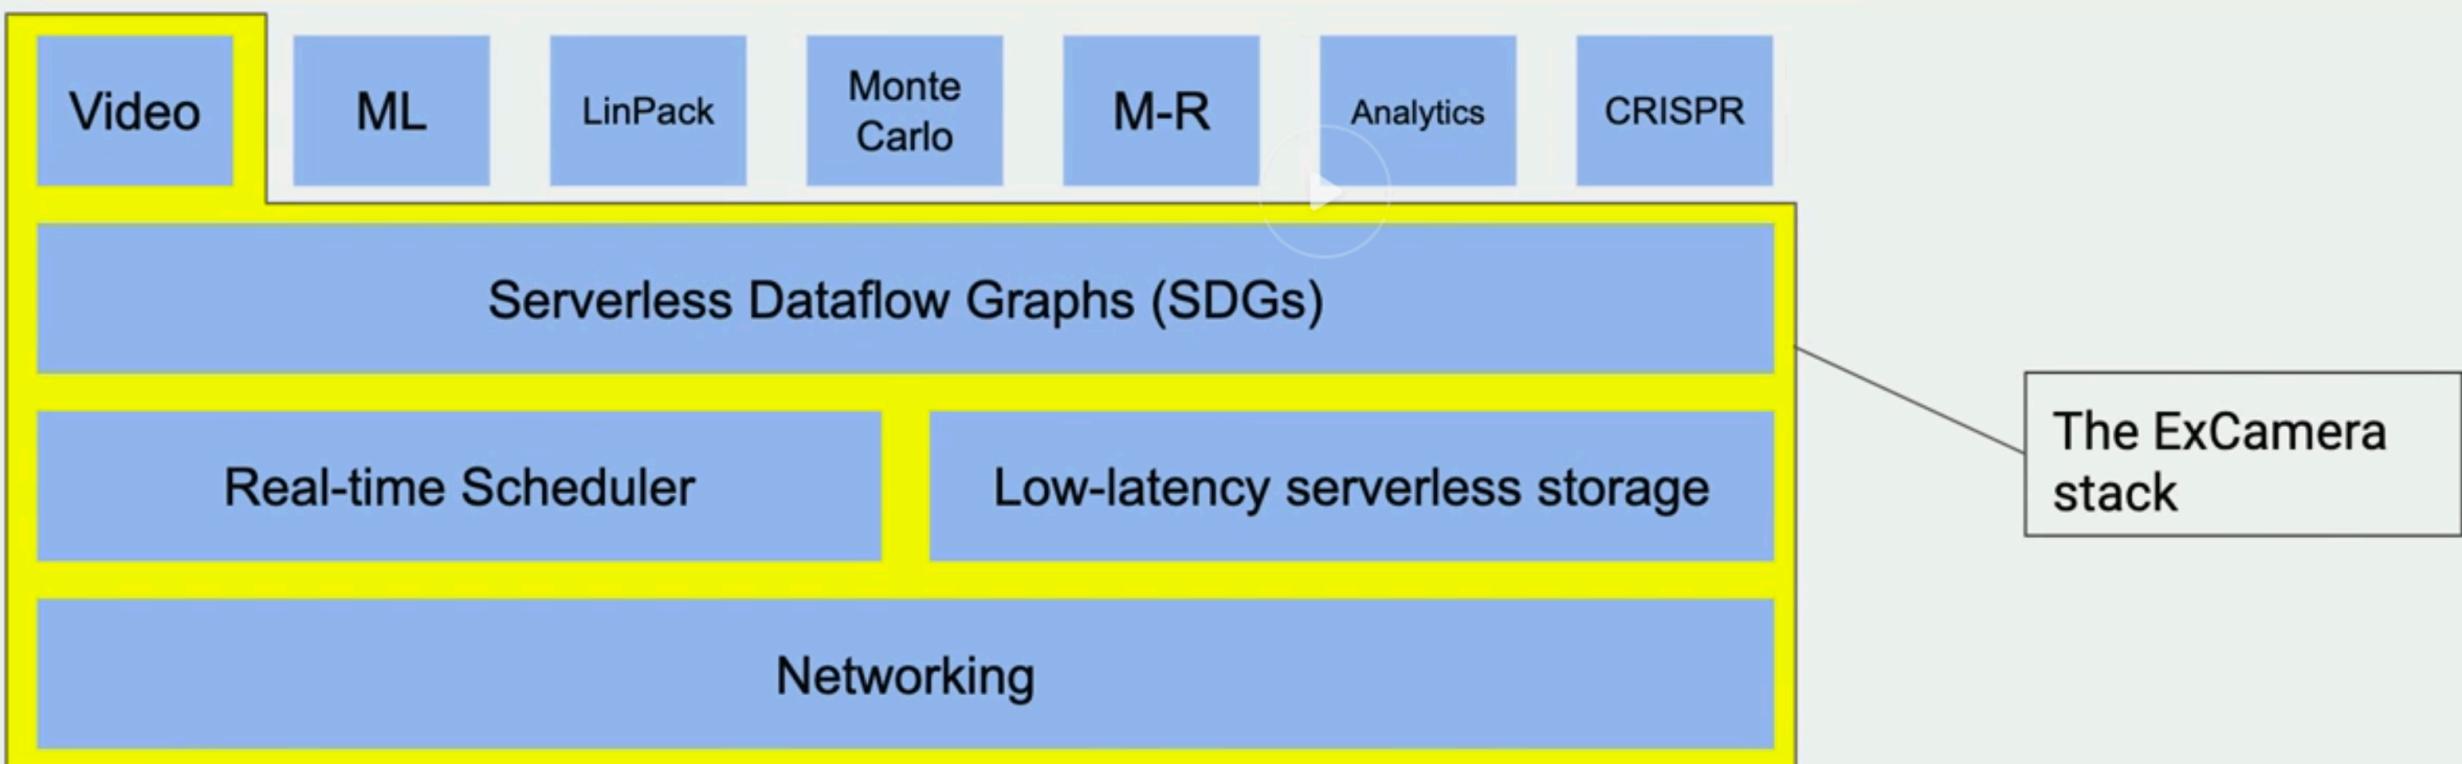


Imperial College  
London



# The Serverless Supercomputer Stack

*Applications: Big data, analytics, ML/AI, simulations, r/t video transcoding, ...*



A CLOUD GURU

Cloud Vendors: Functions, APIs, Managed Service Portfolios

## SERVERMIX APPLICATIONS

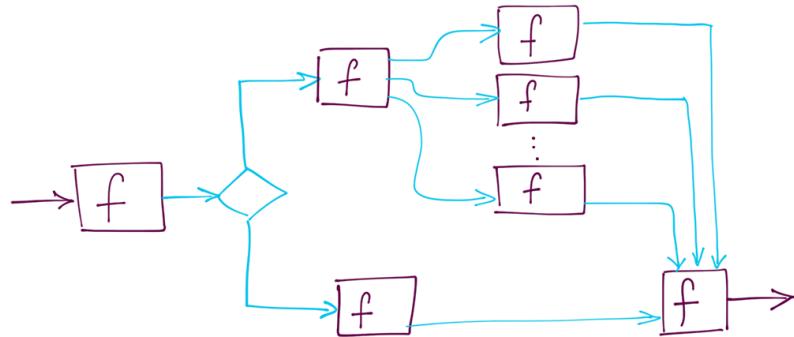
Systems	Components	
	<i>Serverful</i>	<i>Serverless</i>
Locus	Scheduler, Redis	Lambda Functions, S3
PyWren	Scheduler	Lambda Functions, S3
IBM PyWren	Scheduler	IBM Cloud Functions, COS, RabbitMQ
ExCamera	Coordinator, Rendezvous	Lambda Functions, S3
Flint	Scheduler	Lambda Functions, S3, SQS
NumPyWren	Provisioner	Lambda Functions, S3
Cirrus	Scheduler, Parameter Servers	Lambda Functions, S3

# SERVERMIX, SERVERLESS, SERVERFUL

**ServerMix: Tradeoffs and Challenges of Serverless Data Analytics**

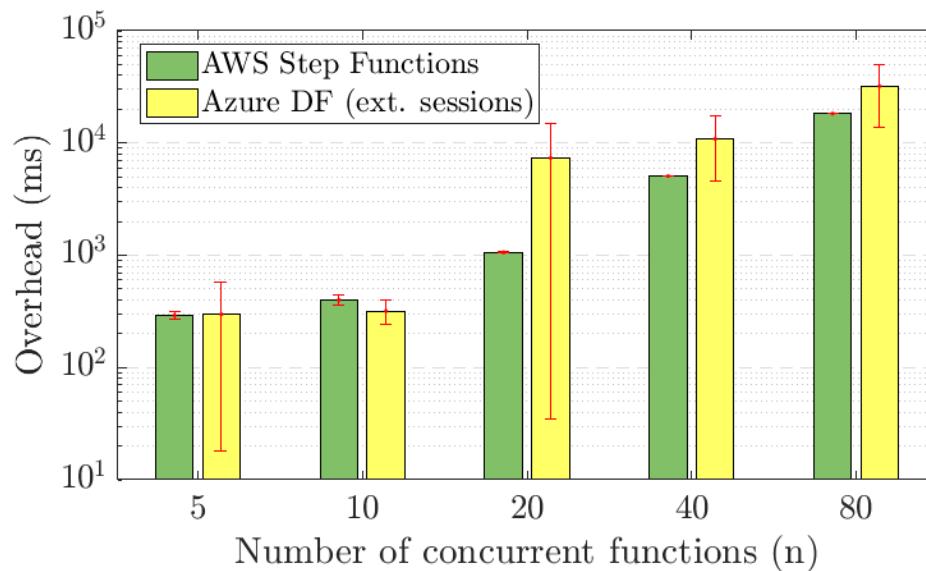


# COMPARISON (WOSC4)

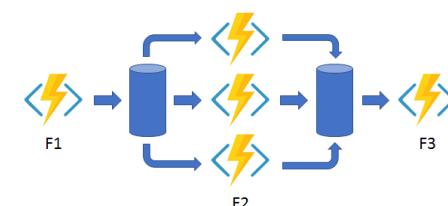


>User Functions

Workflow Framework



AWS Step Functions

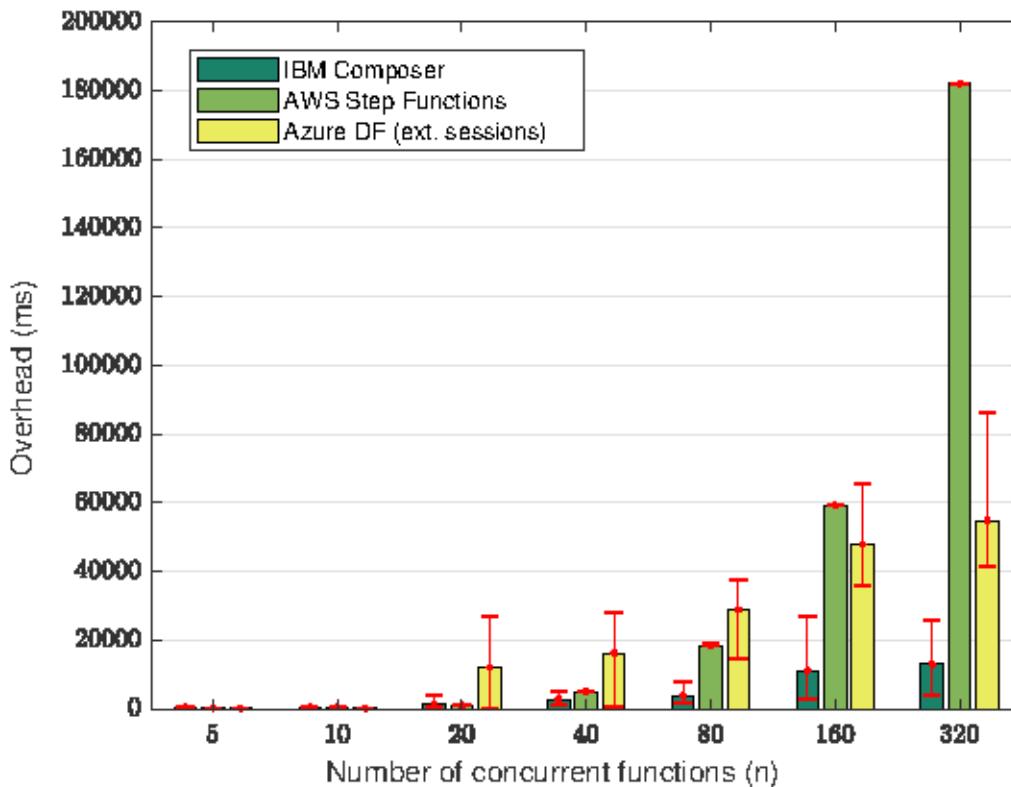


Azure Durable Functions



IBM Function Composer

# COMPARISON (WOSC5)

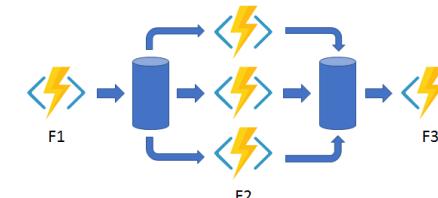


Is IBM Composer the answer now ?



State Machine

AWS Step Functions



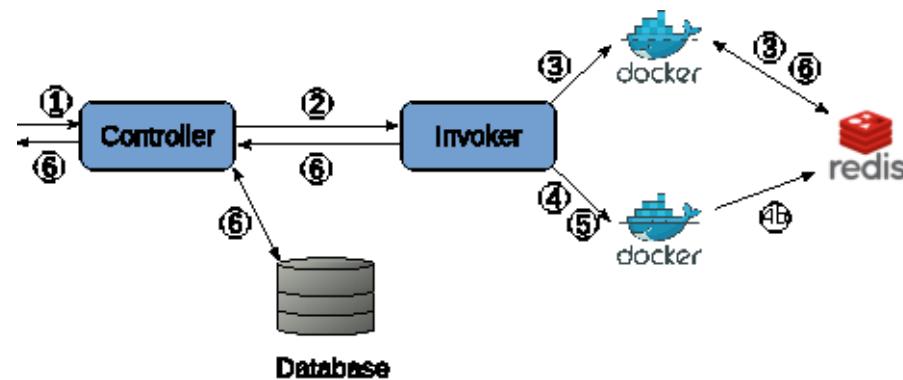
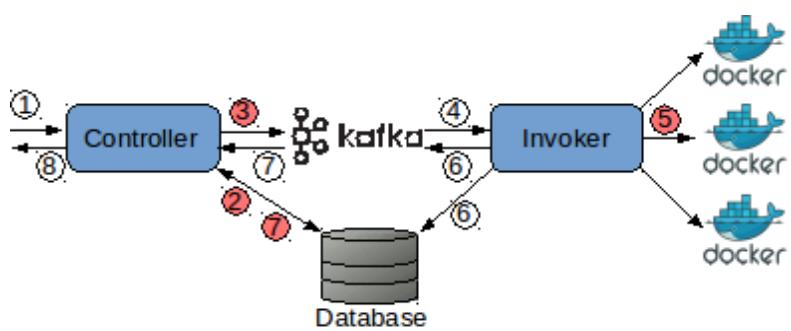
Event Sourcing



AST  
continuations



# OPENWHISK AND COMPOSER

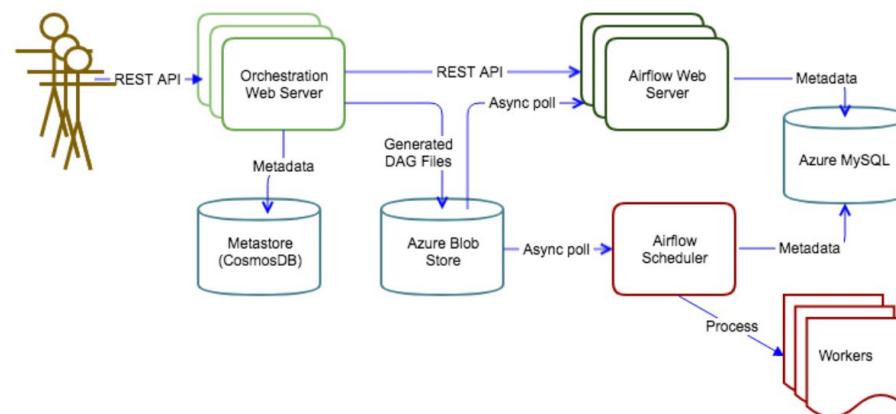
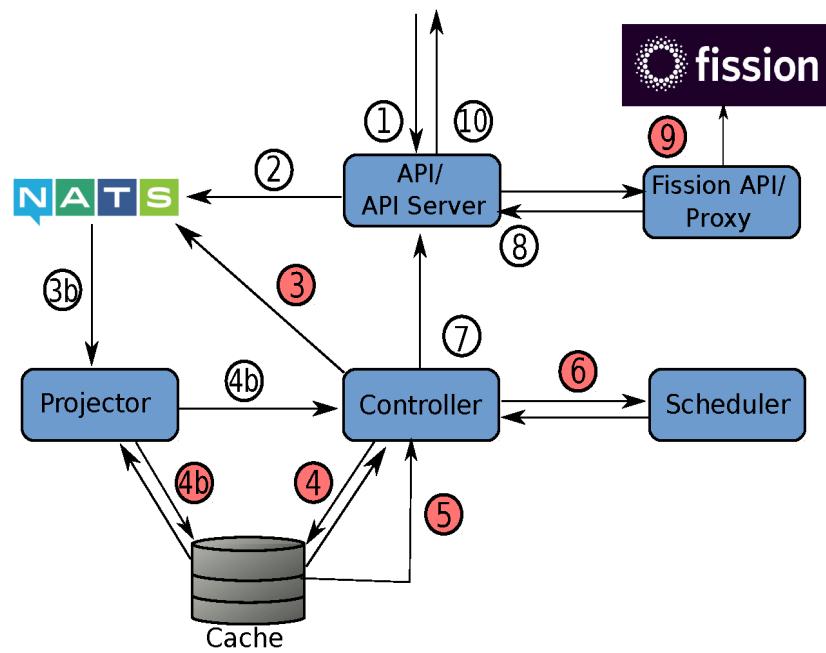


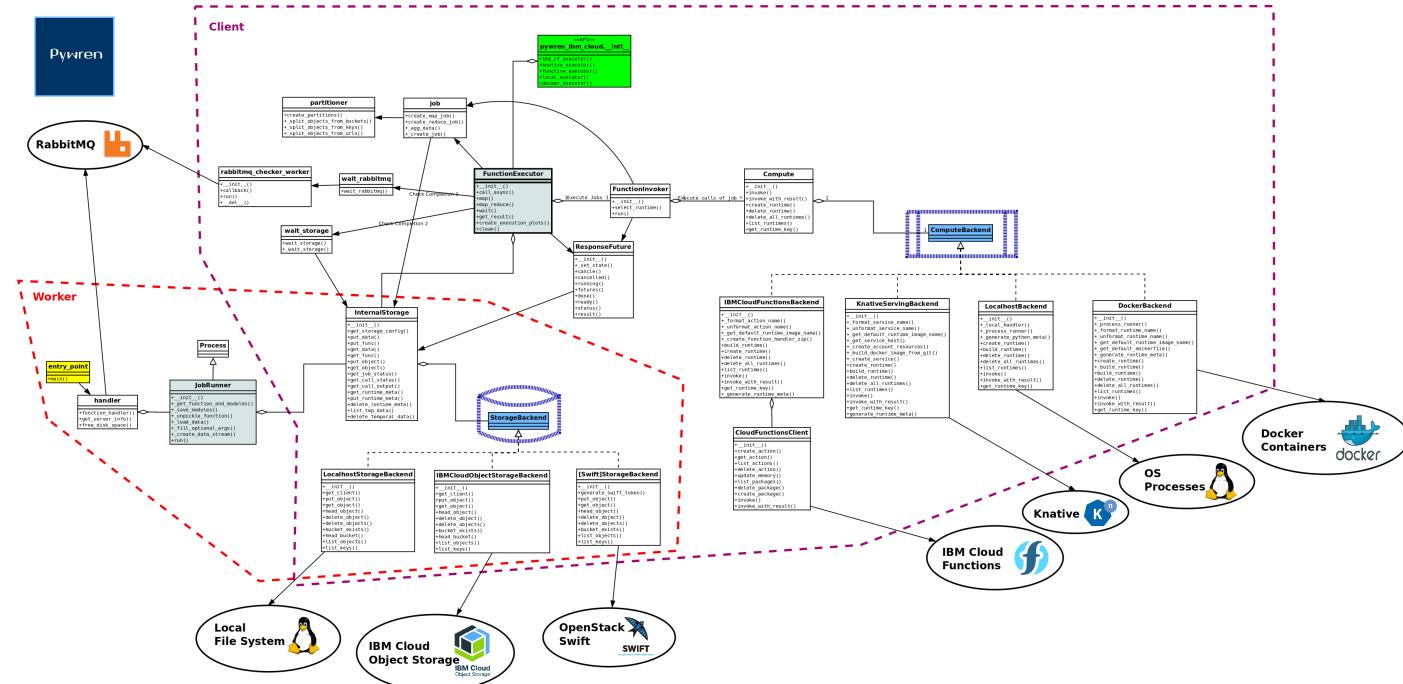
# COMPOSER PROBLEMS

- Coupled Design with OpenWhisk Conductor Actions
- Resource inefficiencies when launching parallel functions
  - High Variability in experiments
- Difficult billing of conductor and secondary actions
- Designed for short-running workflows
- Limited fault-tolerance
- It relies on a user-provided **serverful** Redis Service for fork-join aggregation



# FISSION WORKFLOWS, AIRFLOW





# SOLUTION NOW ? USE YOUR EXTERNAL SCHEDULER LIKE IBM PYWREN



# CONCLUSIONS

- **Innovation in the Cloud providers is needed !!!**
- Existing FaaS Orchestration systems are not currently designed for supporting parallel concurrent jobs
- Messages to Cloud providers:
  - Cloud providers must offer the required tools to orchestrate parallel Data Analytics tasks
  - Cloud providers must offer event-based building blocks to enable the construction of third-party schedulers (AWS Destinations)
  - Multi-tenancy, cost, and resource efficiency will be of paramount importance for these services

