


Scaling a Variant Calling Genomics Pipeline with FaaS

9th International Workshop on Serverless Computing

Part of ACM/IFIP International Middleware Conference, December 11–15, 2023 in DAMSLab, Department of Arts, University of Bologna, Italy



UNIVERSITAT
ROVIRA I VIRGILI



Aitor Arjona^a

Arnau Gabriel-Atienza^a

Sara Lanuza-Orna^a

Xavier Roca-Canals^a

Ayman Bourramouss^a

Tyler K. Chafin^b

Lucio Marcello^b

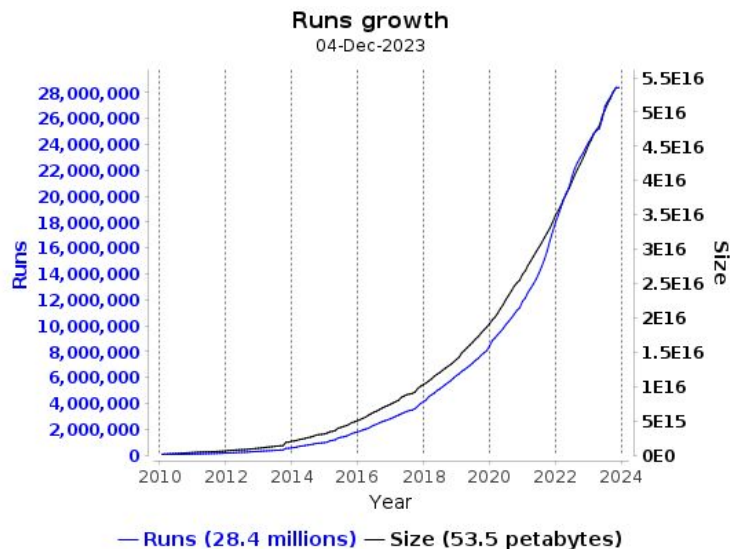
Paolo Ribeca^b

Pedro García-López^a

*a. Universitat Rovira i Virgili,
Tarragona, Spain*


*b. Biomathematics and Statistics Scotland,
Edinburgh, United Kingdom*

Genomics Workloads



- Genomics is a **compute- and data intensive-** task.
- Exponential growth in data size and complexity.
- Biomedical institutions with HPC **struggle to keep up.**


Cloud IaaS for Genomics



The Cloud elasticity is key for scaling genomics workloads using short-term resources.



Cloud IaaS for Genomics



The Cloud elasticity is key for scaling genomics workloads using short-term resources.

- **Complexity for bioinformatic users**
- ◆ Capacity/VM size for processing X GB of data?
 - ◆ Auto-scaling?
 - ◆ Hidden costs?

Configuring, deploying and scaling genomics workloads is challenging for bioinformatics users.

Going serverless

Serverless (FaaS)



1. **Pay only for resources used** at millisecond granularity, scale down to zero when not used
2. **Instant scalability** (~200 ms cold start, thousands of parallel functions)
3. **Completely managed:** Scaling, security...

- Why serverless for genomic pipelines?
- ◆ No servers to manage!
 - ◆ *Less friction to the Cloud for less experienced (bioinformatic) users*
 - ◆ Allows to **massively and effortlessly scale highly-parallel genomics workloads.**


Serverless Genomic Variant Calling

Objective: Adapt a **single-node HPC** variant calling genomics application to **serverless** in order to **scale in parallelism, process larger datasets** and **decrease runtime**.

Serverless Genomic Variant Calling

Objective: Adapt a **single-node HPC** variant calling genomics application to **serverless** in order to **scale in parallelism**, **process larger datasets** and **decrease runtime**.

- **Variant Calling:** detect differences (mutations, variants) in a sampled genome compared to a reference genome.



Distributing a Variant Calling pipeline

Sampled Genome (FASTQ)

```
@SEQ_ID1
CGGTAGCCAGCTGCGTTCAGTATGGAAGATTGATT
+
+&&-&%$%$%$%$%)33&0$&%$' '* '%$%$%$#+-5
@SEQ_ID2
TTCAGTTTATGGGTGCGGGTGTATGTGACAAGAAAG+
"###"$%$%$%) , +) +&' ( , "###"$&0$&%$' '*&0
@SEQ_ID3
GCATGACCATACCGTGACAAGAAAGTCACCGCCCGTC
+
!' '* ((( (**++) ) %%%) ) '%$%$%$ ( '%#
@SEQ_ID4
CGGTAGCCAGCTGCGTTCAGTATGGAAGATTGATT
+
+&&-&%$%$%$%$%)33&0$&%$' '* '%$%$%$#+-5
@SEQ_ID5
TTCAGTTTATGGGTGCGGGTGTATGTGACAAGAAAG+
"###"$%$%$%) , +) +&' ( , "###"$&0$&%$' '*&0
```

Reference Genome (FASTQ)

```
>SEQUENCE_1
MTEITAA MVKELRESTGAGMMDCKNALSETNGDGLEKK
TEDFAAEVAAQLGLEKKTEDFAAEVAAQLFDKAVQLLR
EMGQFYVMDDKKTVEQVIAEKEKEFGGKIKIV
>SEQUENCE_2
SATVSEINSETDFVAKNDQFIALTKDTTAHIQSNLSQS
VEELHSSTINGVKFEEYLSQIATIGENLVRRFATLK
AGANGVVGYYIHTNGRGVVIAAACDSAE
>SEQUENCE_3
MTEITAA MVKELRESTGAGMMDCKNALSETNGDGLEKK
TEDFAAEVAAQLGLEKKTEDFAAEVAAQLFDKAVQLLR
EMGQFYVMDDKKTVEQVIAEKEKEFGGKIKIV
>SEQUENCE_4
SATVSEINSETDFVAKNDQFIALTKDTTAHIQSNLSQS
VEELHSSTINGVKFEEYLSQIATIGENLVRRFATLK
AGANGVVGYYIHTNGRGVVIAAACDSAE
>SEQUENCE_5
MTEITAA MVKELRESTGAGMMDCKNALSETNGDGLEKK
TEDFAAEVAAQLGLEKKTEDFAAEVAAQLFDKAVQLLR
EMGQFYVMDDKKTVEQVIAEKEKEFGGKIKIV
```

- Process: **String similarity search** - Compare all sample genome to all reference genome

Distributing a Variant Calling pipeline

Sampled Genome (FASTQ)

```
@SEQ_ID1
CGGTAGCCAGCTGCGTTTCAGTATGGAAGATTGATT
+
+&&-&%$%$%$%$#)33&0$&%$' '* '%$#%$%#+-5
@SEQ_ID2
TTCAGTTTATGGGTGCGGGTGTATGTGACAAGAAAG+
"###" "$%$#) %, +) +&' (, "###" "&0$&%$' '*&0
```

```
@SEQ_ID3
GCATGACCATACCGTGACAAGAAAGTCACGCCCGTC
+
! '* (((***)) %%%+) (%%%) '%%##$( '%#
@SEQ_ID4
CGGTAGCCAGCTGCGTTTCAGTATGGAAGATTGATT
+
+&&-&%$%$%$%$#)33&0$&%$' '* '%$#%$%#+-5
```

Reference Genome (FASTQ)



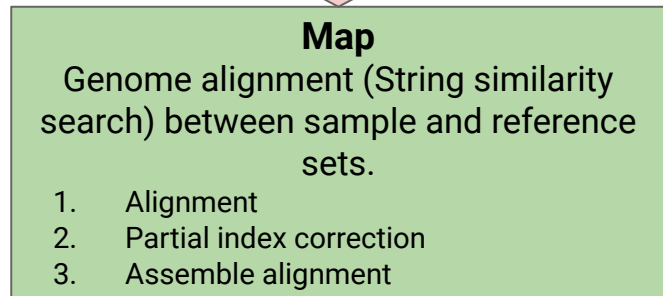
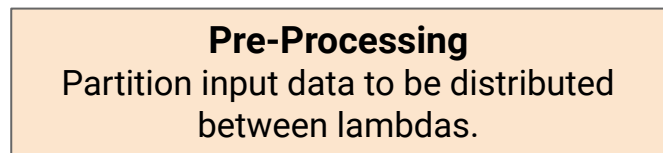
```
>SEQUENCE_1
MTEITAAMVKELRESTGAGMMDCKNALSETNGDGLEK
>SEQUENCE_2
SATVSEINSETDFVAKNDQFIALTKDTTAHIQSNLSQS
VEELHSSTINGVKFEEY
```

```
>SEQUENCE_3
MTEITAAMVKELRESTGAGMMDCKNALSETNGDGLEKK
TEDFAAEVAAQLGLEK
>SEQUENCE_4
MTEITAAMVKELRESTGAGMMDCKNALSETNGDGLEK
```


```
>SEQUENCE_5
MTEITAAMVKELRESTGAGMMDCKNALSETNGDGLEKK
TEDFAAEVAAQLGLEKTTEDFAAEVAAQLFDKAVQLLR
EMGQFYVMDDKKTVEQVIAEKEKEFGGKIKIV
```

- Process: **String similarity search** - Compare all sample genome to all reference genome
- How to distribute the workload?
 - Partition the dataset and perform an all to all comparison
 - Cartesian product → **Lots of parallel tasks**
- Partitioning the dataset implies:
 - Partial correction process
 - Merging all partial results to produce the final output


Serverless Variant Calling Pipeline Architecture



How to distribute the workload?
→ Partition the dataset

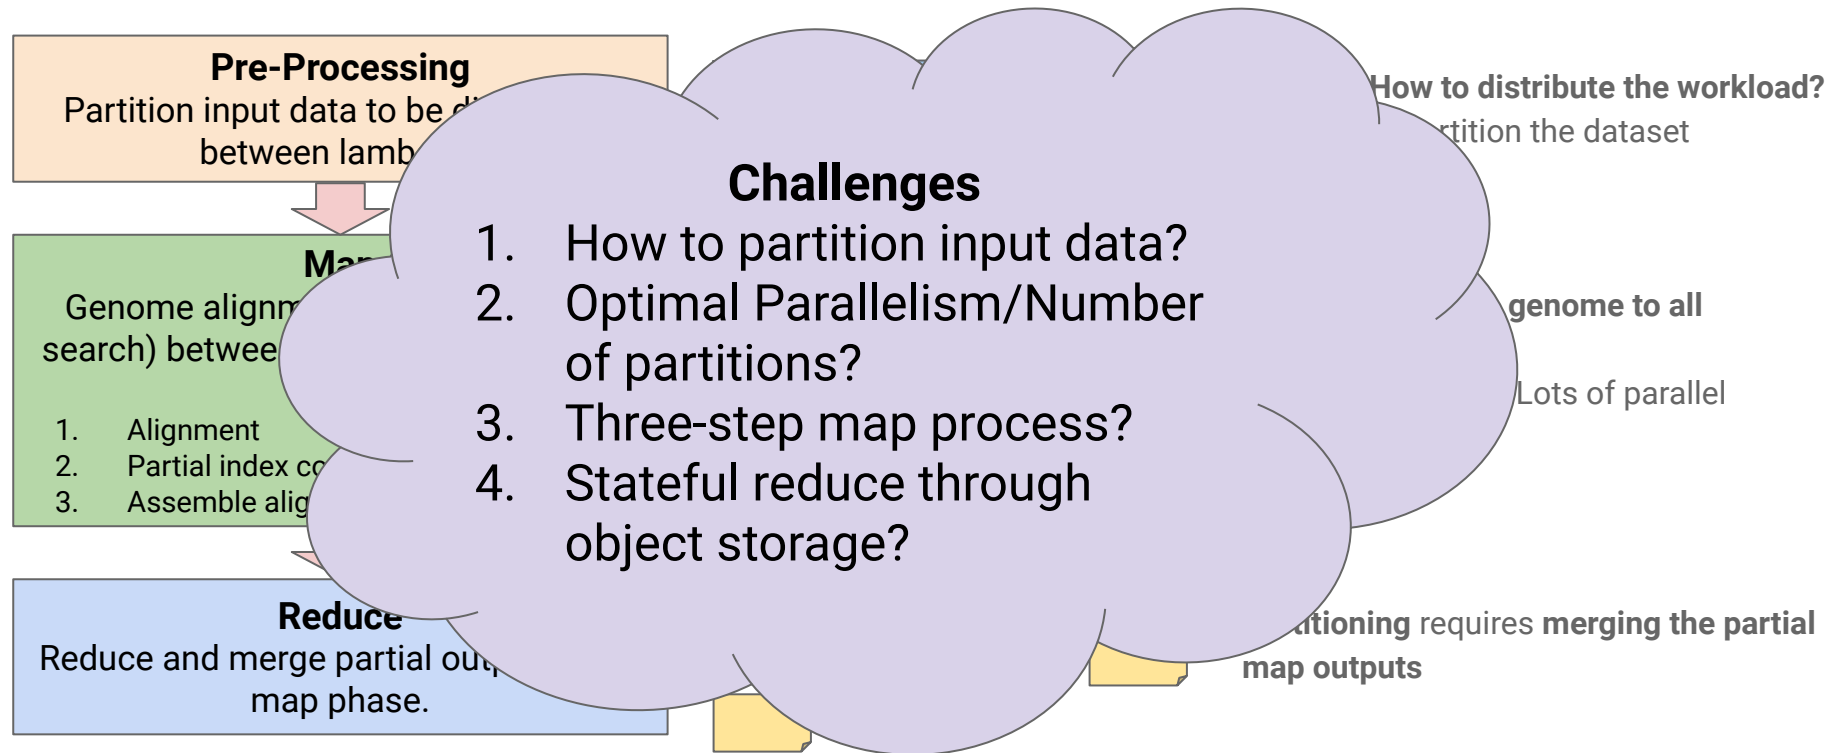


Compare all sample genome to all reference genome
→ Cartesian product: Lots of parallel tasks




Partitioning requires merging the partial map outputs

Serverless Variant Calling Pipeline Architecture



Challenge 1

Input Data Partitioning



- **Objective:** **partition** input data
- Partial reads with HTTP GET byte-range requests

- Arbitrary **byte-ranges** breaks the genome file
- We require more **metadata** for each partition (sequence identifier and offset)

- Indexing to **locate and identify** each sequence for **any arbitrary byte-range**
- Lookup index for partition metadata


Challenge 2

Data dependencies from synchronous HPC code

- **Data dependencies**→**Functions calling functions**: Stop the process, call another task, synchronize (wait), get result, then resume.
- **No preemption** in serverless →Blocked tasks **occupy concurrency slot**
- Can provoke deadlocks and limit scalability
- Blocking code must be **split** into many **asynchronous tasks** that can be **scheduled independently**
- **Data dependencies** must be passed through **object storage**

Challenge 3


Optimal degree of parallelism



- How many tasks to launch?

- **More tasks** → More parallelism → **Less scalable** (concurrency limit)
- **More tasks** → Larger data chunks → Less efficient

- Leverage **intra-function parallelism** to launch **less tasks** without sacrificing **performance**




Evaluation

HPC vs Serverless

Stage	HPC	Serverless
Genome Indexing	0 min 14.20 s	0 min 9.81 s
Alignment	0 min 14.20 s	0 min 48.10 s
Index correction	-	0 min 7.63 s
Generate mpileup	51 min 15.79 s	1 min 6.55 s
Index shuffle	-	0 min 10.73 s
Variant Calling	54 min 5.04s s	0 min 27.82 s
Total	106 min 8.21 s	2 min 50.64 s

Large scale execution



Insights and conclusions

1. **Optimizing performance-cost** with **multi-CPU** functions requires **balancing intra- and inter-function parallelism**
2. Effective **data partitioning** is vital to exploit the massive parallelism offered by FaaS.
3. **Asynchronous and non-blocking** code is strictly necessary to guarantee scalability and to **avoid the concurrency limits** of serverless platforms.
4. Although **object storage** offers scalability for intermediate workflow data storage, it suffers from **slow performance and elevated costs**.
5. **Serverless data processing services** for object storage, like S3 SELECT, provide a nice **synergy** with serverless FaaS workloads.

Thank you!
Any questions?

You can find me at aitor.arjona@urv.cat

Challenge 4

Stateful data movements

- Reduce stage → Stateful data movements through object storage

- **I/O time** from Lambda to S3 is **expensive**
- We want to **simplify** the pipeline

- Delegate **shuffle logic** to S3 SELECT
- **S3 SELECT** allows to define simple SQL queries over semi-structured data
- **Cheaper, more simple, and less error-prone** than doing **ad-hoc shuffling**