



Serverless Confidential Containers: Challenges and Opportunities

Carlos Segarra

(w/ Tobin Feldman-Fitzthum and Daniele Buono)

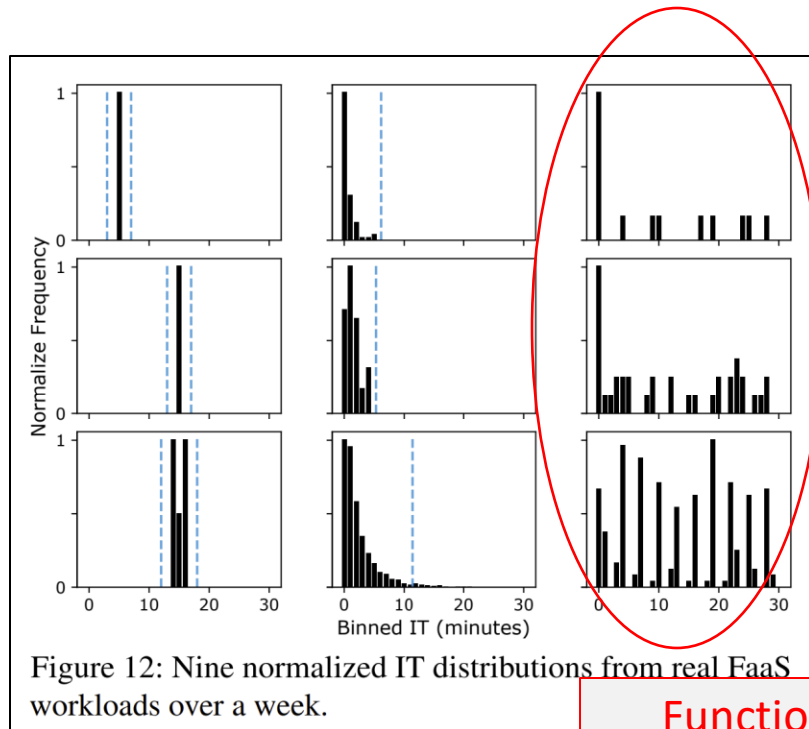
Large-Scale Data & Systems (LSDS) Group - Imperial College London

Visiting IBM TJ Watson (Sep'23 – Nov'23)


<https://carlossegarra.com>
<cs1620@ic.ac.uk>



Introduction: Characterizing Serverless Functions



Functions are bursty!



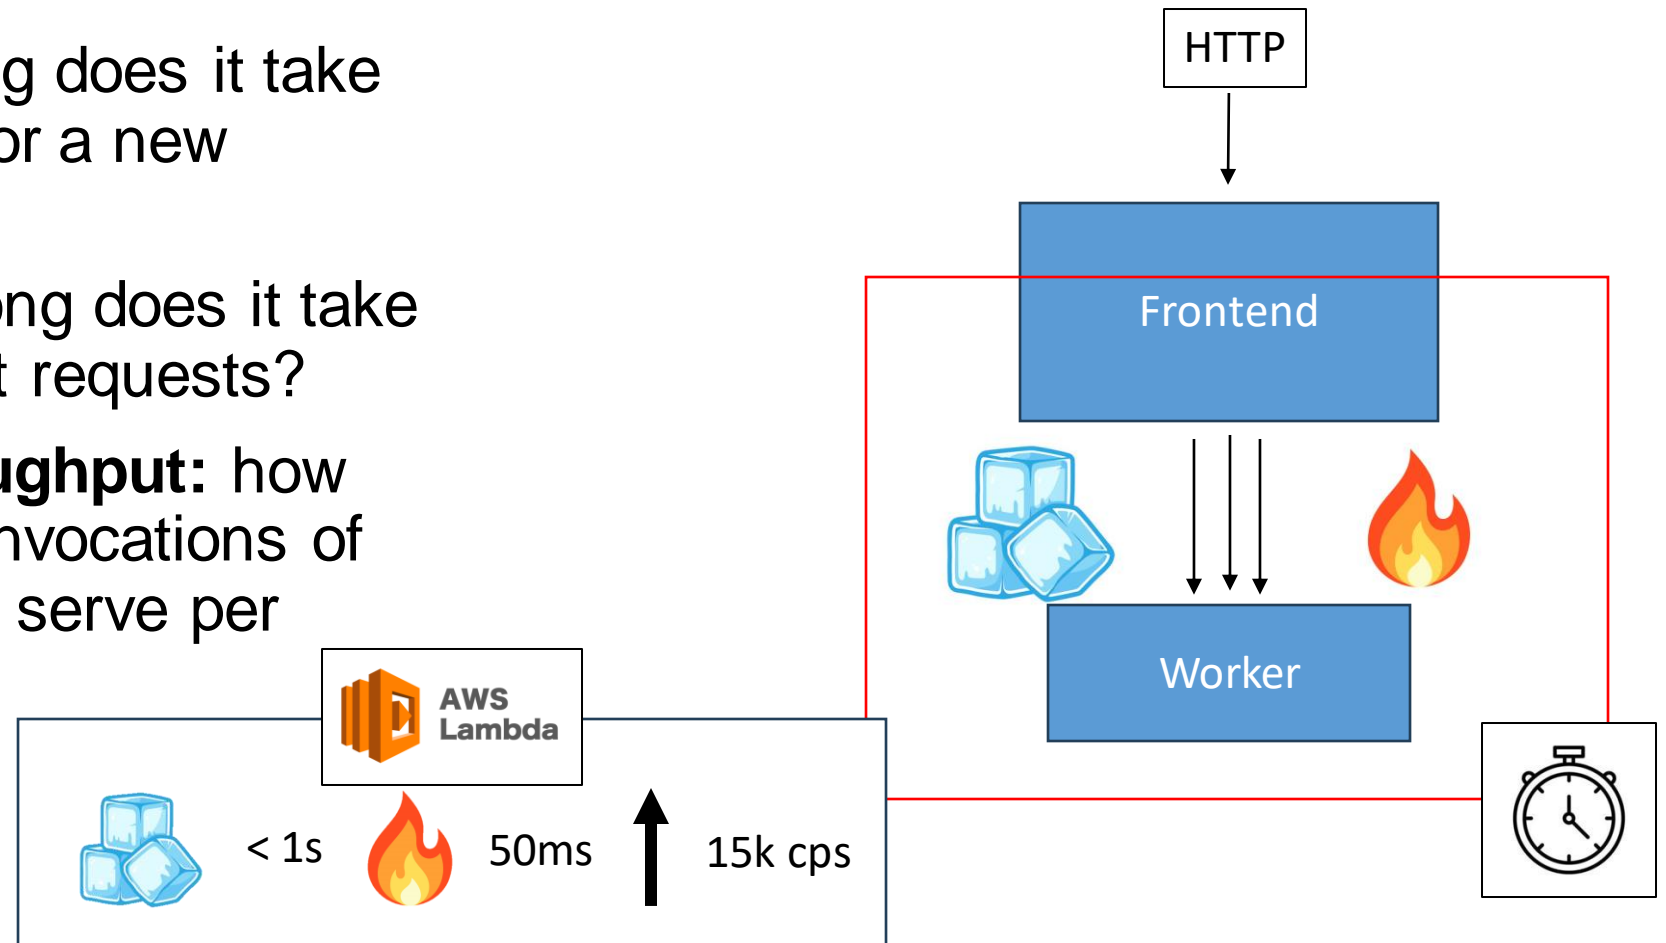
[ATC'20] Serverless in the Wild: Characterizing and Optimizing the Serverless Workload at a Large Cloud Provider

Introduction: Problems in Serverless

Cold-Start: how long does it take to serve a request for a new function?

Warm-Start: how long does it take to serve subsequent requests?

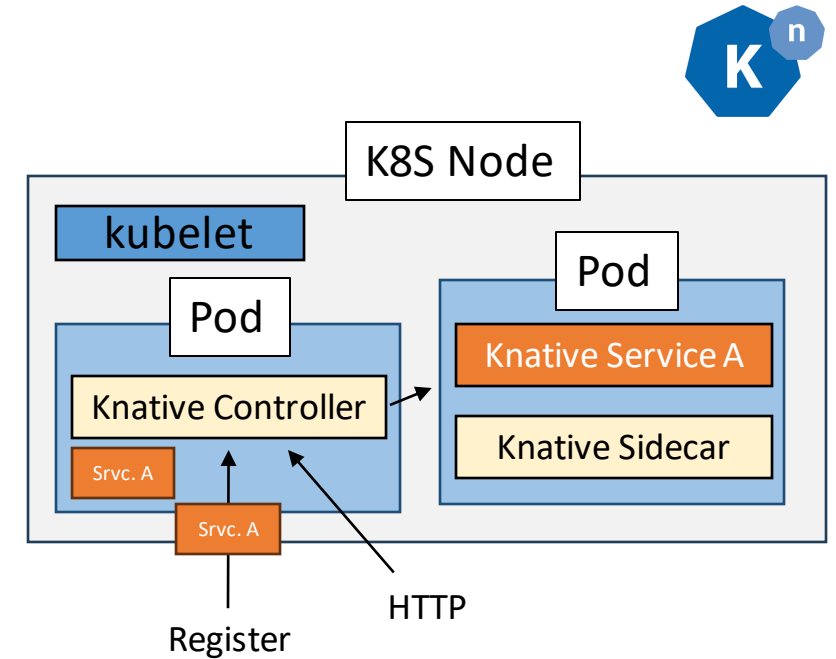
Instantiation Throughput: how many (concurrent) invocations of this function can we serve per second?



[ATC'23] On-demand Container Loading in AWS Lambda

Introduction: More Problems in Serverless!

Inter-function isolation is fine, but not enough!

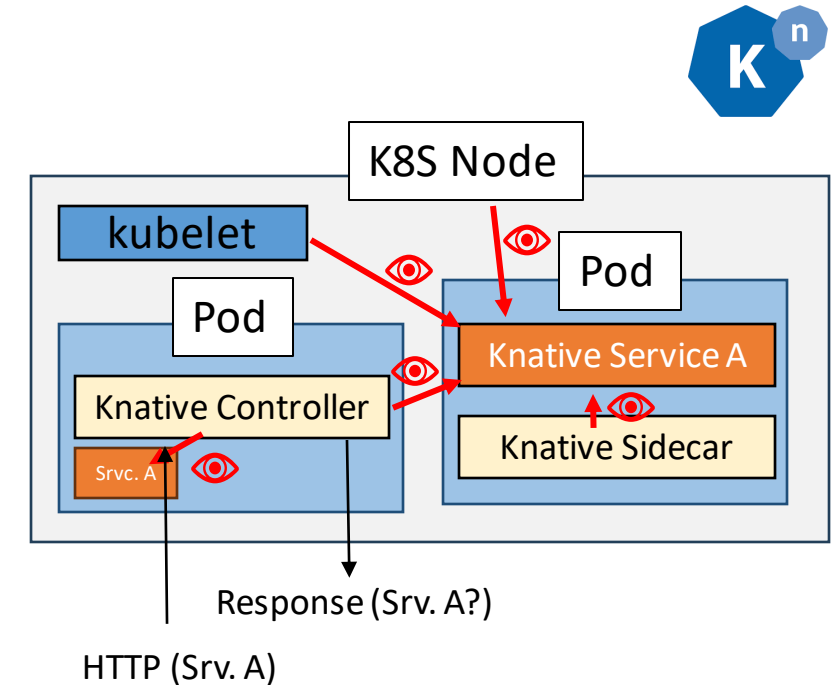


Introduction: More Problems in Serverless!

Inter-function isolation is fine, but not enough!

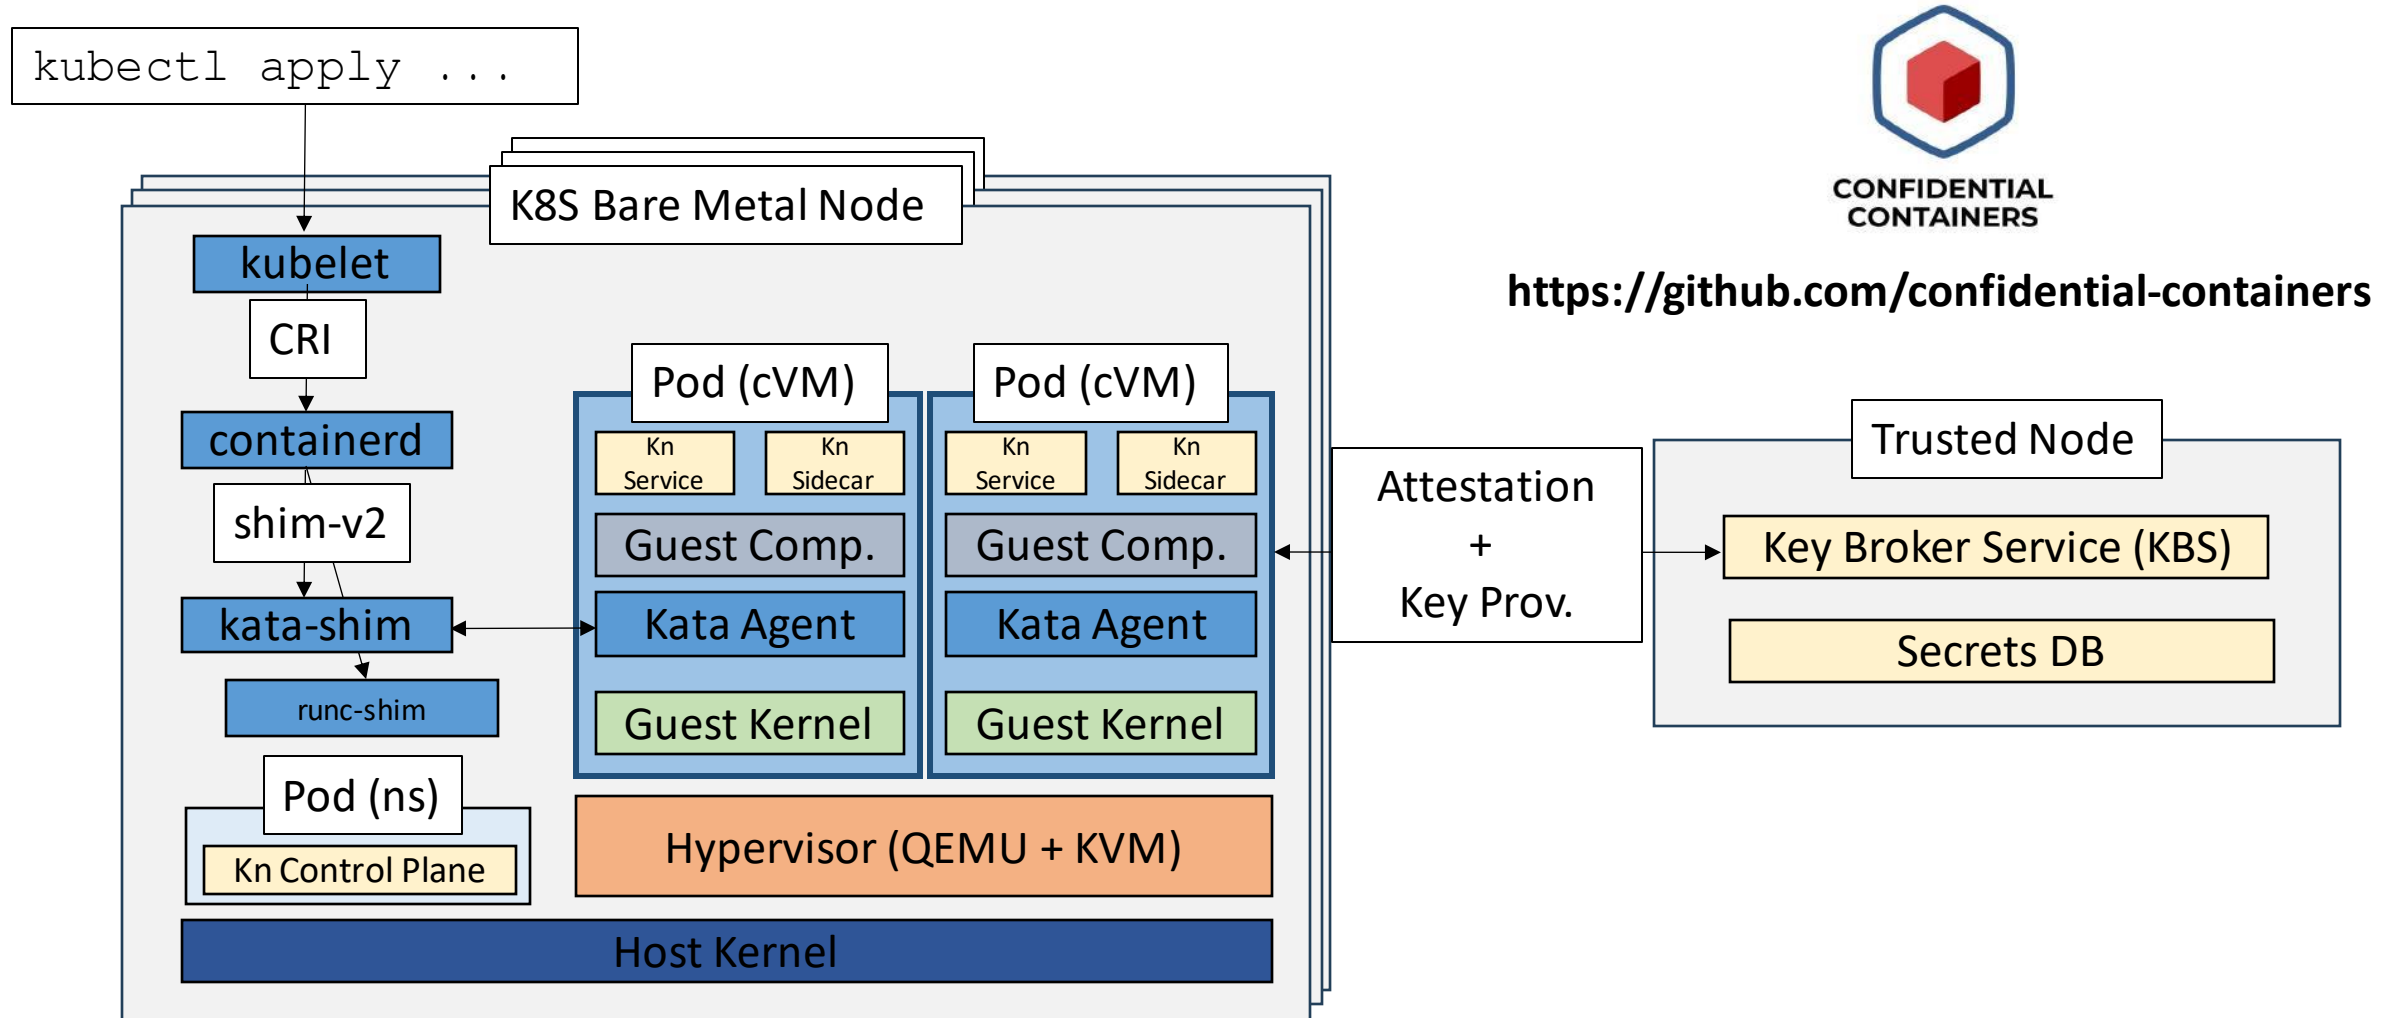
We need isolation from the host environment to guarantee...

- Data Confidentiality
- Code Confidentiality
- Execution Integrity






Confidential Computing

PoC: Knative on Confidential Containers




Evaluation

We want to evaluate the feasibility of our PoC according to the three key metrics we identified for serverless:

	1. Cold Start Times	2. Warm Start Times	3. Instantiation Throughput
 K8S RUNC	6s	1s	1 fps
 kata containers	7s	2s	0.5 fps
 CONFIDENTIAL CONTAINERS	??	??	??

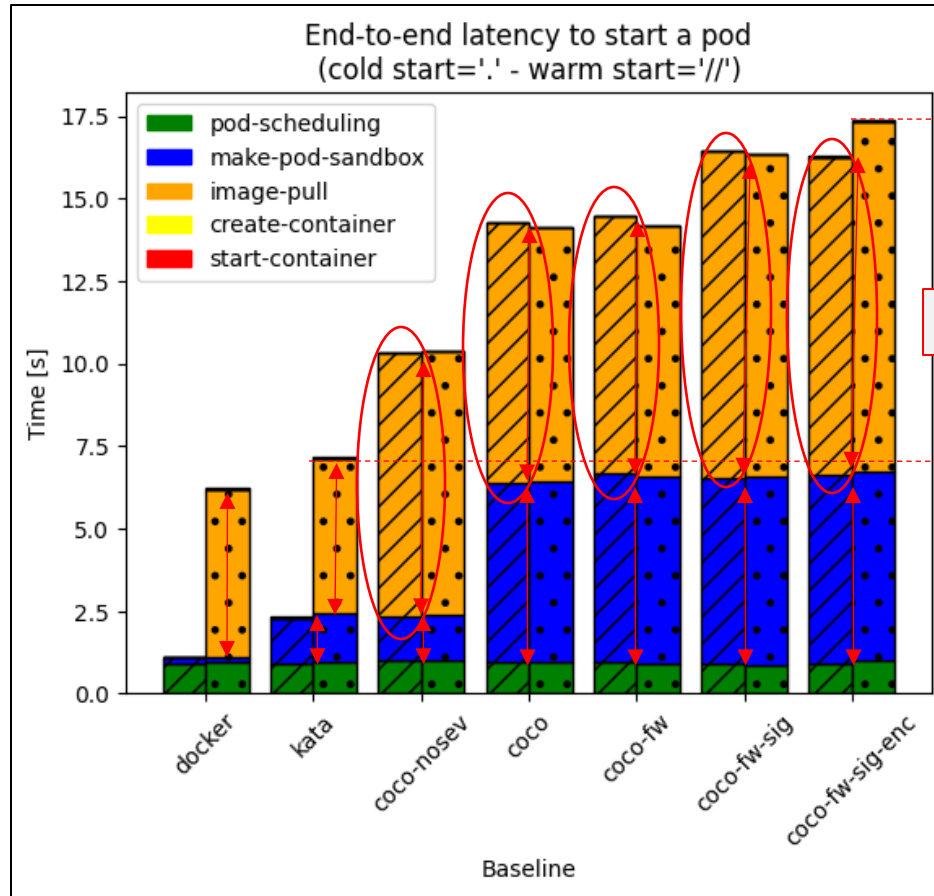
Evaluation: Baselines



- 0. **docker (i.e. runc)**: no VMs
- 1. **kata**: VMs
- 2. **coco-nosev**: + pull in guest
- 3. **coco-nosev-ovmf**: + OVMF
- 4. **coco**: + SEV
- 5. **coco-fw**: + HW att
- 6. **coco-fw-sig**: + image signature
- 7. **coco-fw-sig-enc**: + image enc.

Knative Service is a simple "Hello World" in Python

Evaluation: Cold/Warm Starts



Observations:

1. Why is VM start-up 4x slower with SEV?
2. Why is image pulling 2-3x slower w.r.t docker?
3. Why are there no warm starts?

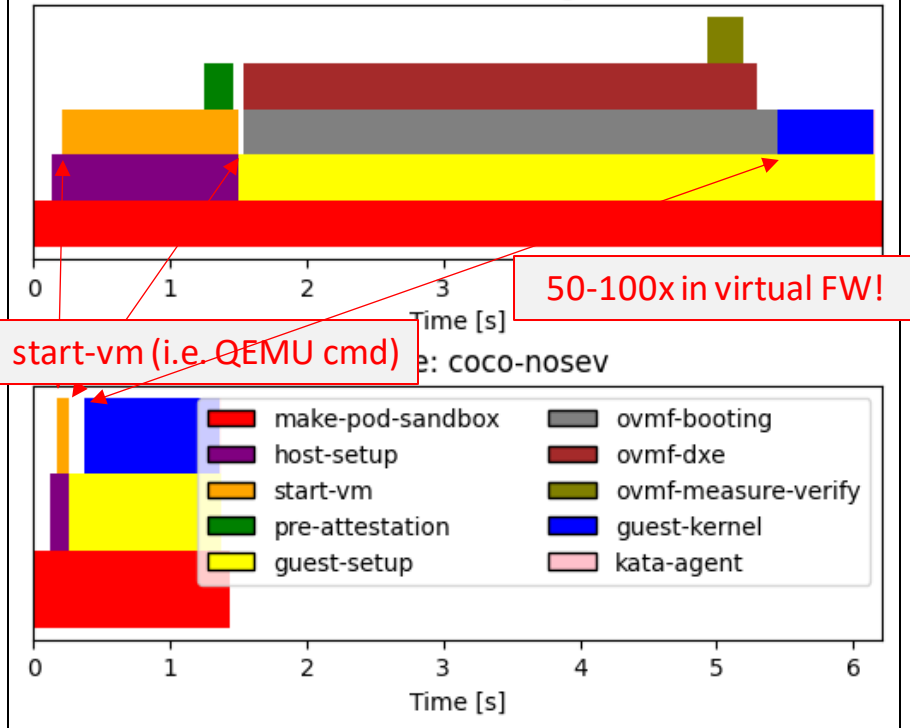
Additional 10 s!

Let us address these questions one by one

Evaluation: VM Start-Up in detail

VM Start-Up with different SEV configurations

Baseline: coco-fw-sig-enc

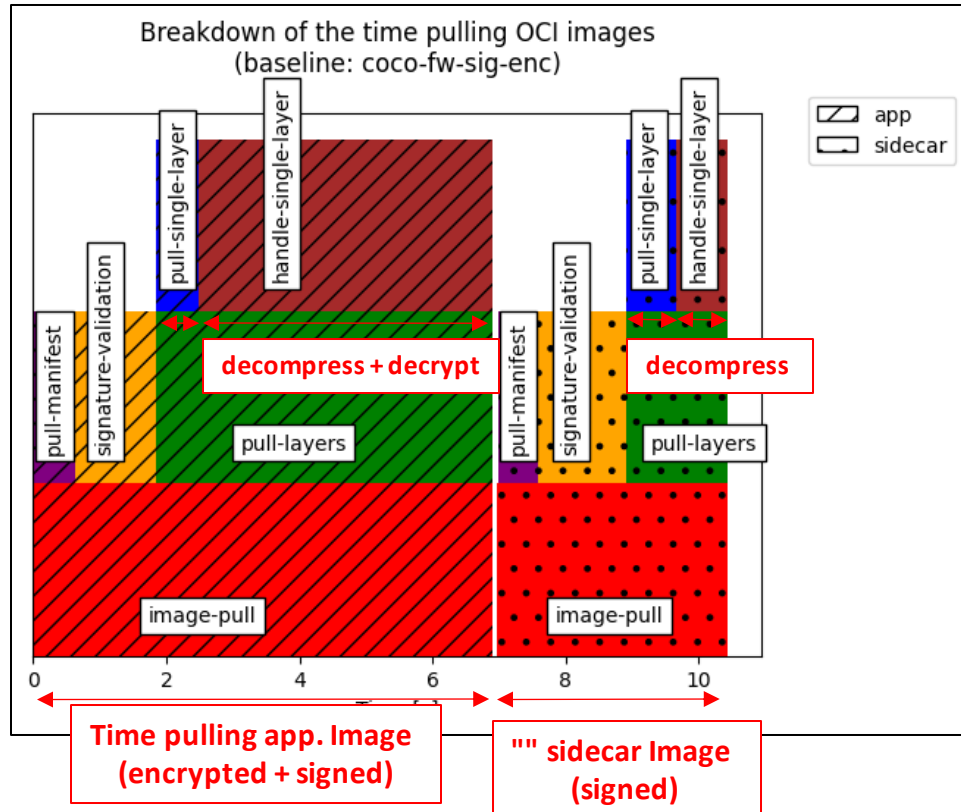


Q1: Why is VM start-up 3x slower with SEV?

start-vm: Provisioning guest memory (pages introduces 1-2 extra seconds (for 2GB of memory)

virtual-fw: OVMF DXE driver initialization introduces 3-4 extra seconds

Evaluation: VM Start-Up in detail




Q2: Why is image-pulling 2x slower w.r.t Docker?

A: containerd's PullImage becomes blocking!


A(ctd): Decrypting image layers is the bottleneck!

Evaluation: Instantiation Throughput



Starting 16 concurrent functions takes > 3' !!

Evaluation: Instantiation Throughput (ctd.)






Q: Why Starting 16 concurrent functions takes > 3'?

A: We are being throttled by the registry!

Evaluation

We want to evaluate the feasibility of our PoC according to the three key metrics we identified for serverless:

	1. Cold Start Times	2. Warm Start Times	3. Instantiation Throughput
 K8S RUNC	6s	1s	1 fps
 kata containers	7s	2s	0.5 fps
 CONFIDENTIAL CONTAINERS	17.5 s	17.5 s	~ 0.1 cps



Serverless Confidential Containers: Challenges and Opportunities

Carlos Segarra

(w/ Tobin Feldman-Fitzthum and Daniele Buono)

Large-Scale Data & Systems (LSDS) Group - Imperial College London

Visiting IBM TJ Watson (Sep'23 – Nov'23)

<https://carlossegarra.com>
<cs1620@ic.ac.uk>

