

GeoFaaS: An Edge-to-Cloud FaaS Platform

M. Malekabbasi, T. Pfandzelter, T. Schirmer, D. Bermbach | Scalable Software Systems

Edge-to-Cloud FaaS & Serverless

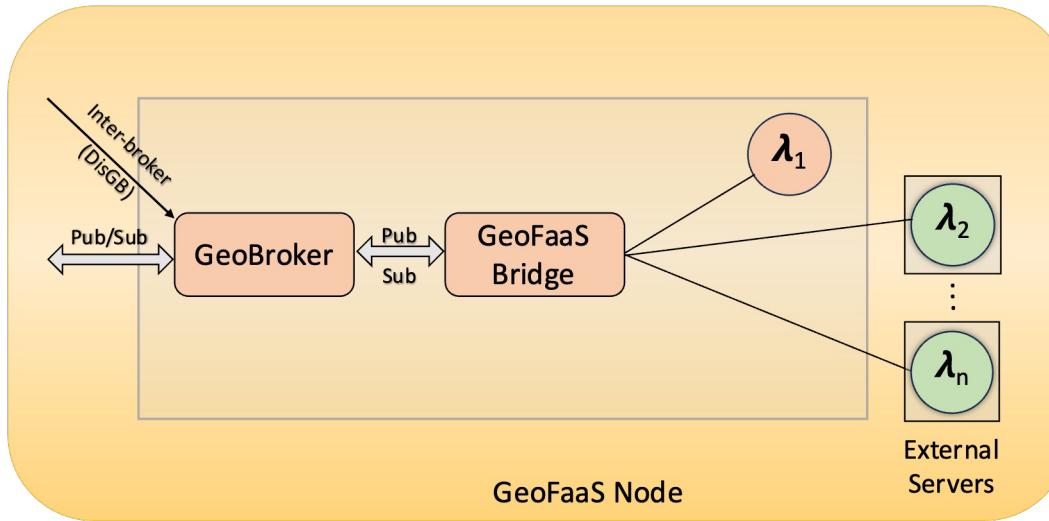
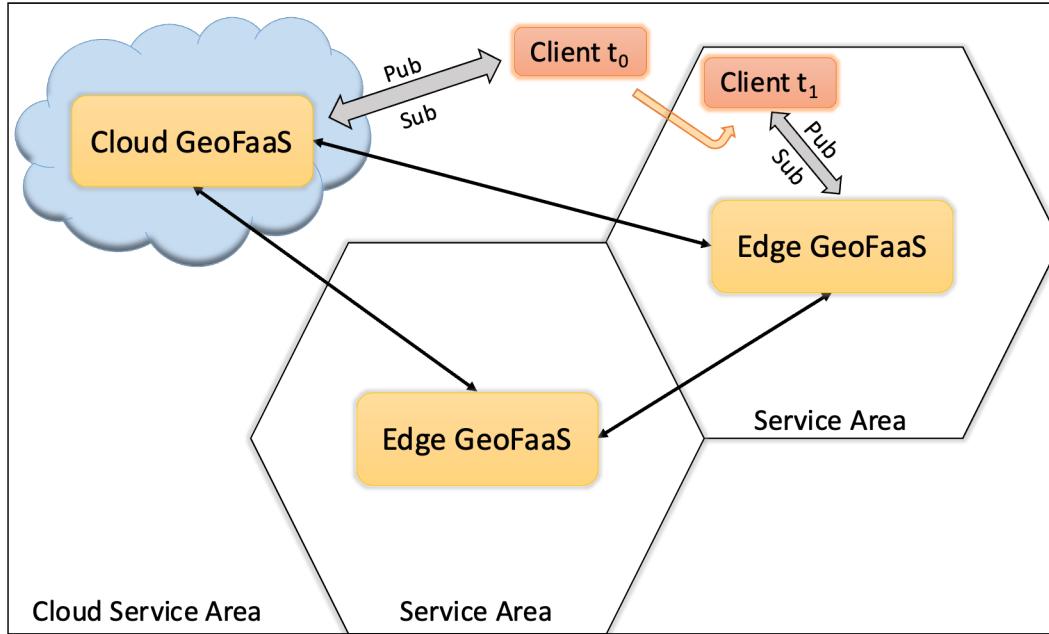
- Imagine (close) future with mobile clients
- Their characteristics:
 - Geographically-distributed
 - limited battery and computation power
 - many latency-sensitive applications
- Problem: edge-to-cloud systems complex dev & management.
- Solution: FaaS abstraction is promising. Now even stateful [11].

- BUT, we are not considering client's location!
 - In a changing network, physical distance effectively approximates latency [13], [14].
 - Client location measurement is cheap (no interaction with the system)



Geo-aware Function-as-a-Service

- A geo-distributed FaaS platform, across e2c continuum
- Transparency for end clients
 - i.e. publish on “ $f_1/call$ ” to call f_1
 - Through a client library for the common serverless abstraction
- GeoFaaS node has three key elements
 - Distributed Geo-aware Message Broker (“DisGB”)
 - FaaS server(s) (local or in the same data center)
 - Bridge, a middleware between the other two



Evaluation

➤ Prototype:

- Kotlin prototype and client library.
- Topics for a sample function f_1 (next table)

➤ Experiment setup:

- 2 *RPi*¹ Edge nodes, and one *GCP*² Cloud node (next figure)
- Each node, running tinyFaaS, GeoBroker and GeoFaaS Bridge instances
- All clients on a *Rpi*, in the same network as the edges

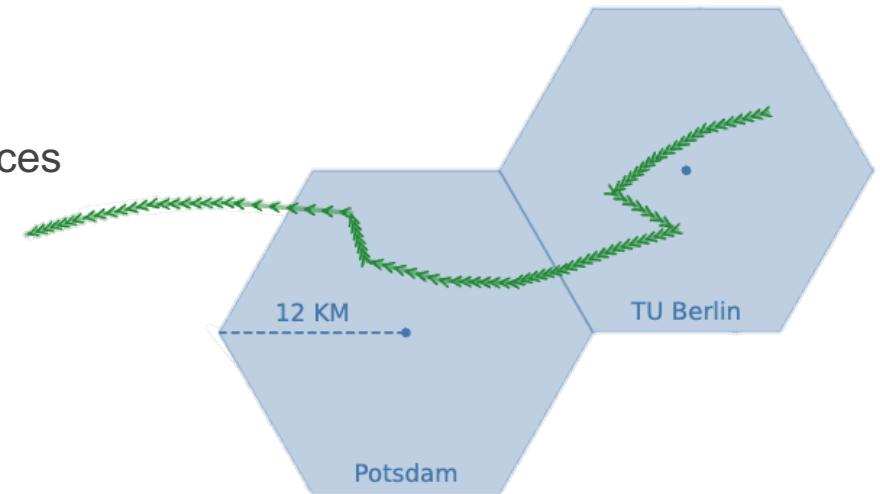
➤ Scenarios:

- 1) “Distance/Latency Change”
- 2) “High Load”, transparent offloading for uninterrupted service
- 3) “Outage”, system’s resilience against GeoFaaS Bridge failure

1. Raspberry Pi 4 B
2. Google Cloud Platform VM

TABLE I
TOPICS CREATED BY *GeoFaaS* FOR THE f_1 FUNCTION.

#	Topic	Explanation
1	/f1/call	Client calls function (Bridge subscribes)
2	/f1/ack	Bridge acknowledges call (client subscribes)
3	/f1/result	Client subscribes for result (Bridge publishes)
4	/f1/nack	Edge Bridge offloads call (Cloud subscribes)
5	/f1/call/retry	Client direct cloud call (Cloud subscribes)



Source & Experiments:
[Github.com/ChaosRez/GeoFaaS](https://github.com/ChaosRez/GeoFaaS)



Results

- FaaS's physical distance impacts response times. GeoFaaS effectively routes clients to nearest servers.
- GeoFaaS offloads requests for transparent client responses under load.
- GeoFaaS with reliable DisGB routes to cloud on edge (internal) failure.

- Full results in our paper (under review)
 - <http://arxiv.org/abs/2405.14413>

mm@3s.tu-berlin.de

References

- [3] Pfandzelter, Tobias, and David Bermbach. "tinyfaas: A lightweight faas platform for edge environments." 2020 IEEE International Conference on Fog Computing (ICFC). IEEE, 2020.
- [5] Russo, Gabriele Russo, et al. "Serverledge: Decentralized function-as-a-service for the edge-cloud continuum." 2023 IEEE International Conference on Pervasive Computing and Communications (PerCom). IEEE, 2023.
- [6] Oliveira, Bárbara, et al. "Function-as-a-Service for the Cloud-to-Thing continuum: a Systematic Mapping Study." 8th International Conference on Internet of Things, Big Data and Security-IoTBDS. 2023.
- [13] Hasenburg, Jonathan, and David Bermbach. "DisGB: Using geo-context information for efficient routing in geo-distributed pub/sub systems." 2020 IEEE/ACM 13th International Conference on Utility and Cloud Computing (UCC). IEEE, 2020.
- [18] Ciavotta, Michele, et al. "DFaaS: Decentralized function-as-a-service for federated edge computing." 2021 IEEE 10th International Conference on Cloud Networking (CloudNet). IEEE, 2021.