


MARCIN COPIK, ALEXANDRU CALOTOIU, GYORGY RETHY, ROMAN BÖHRINGER, RODRIGO BRUNO, TORSTEN HOEFER



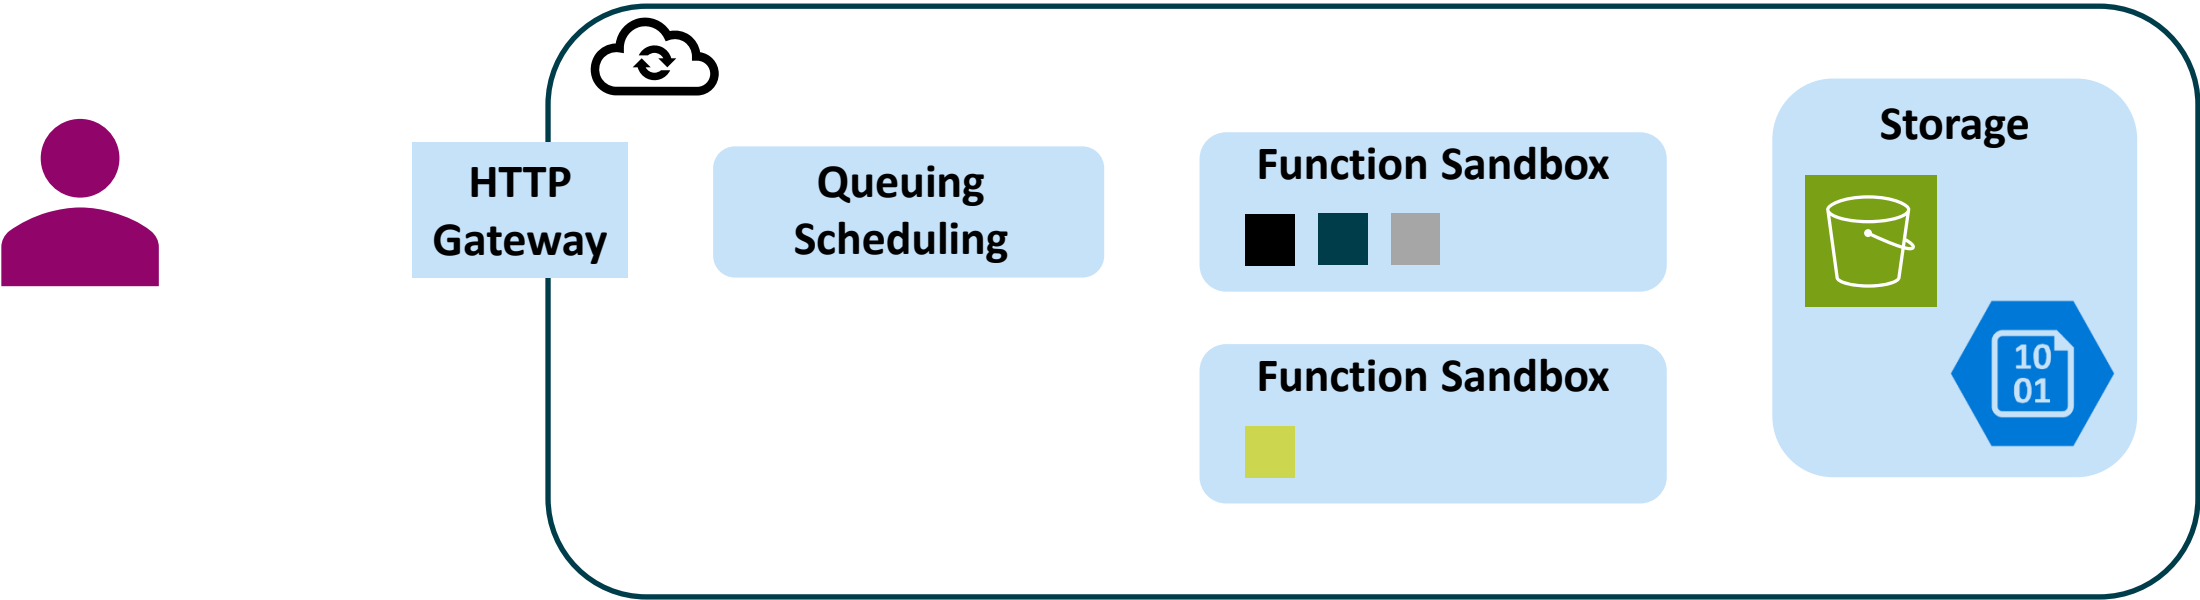
ACM Symposium  
on Cloud Computing

# Process-as-a-Service: Unifying Elastic and Stateful Clouds with Serverless Processes



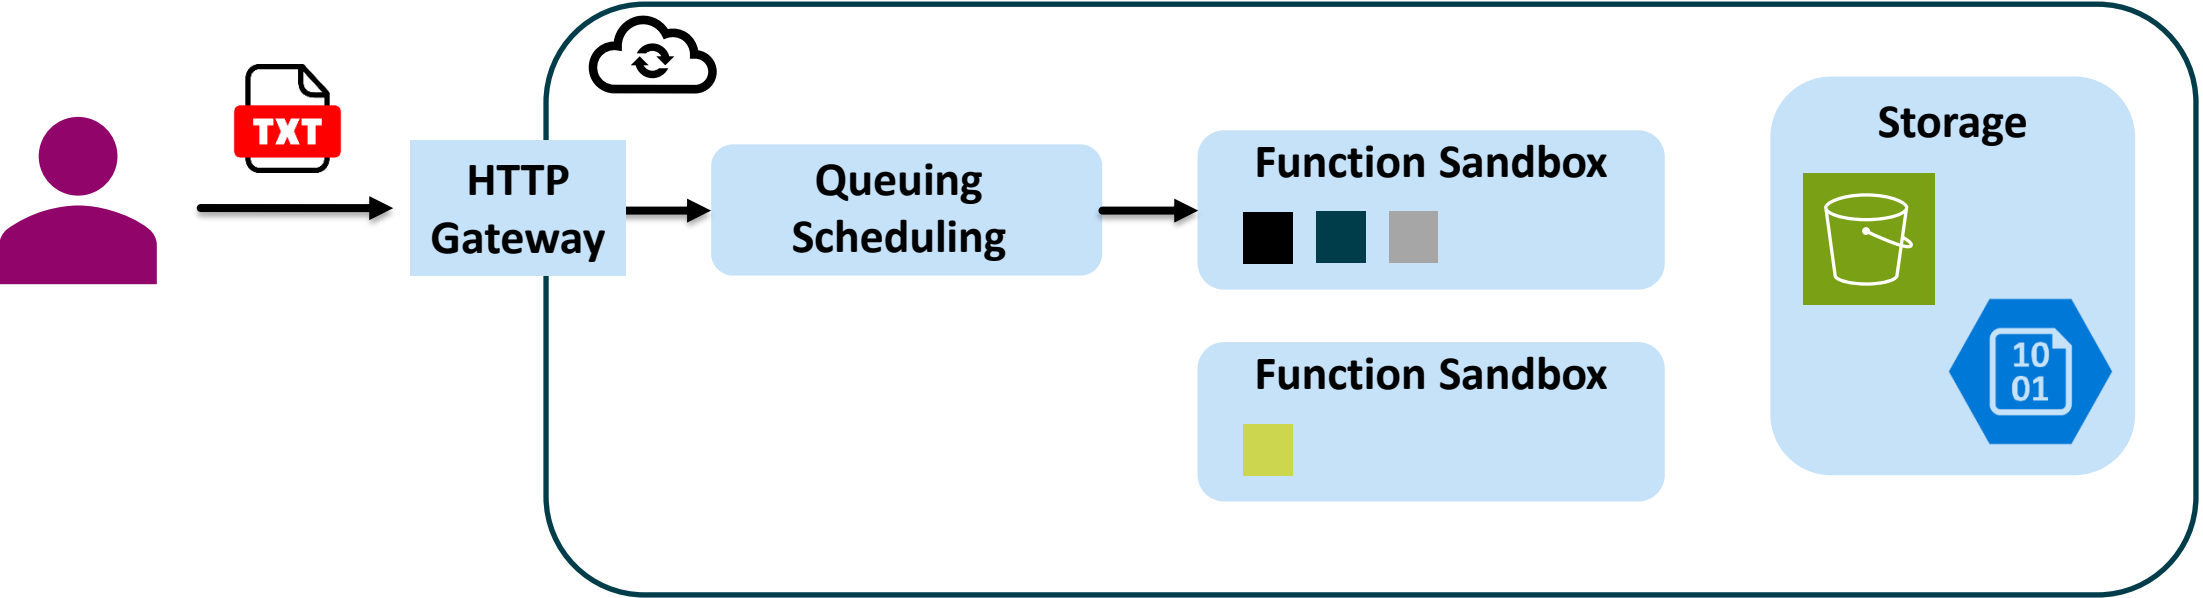
TÉCNICO  
LISBOA

# LaTeX Microservice in Serverless



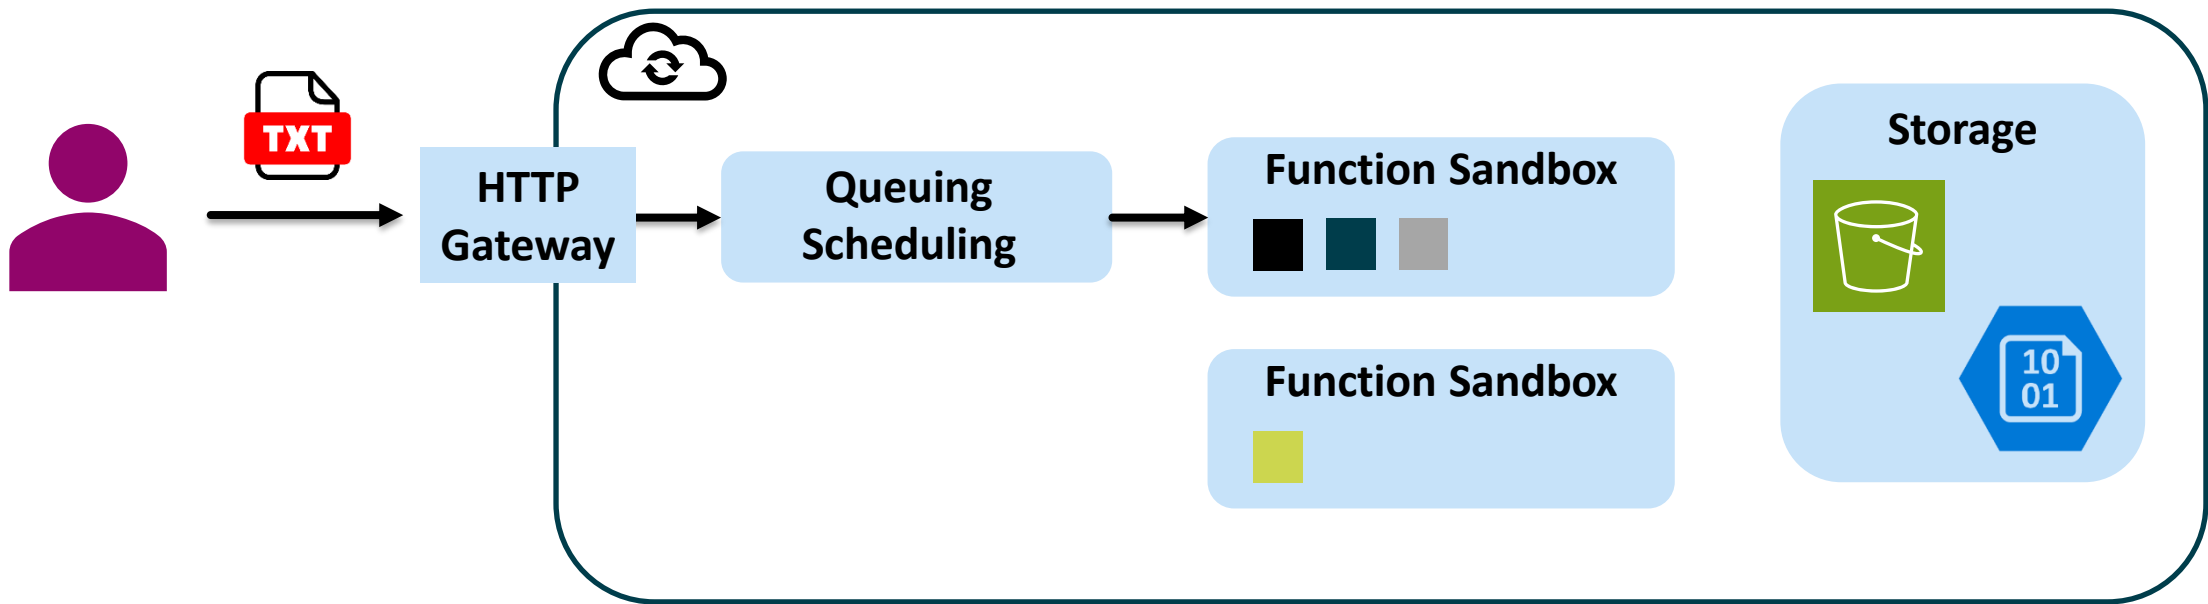


# LaTeX Microservice in Serverless



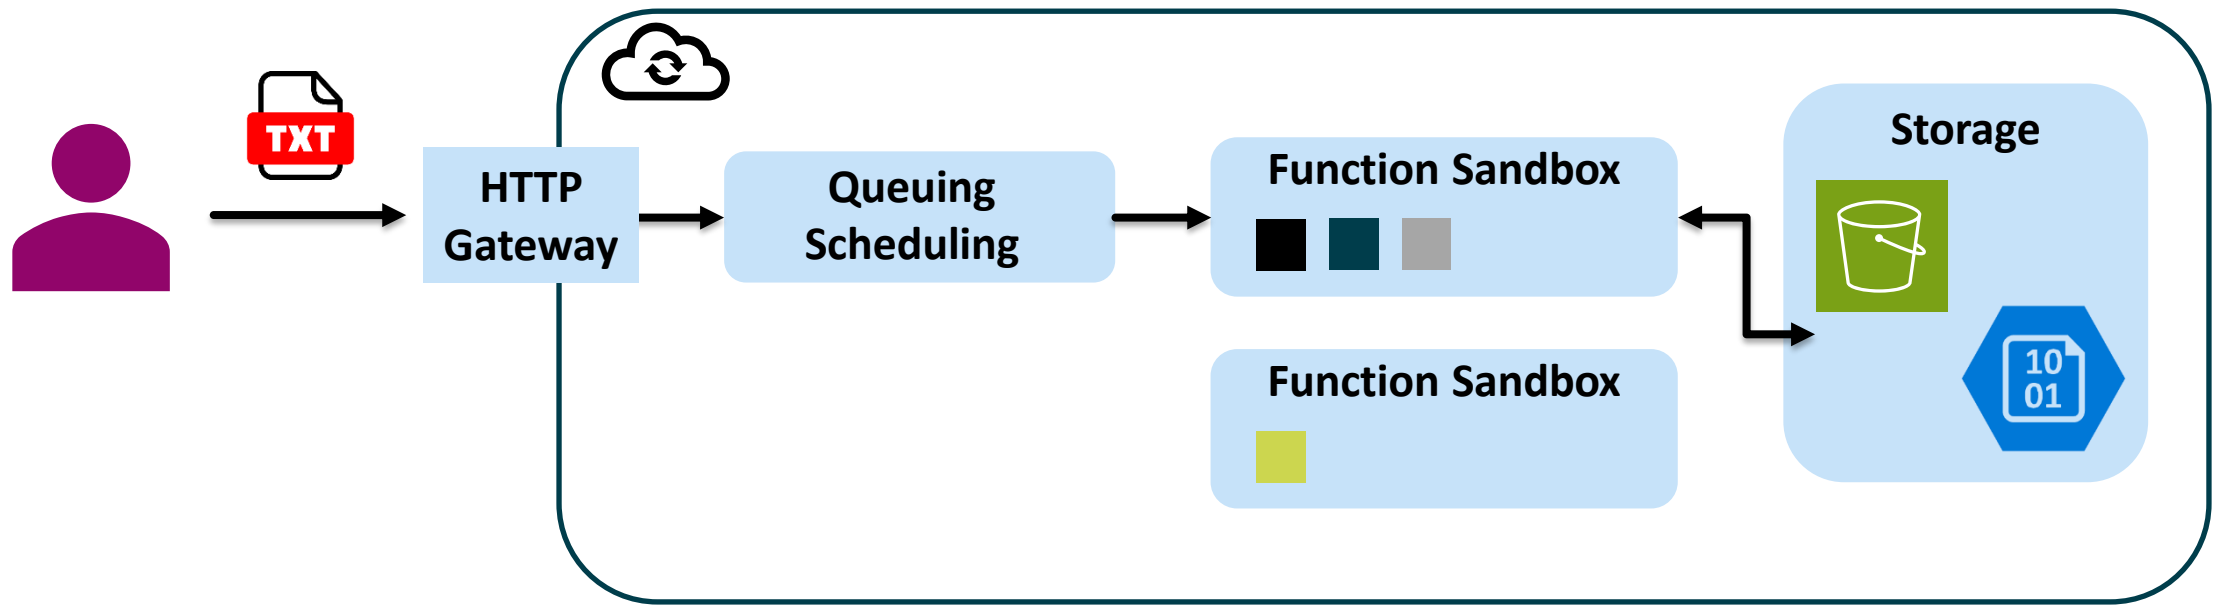
# LaTeX Microservice in Serverless

Is cache up to date?  
Did other function edit file of this client?




# LaTeX Microservice in Serverless

Is cache up to date?  
Did other function edit file of this client?



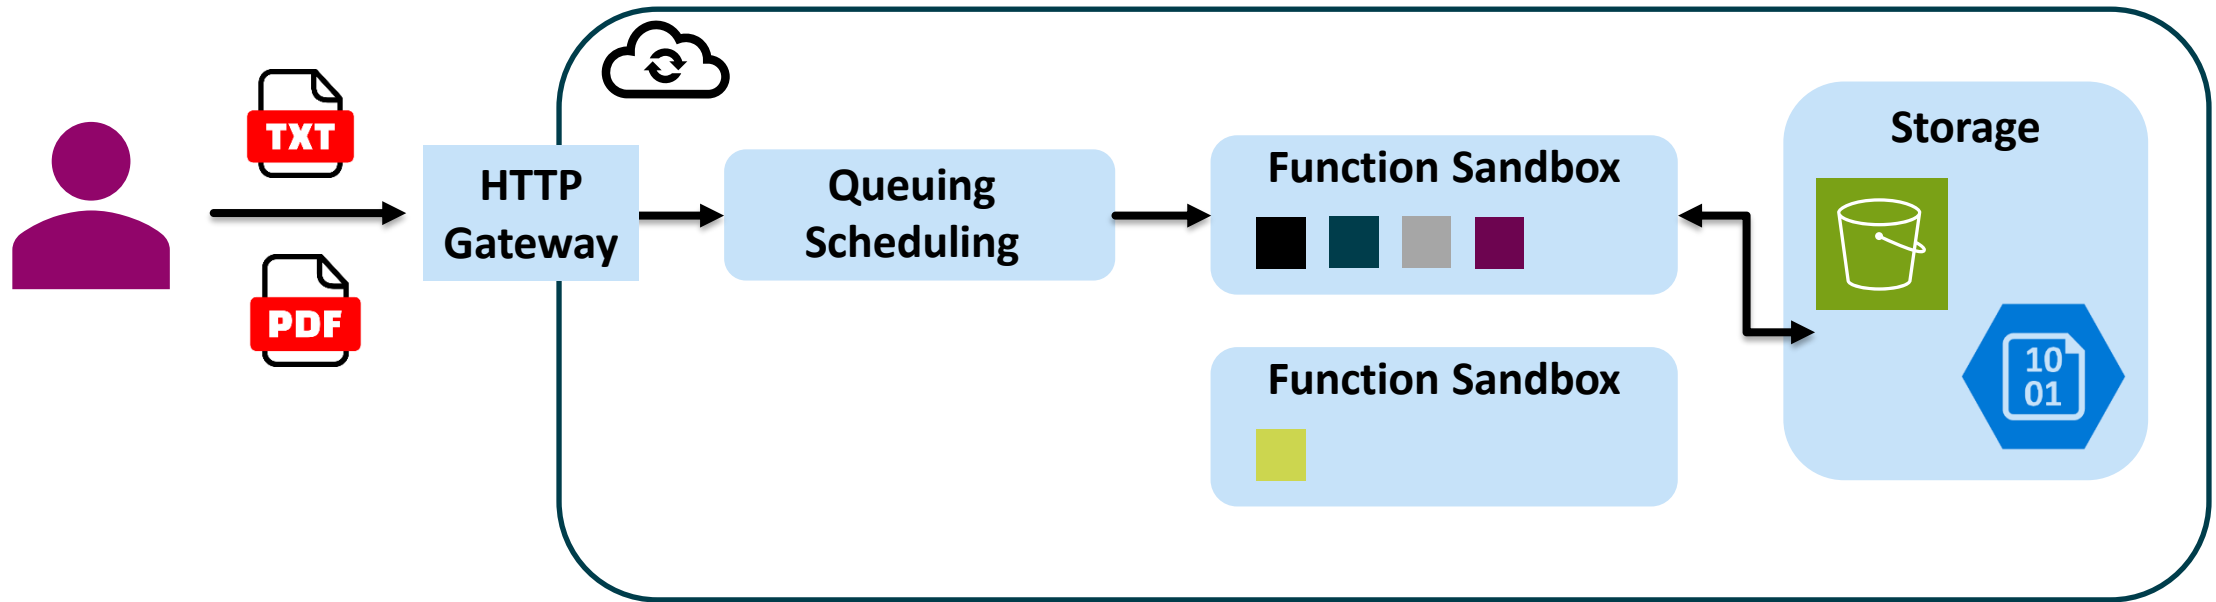
# LaTeX Microservice in Serverless



## Serverless Service

✗ Cached function state is not reliable

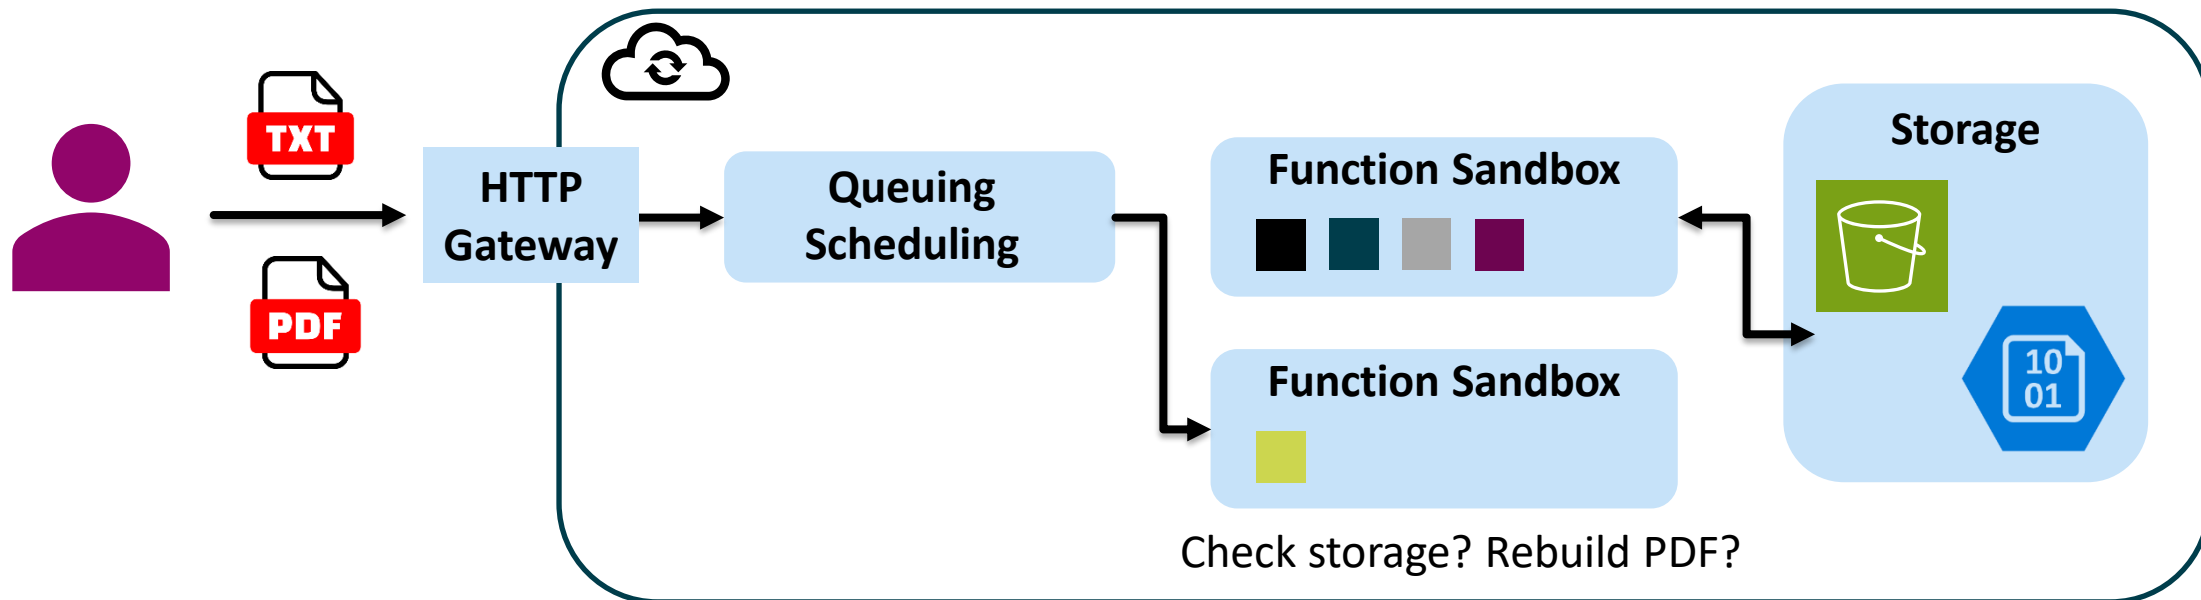
# LaTeX Microservice in Serverless



## Serverless Service

✗ Cached function state is not reliable

# LaTeX Microservice in Serverless

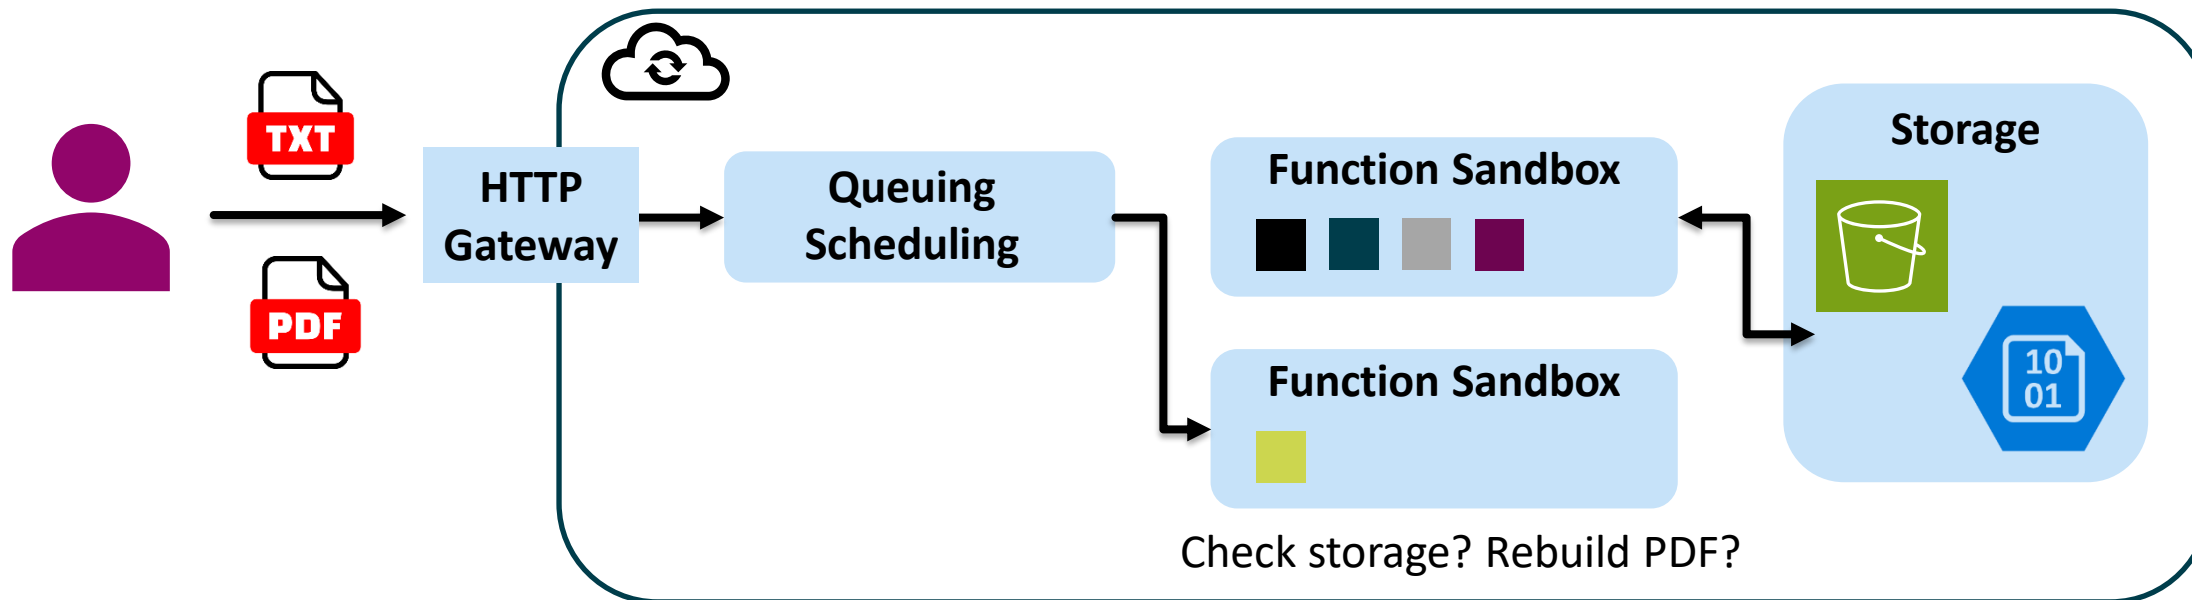


## Serverless Service

✗ Cached function state is not reliable



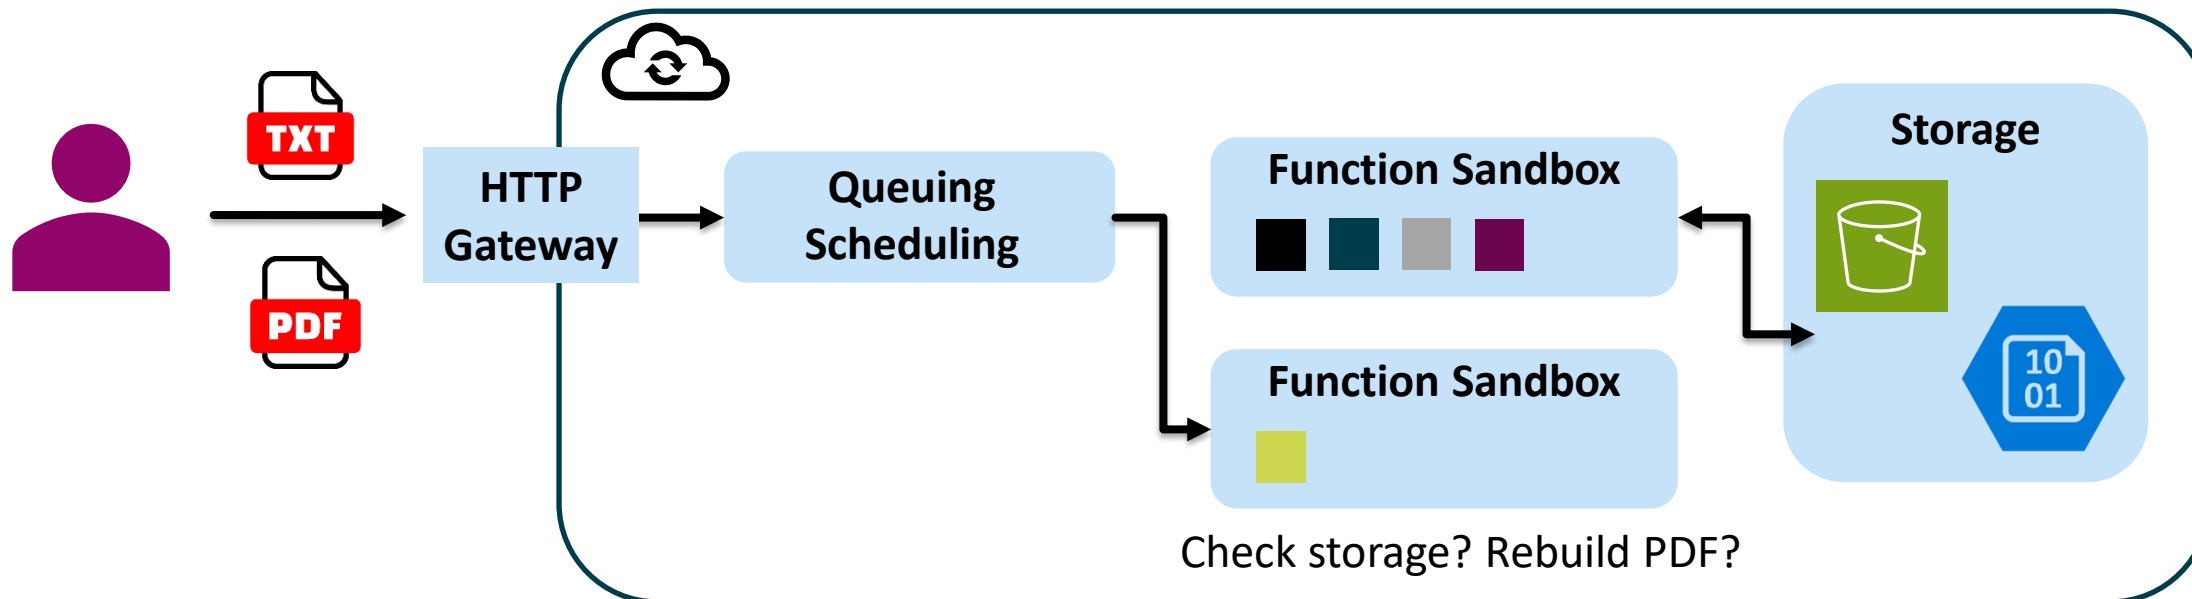
# LaTeX Microservice in Serverless



## Serverless Service

- ❌ Cached function state is not reliable
- ❌ Request context has no influence on scheduling

# LaTeX Microservice in Serverless



## Serverless Service

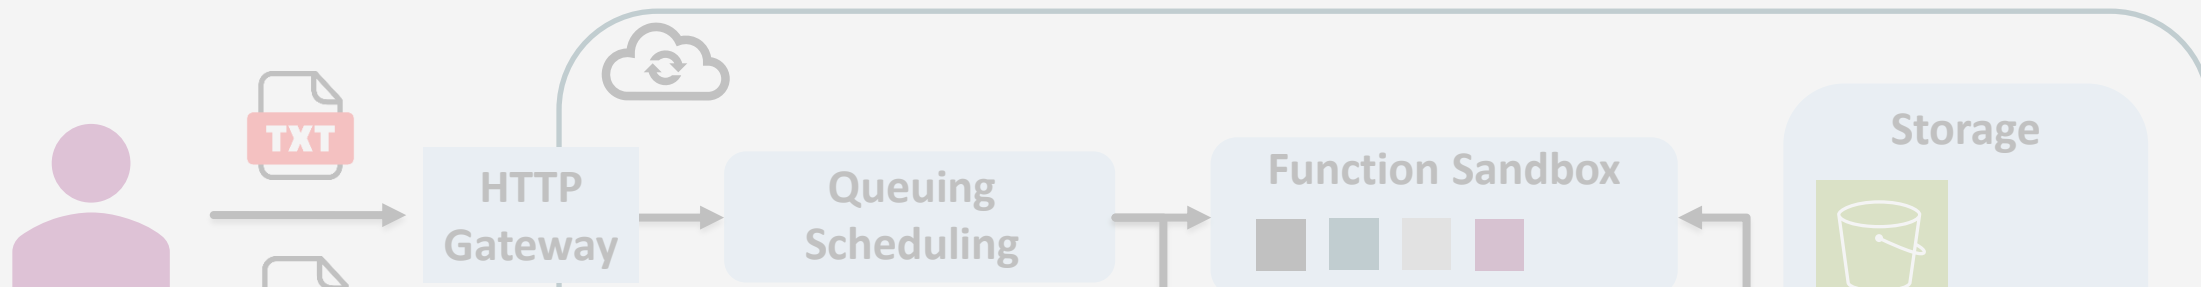
- ✗ Cached function state is not reliable
- ✗ Request context has no influence on scheduling

## Standard Microservice

- ✓ Extensive caching
- ✓ Load balancing with sticky sessions

# LaTeX Microservice in Serverless

Is cache up to date?  
Did other function edit file of this client?



Should we fix this problem with dedicated optimizations?

Maybe functions are not enough and need a better abstraction?


## Serverless Service

- ✗ Cached function state is not reliable
- ✗ Request context has no influence on scheduling

## Standard Microservice

- ✓ Extensive caching
- ✓ Load balancing with sticky sessions


# Cloud Models



# Cloud Models

## IaaS

## CaaS



**Persistent**

**Persistent**



**Provisioned**

**Provisioned**



**Persistent,  
local**

**Persistent,  
local**



**Months**

**Hours, days**

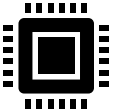






**Manual,  
minutes**

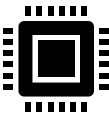




**Semi-automatic,  
minutes & seconds**



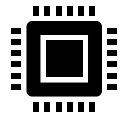




# Cloud Models

	IaaS	CaaS	FaaS
	Persistent	Persistent	Ephemeral
	Provisioned	Provisioned	Pay-as-you-go
	Persistent, local	Persistent, local	Ephemeral, remote
	Months	Hours, days	Minutes, seconds
	Manual, minutes	Semi-automatic, minutes & seconds	Automatic, seconds

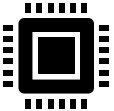




# Cloud Models

	IaaS	CaaS	FaaS
	Persistent	Persistent	Ephemeral
	Provisioned	Provisioned	Pay-as-you-go
	Persistent, local	Persistent, local	Ephemeral, remote
	Months	Hours, days	Minutes, seconds
	Manual, minutes	Semi-automatic, minutes & seconds	Automatic, seconds

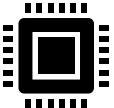




# Cloud Models

	IaaS	CaaS	FaaS
	Persistent	Persistent	Ephemeral
	Provisioned	Provisioned	Pay-as-you-go
	Persistent, local	Persistent, local	Ephemeral, remote
	Months	Hours, days	Minutes, seconds
	Manual, minutes	Semi-automatic, minutes & seconds	Automatic, seconds

# Cloud Process

	IaaS	CaaS	FaaS
	Persistent	Persistent	Ephemeral
	Provisioned	Provisioned	Pay-as-you-go
	Persistent, local	Persistent, local	Ephemeral, remote
	Months	Hours, days	Minutes, seconds
	Manual, minutes	Semi-automatic, minutes & seconds	Automatic, seconds

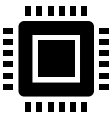




# Cloud Process

	IaaS	CaaS	PaaS	FaaS
	Persistent	Persistent		Ephemeral
	Provisioned	Provisioned		Pay-as-you-go
	Persistent, local	Persistent, local		Ephemeral, remote
	Months	Hours, days		Minutes, seconds
	Manual, minutes	Semi-automatic, minutes & seconds		Automatic, seconds



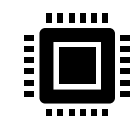




# Cloud Process

Durable state of a container


	IaaS	CaaS	PraaS	FaaS
	Persistent	Persistent		Ephemeral
	Provisioned	Provisioned		Pay-as-you-go
	Persistent, local	Persistent, local	Persistent, local	Ephemeral, remote
	Months	Hours, days	Minutes, hours	Minutes, seconds
	Manual, minutes	Semi-automatic, minutes & seconds	Automatic, seconds	Automatic, seconds

# Cloud Process


Durable state of a container and **ephemeral compute** of serverless.

	IaaS	CaaS	PaaS	FaaS
	Persistent	Persistent	Ephemeral	Ephemeral
	Provisioned	Provisioned	Pay-as-you-go	Pay-as-you-go
	Persistent, local	Persistent, local	Persistent, local	Ephemeral, remote
	Months	Hours, days	Minutes, hours	Minutes, seconds
	Manual, minutes	Semi-automatic, minutes & seconds	Automatic, seconds	Automatic, seconds

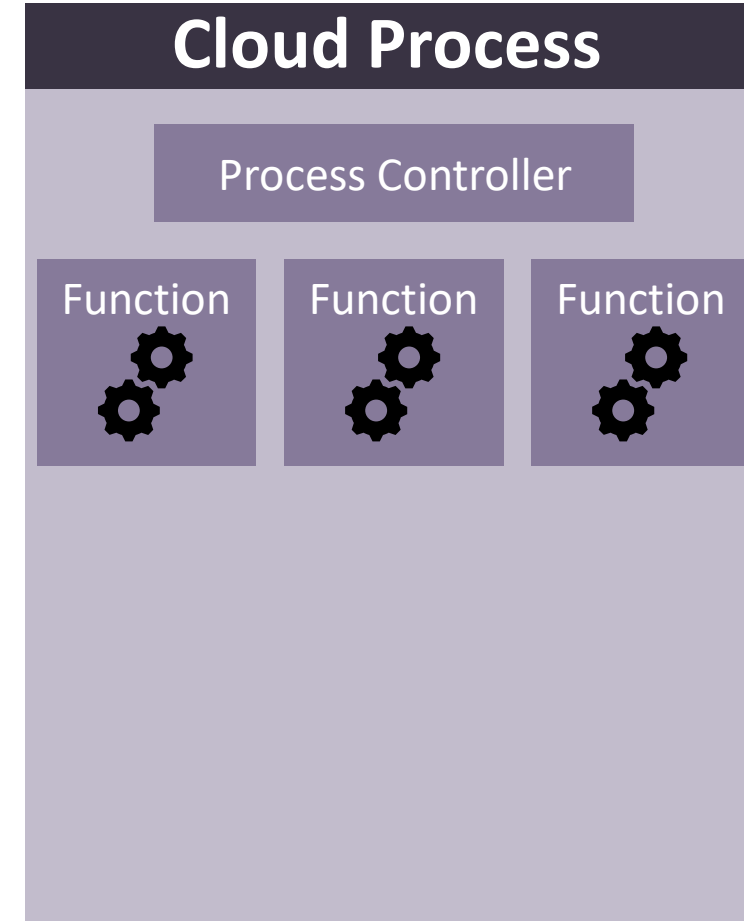
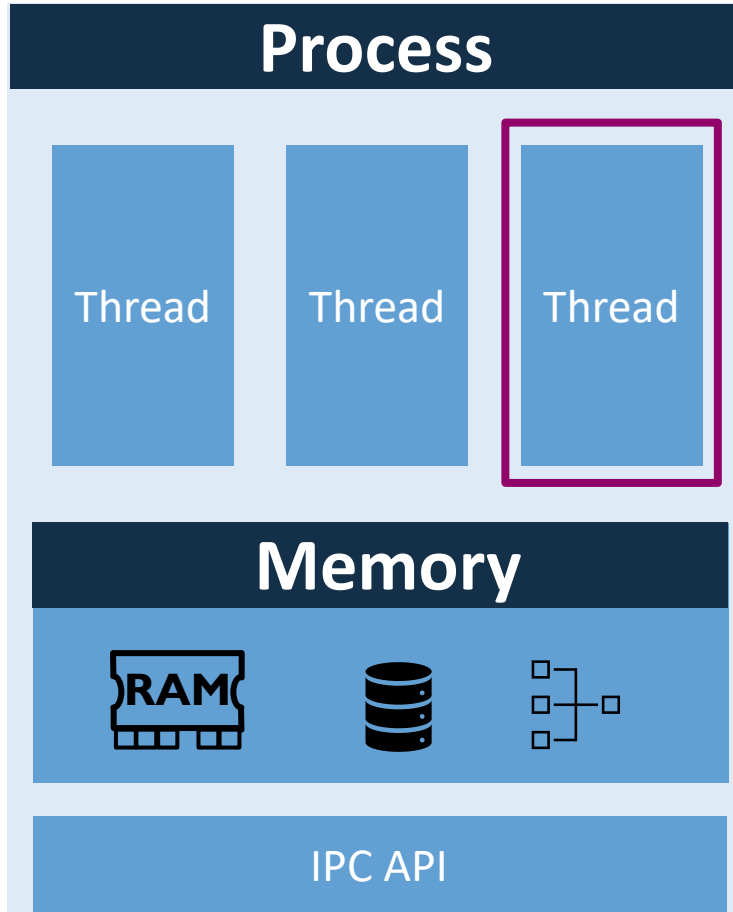
# Cloud Process Model



# Cloud Process Model



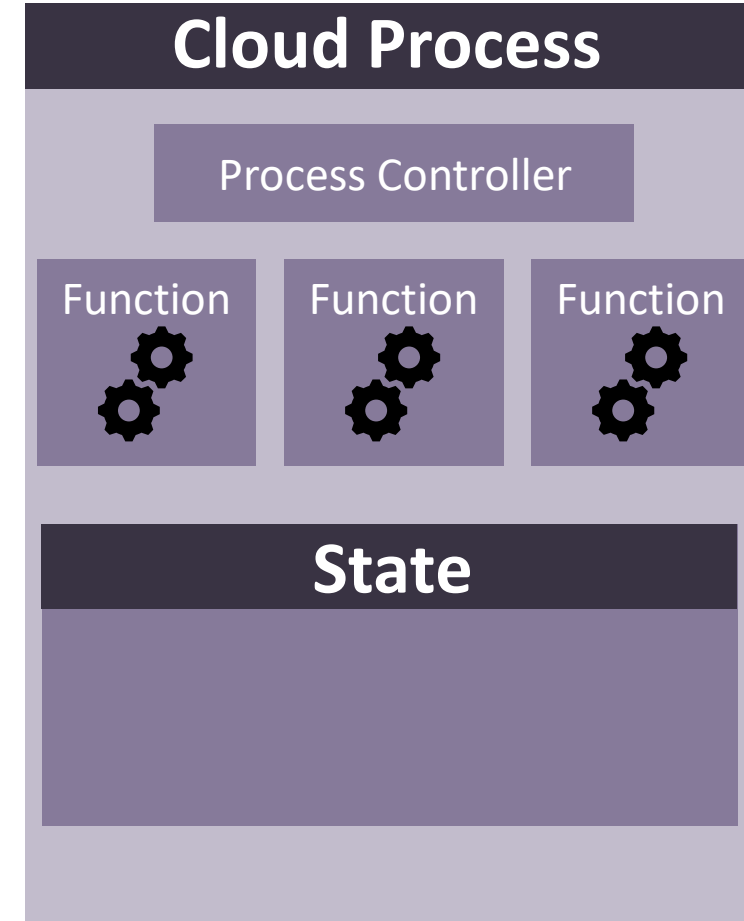

# Cloud Process Model



**Functions** represent computational threads within process.





# Cloud Process Model





**State** is represented as a collection of generic objects

# Cloud Process Model




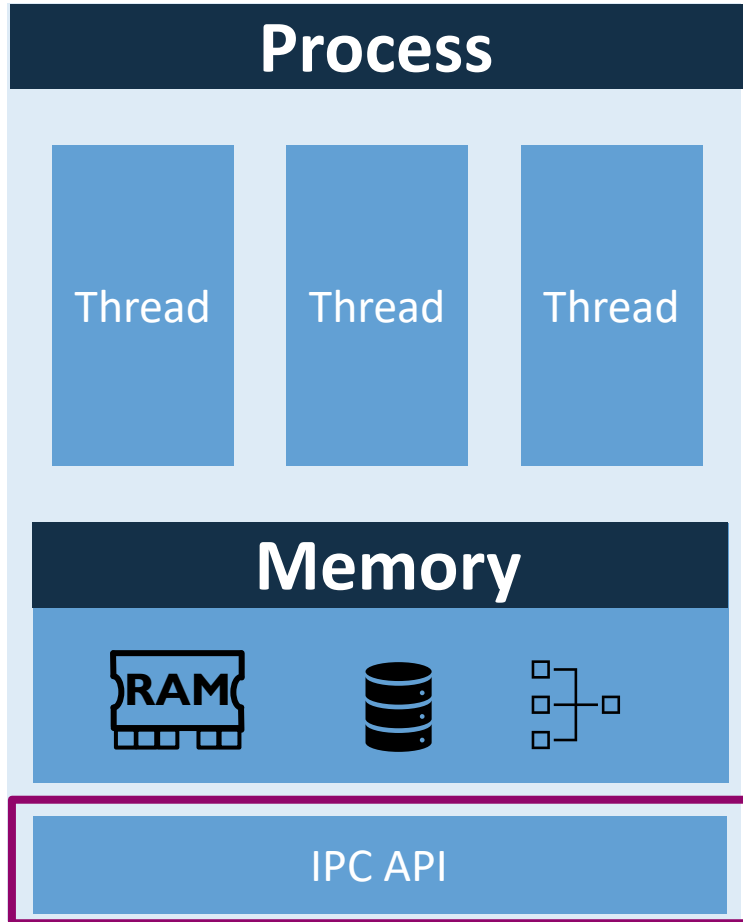
**State** is represented as a collection of generic objects

# Cloud Process Model




**Channels** are an abstraction for actual data movement.

# Cloud Process Model




**Channels** are an abstraction for actual data movement.

# Inter-Process Communication






# Inter-Process Communication




# Inter-Process Communication




```
invoke(pid_destination, function, data)
```

# Inter-Process Communication



```
invoke(pid_destination, function, data)
```


# Inter-Process Communication



```
invoke(pid_destination, function, data)
```

```
state(key, data)  
data = state(key)
```


# Inter-Process Communication



```
invoke(pid_destination, function, data)
```

```
state(key, data)  
data = state(key)
```

# Inter-Process Communication




```
invoke(pid_destination, function, data)
```


```
state(key, data)  
data = state(key)
```

```
send(pid_destination, key, data)  
data = recv(pid_source, key)
```

# Benchmark: LaTeX Microservice




# Benchmark: LaTeX Microservice





# Benchmark: LaTeX Microservice





**Marcin Copik**  
ETH Zurich



**Alexandru Calotoiu**  
ETH Zurich



**Gyorgy Rethy**  
Oracle Labs



**Roman Böhringer**  
OpenCore GmbH





**Rodrigo Bruno**  
INESC-ID




**Torsten Hoefler**  
ETH Zurich

## More of SPCL's research:

 [youtube.com/@spcl](https://youtube.com/@spcl) **210+ Talks**

 [twitter.com/spcl\\_eth](https://twitter.com/spcl_eth) **1.6K+ Followers**

 [github.com/spcl](https://github.com/spcl) **5,5K+ Stars**

... or [spcl.ethz.ch](https://spcl.ethz.ch)



**spcl/PraaS**



**SoCC'24  
Paper**